

# Anti-Pattern

# Reinventing the wheel



# Dare To Be Different



# Monolith database



# Hardcoded hell



# Manual Configurations Management



# Flying Blind



# Unknown caller



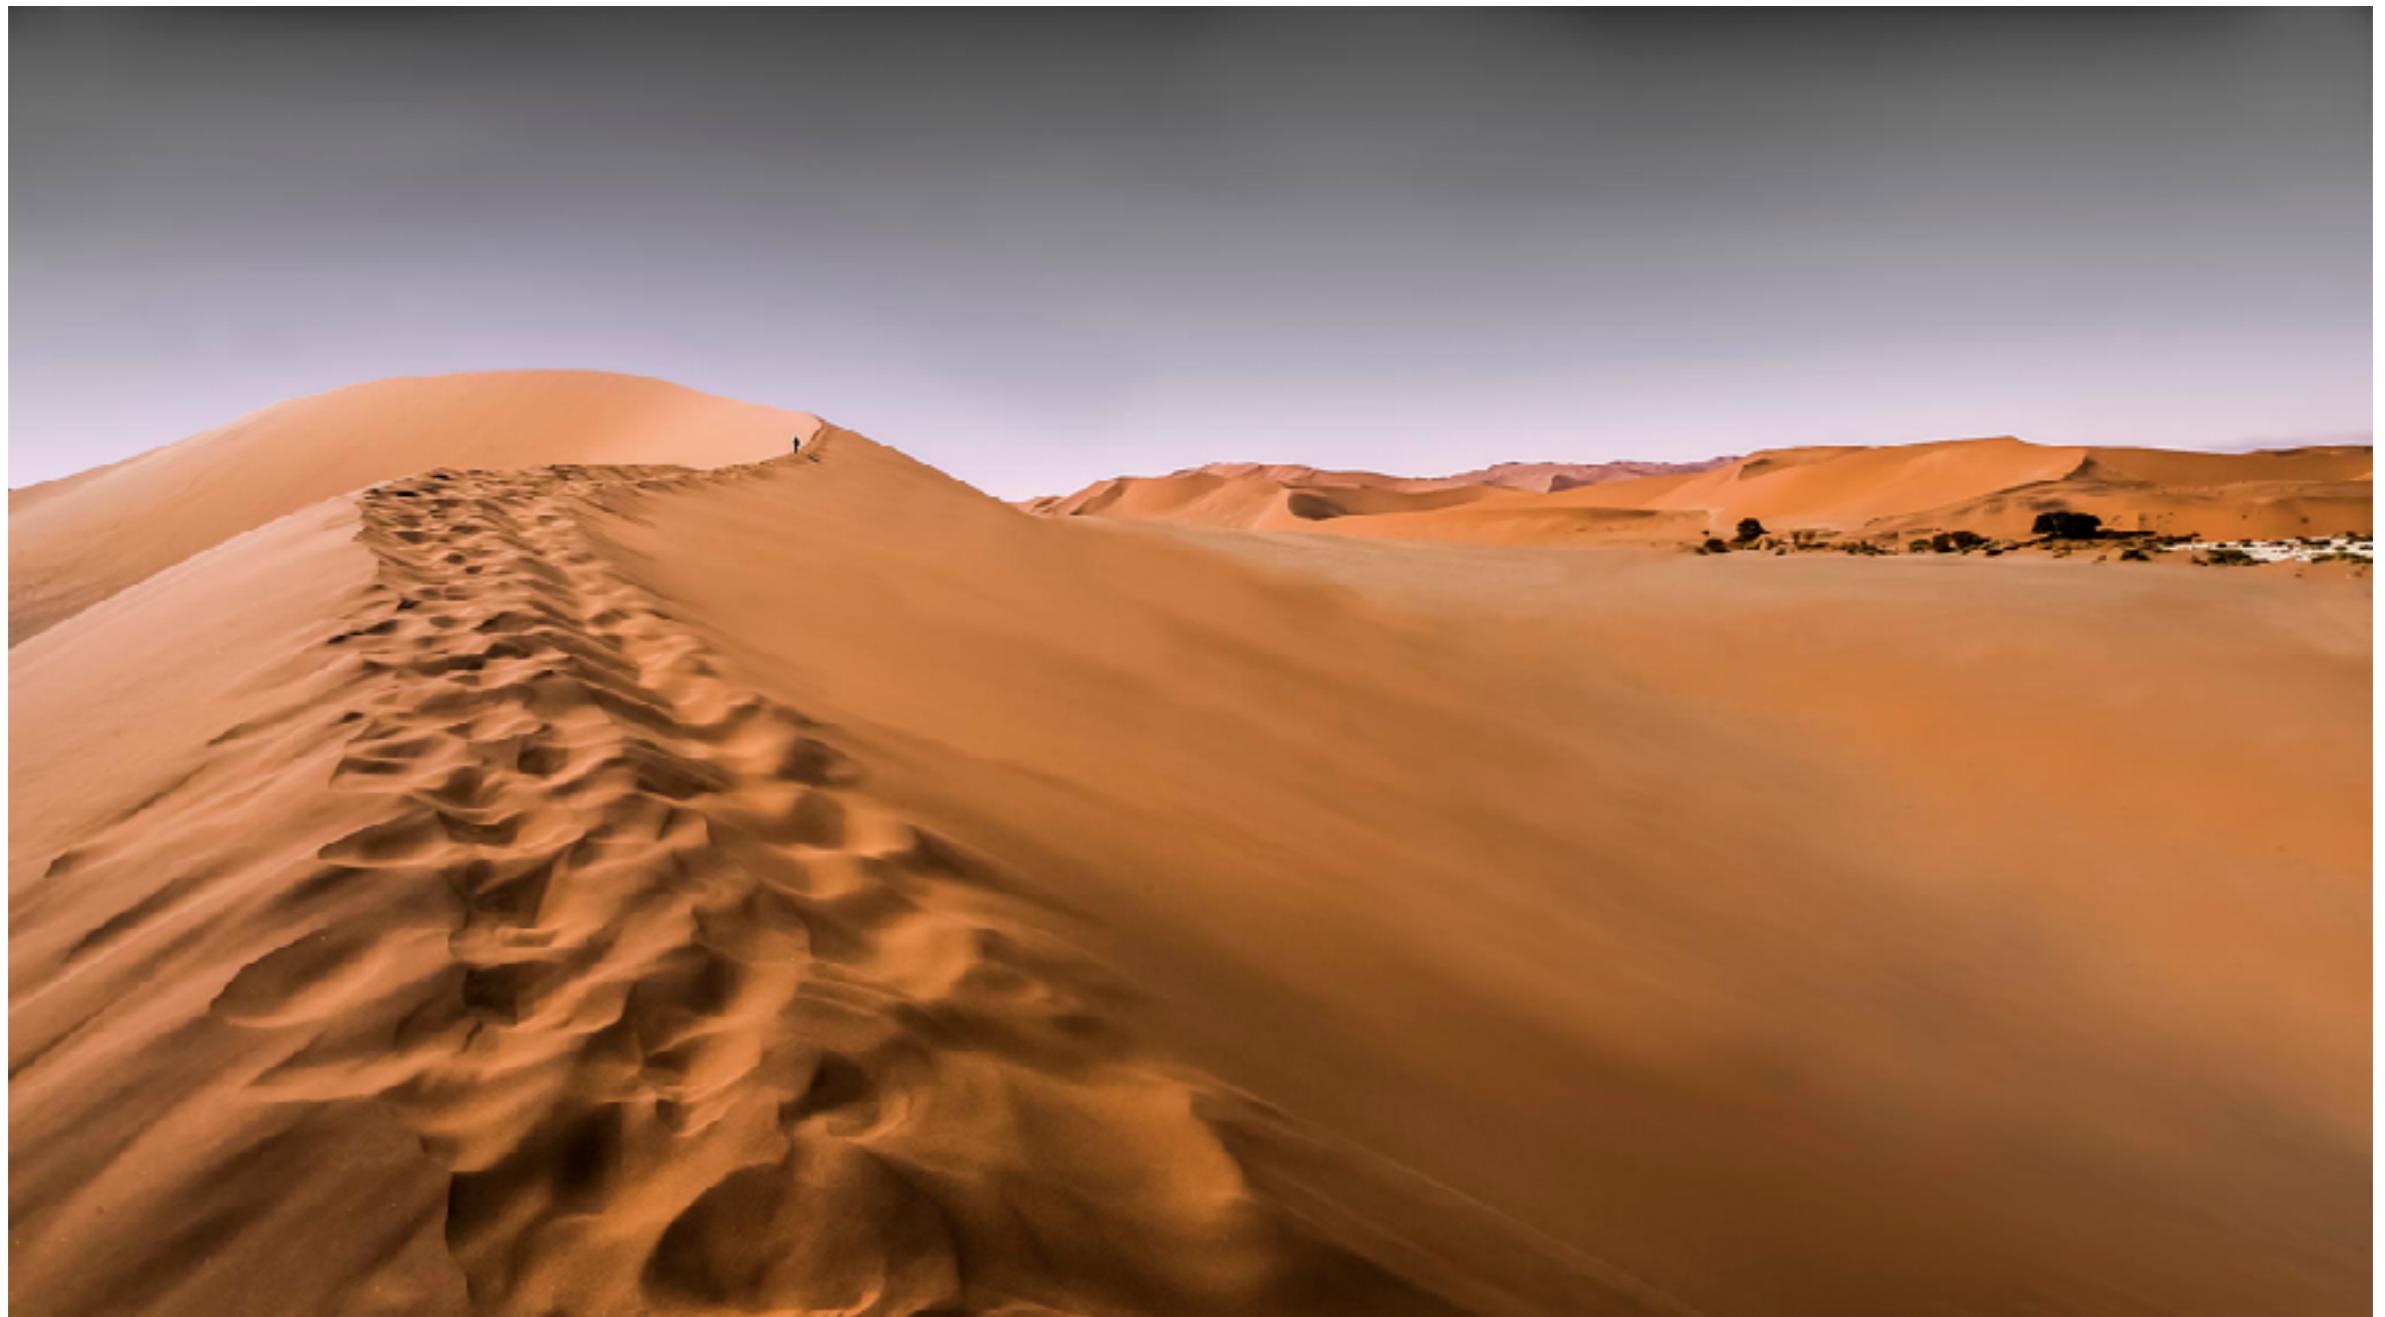
# Dogpiles



# Synchronous world



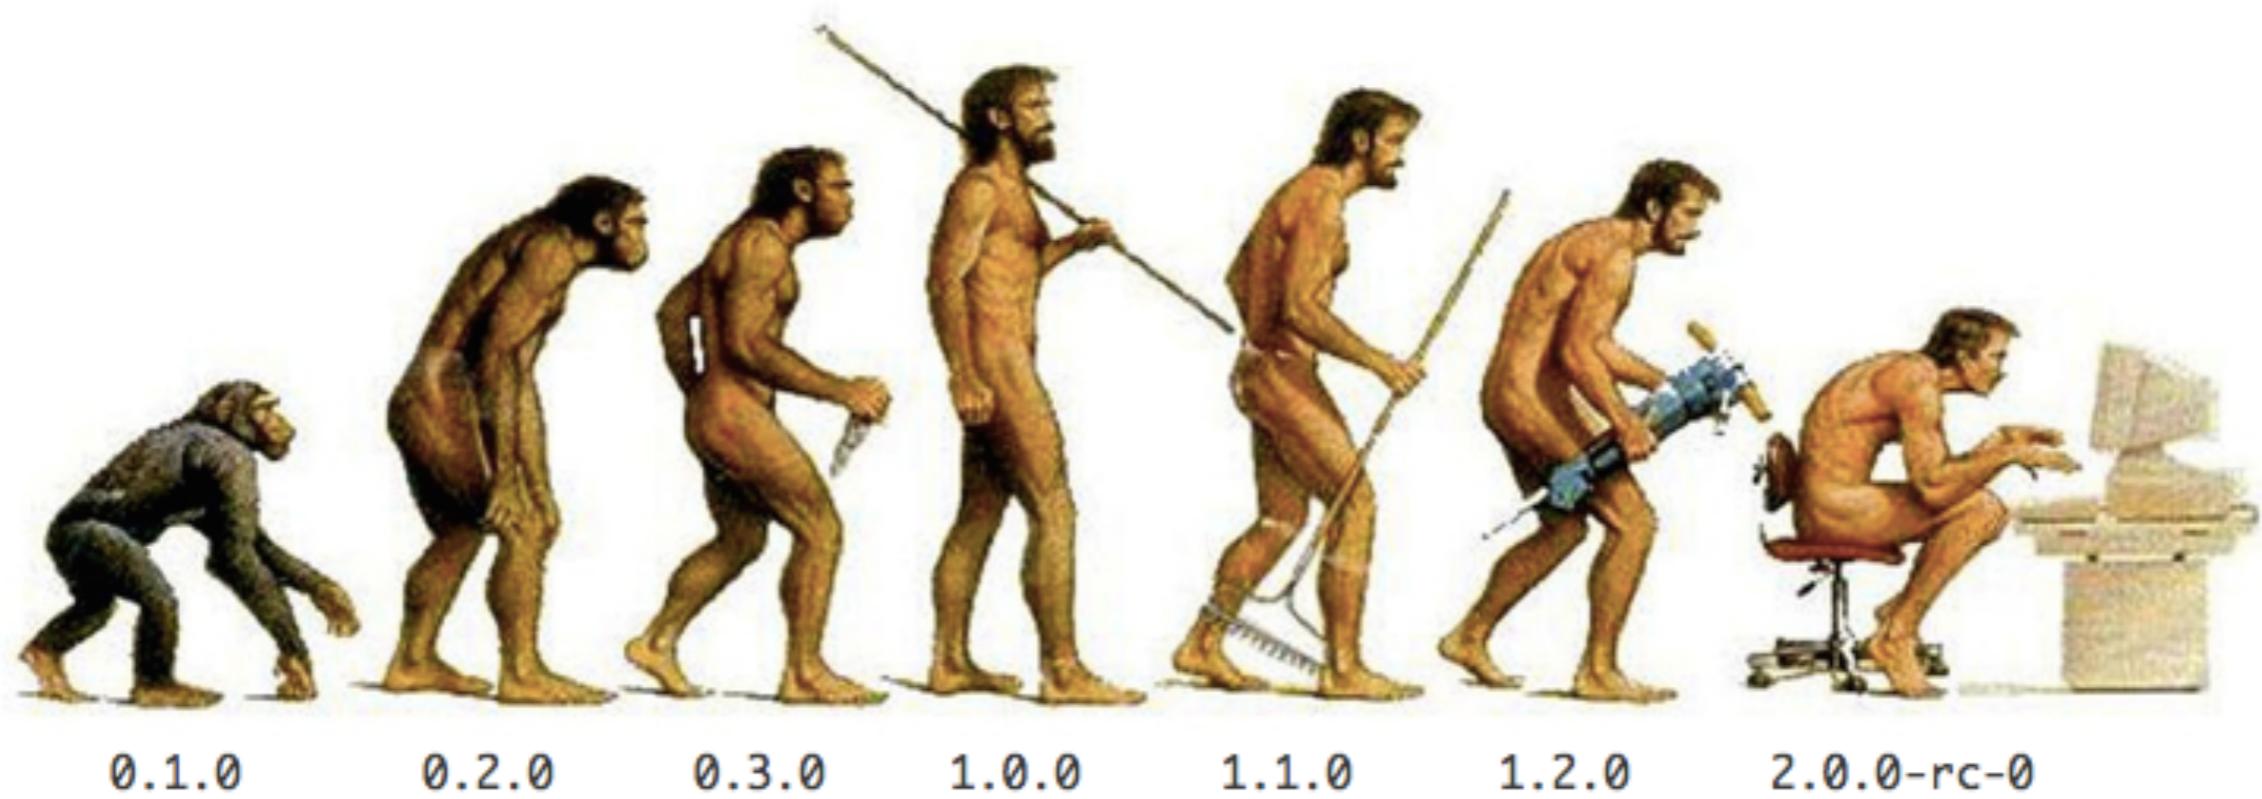
# Grains of Sand



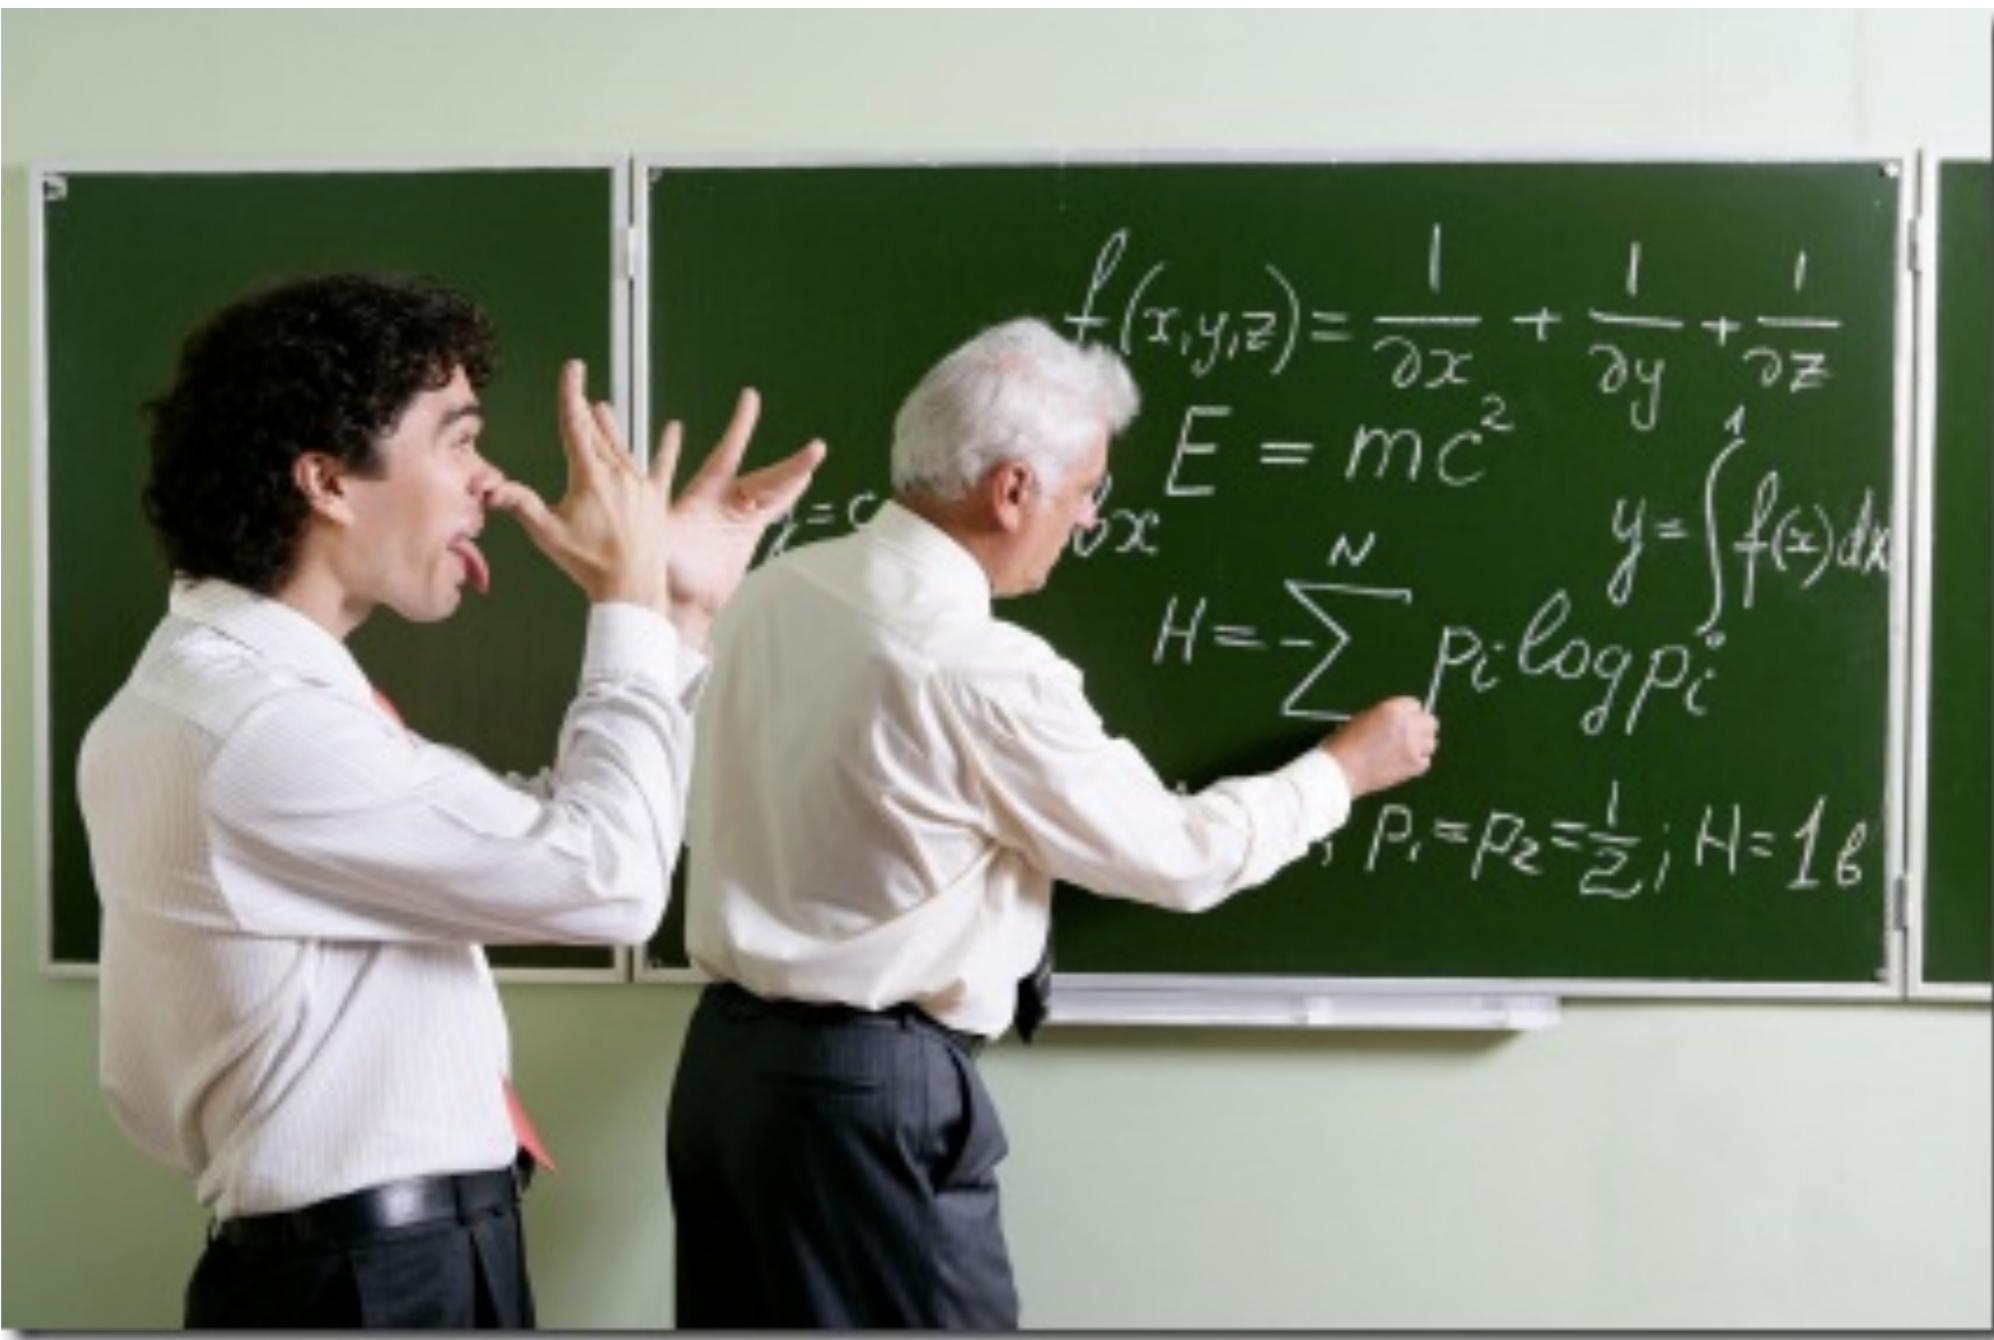
# Jump On The Bandwagon



# Versioning Avoidance



# Missing mock servers



# Give It A REST



# Broken Gate



# Layered Services Architecture



# Lipstick on the Pig

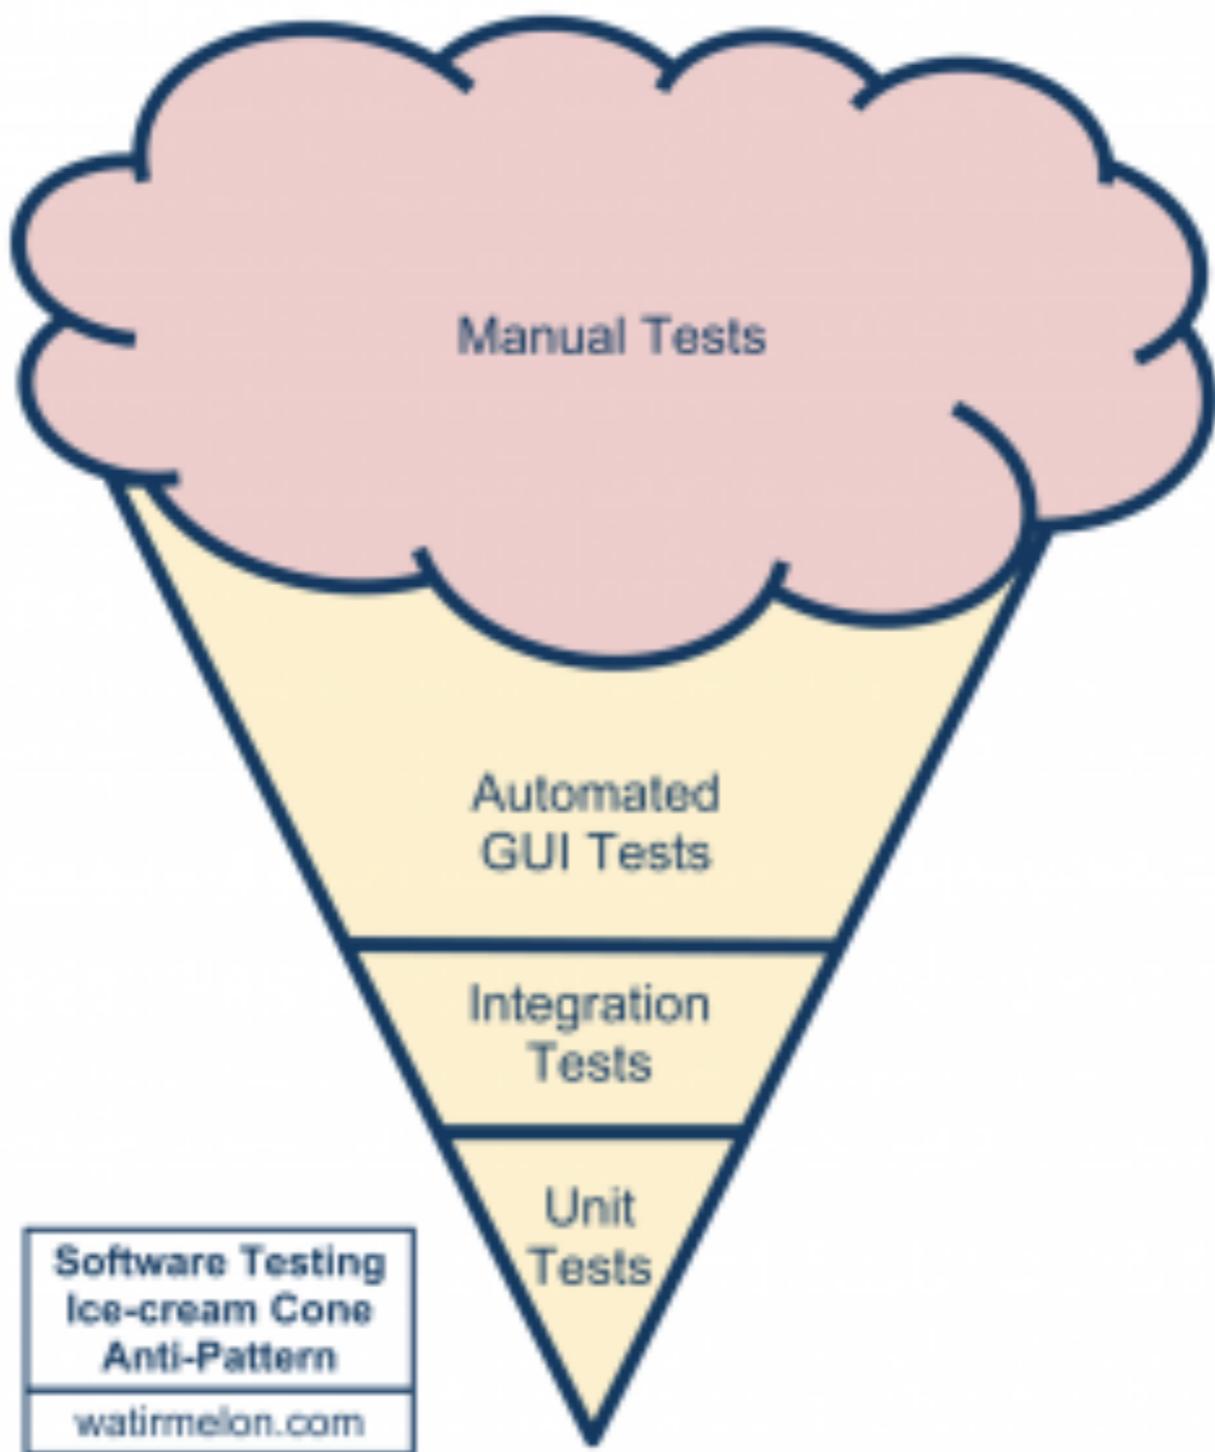


# Backend Micrservice

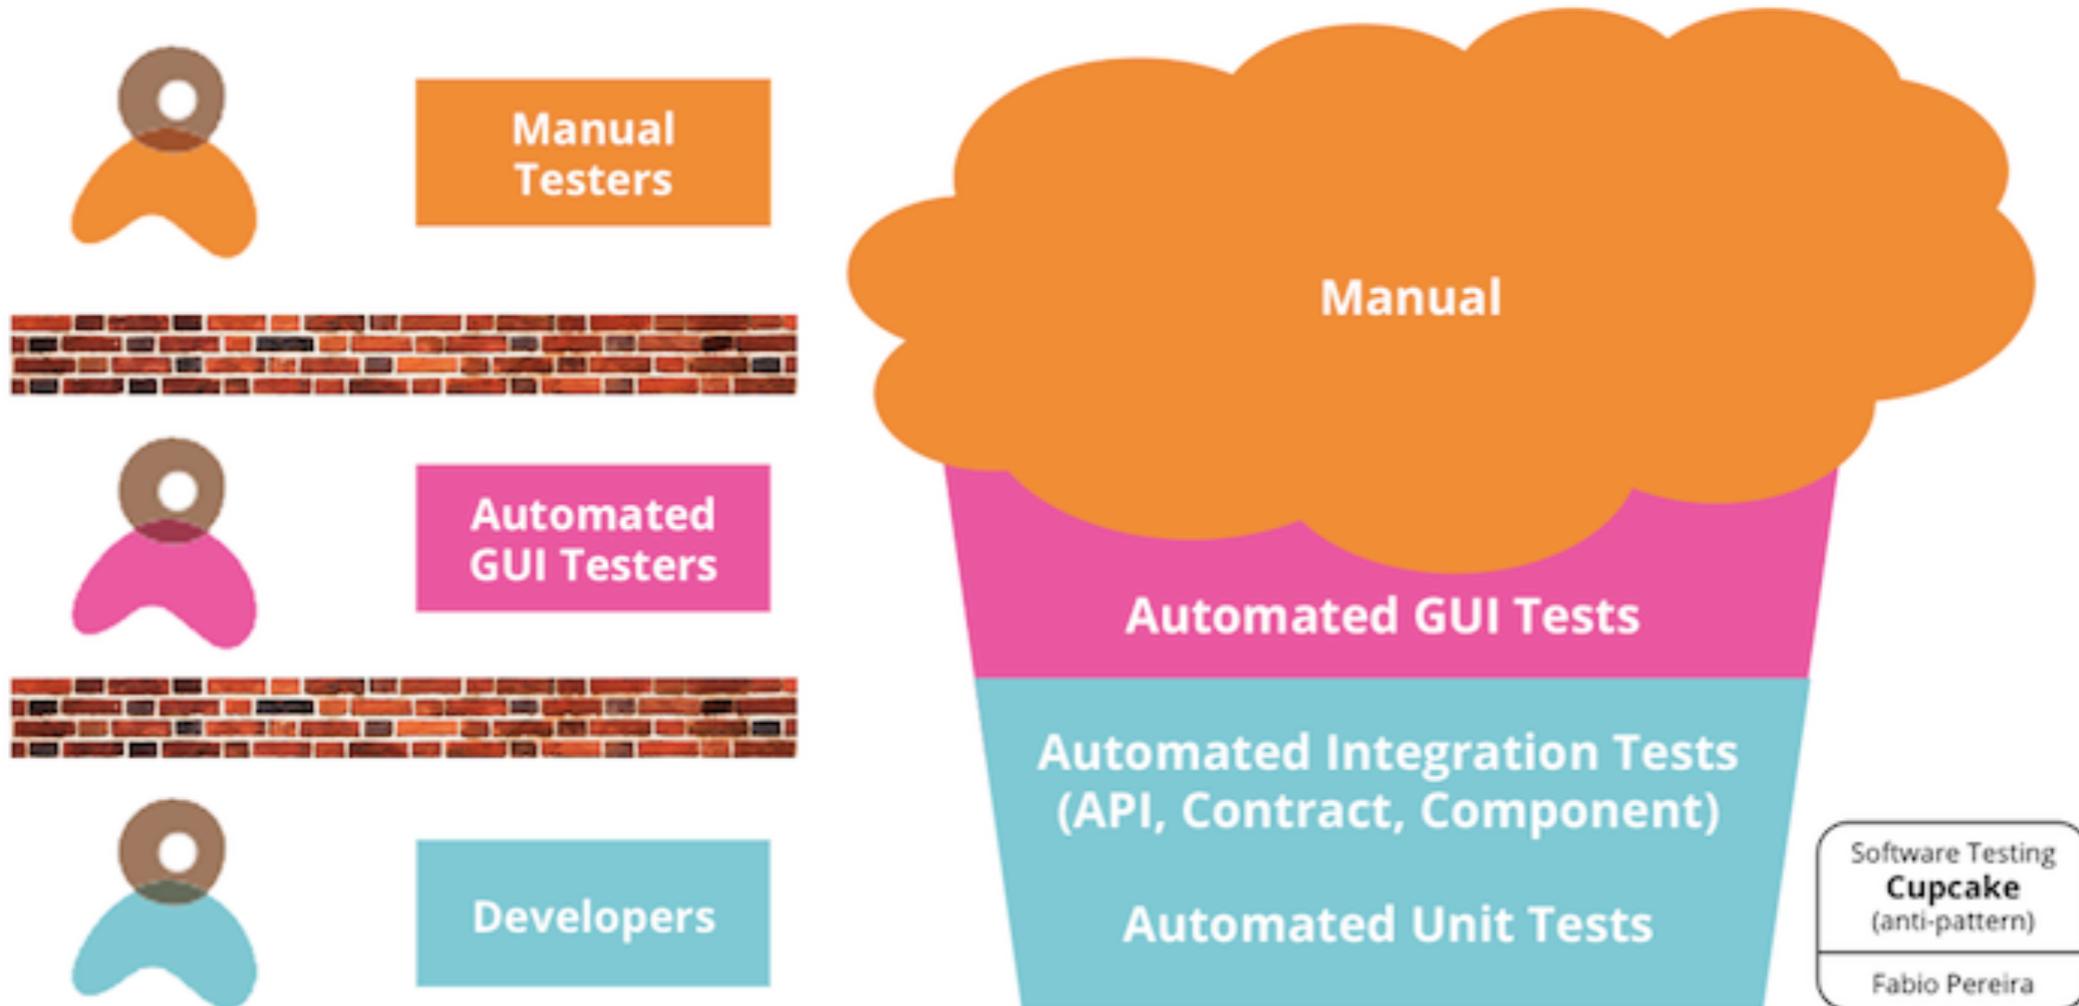


NETFLIX

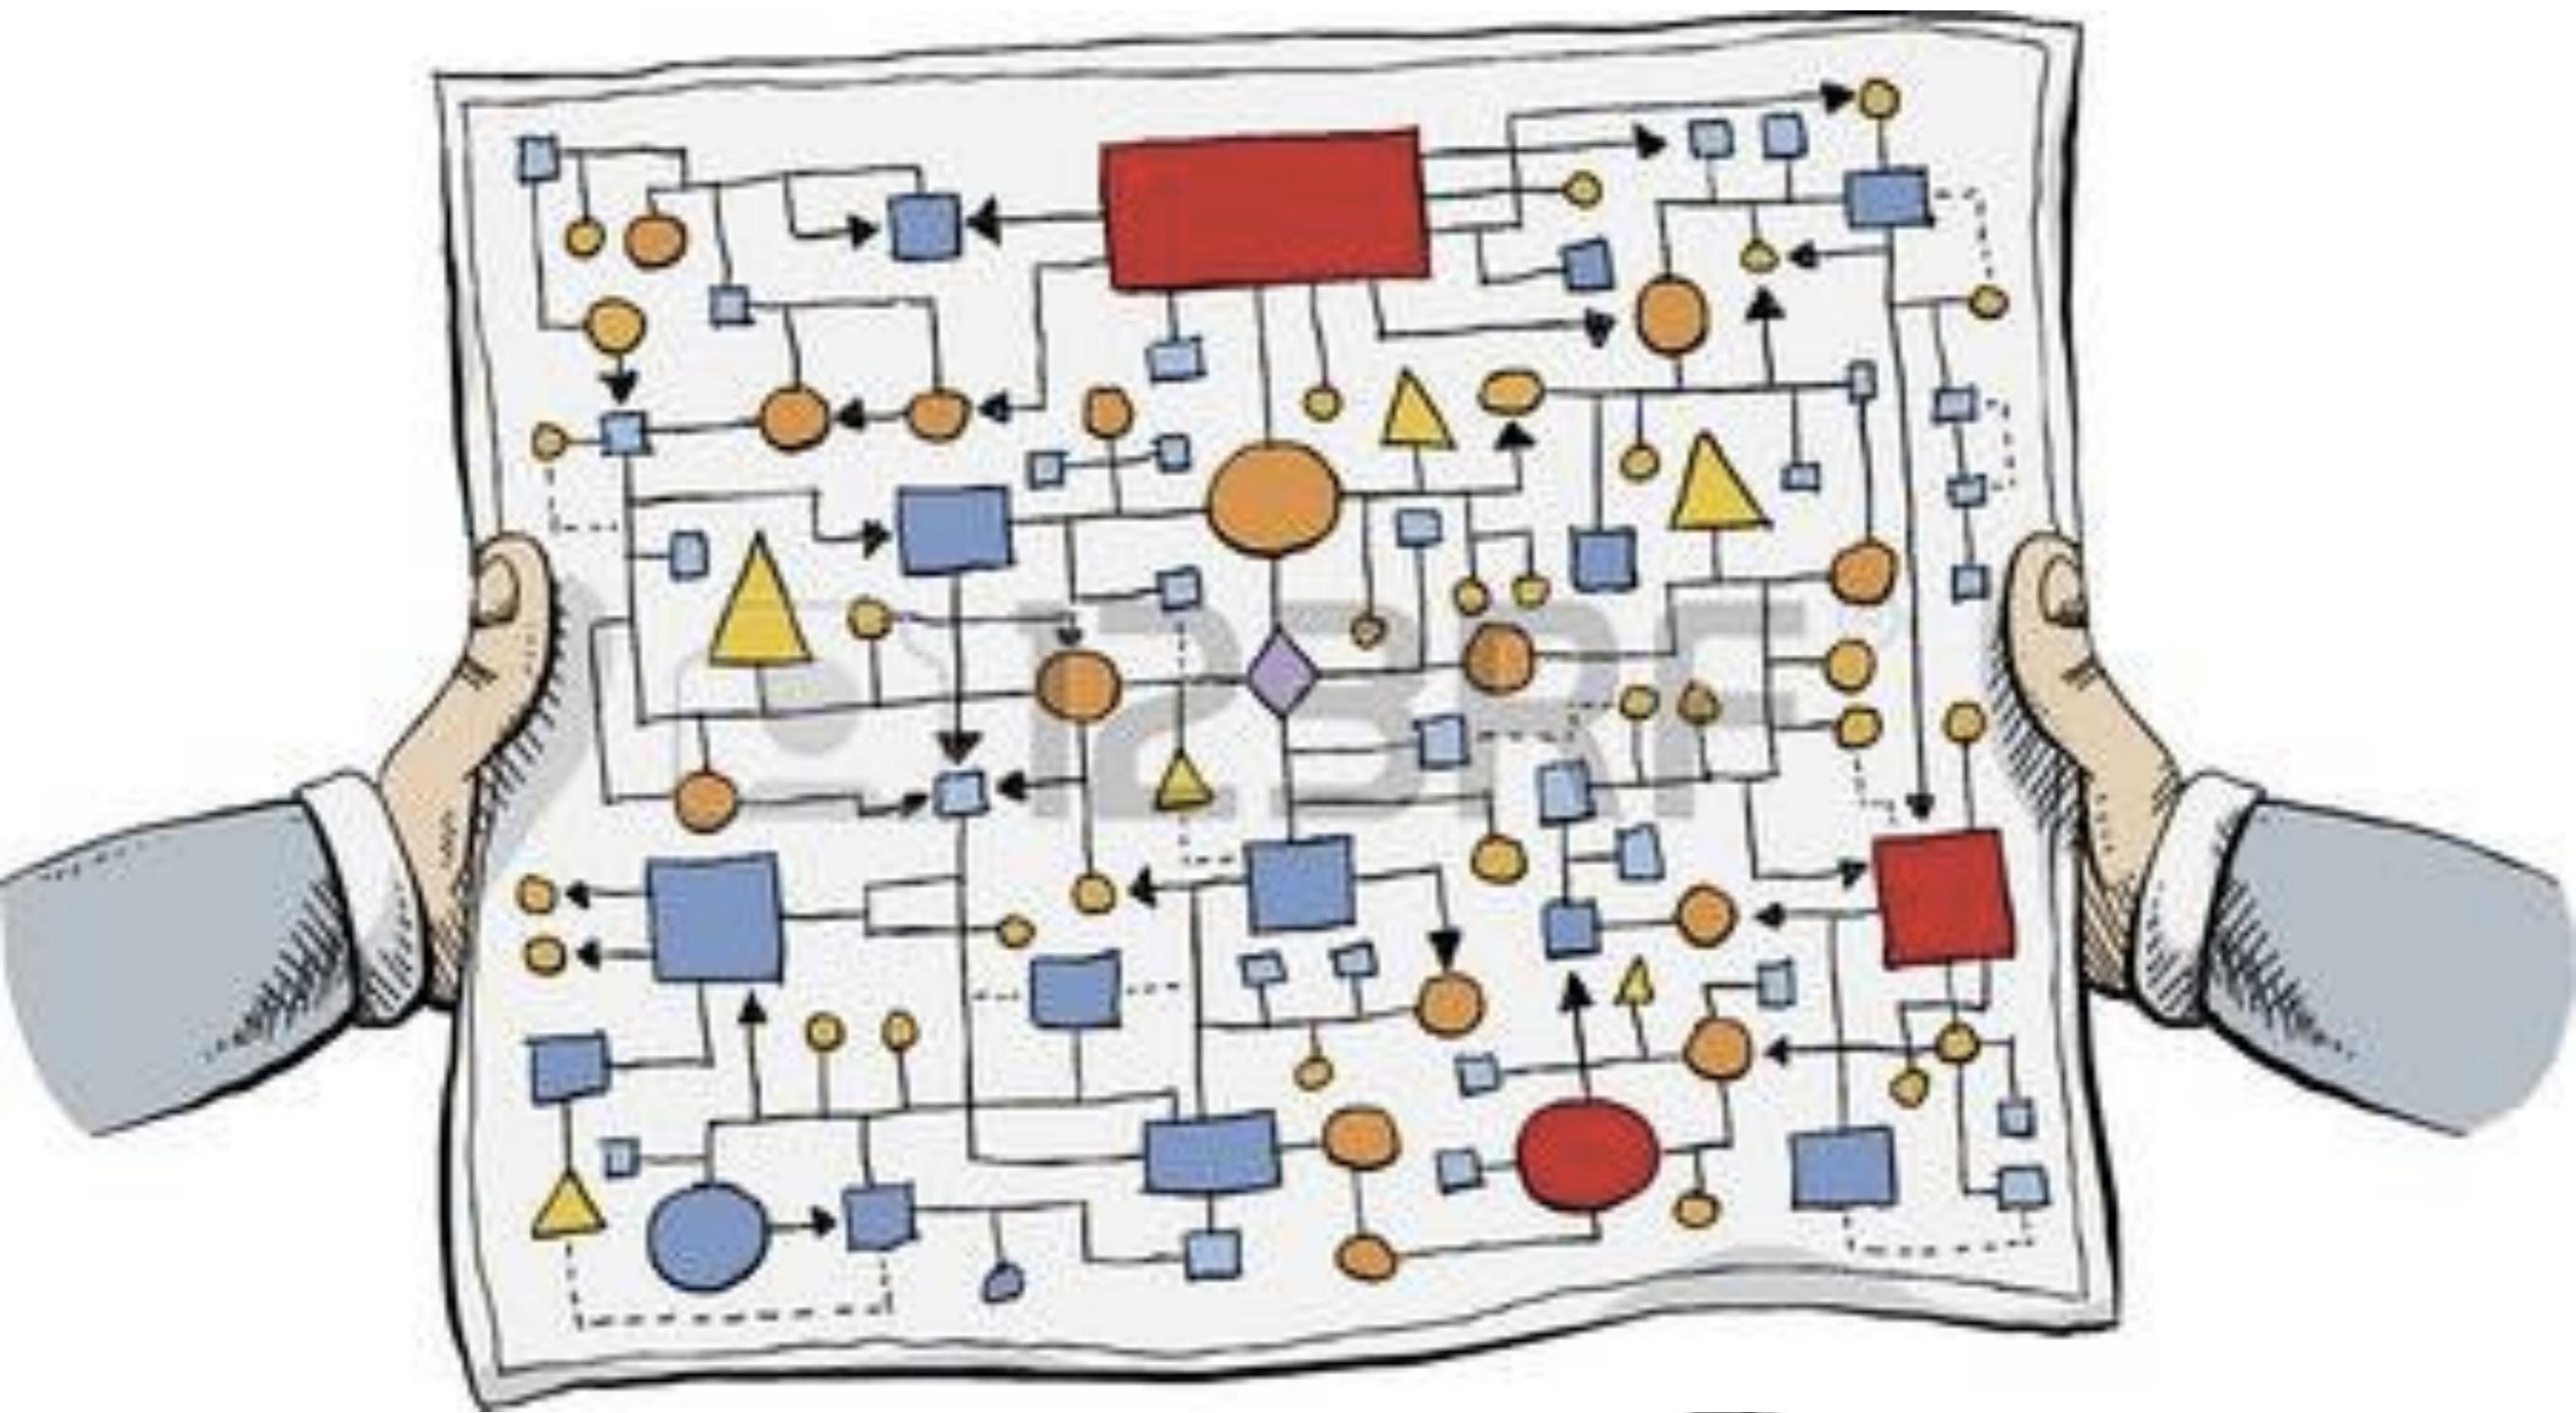
# Ice Cream Cone Test



# Cup Cake Anti Pattern



# Addicted to complexity



# Politics Oriented Architecture



# The Unteachables

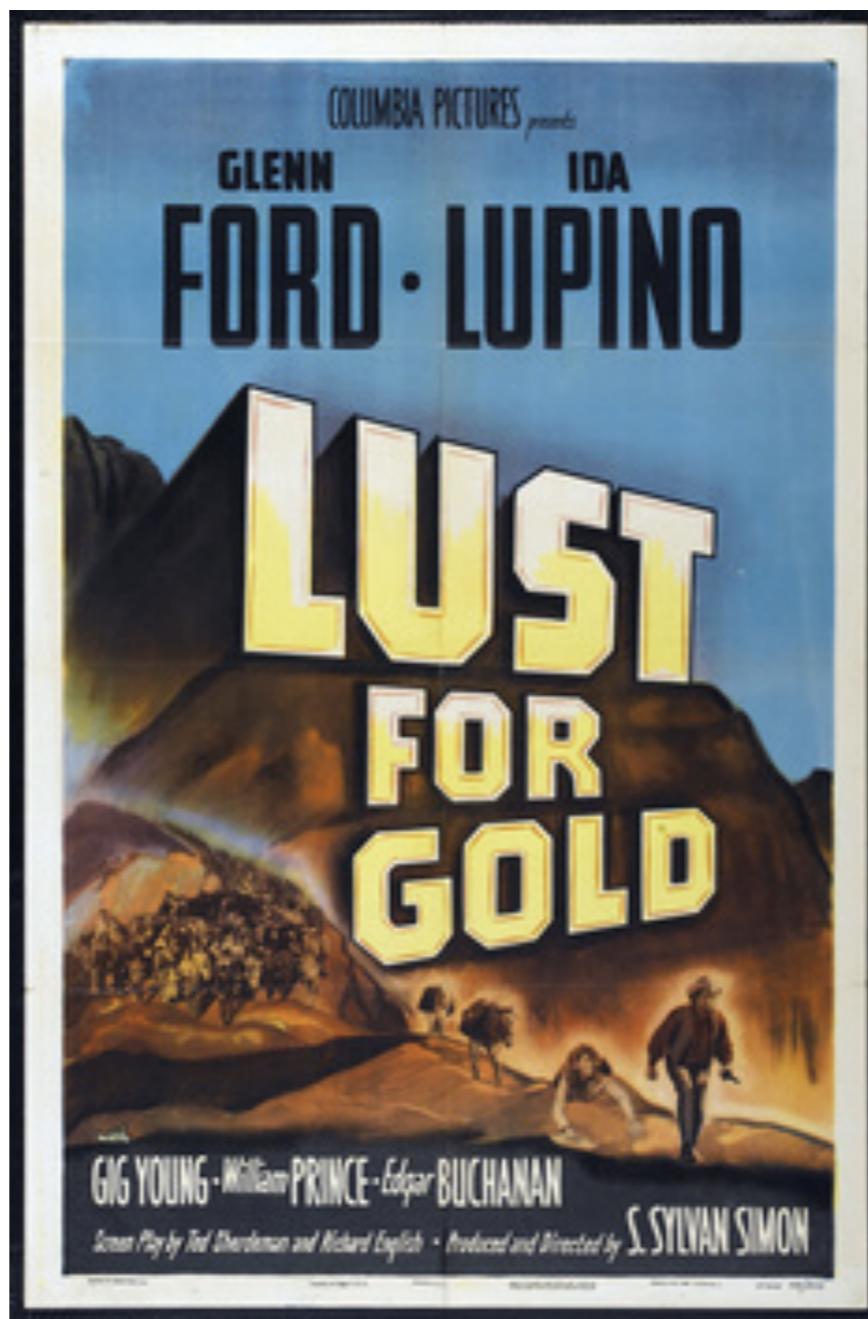


"If you can't understand it, don't invest in it.“

- 

Warren Buffet

# Lust



Do you really need it ?

# Appendix

# Flying Blind



there are no healthchecks, no monitoring and no metrics, so you assume everything is well in your systems, all the time.

- you need each service to be able to produce a self healthcheck and metrics, possibly through an HTTP endpoint, and surround it by some external monitoring.
- dropwizard, spring boot

# Monolith database



microservices expose nice different REST endpoint, but are  
are sharing the same database

- moving on from the big SQL joins and welcoming asynchronous completion of tasks, a user interface that progressively updates itself

# Unknown caller



there's no correlation between calls, so each call is a new one, completely unrelated to the source of it.

- each server must inject a call id if missing , each server must propagate the call id if present, each server should log each call including the id. Thus providing a clear chain of invocations between the services

# Hardcoded hell



All the address (endpoints) of services are hardcoded somewhere, sometimes directly in the code.

- introduce a discovery mechanism like eureka, zookeeper, consul, msnos

# Synchronous world



Every call in your systems is synchronous, and you wait for things to actually happen before returning to the caller

- using queues on top of your receivers and implementing asynchronous protocols from the start.

# Grains of Sand



defining the size of a microservice using strict line of code counts

- As a rule of thumb you should start out with bigger, more coarse-grain services. It will always be easier to further decompose a single service than it will be to combine multiple services.

# Jump On The Bandwagon



starts building microservices before the tools and processes  
are in place to support them.

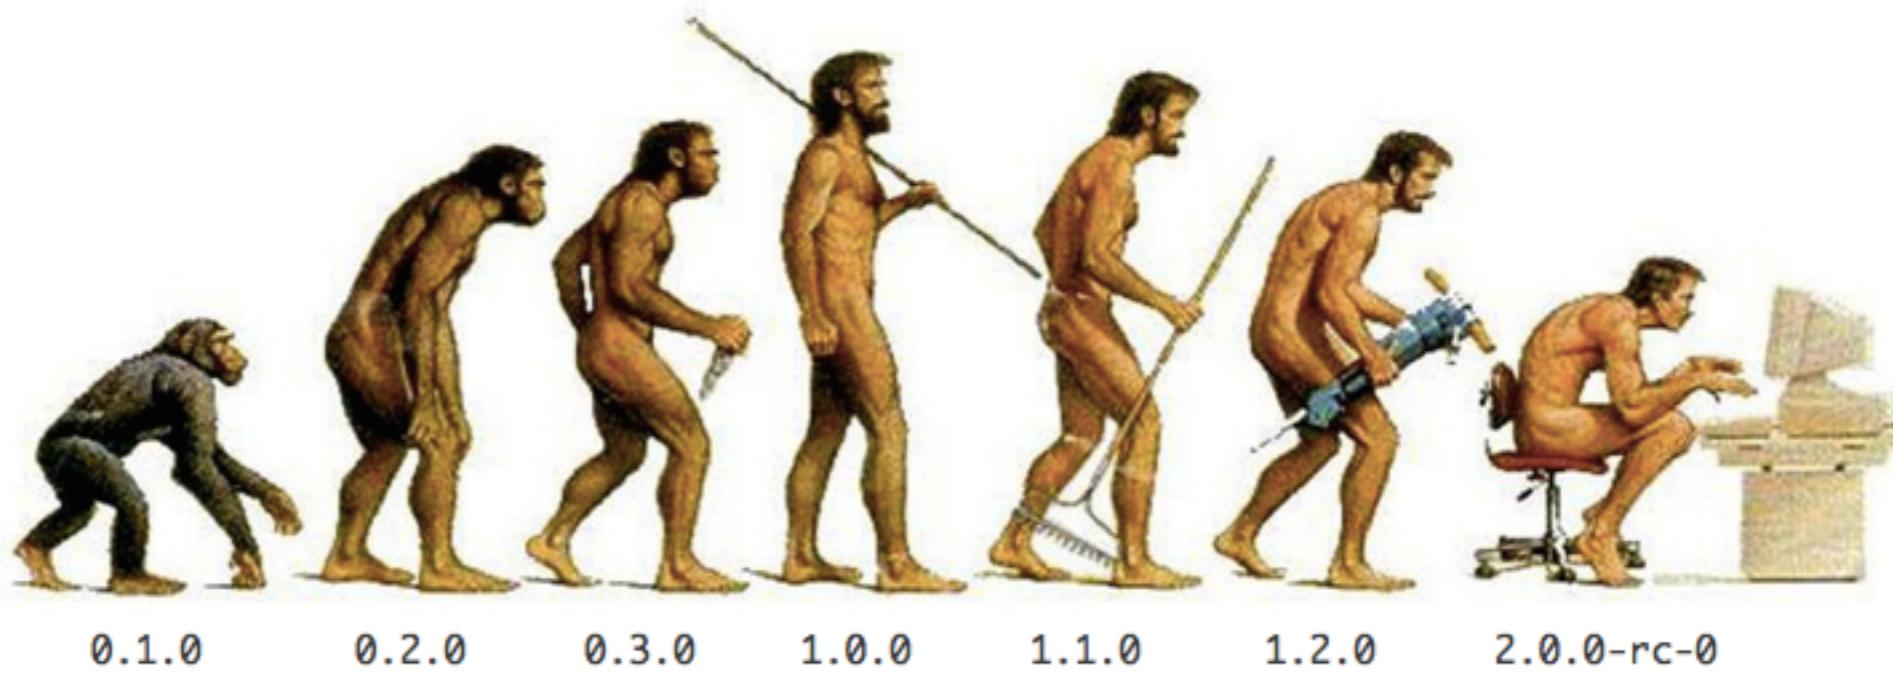
- if your organization is not ready or willing to change its testing, monitoring and deployment mentality, then you likely won't get all the benefits of microservices.

# Dare To Be Different



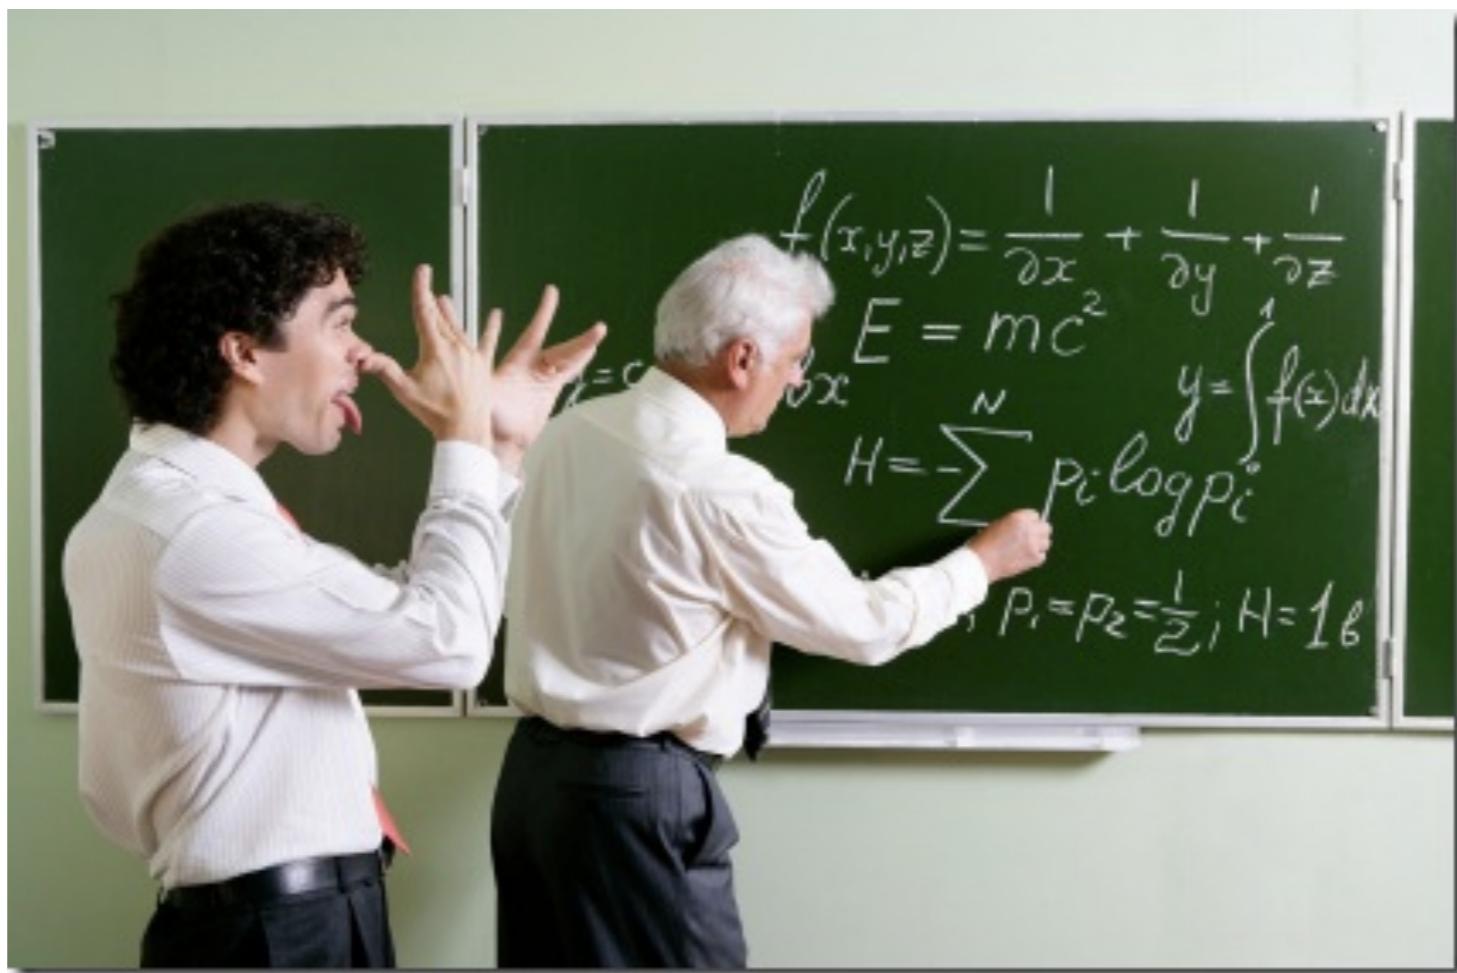
- Each of these frameworks provide their own mechanisms for security, logging, configuration, embedded HTTP server, UI template engine, dependency injection, and more.
- one of your first tasks, before you build a single microservice, should be to standardize the technology stack for everyone.

# Versioning Avoidance

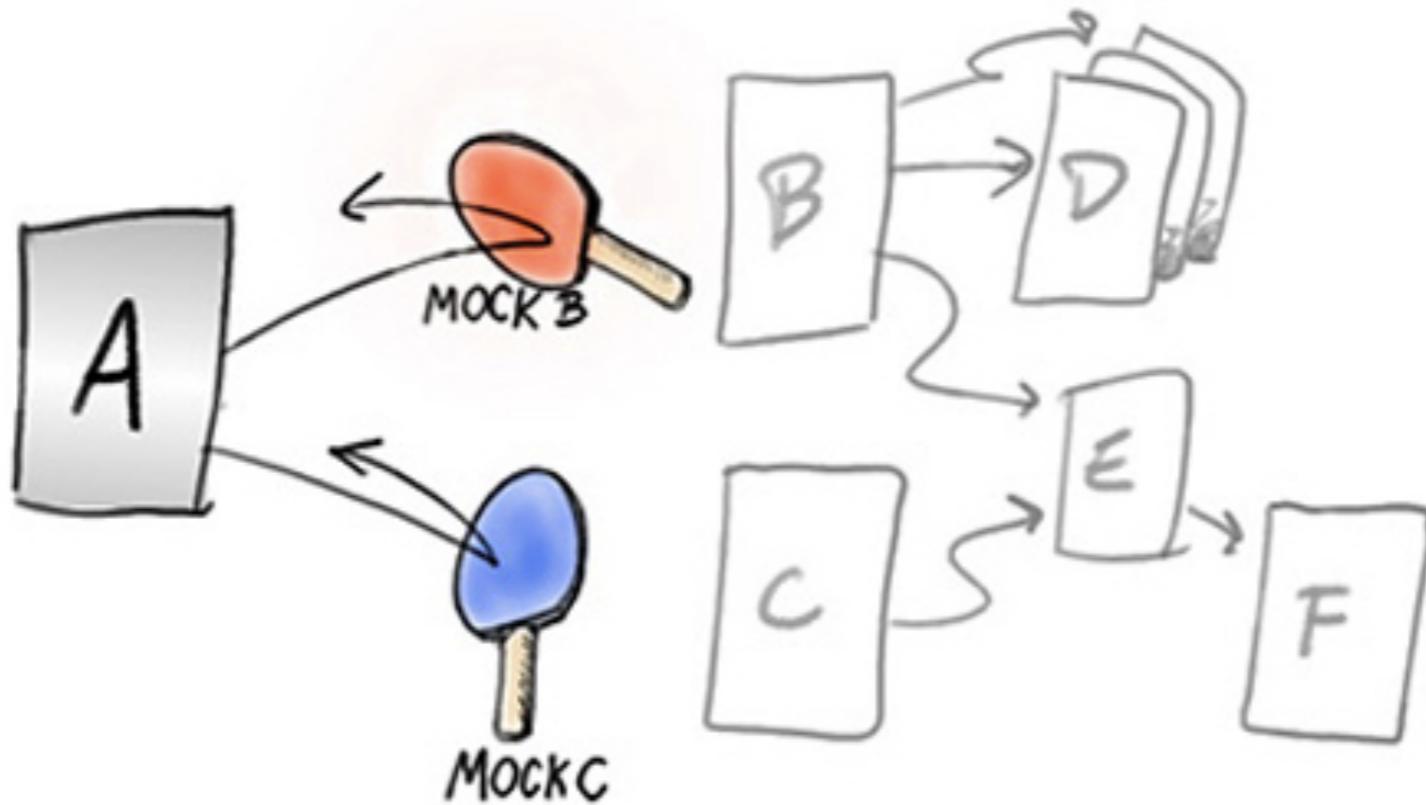


Thought it would be only need one version of the service.

# Missing mock servers



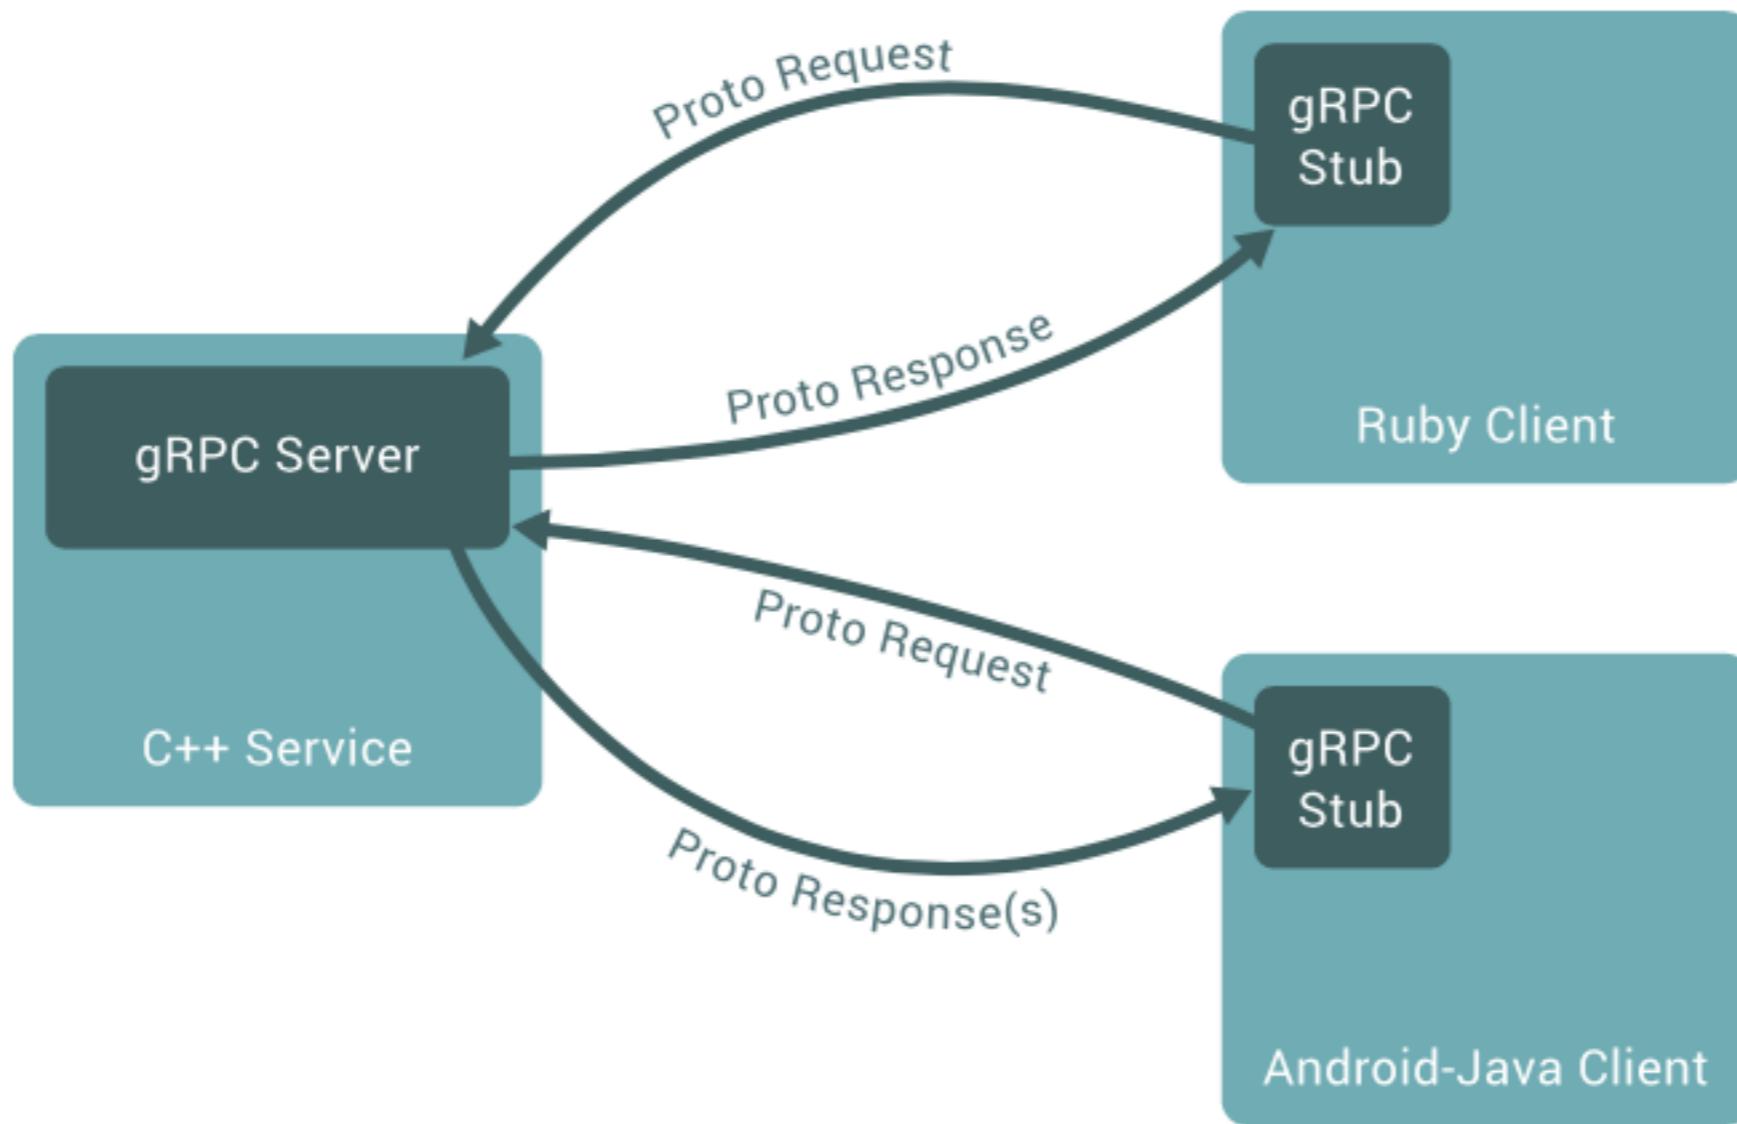
each consuming team of a service has to write their own mocks and stubs, which is wasteful.



# Give It A REST



Relying heavily on RESTful APIs when other messaging paradigms may be more appropriate.



REST  
gRPC  
Apache Thrift

- Have a versioning strategy that can allow the consumers a graceful migration and assure providers can transparently deploy changes without affecting anyone. Limit the number of side-by-side major versions in the production and govern them.

# Manual Configurations Management



managing the configurations for each service

# Layered Services Architecture

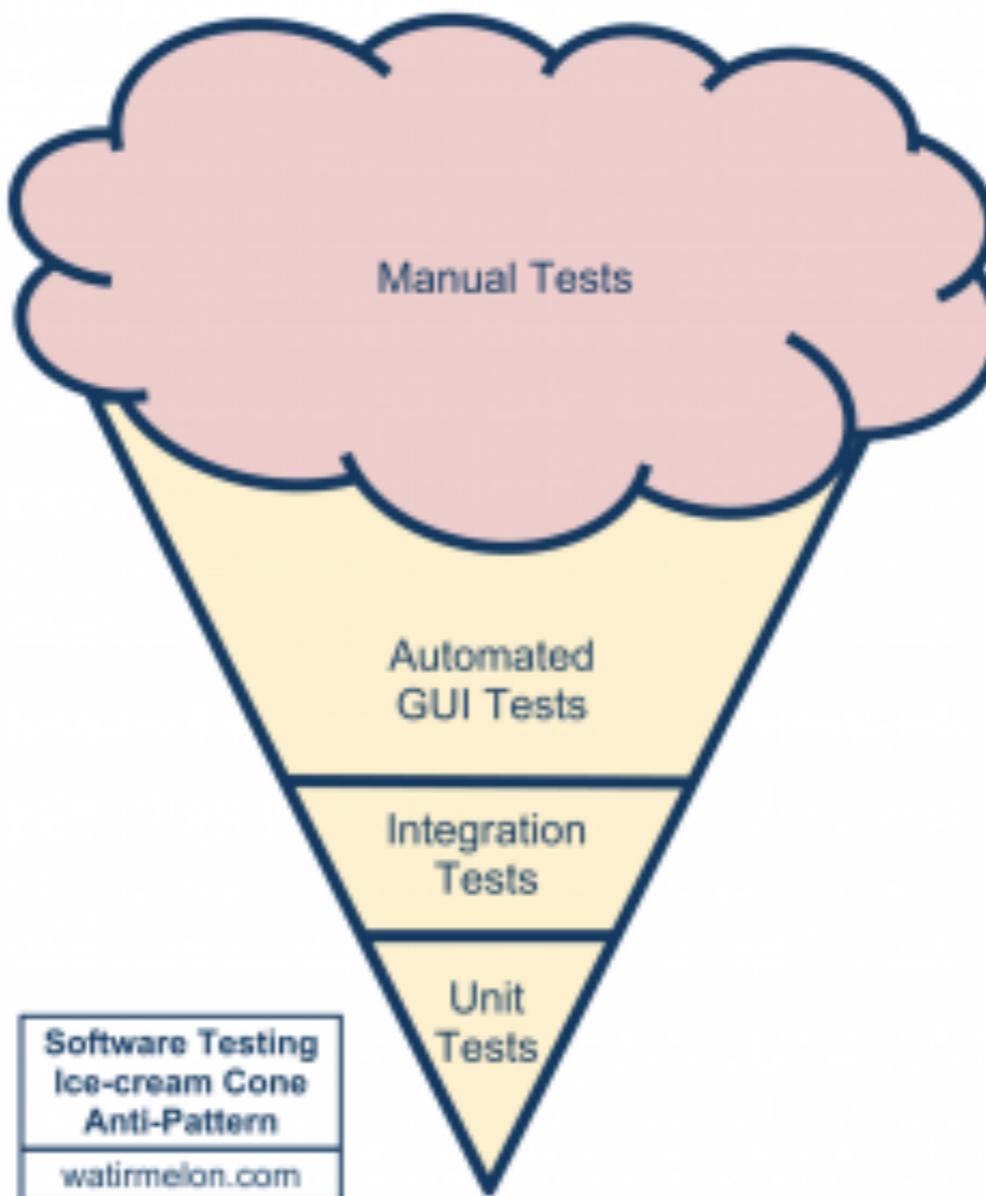


Teams mostly focused on technical cohesion rather than functional regarding reusability. For example, several services functioned as a data access layer to expose tables as services;

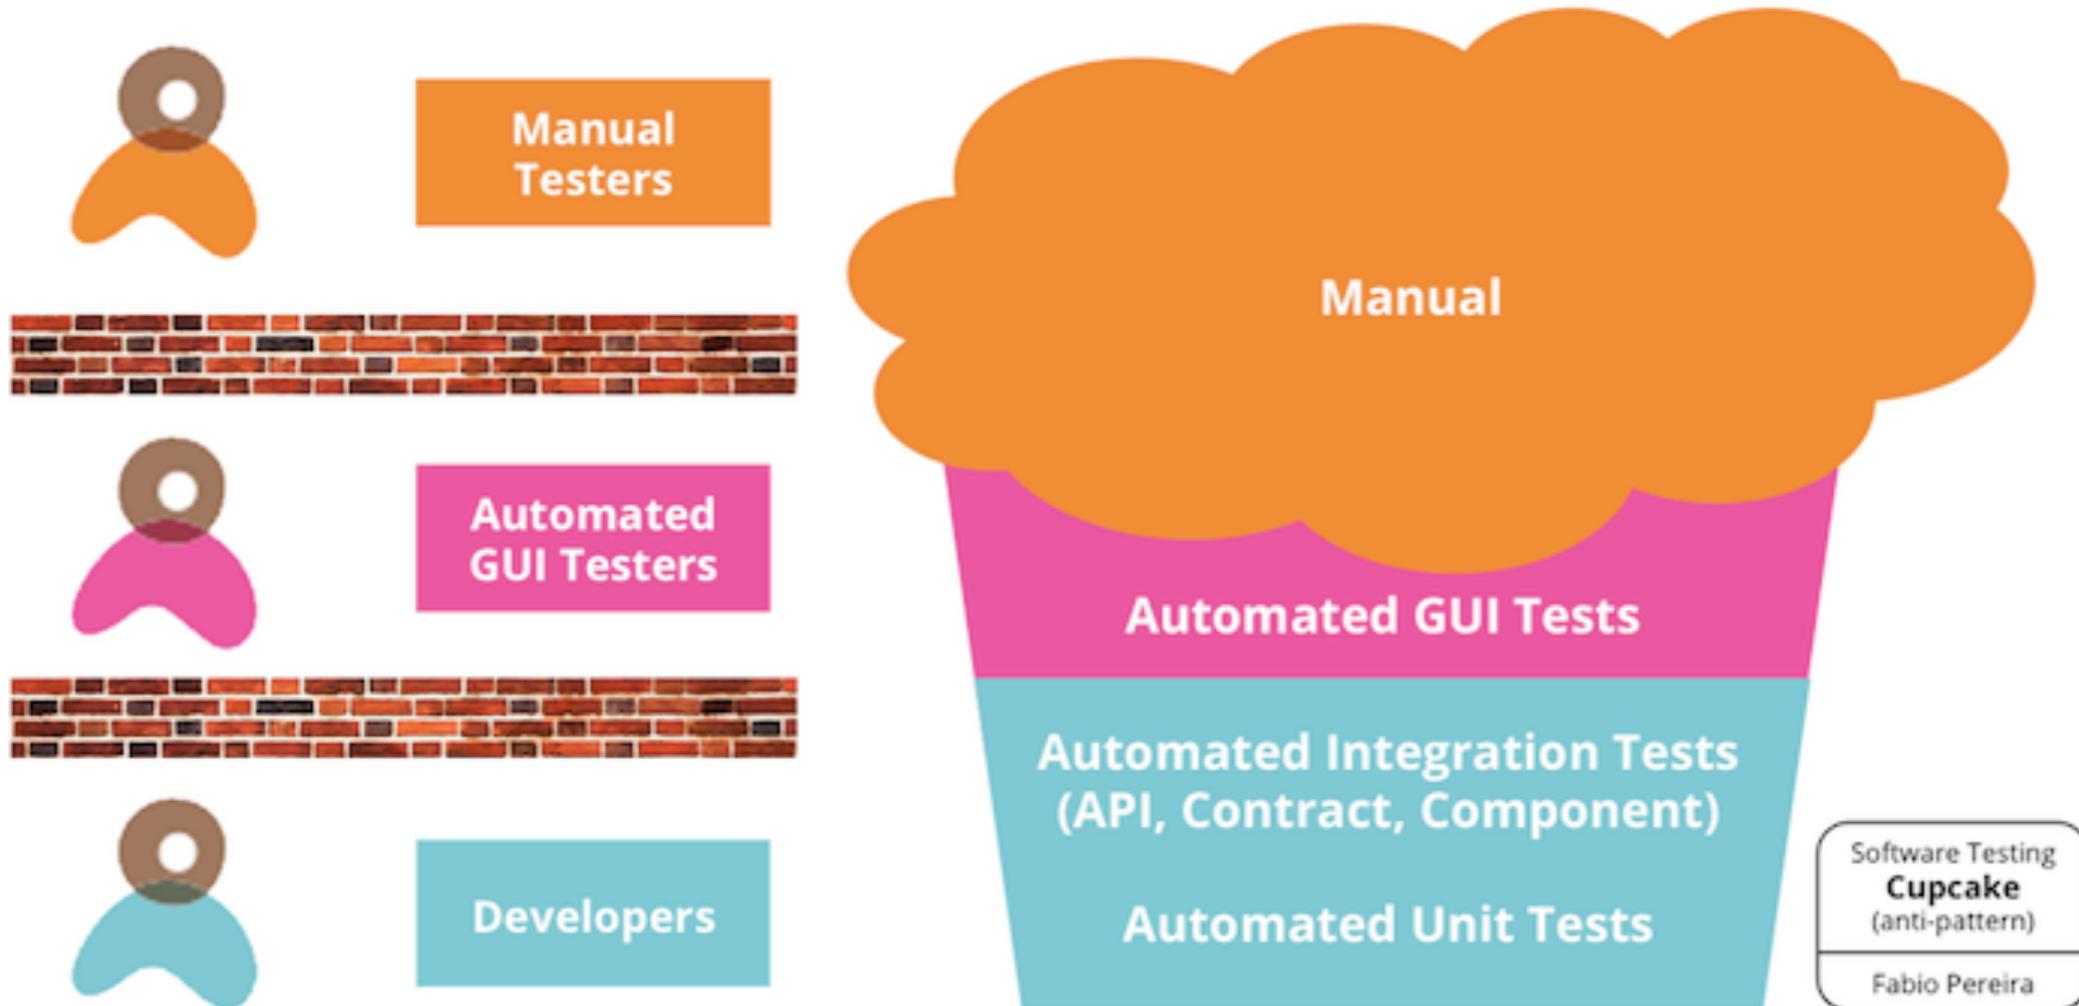
# Not using a Gate



# Ice Cream Cone Test



# Cup Cake Anti Pattern



# Appendix

- use application configuration management tools

- Logical separation of layers within a service is fine, however, there should not be any out of process calls. Try to look at a service as one atomic business entity, which must implement everything to achieve the desired business functionality.

Invest in API Management solutions to centralize, manage and monitor some of the non-functional concerns and which would also eliminate the burden of consumer's managing several microservices configurations.

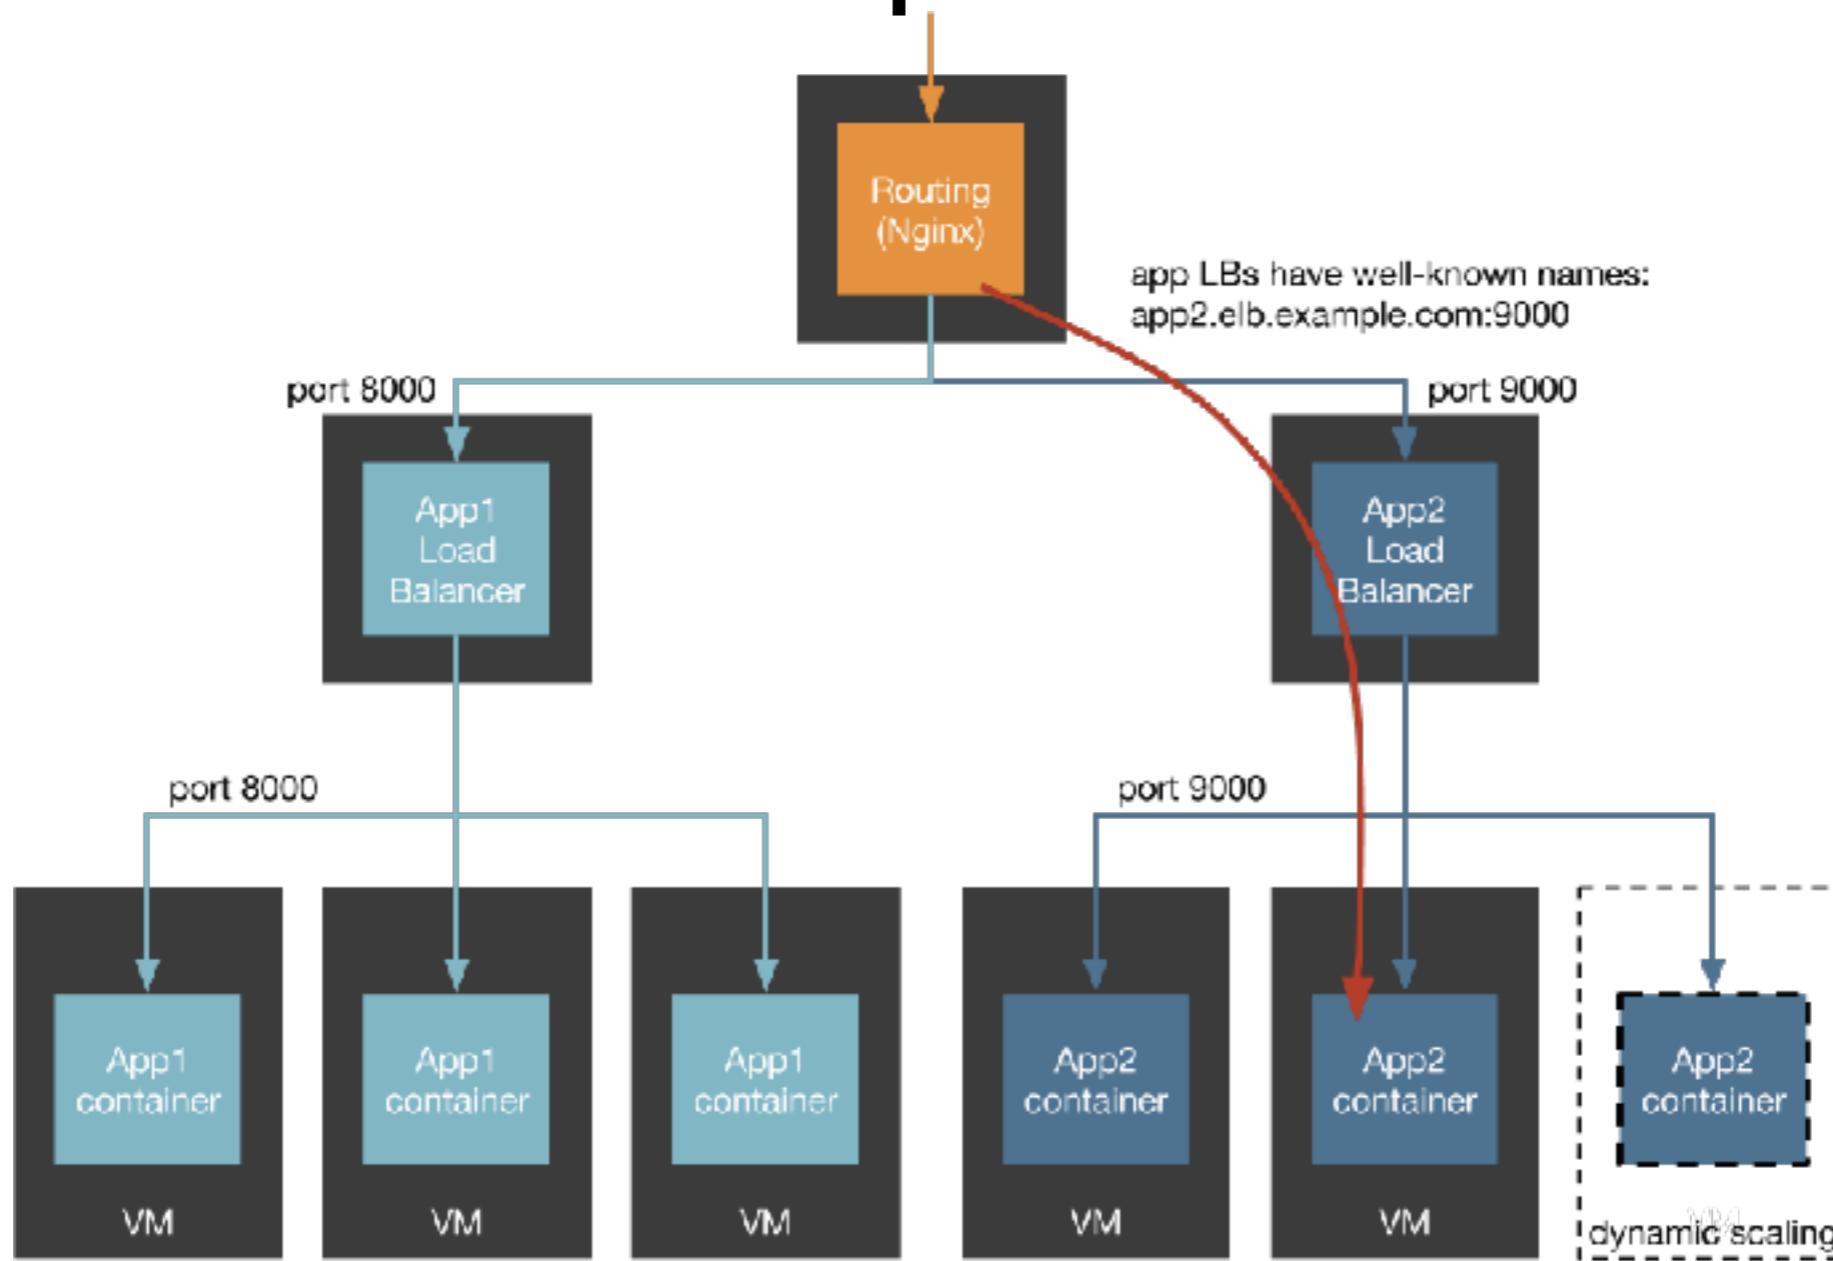
# Netflix Architecture

# Google Architecture

- happens when there is not enough low-level testing (unit, integration and component), too many tests that run through the Graphical User Interface (GUI) and an even larger number of manual tests.

- Writing tests usually spans 3 teams:
  - Developers tests.
  - UI Automation.
  - Manual testers.
- This results in duplication.

# Local proxies



a load balancing service fronting each stateless microservice.

- If an application instance is unable to respond to remote health checks, then the load balancer will remove that application instance from the pool.

# Local proxies - Anti pattern

- Health checking: If an application instance is unable to respond to remote health checks, then the load balancer will remove that application instance from the pool. But if the application instance is unable to respond because it is under heavy load or is in a blocking state (ex. with Node.js performing computationally-heavy work or Rails unicorn workers waiting on long-running queries), then the load balancer may be removing instances just when the service needs more instances in the pool. Scaling up additional VM capacity takes on the order of minutes, and any new instances will be coming live in a state where there are still not enough nodes marked as healthy to serve queued requests. There is actually enough capacity available to run the service but the LB is "helpfully" not leaving instances marked healthy long enough to take up the load and the service is dead in the water until the health checks are turned off temporarily by a bleary-eyed operator at 2AM.

- Networking Inefficiency: We've introduced an extra network hop for every single internal request.
- Utilization: The hidden assumption of this architecture is that we have one container per VM. If we provision multiple containers per VM to improve utilization, then we need to configure the LB to listen for multiple ports and the question of how health checks are supposed to work becomes complicated. Do we remove the entire node because one of its services dies? Do we need an instance of every service on each VM or can we route some ports to only some instances? Do we even have that option if we're using a managed LB?

- pushes the responsibility for the application topology away from the network infrastructure and into the application itself where it belongs -- only the application can have the full awareness of its state necessary to manage itself
- <https://www.joyent.com/blog/container-native-discovery>

