

DAAAAAAD, I lost
one of my
microServices
I can't play
without it.

I told you he is
not mature enough
for it. We should
have bought him a
monolith.

I believe you.





- » Skan.ai - chief Architect
- » Ai.robotics - chief Architect
- » Genpact - solution Architect
- » Welldoc - chief Architect
- » Microsoft
- » Mercedes
- » Siemens
- » Honeywell



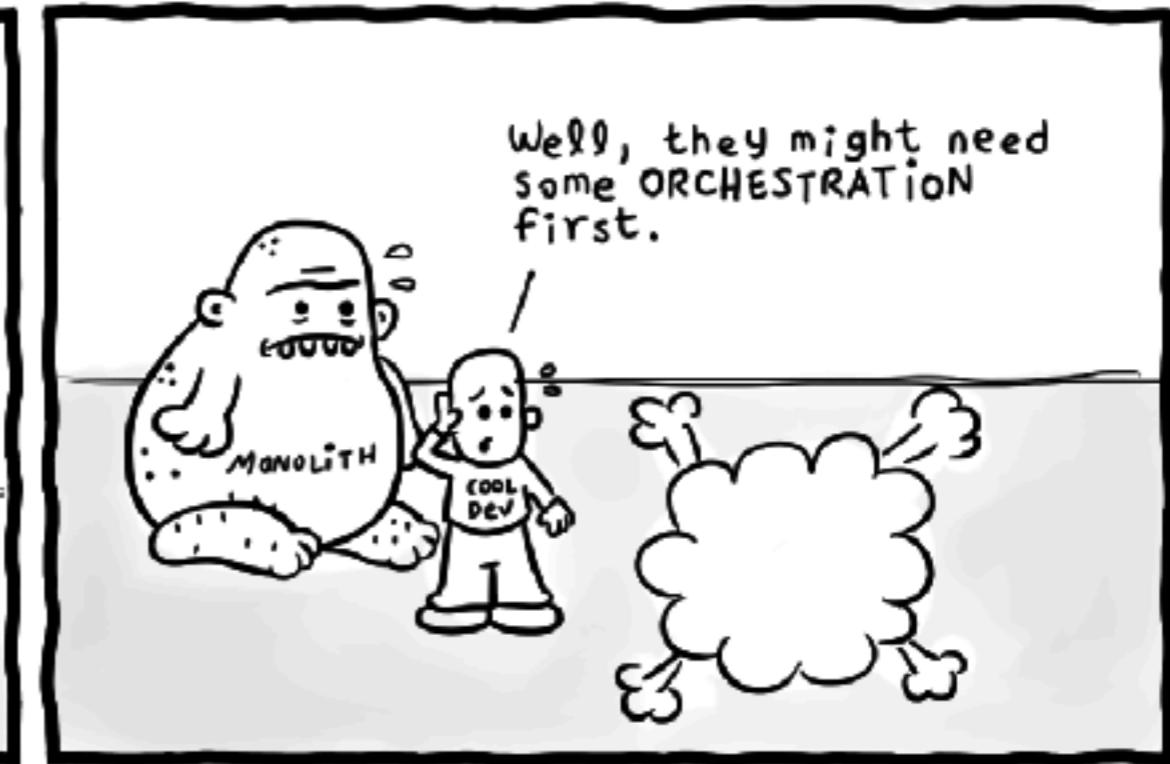
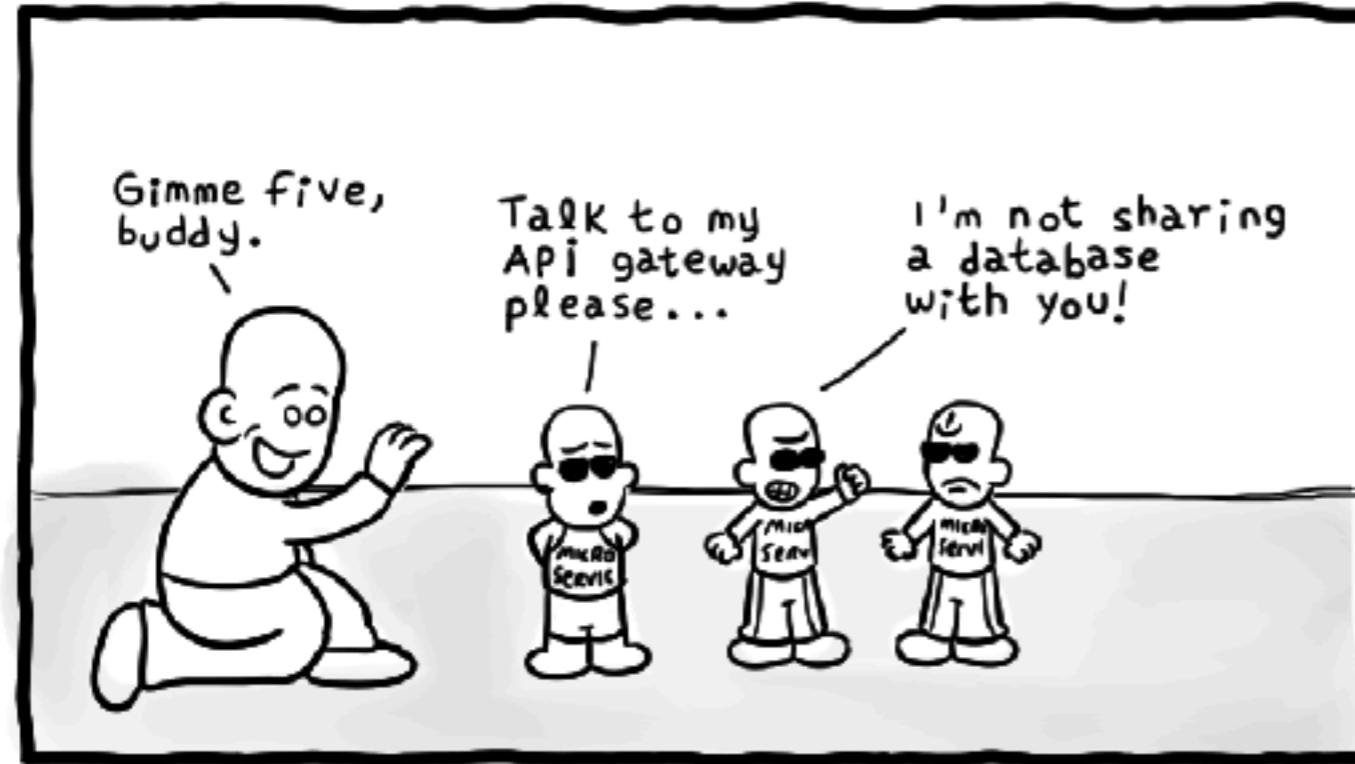
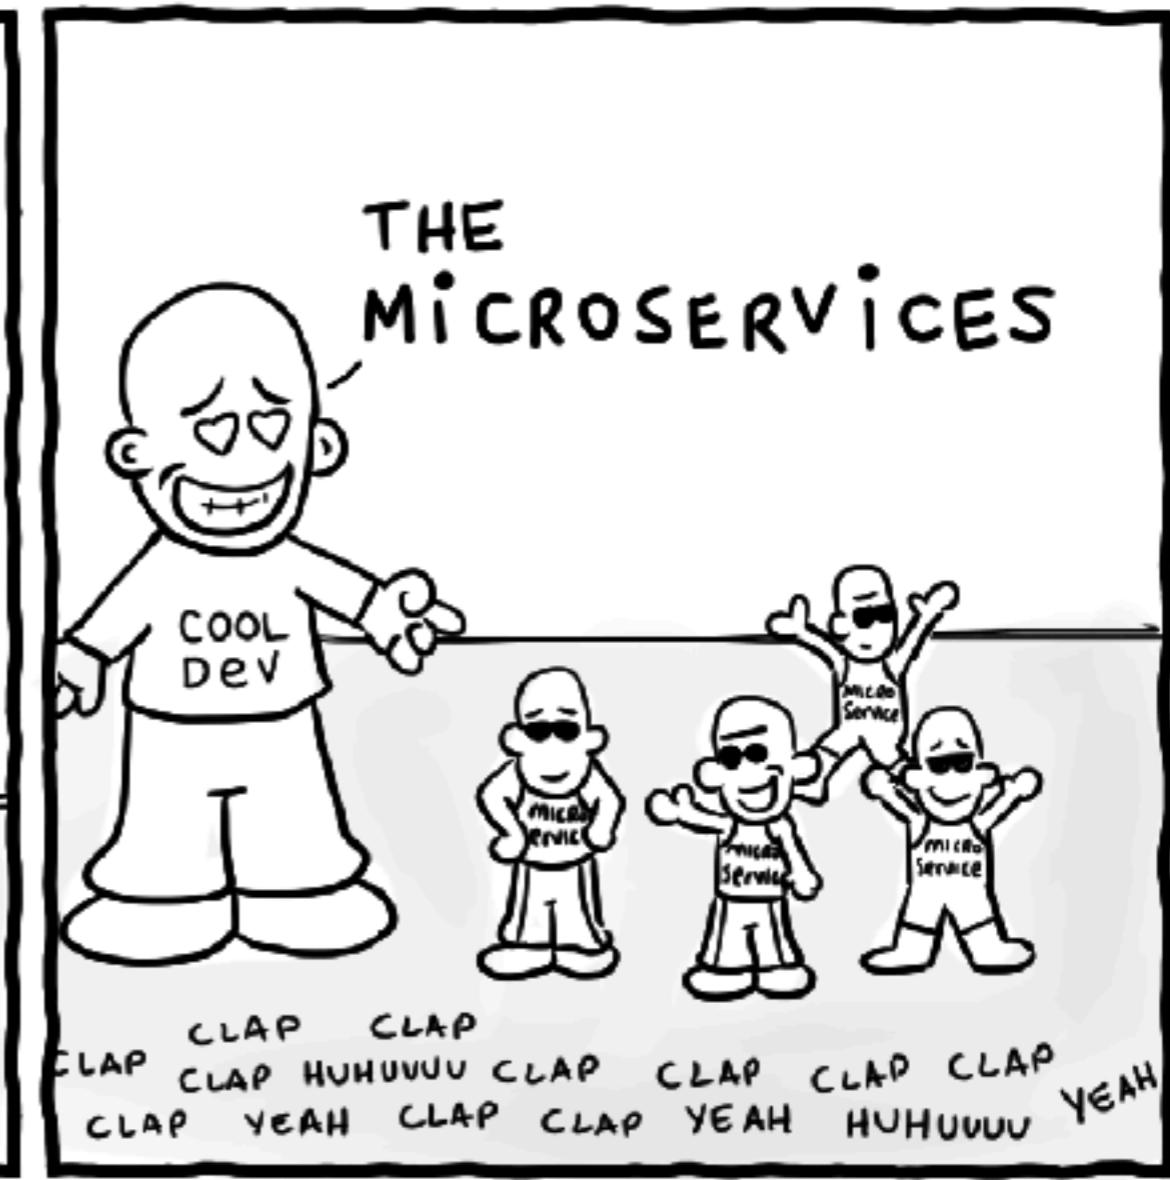
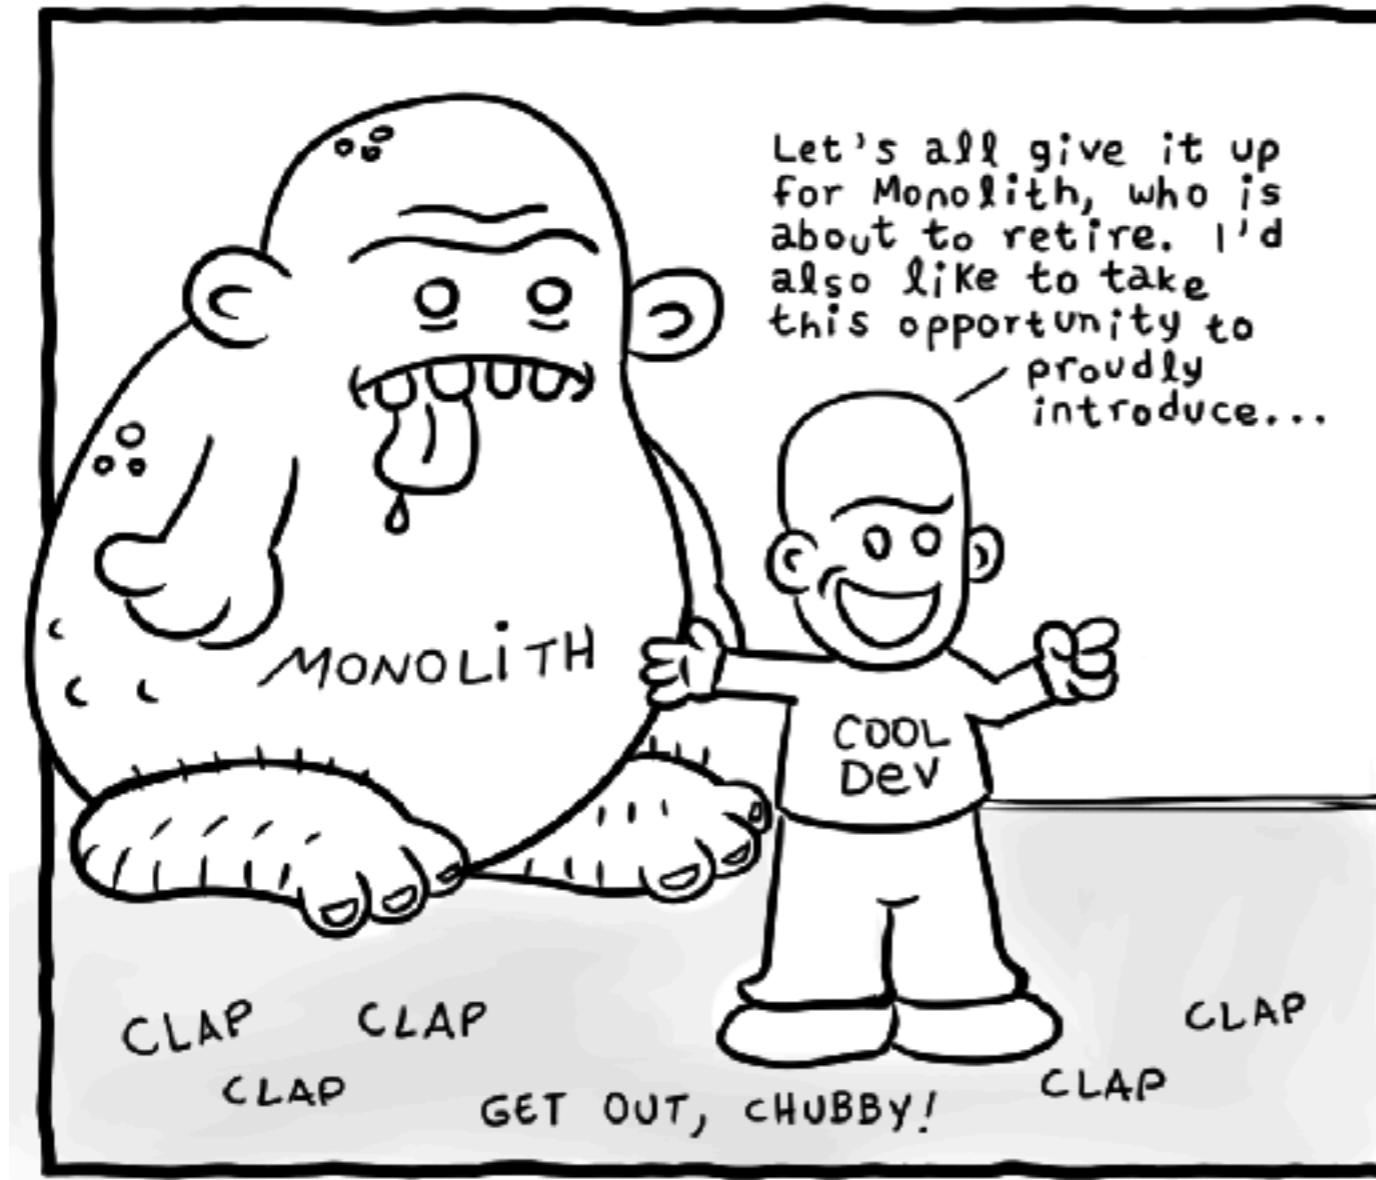
Mubarak



- » How to build Microservice
- » Patterns
- » Anti Patterns
- » Hands on Demos
- » Case Studies

what do
YOU
expect?

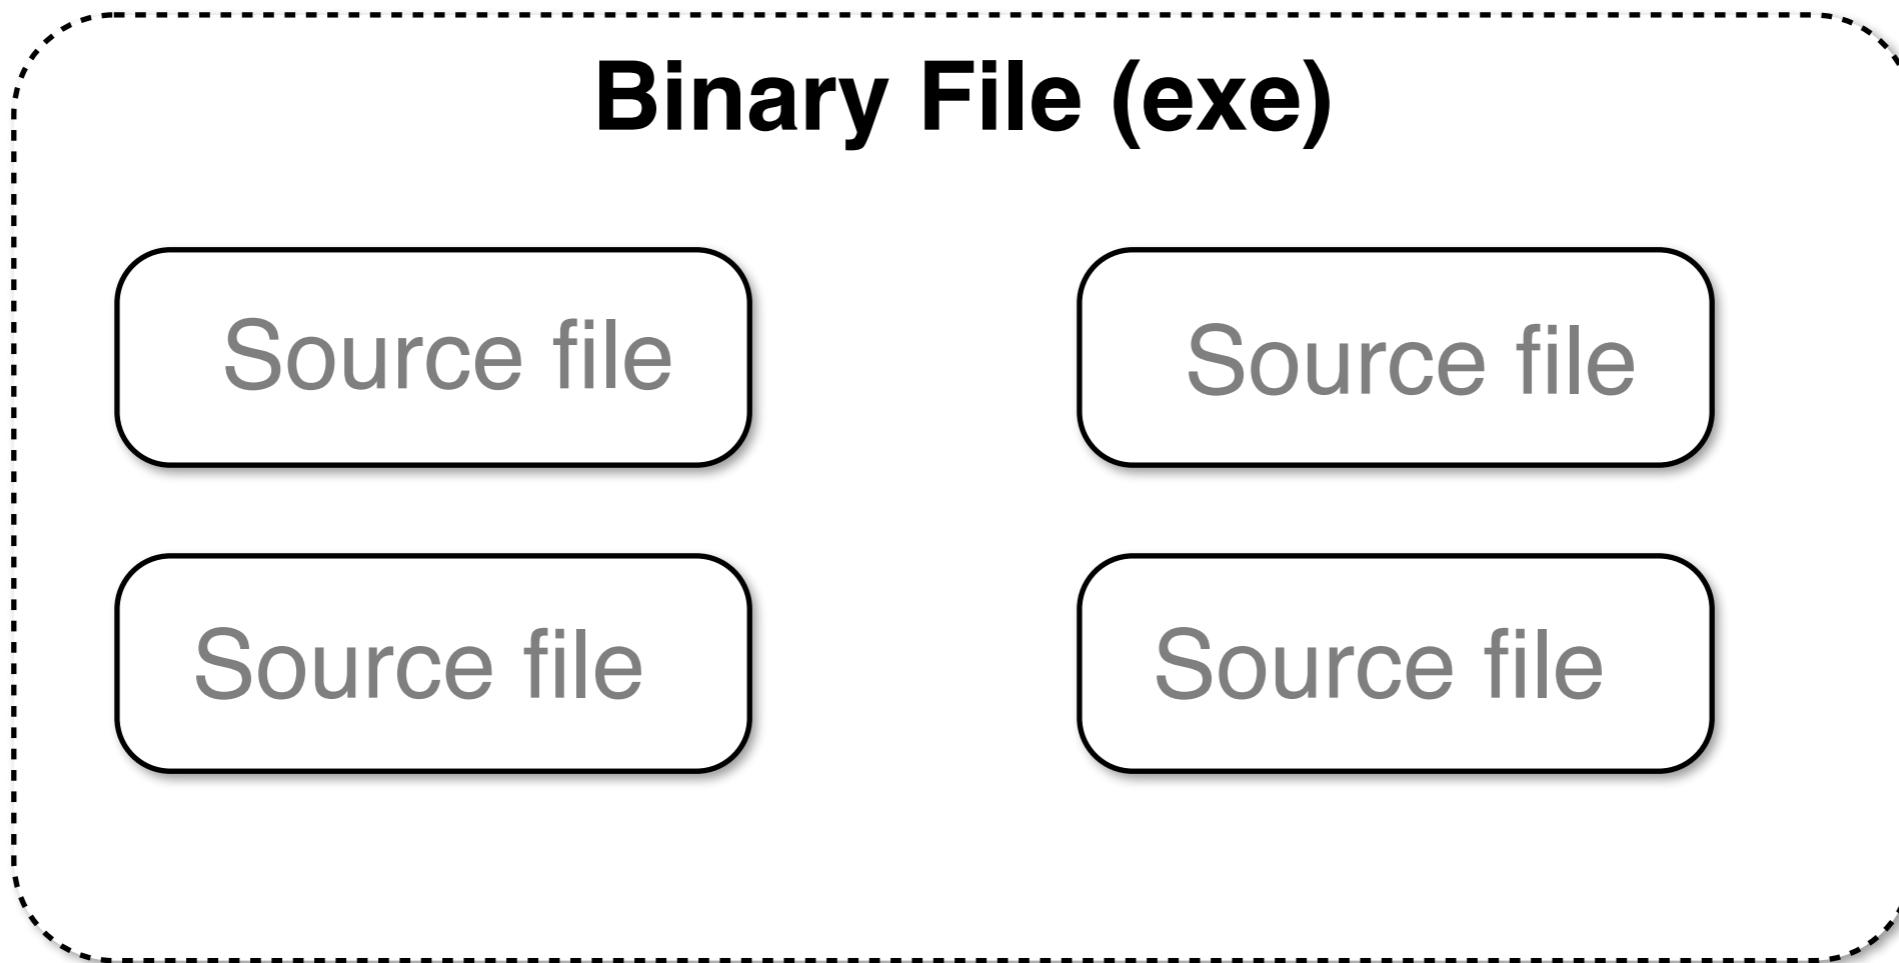
- > Your Technology stack
- > Total Years of experience



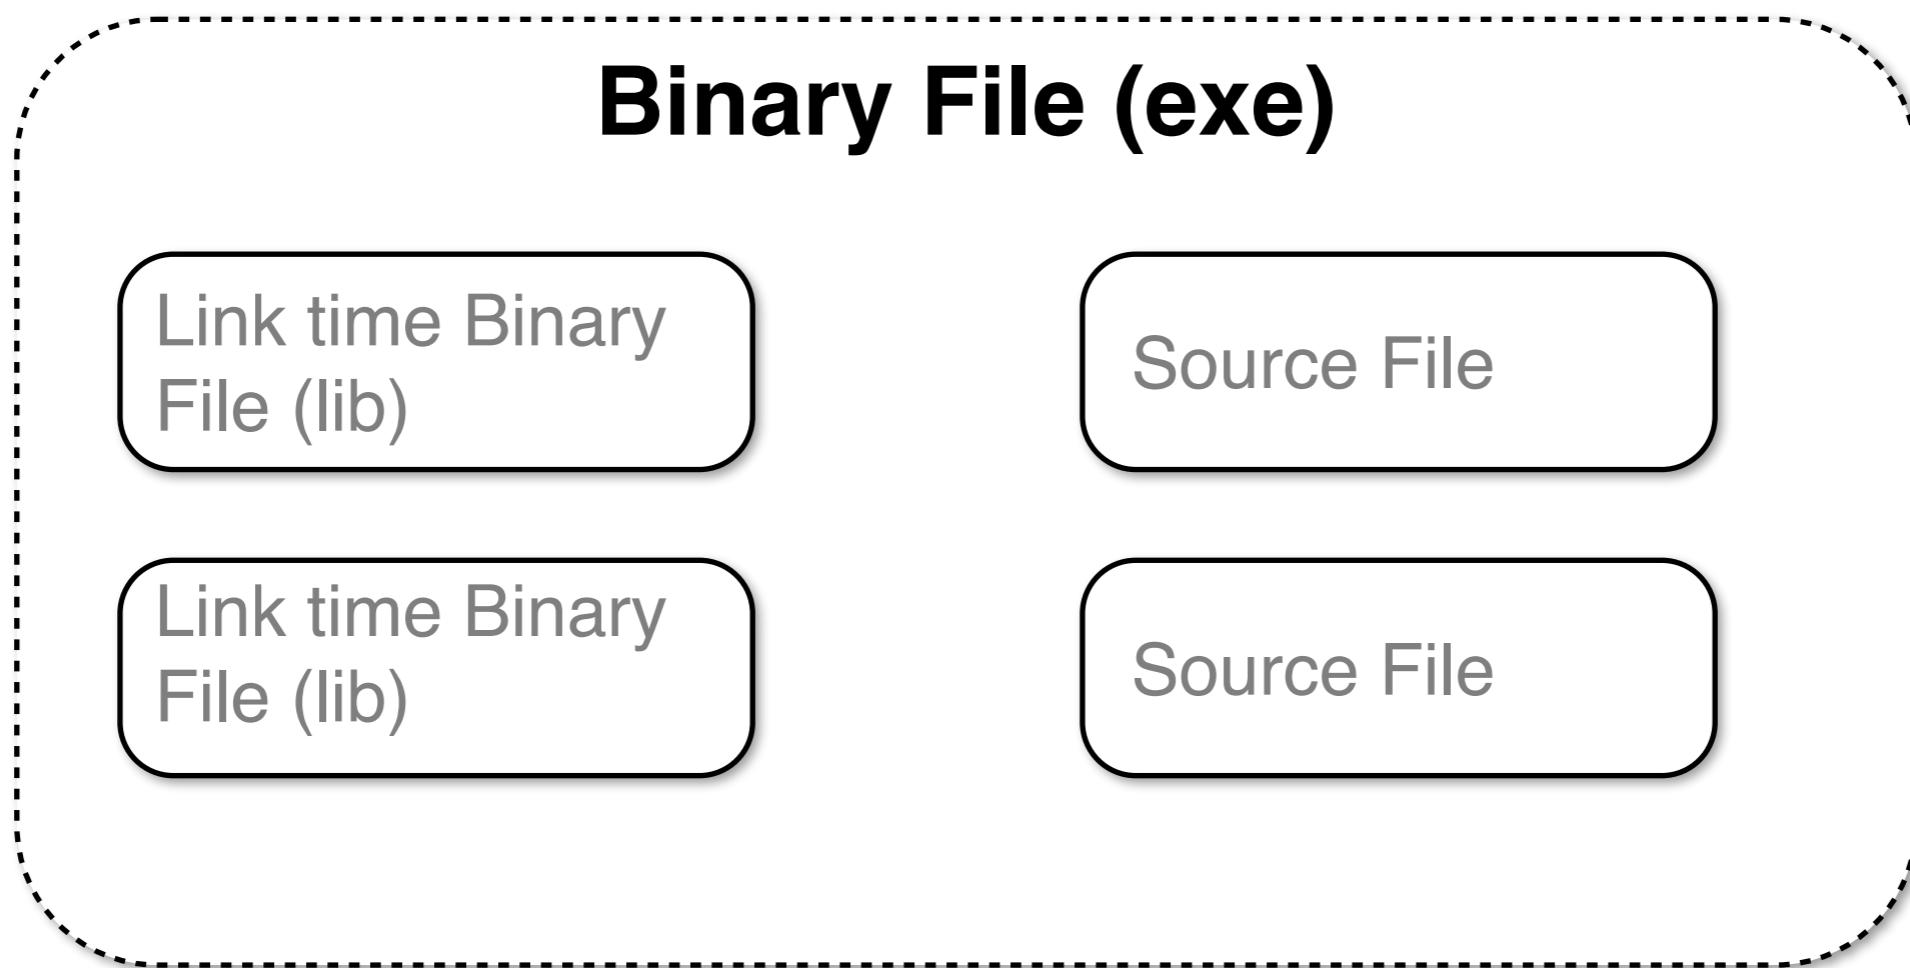
Day 1

What is Microservice ?

Completetime Monolithic



Linktime Monolithic / Monolithic Deployment



Runtime Monolithic

Monolithic Process

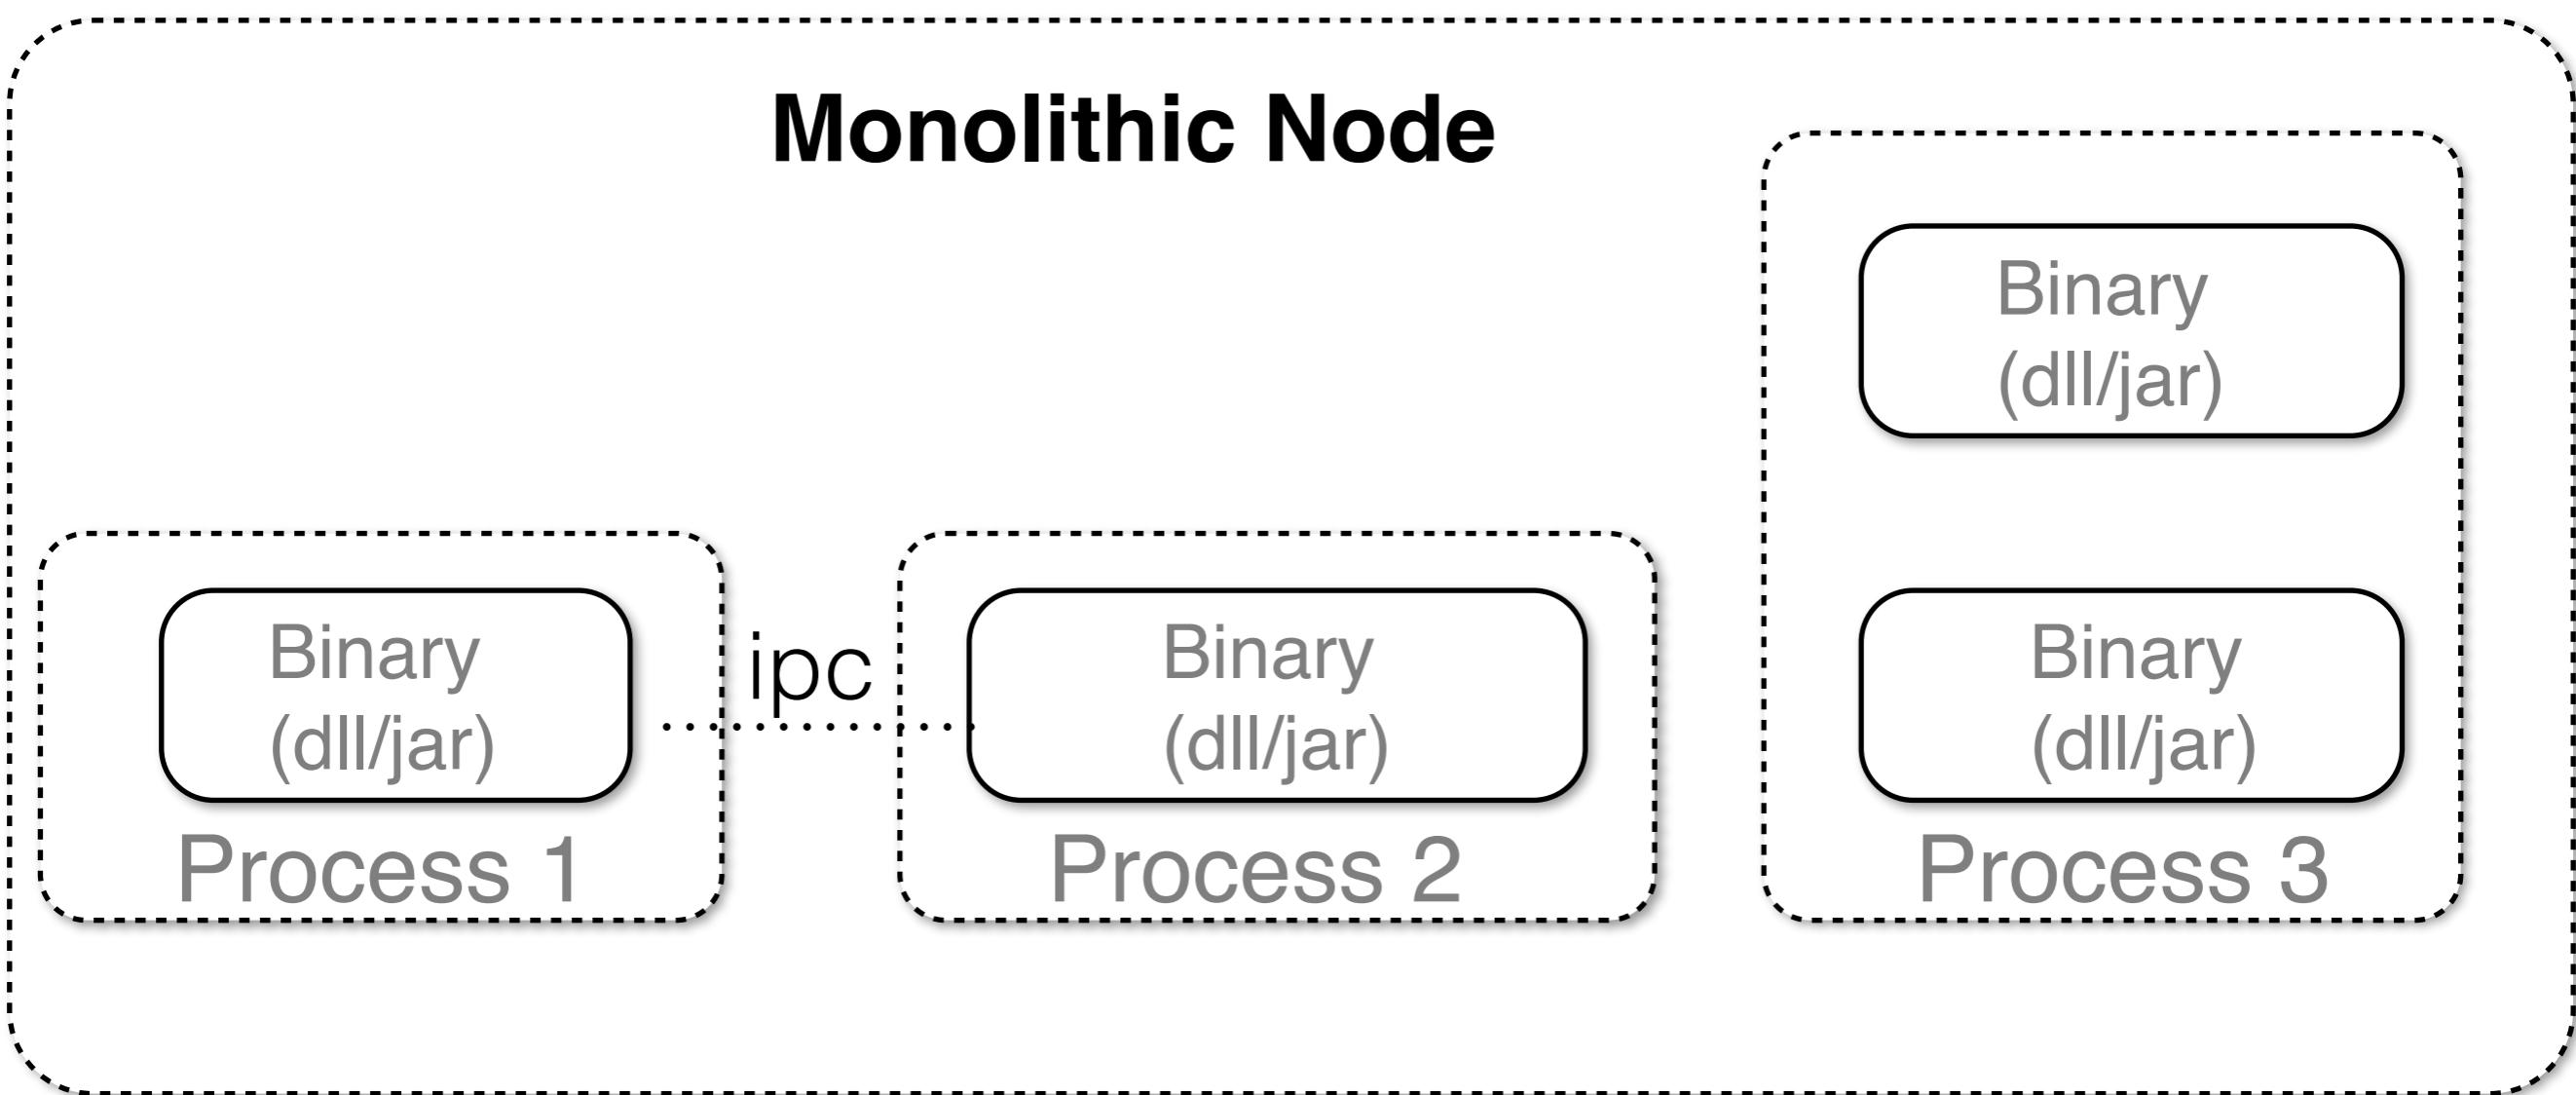
Binary (dll/jar)

Binary (dll/jar)

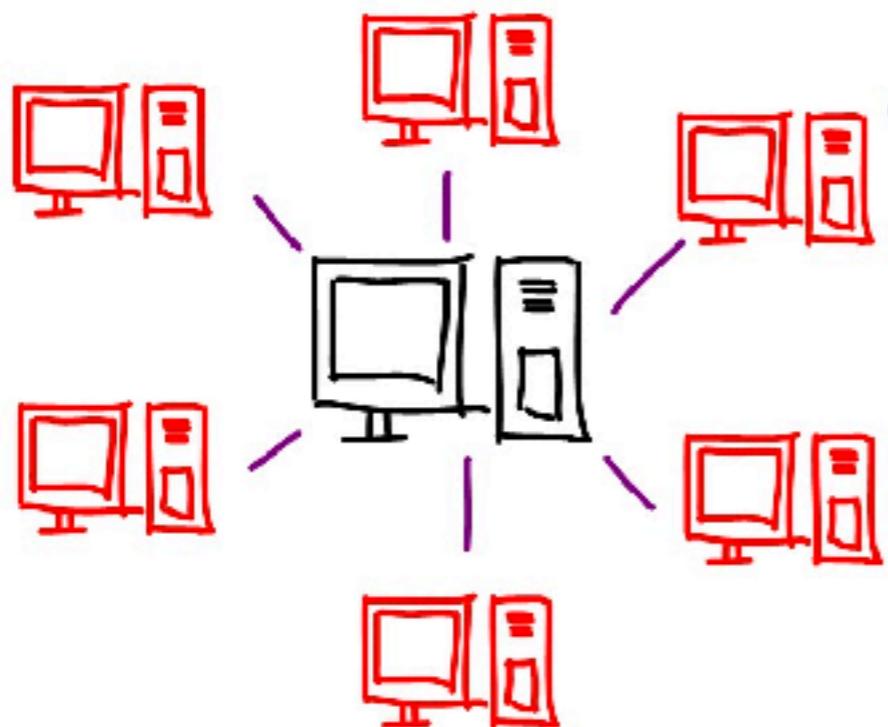
Binary (dll/jar)

Binary (dll/jar)

Multi Process Application

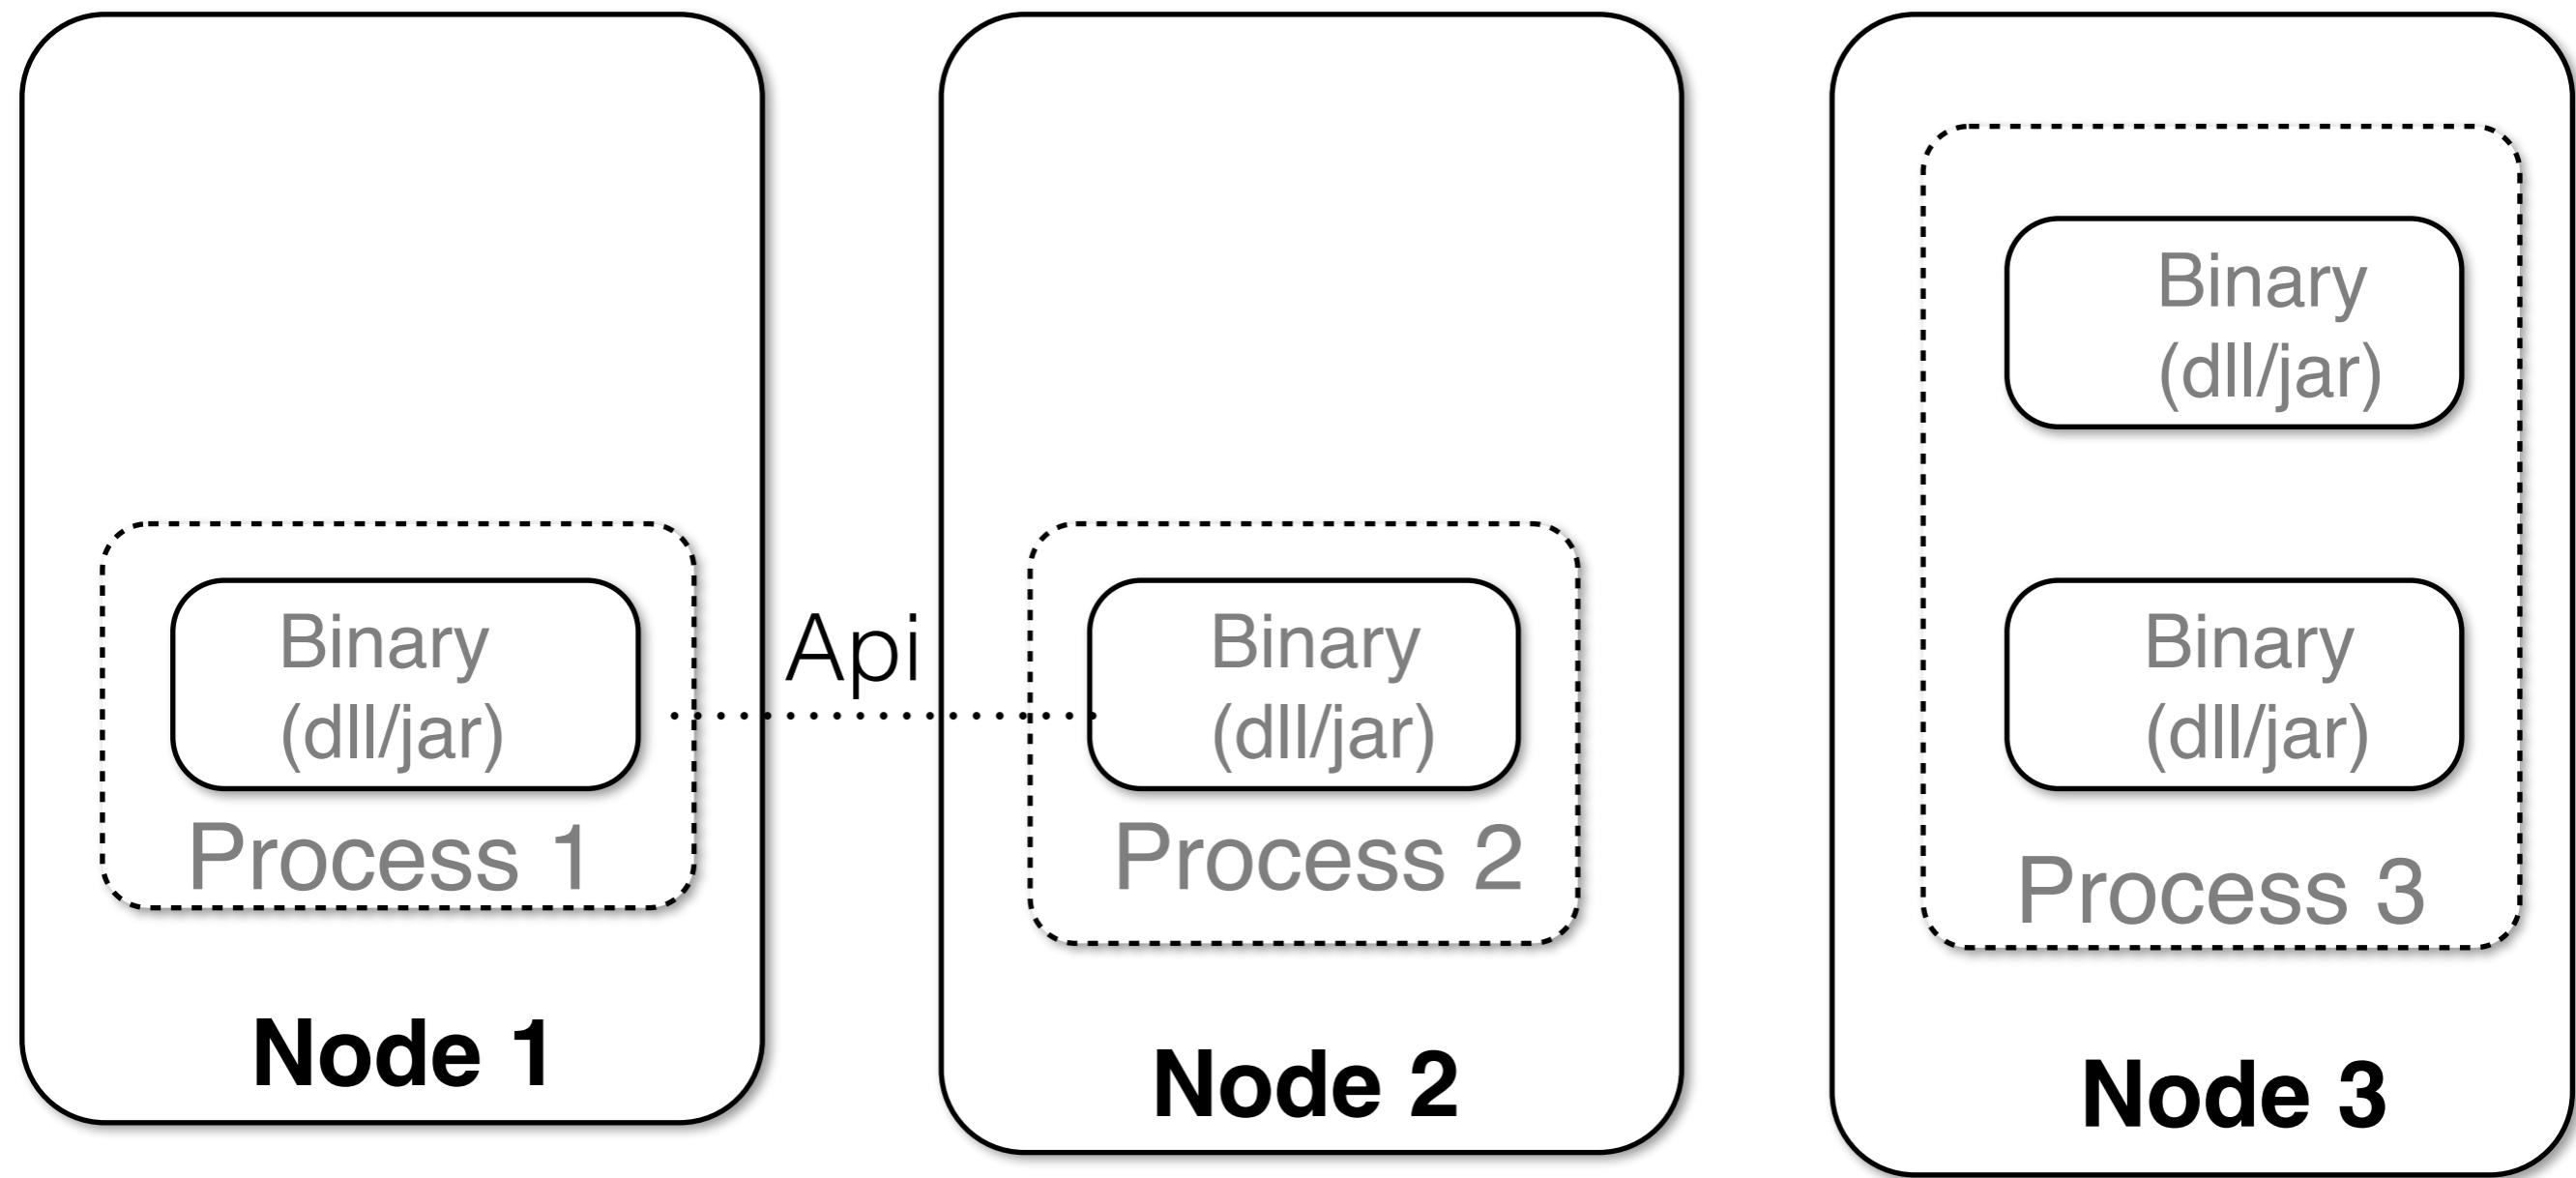


Distributed Computing

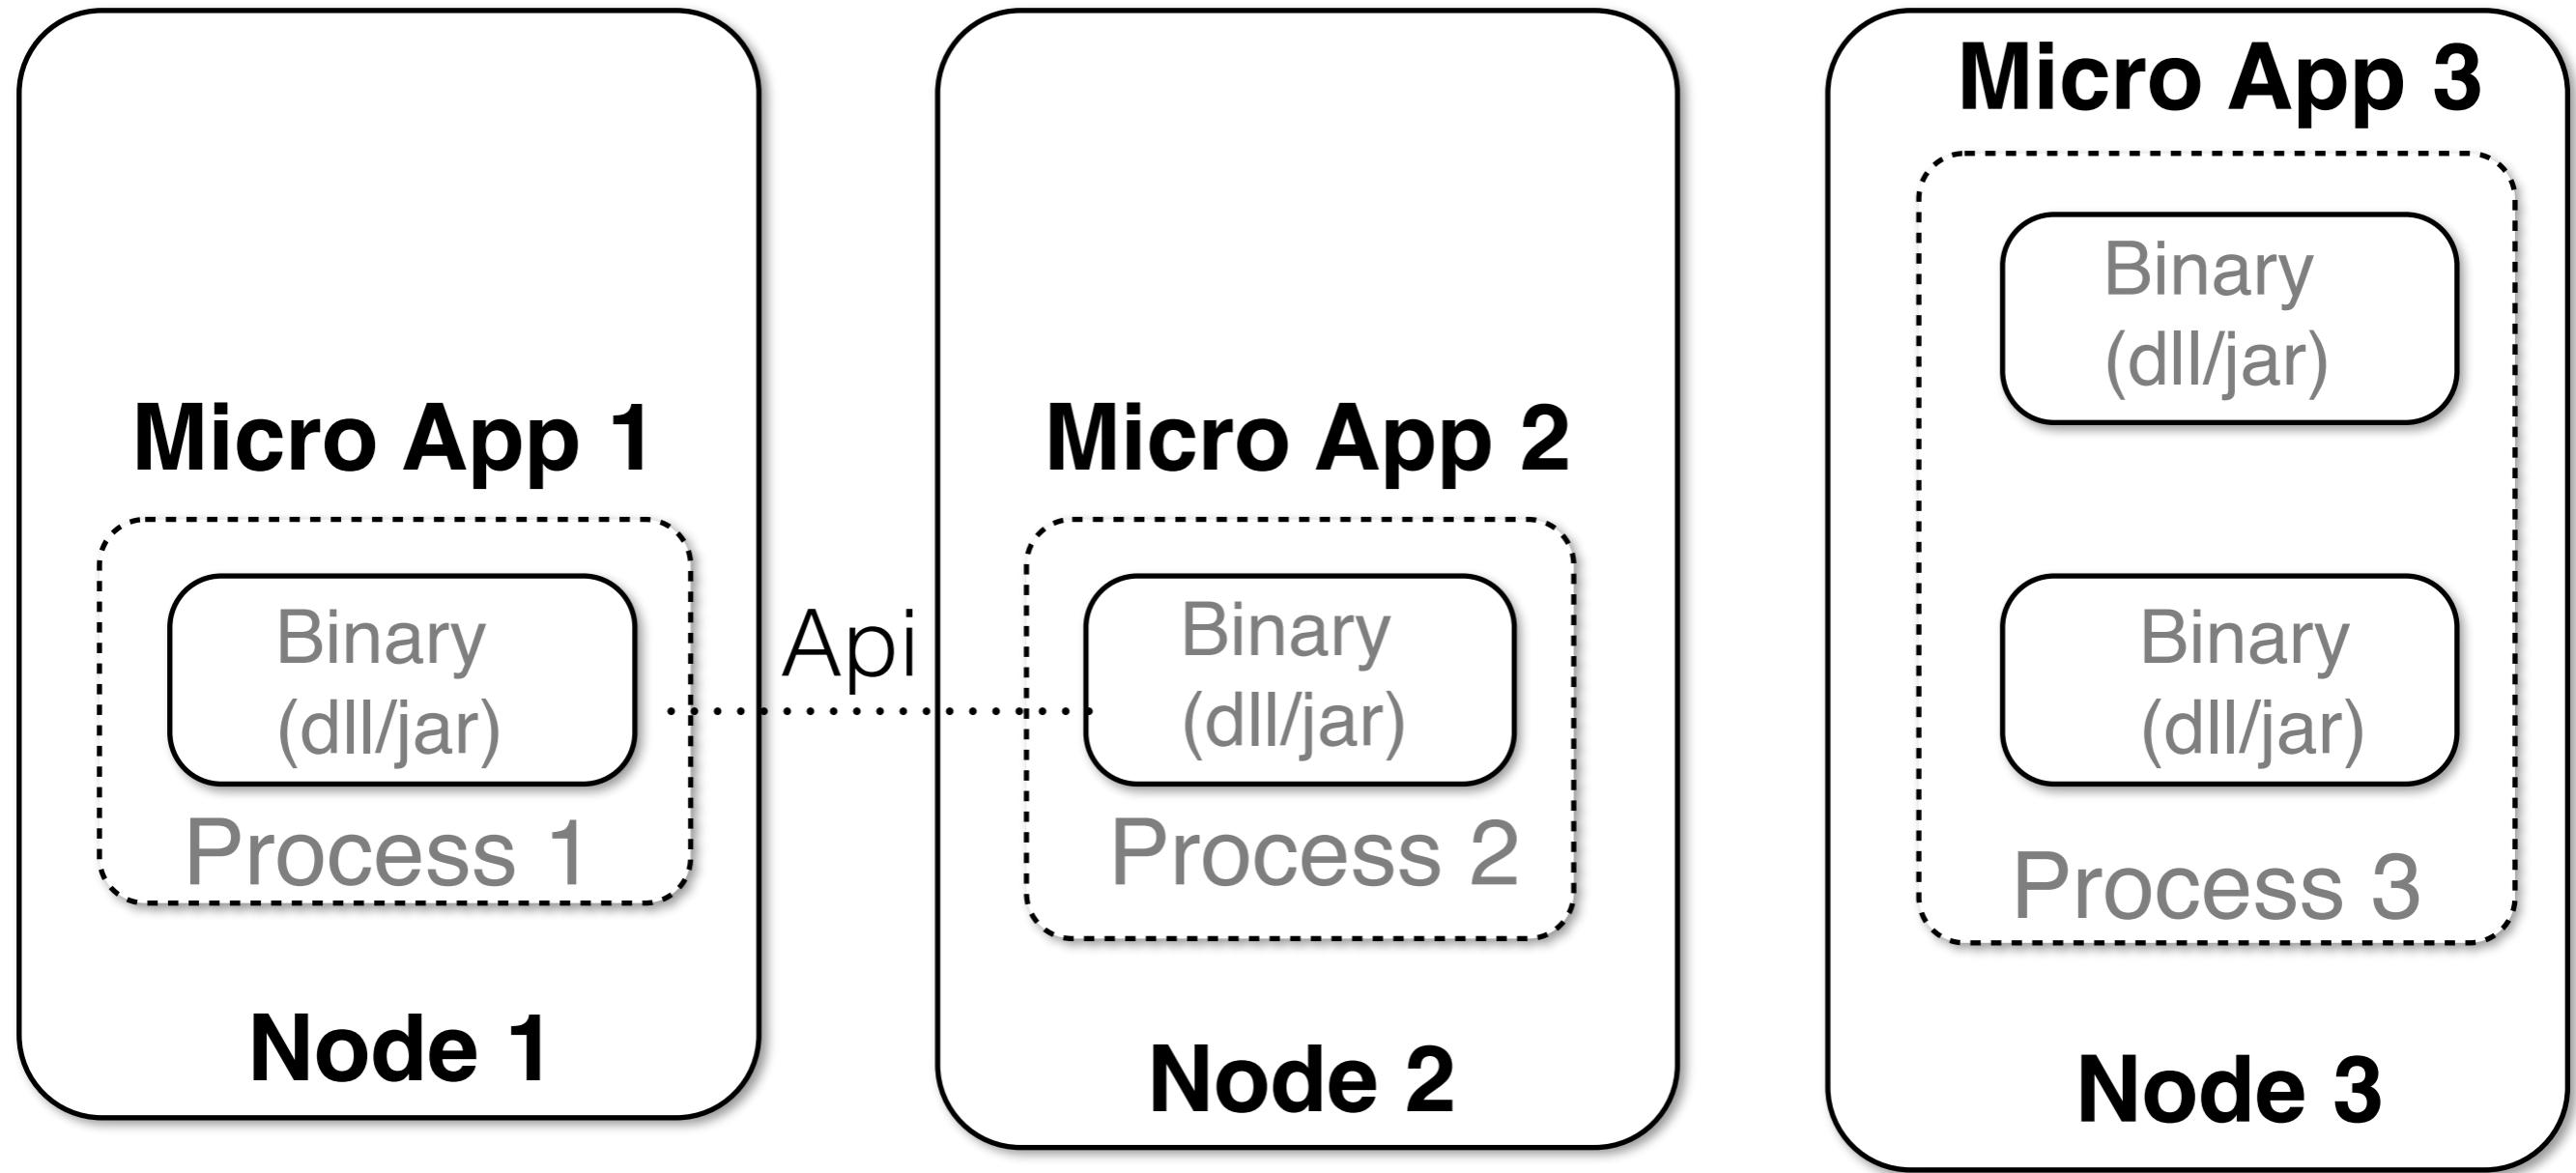


- DCOM
- CORBA
- SOAP
- Message Queue
- REST
- GRPC
- THRIFT

Distributed Application



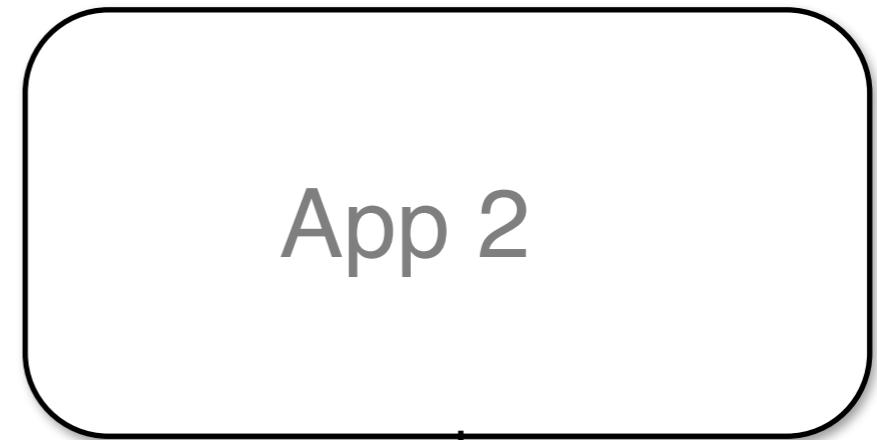
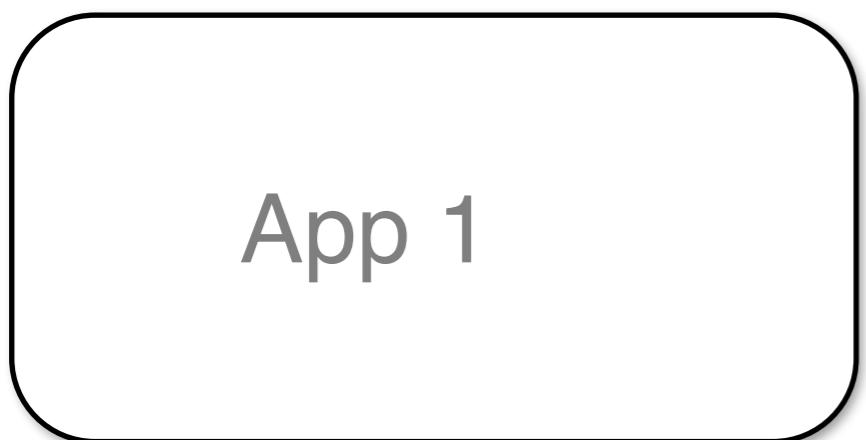
Micro Distributed Application





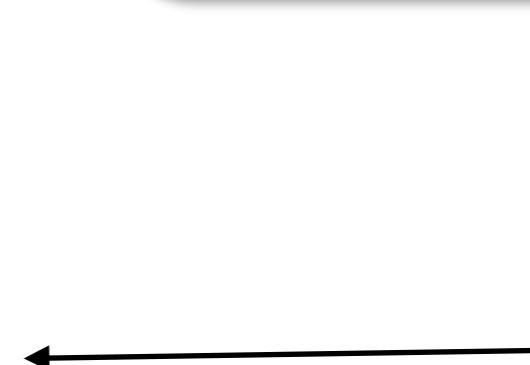
Traditional Distributed Application
vs
Micro Distributed Application

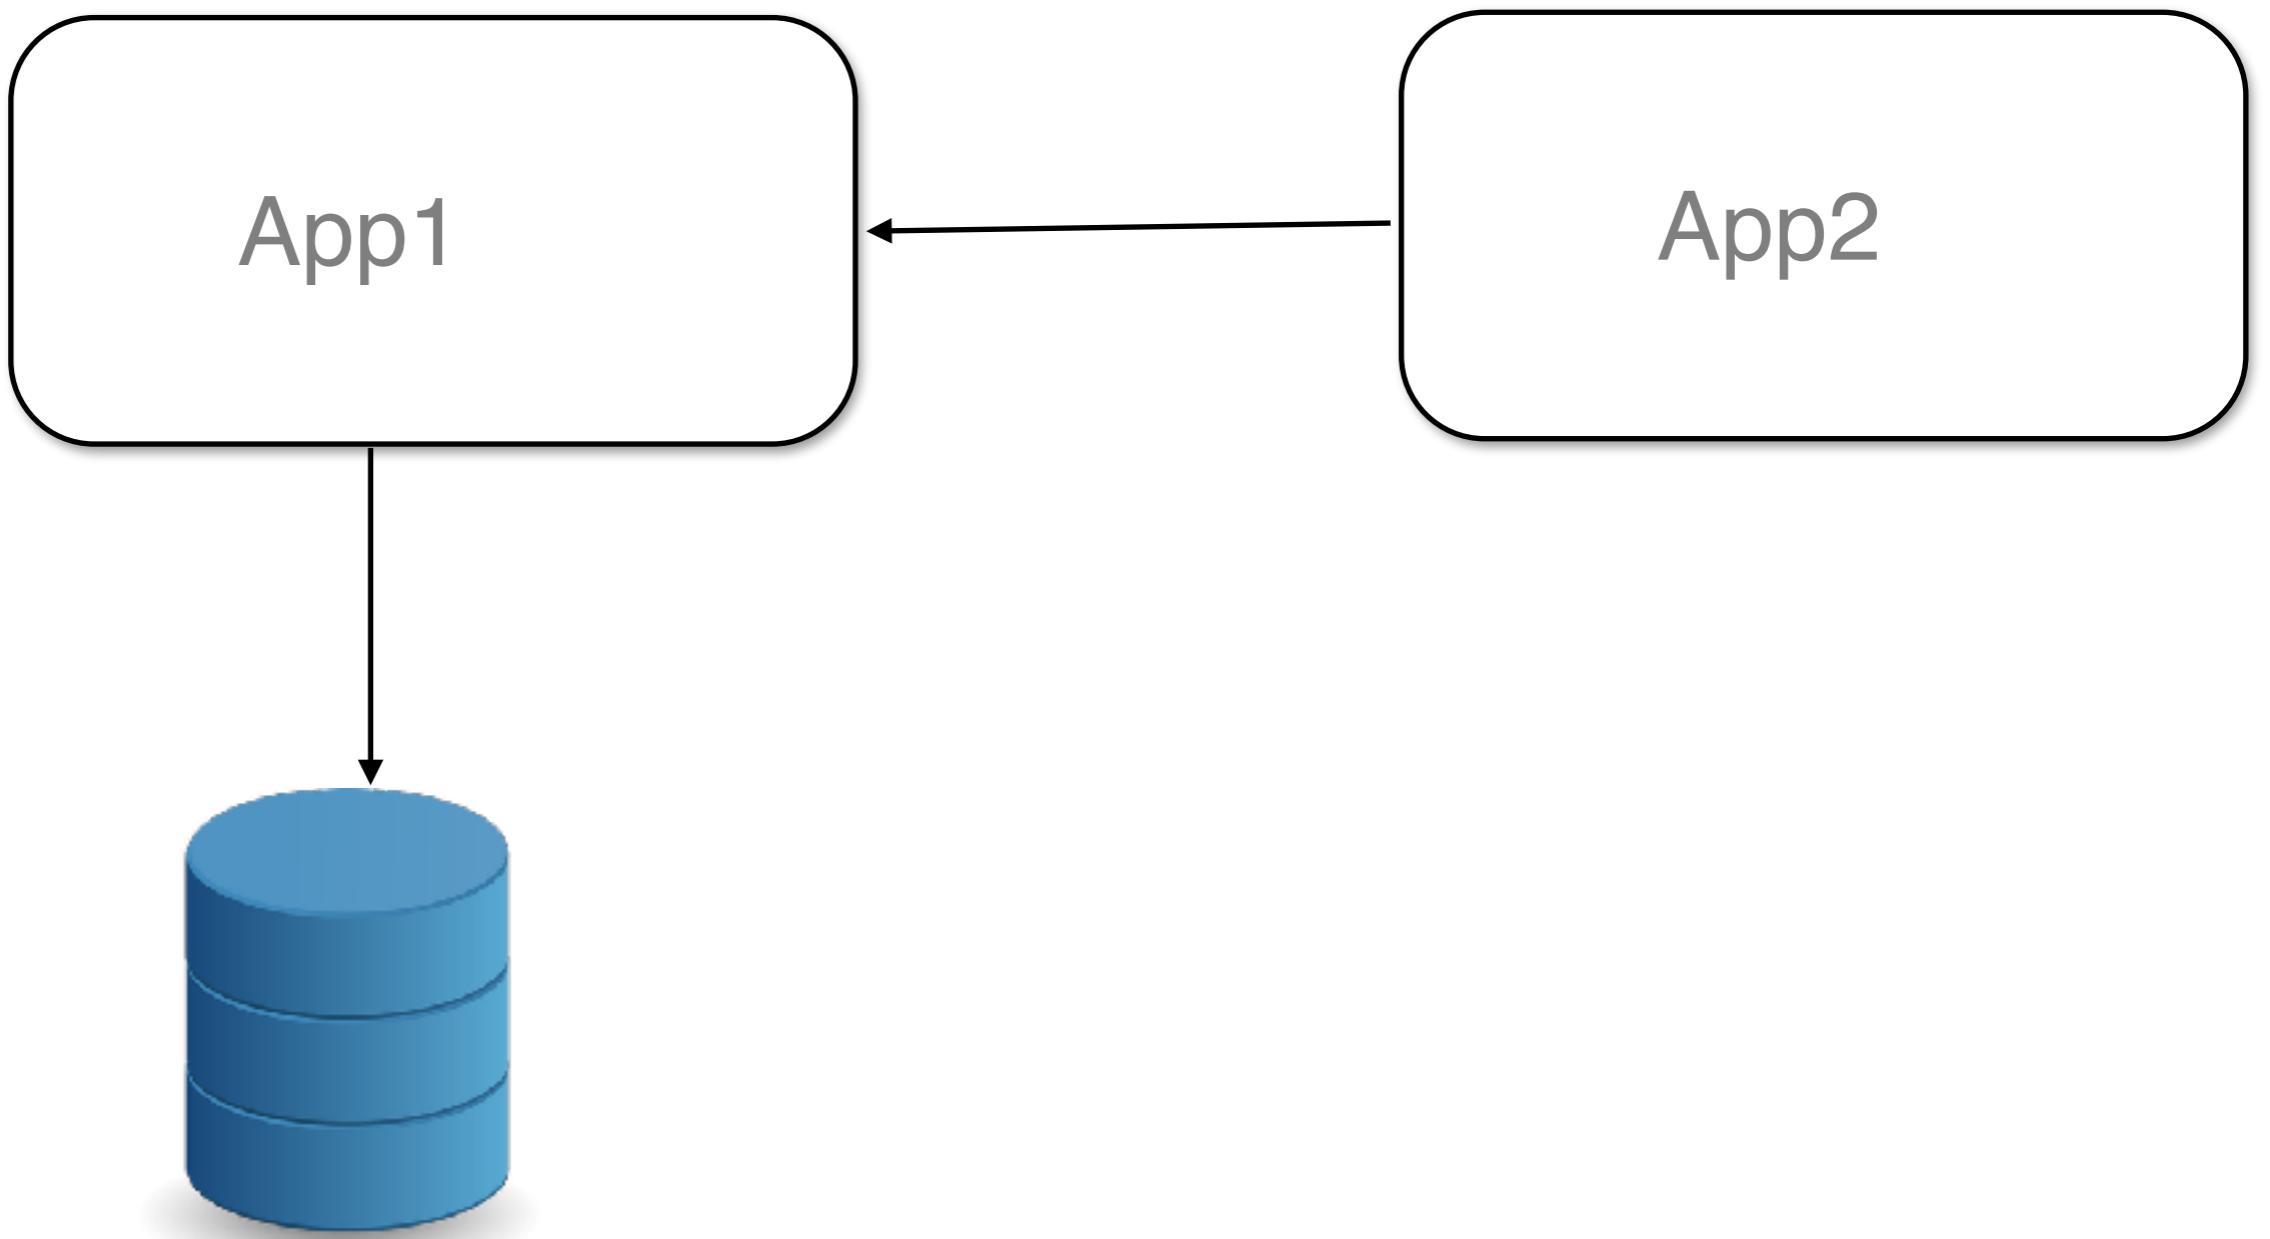
	Monolithic	Micro service
Database / Storage	Shared	Not shared
Infra (Hosting)	Shared	Not Shared
CI/CD (Build Server)	Shared	Not Shared
Fun Requirements	Shared	Not Shared
SCRUM Team / Sprint	Shared	Not Shared
Test Cases	Shared	Not Shared
Sorce Control	Shared	Not Shared
Architecture	Shared	Not Shared
Technology Stack / Fwks	Shared	Not shared



App 1

App 2





Virtual Machine

App 1

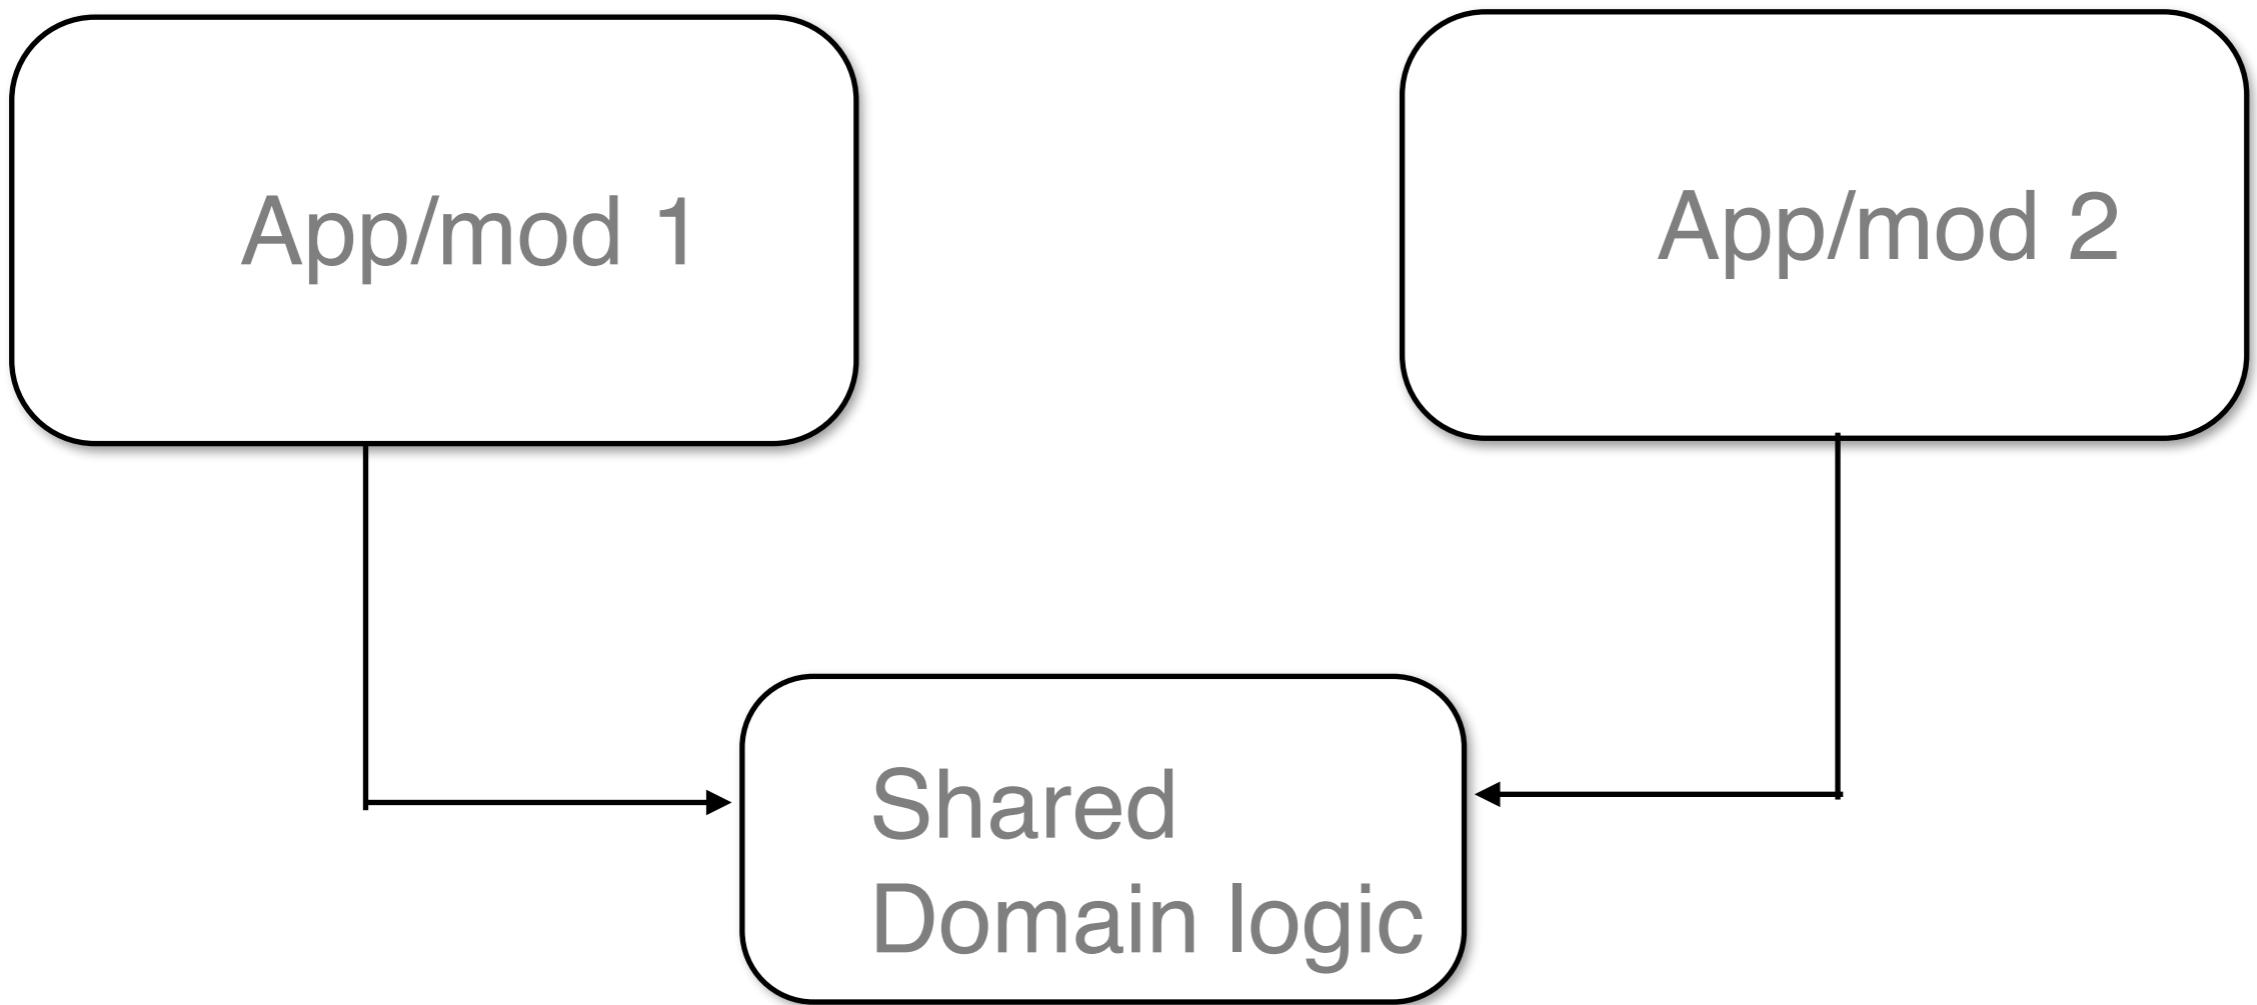
App 2

Virtual Machine

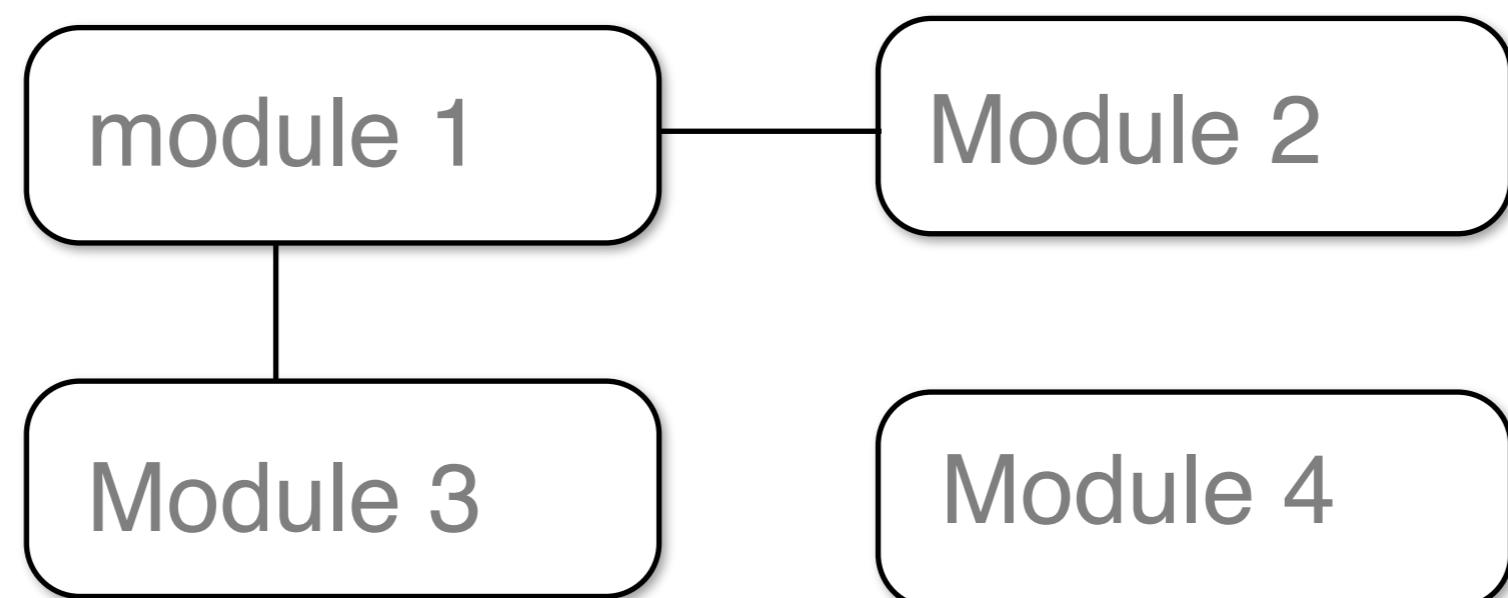
App 1

Virtual Machine

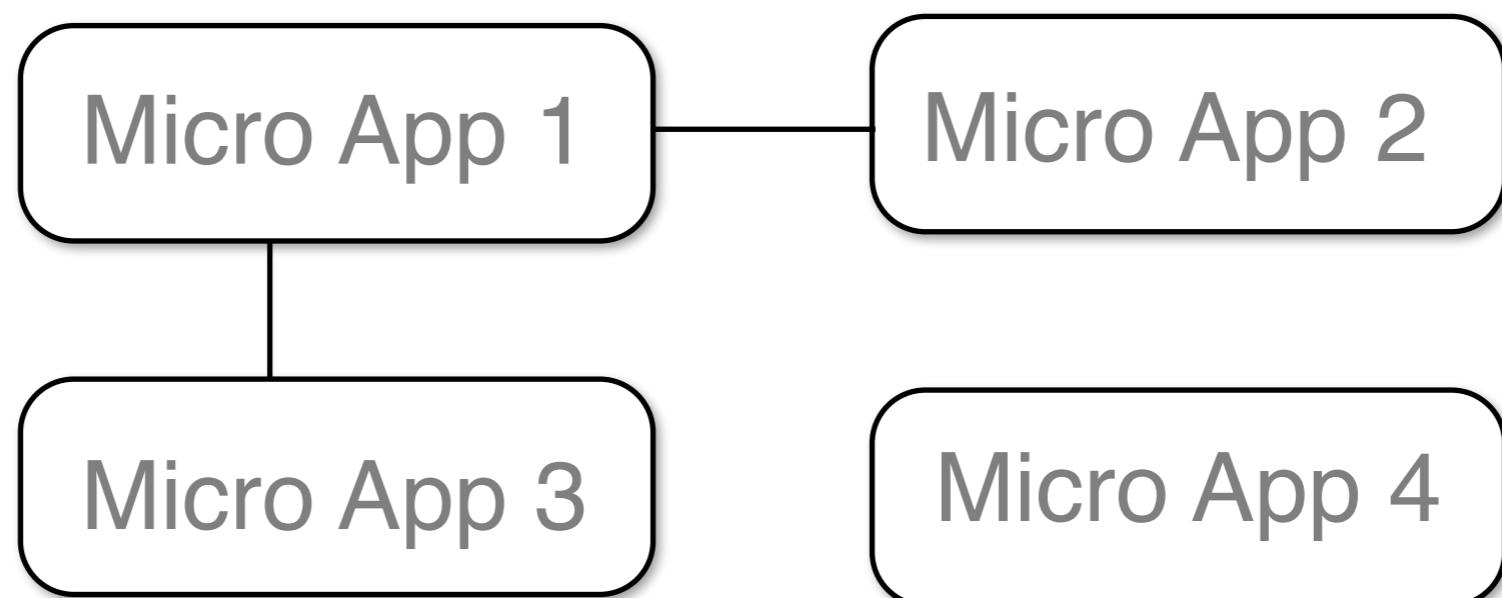
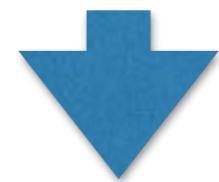
App 2



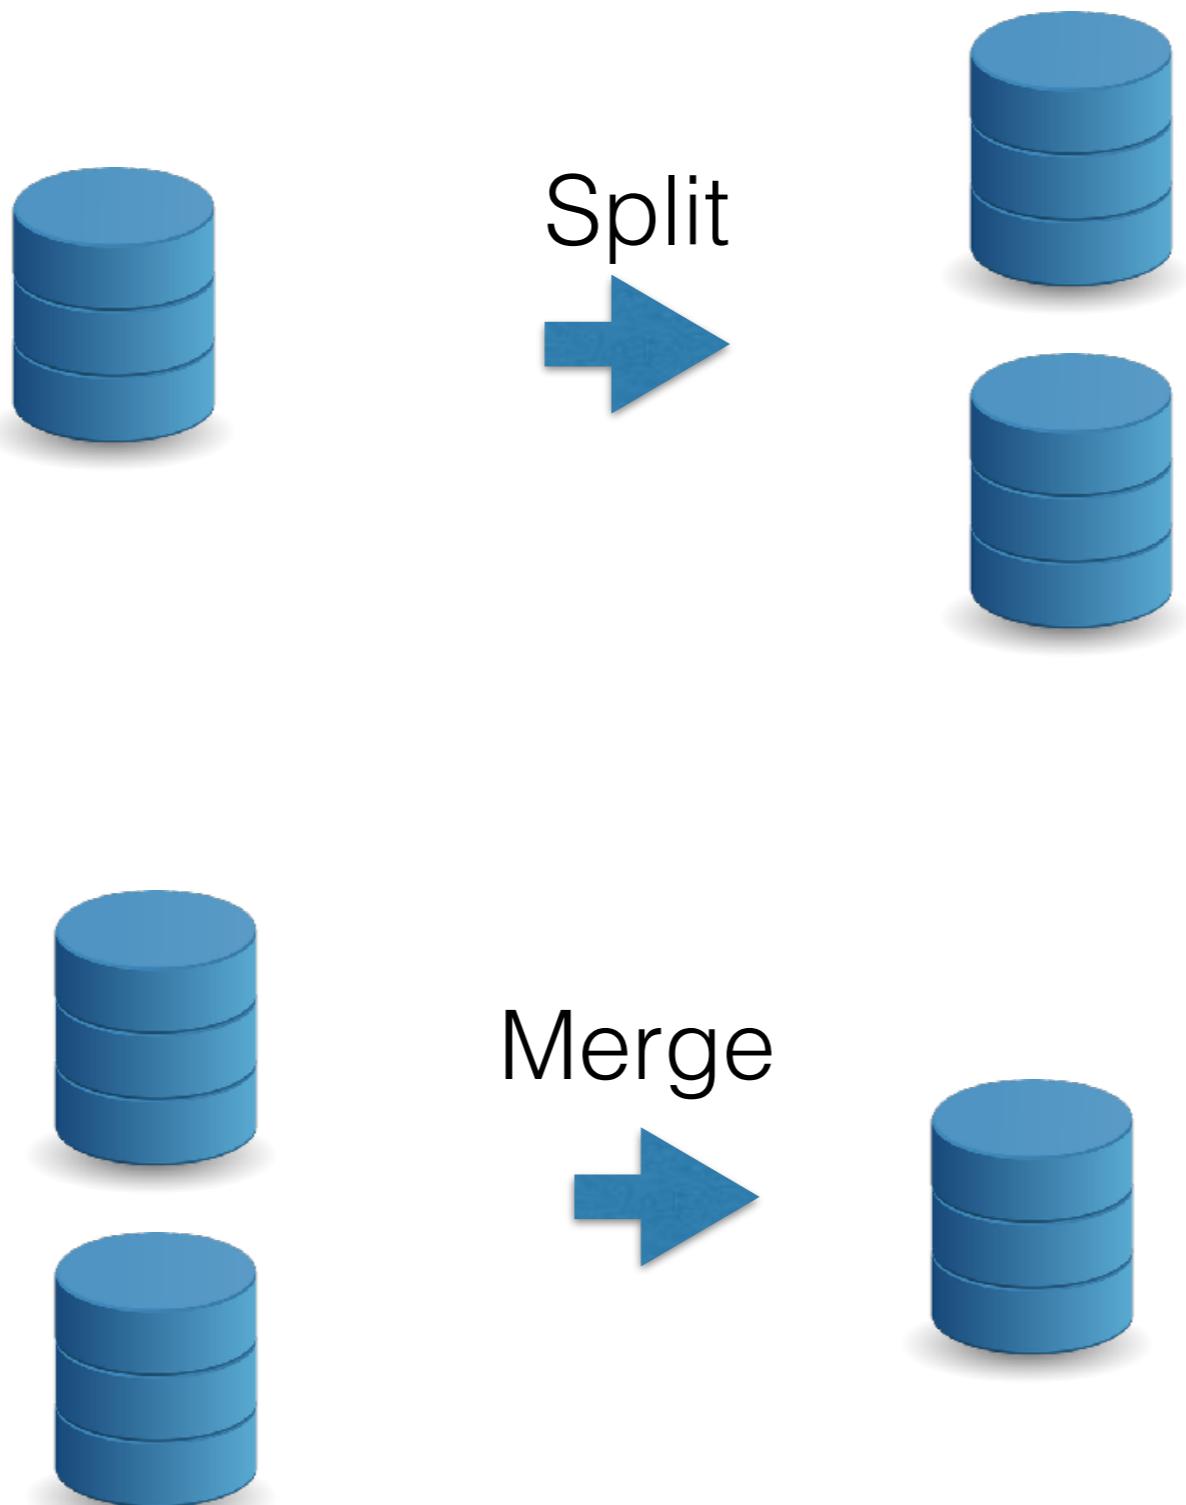
Application



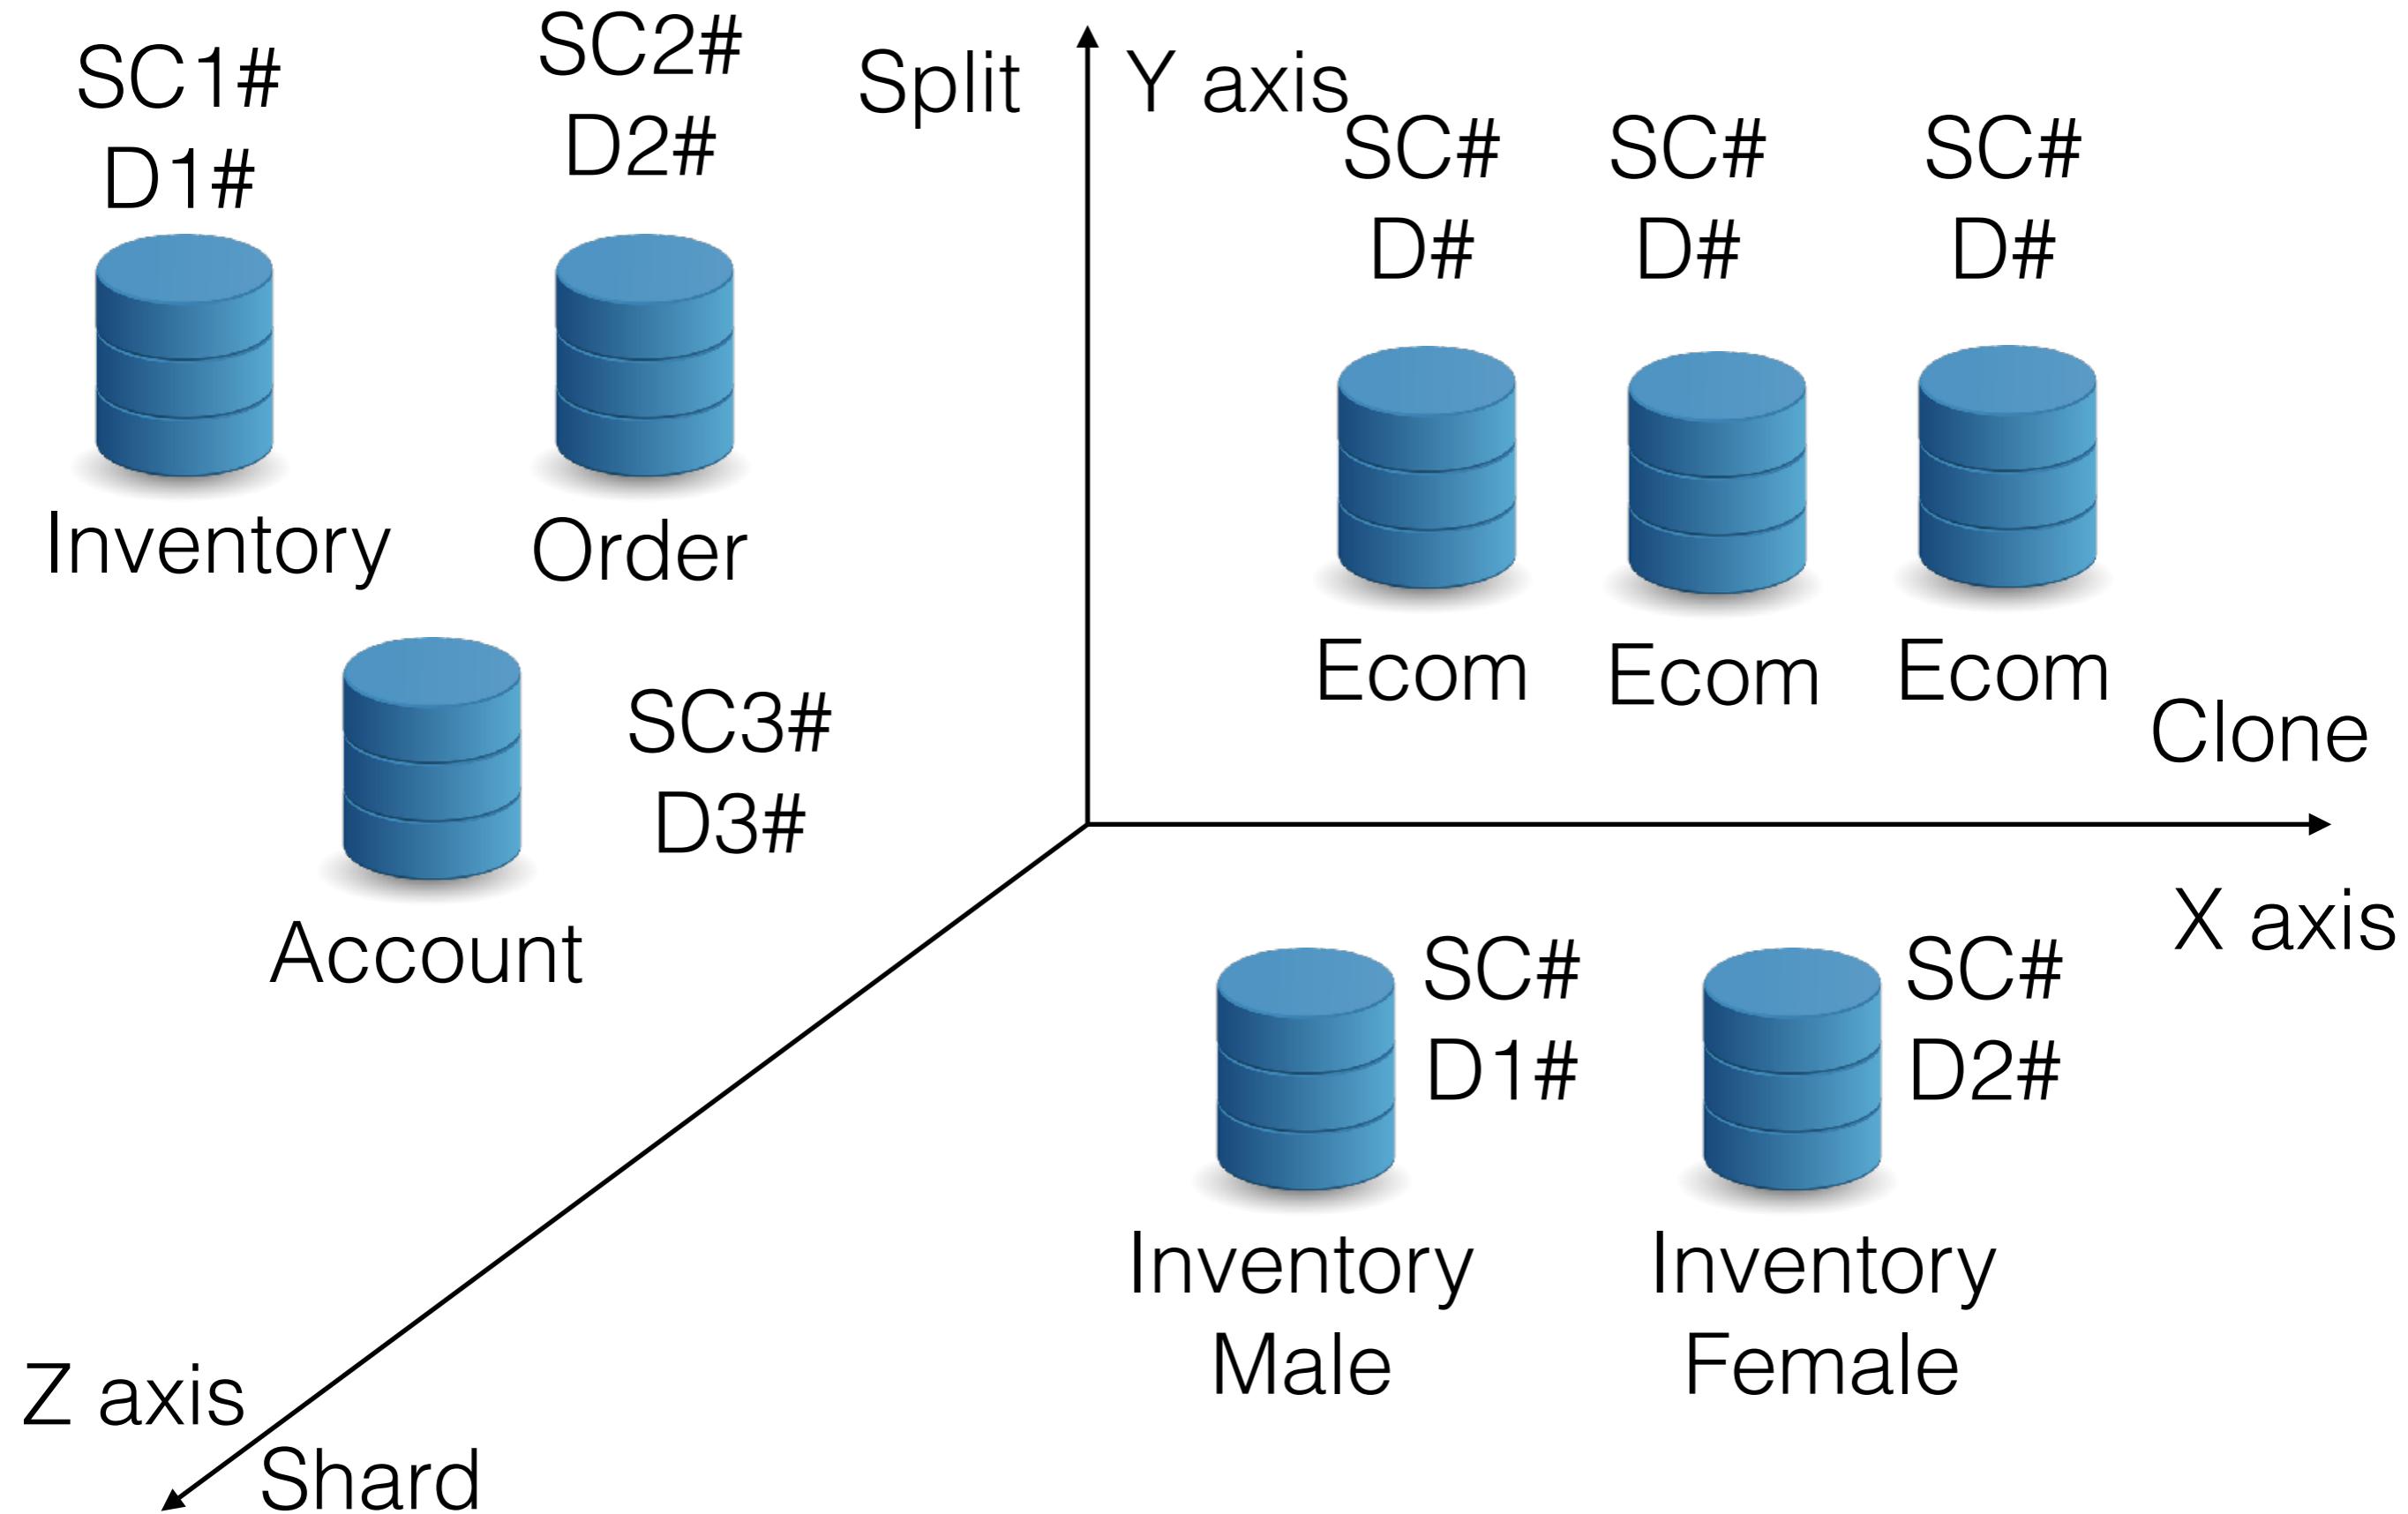
Application



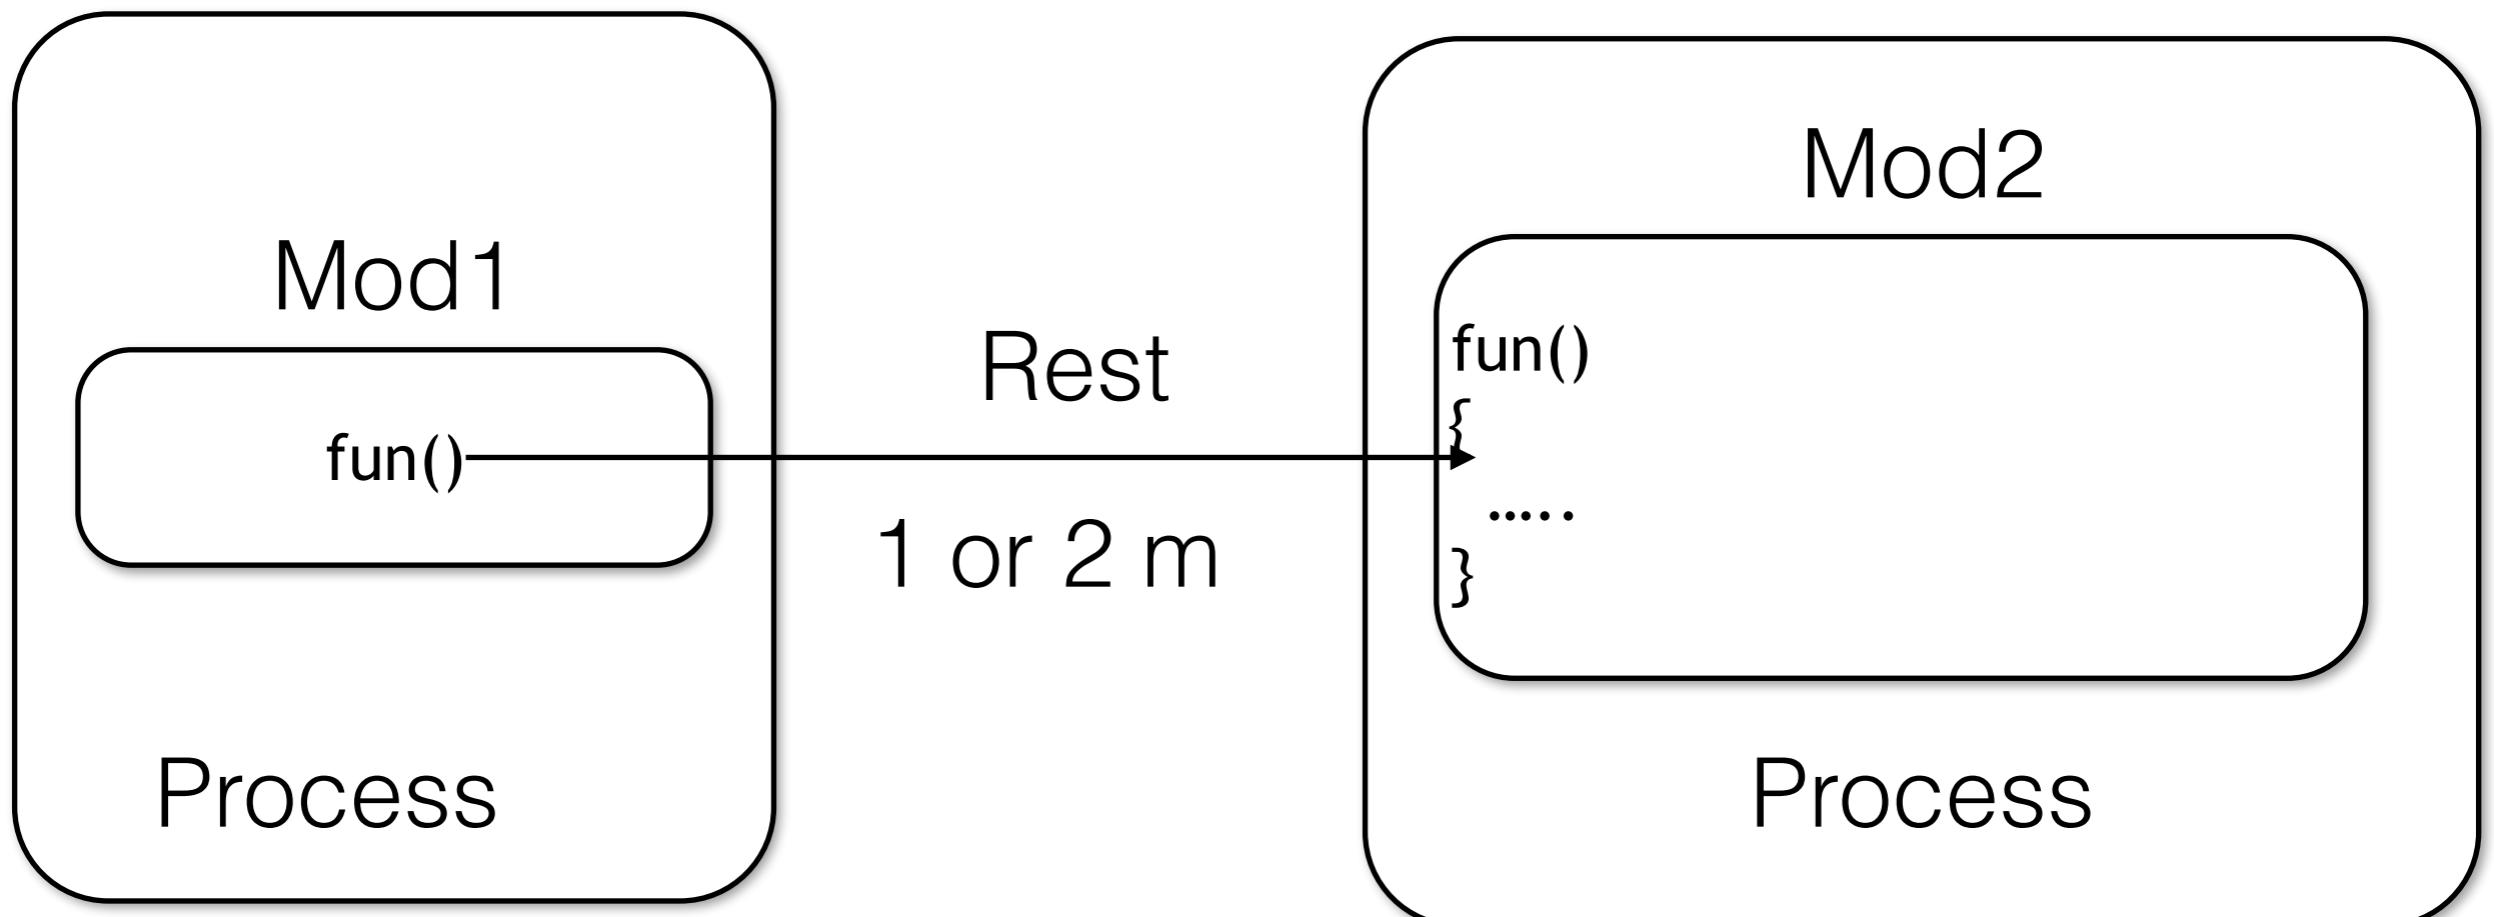
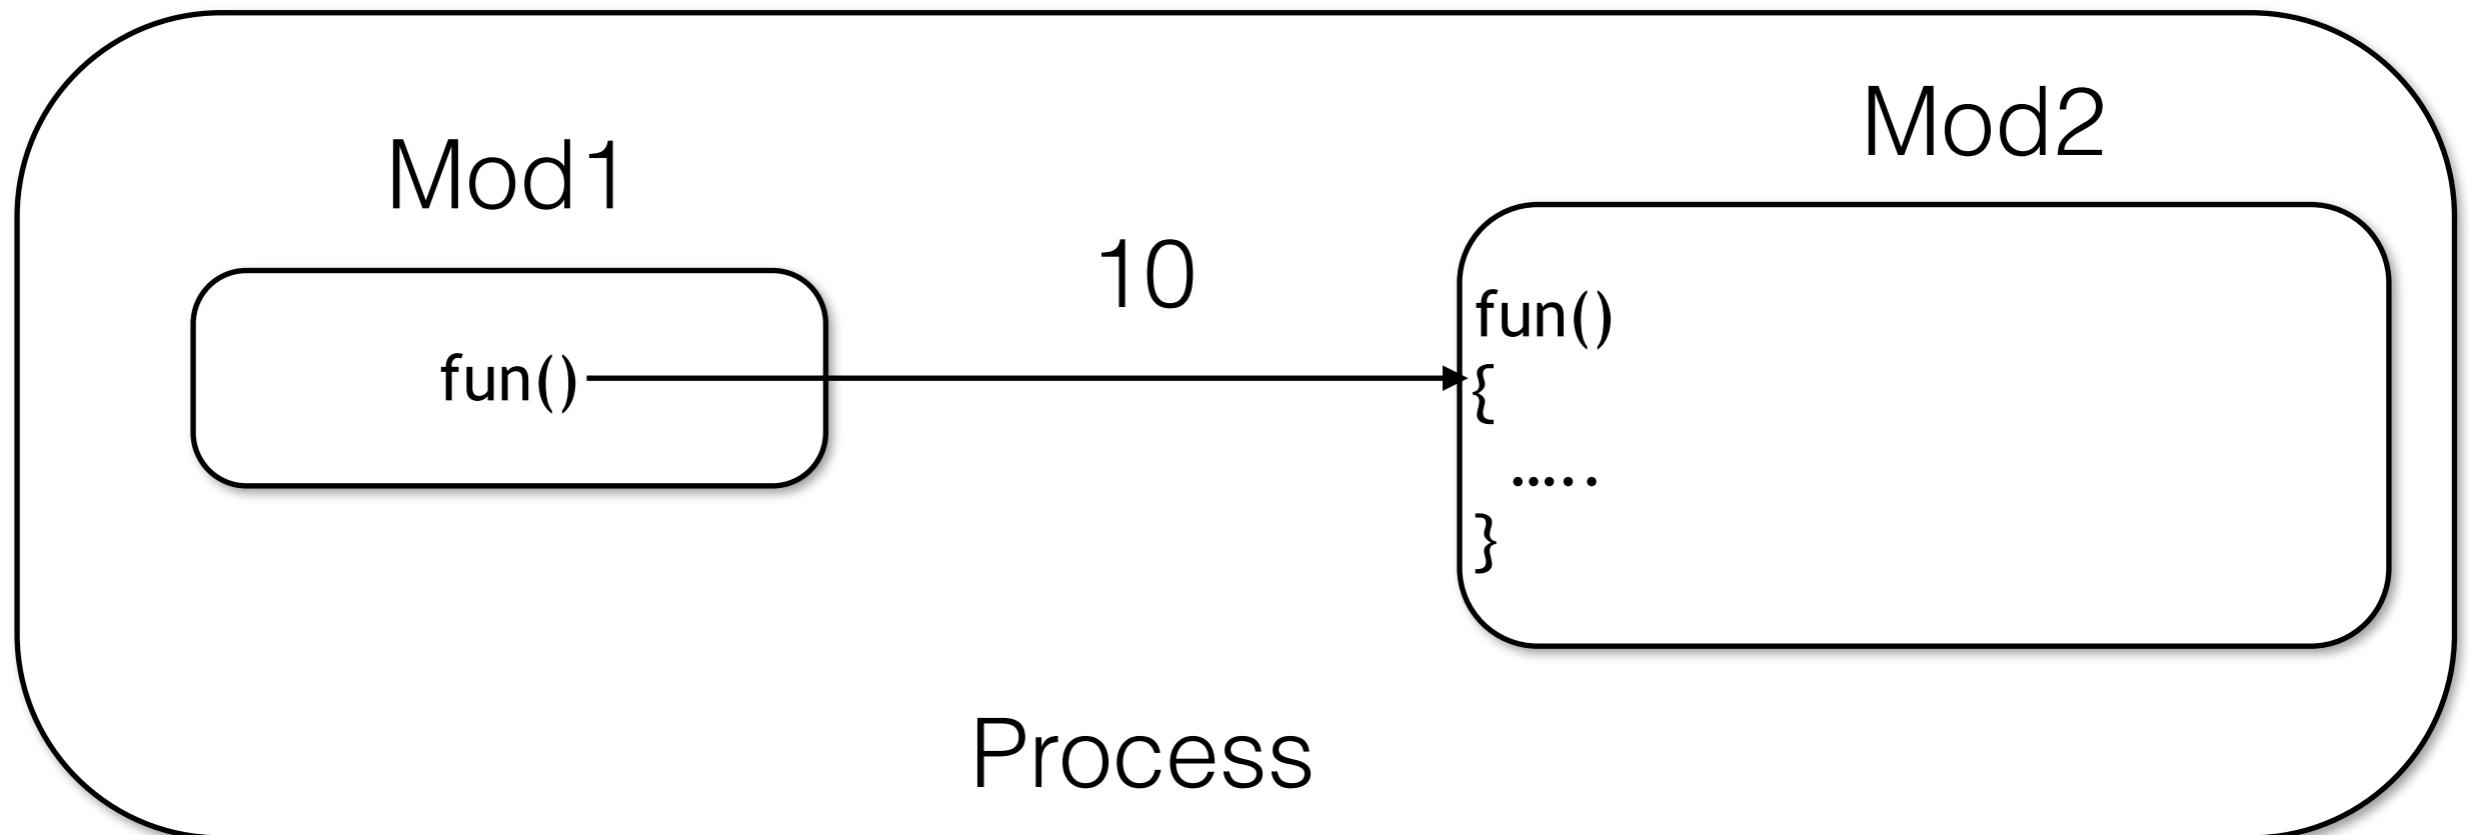
	Pros/ Cons
Performance	- -
Transaction Management	- -
Joins / Report / Materialized view	- -
Development time (1:3)	- -
Learning Curve	- -
Infra Cost	- -
Debugging (Finding bugs)	- -
Integration Test	- -
Log Mgmt	- -
Config Mgmt	- -
Authentication	- -
Authorization	- -
Monitoring / Alerting	- -
CI/CD	- -
Agile Architecture (Incremental arch)	+ +
Feature Shipping	+ +
Scalability	+ +
Resilient	+ +
Polygot	+ +
Maintainability	+ +



Scalability Cube - 50 rules for high Scalability

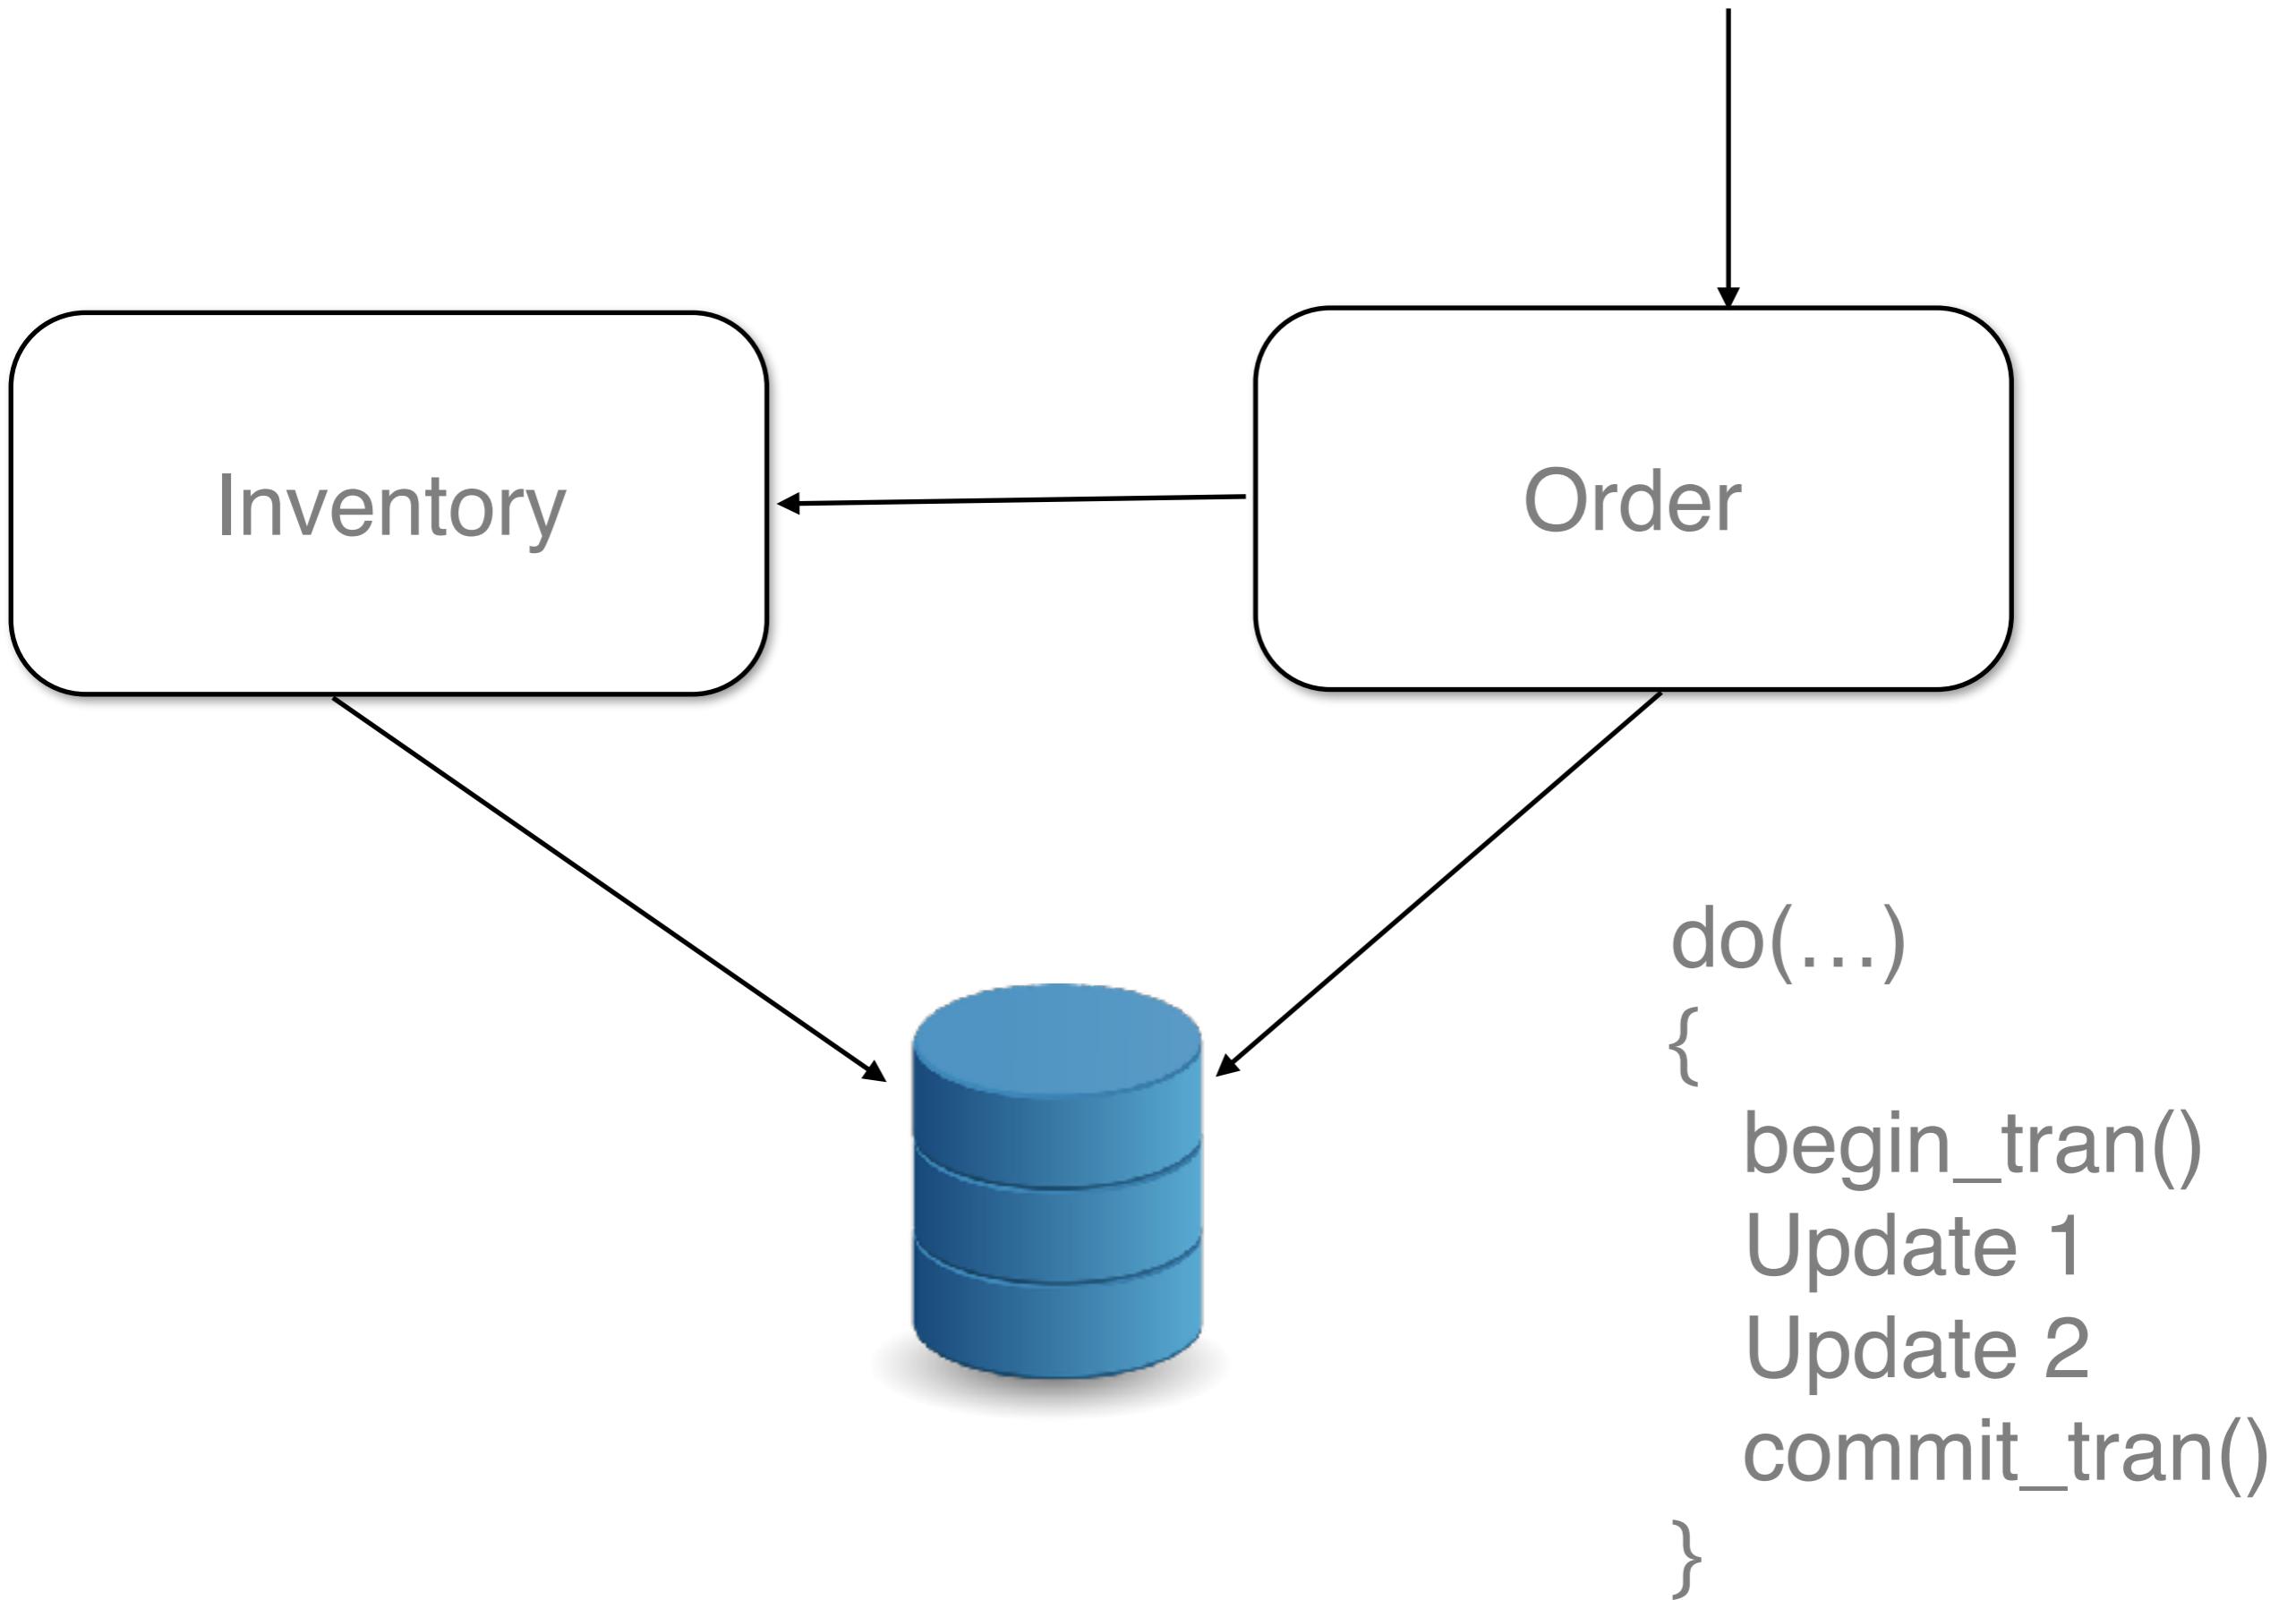


YAGNI - you aren't gonna need it
KISS
Solve tomorrow's problem tomorrow

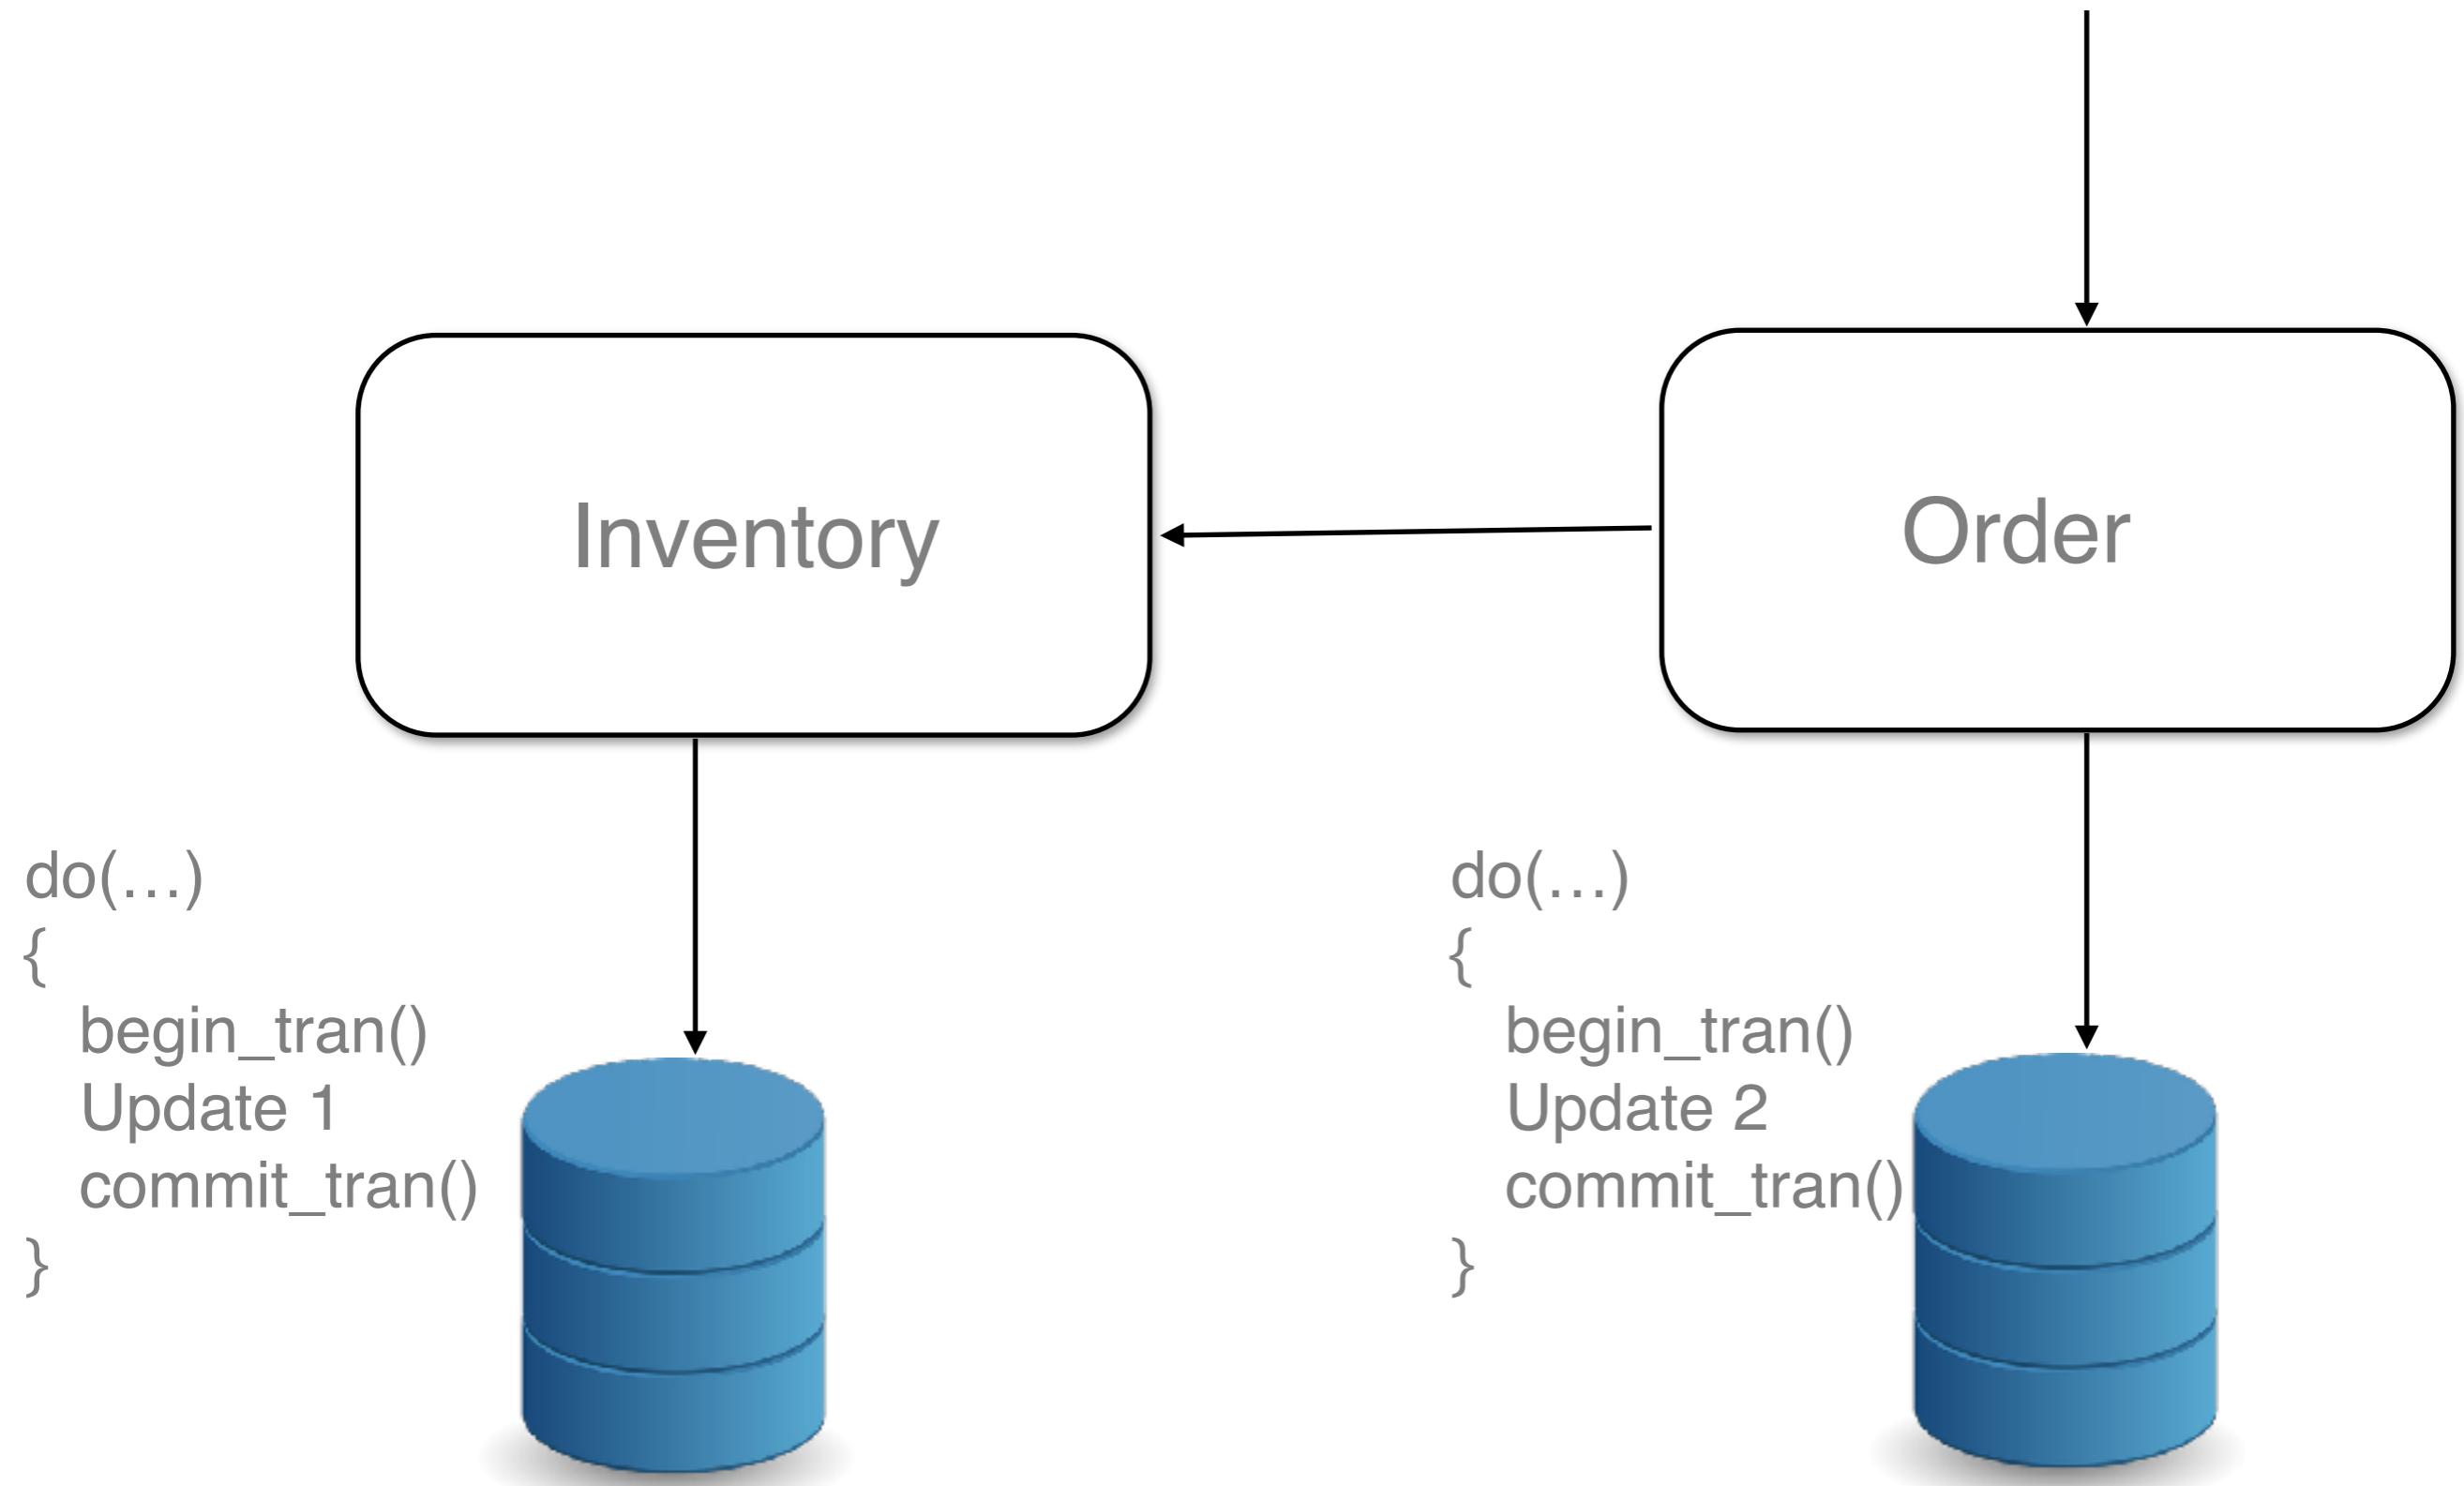


Operation	Cpu Cycles
• $10 + 12$	5
• Calling a in-memory Method	10
• Create Thread	2,00,000
• Destroy Thread	1,00,000
• Database Call	40,00,000
• Distributed Fun Call	20,00,000

Transaction



Transaction



2 phase commit
JTX , MSDTC, ...

msdtc.begin_tran()

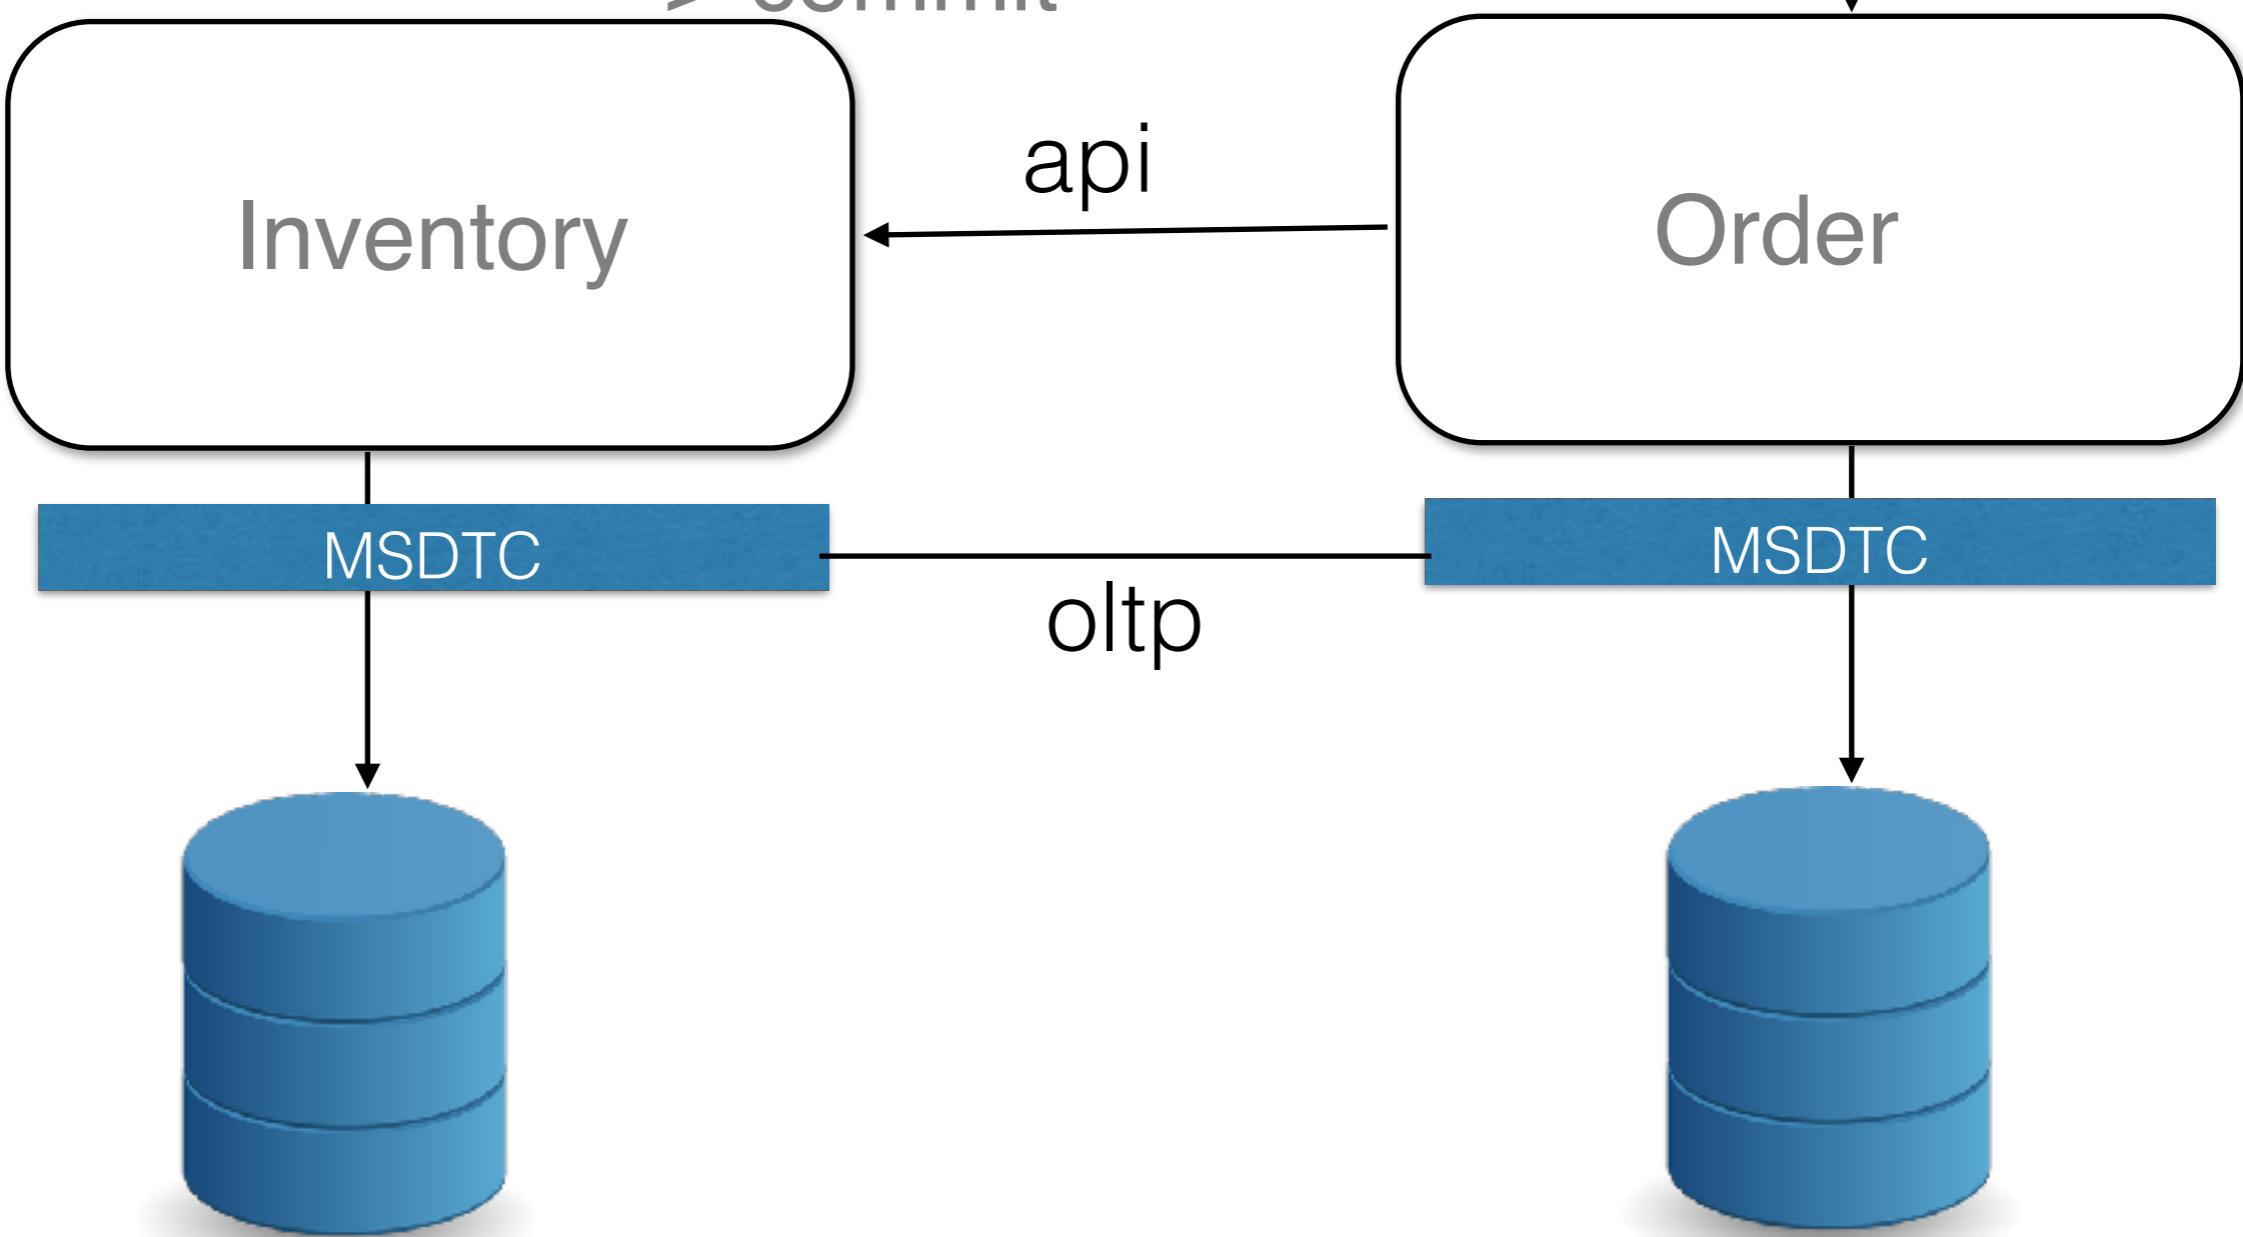
Update 1

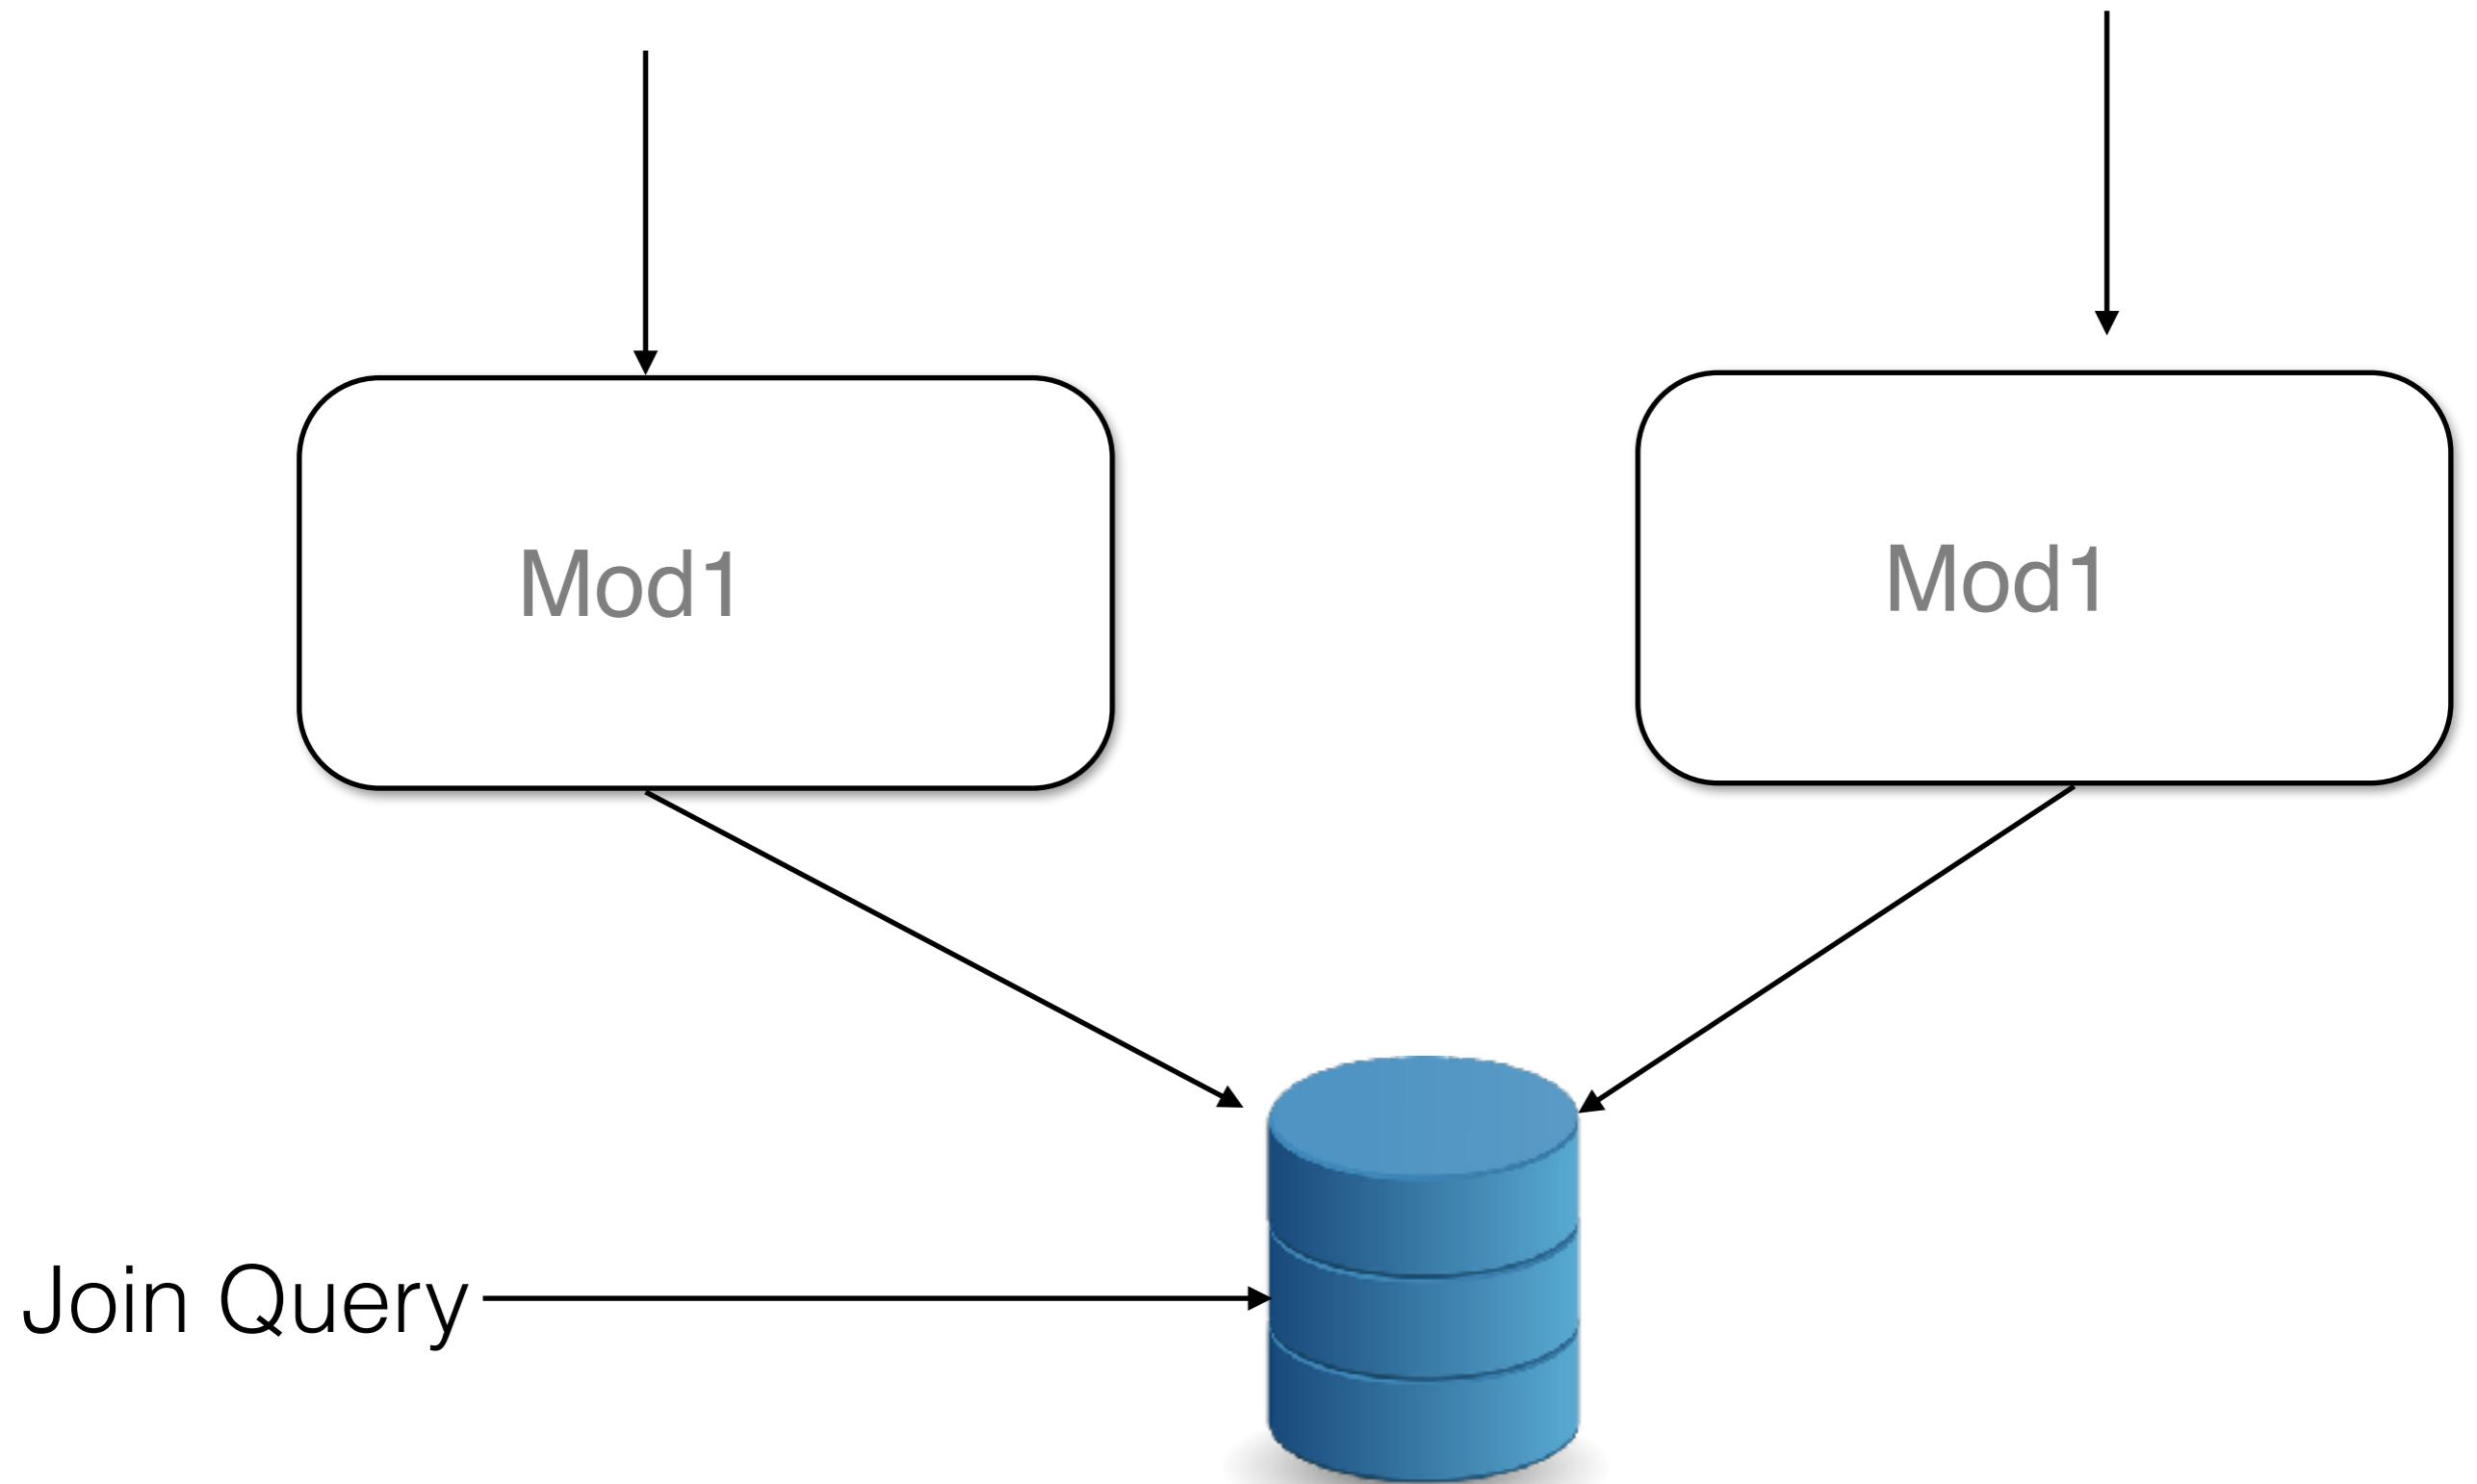
Update 2

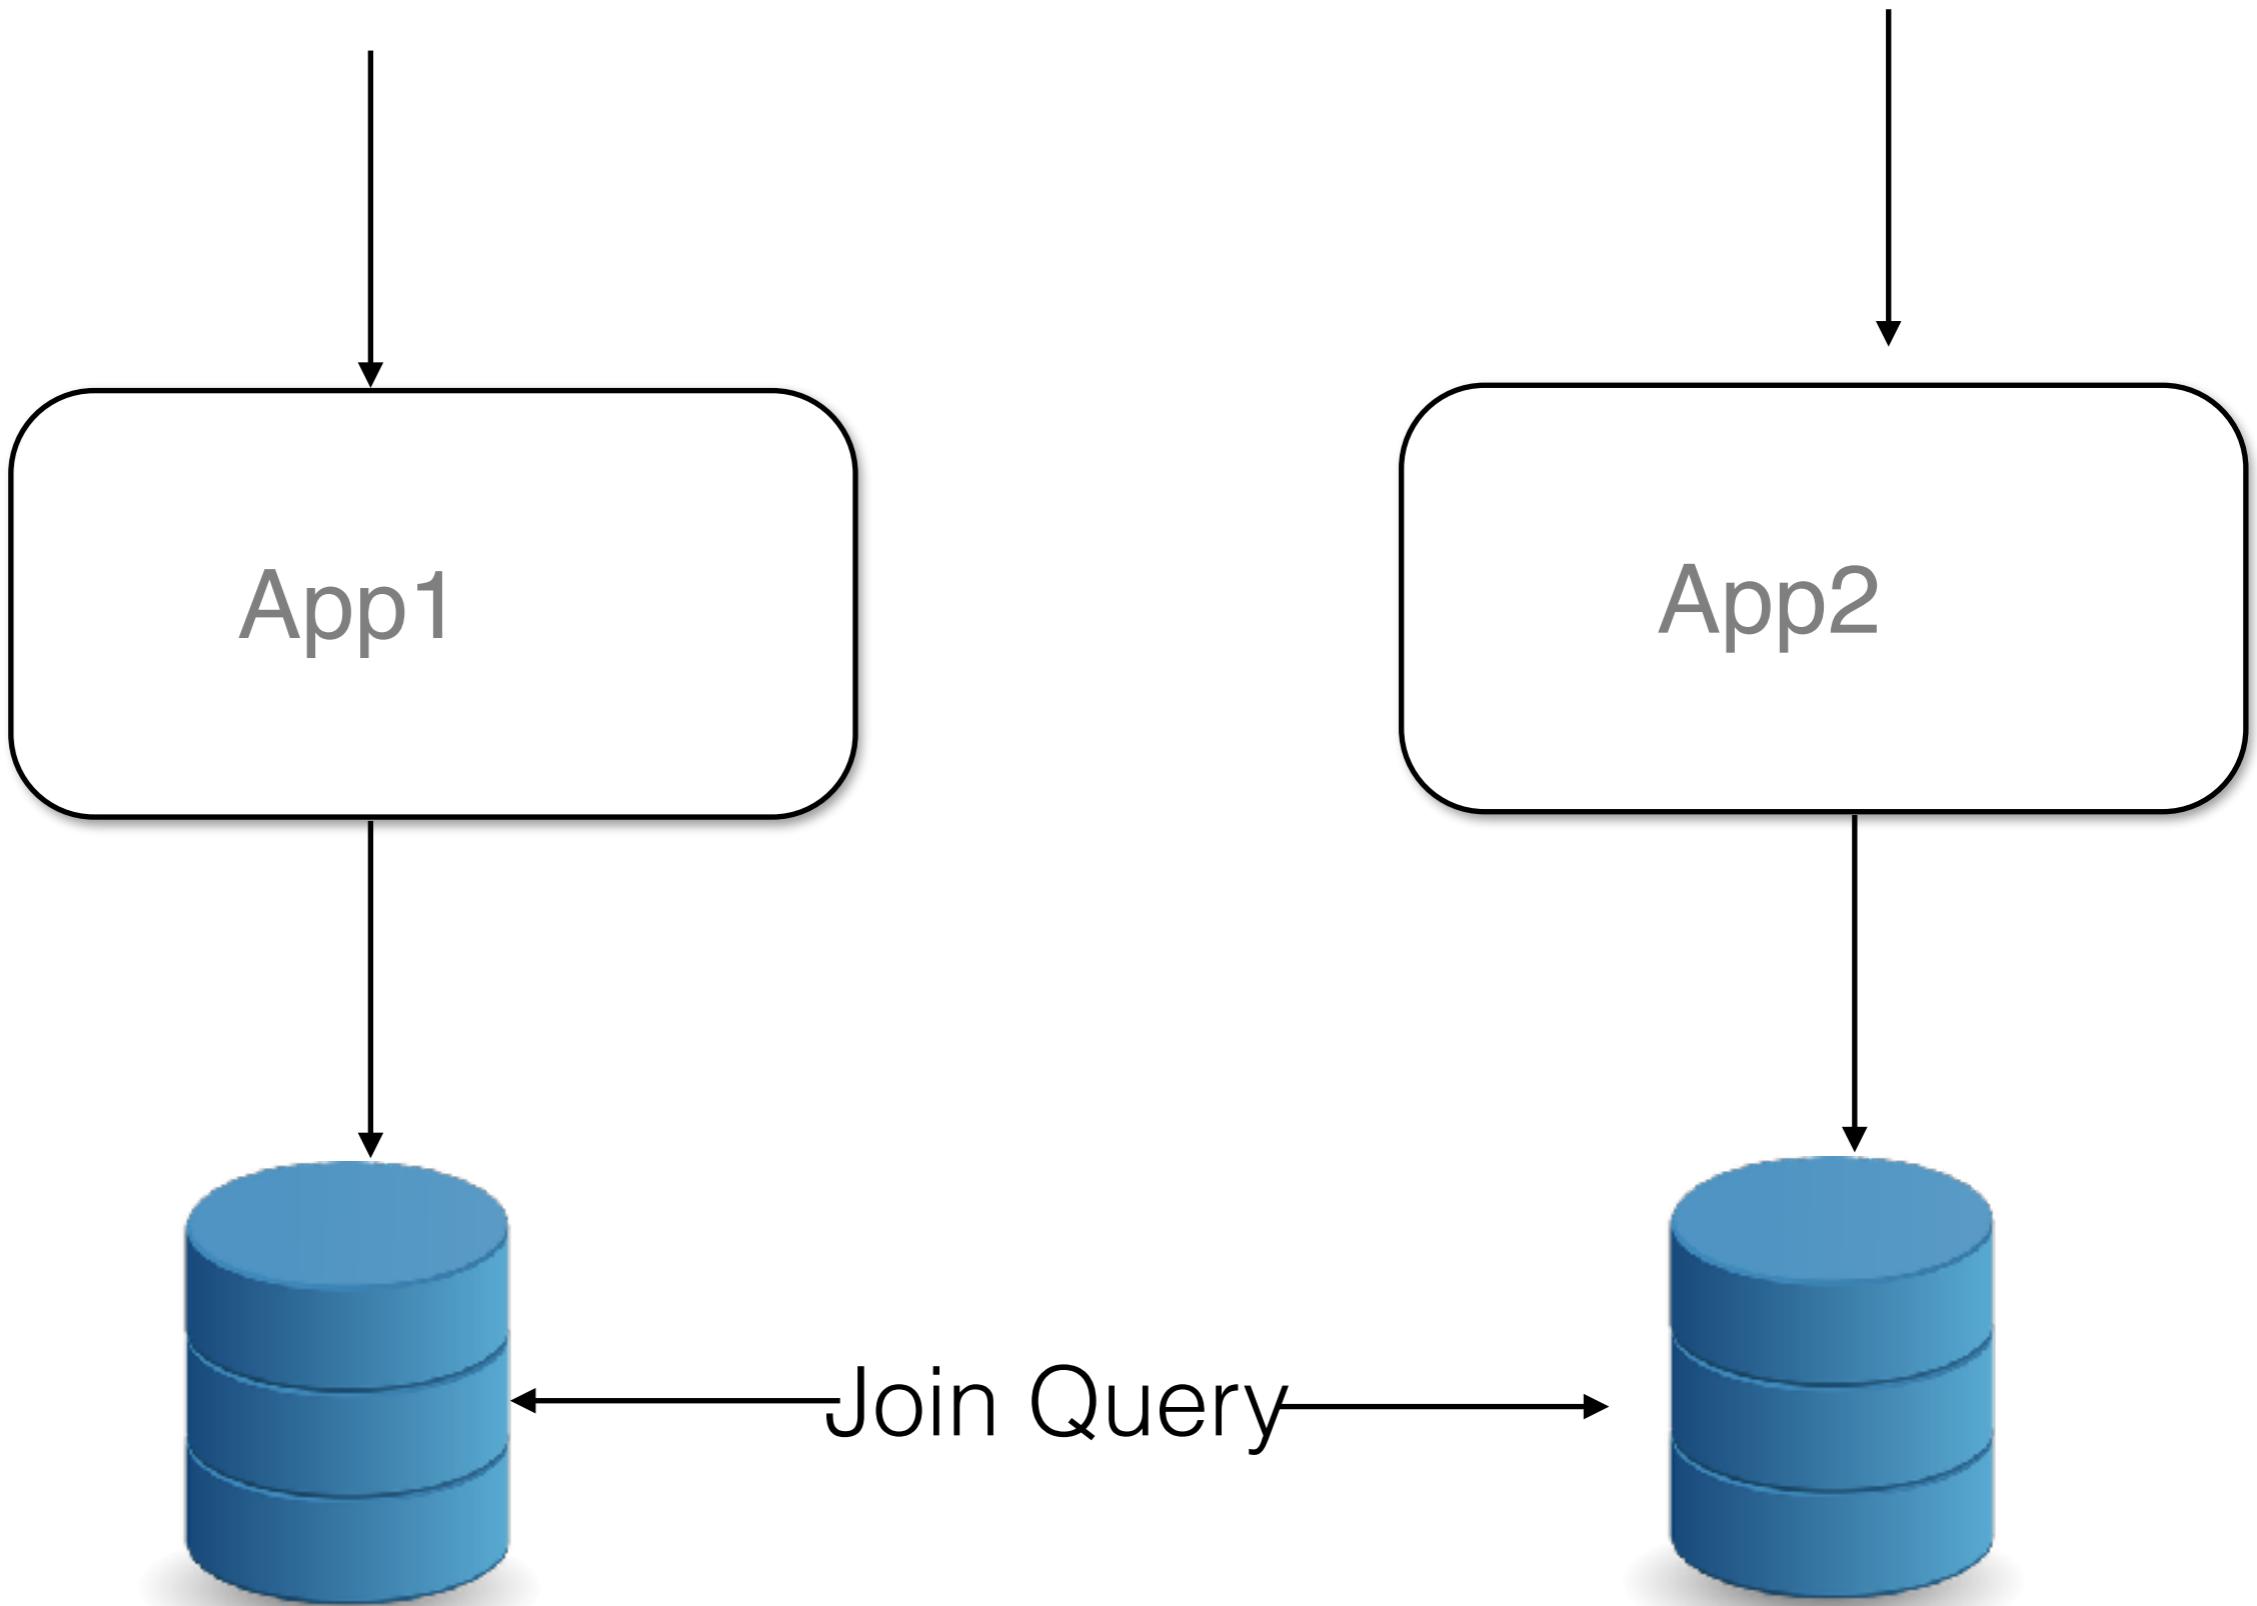
mstdc.commit_tran()

-> take votes

-> commit









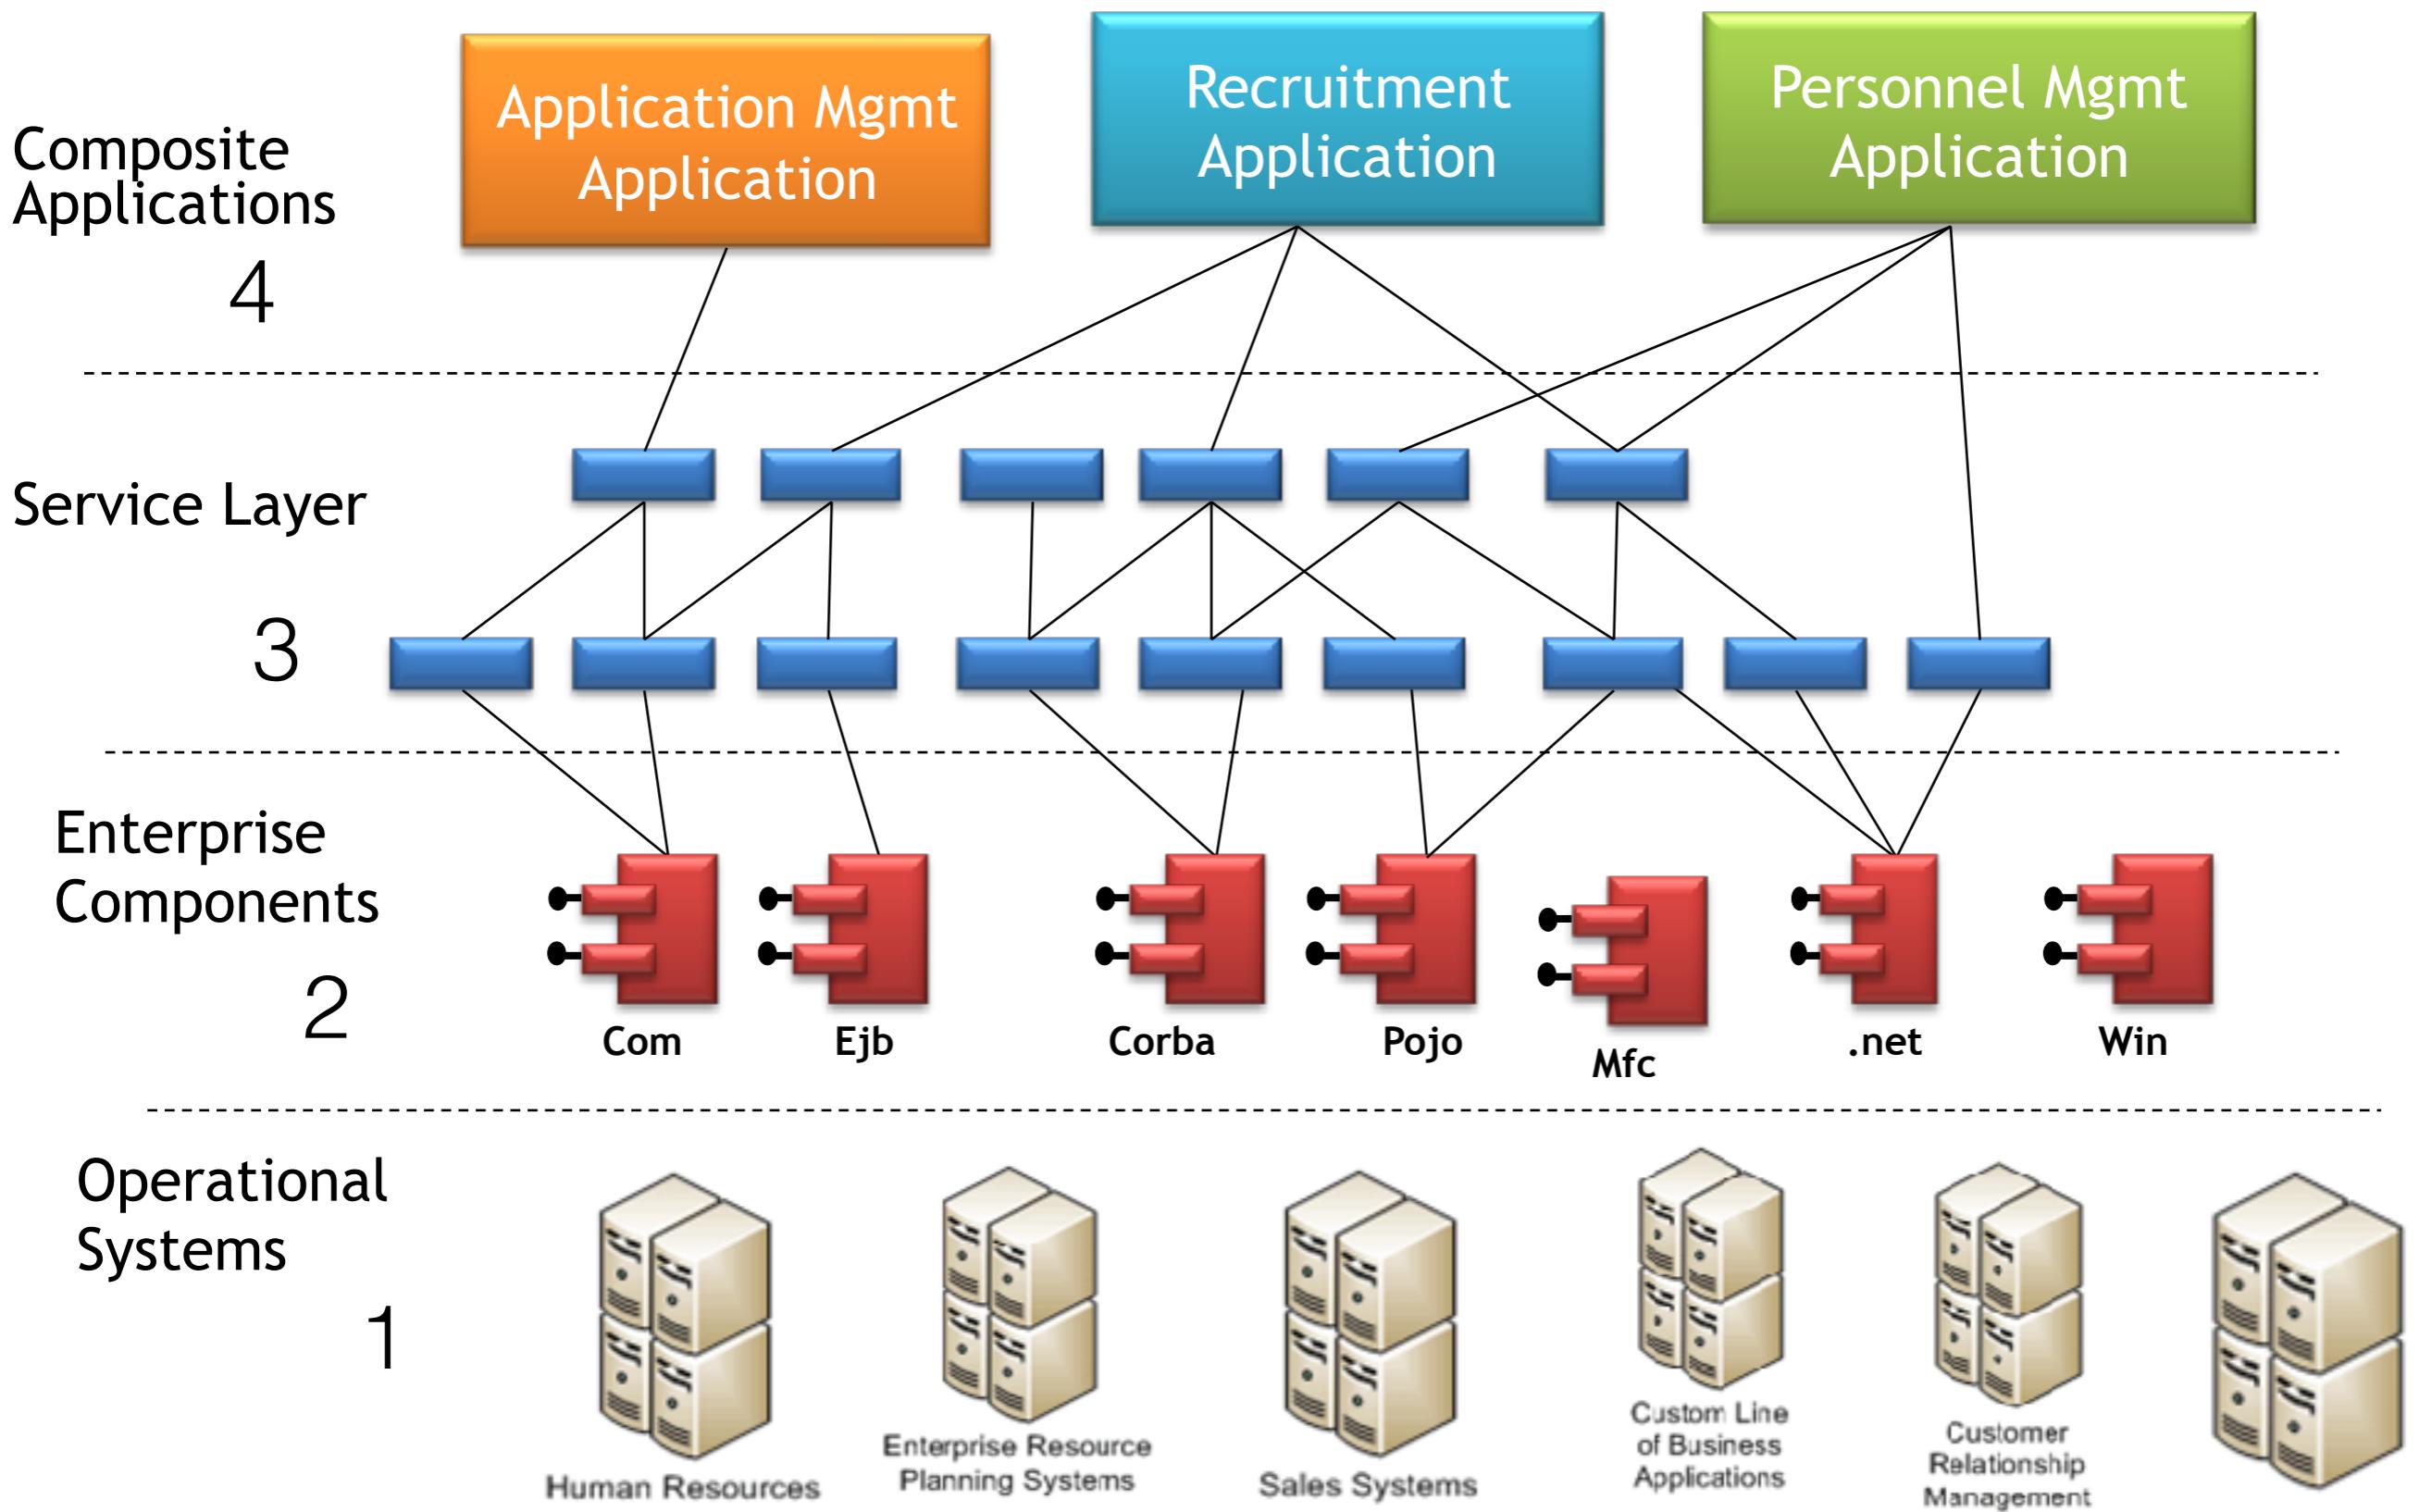
MICROSERVICES

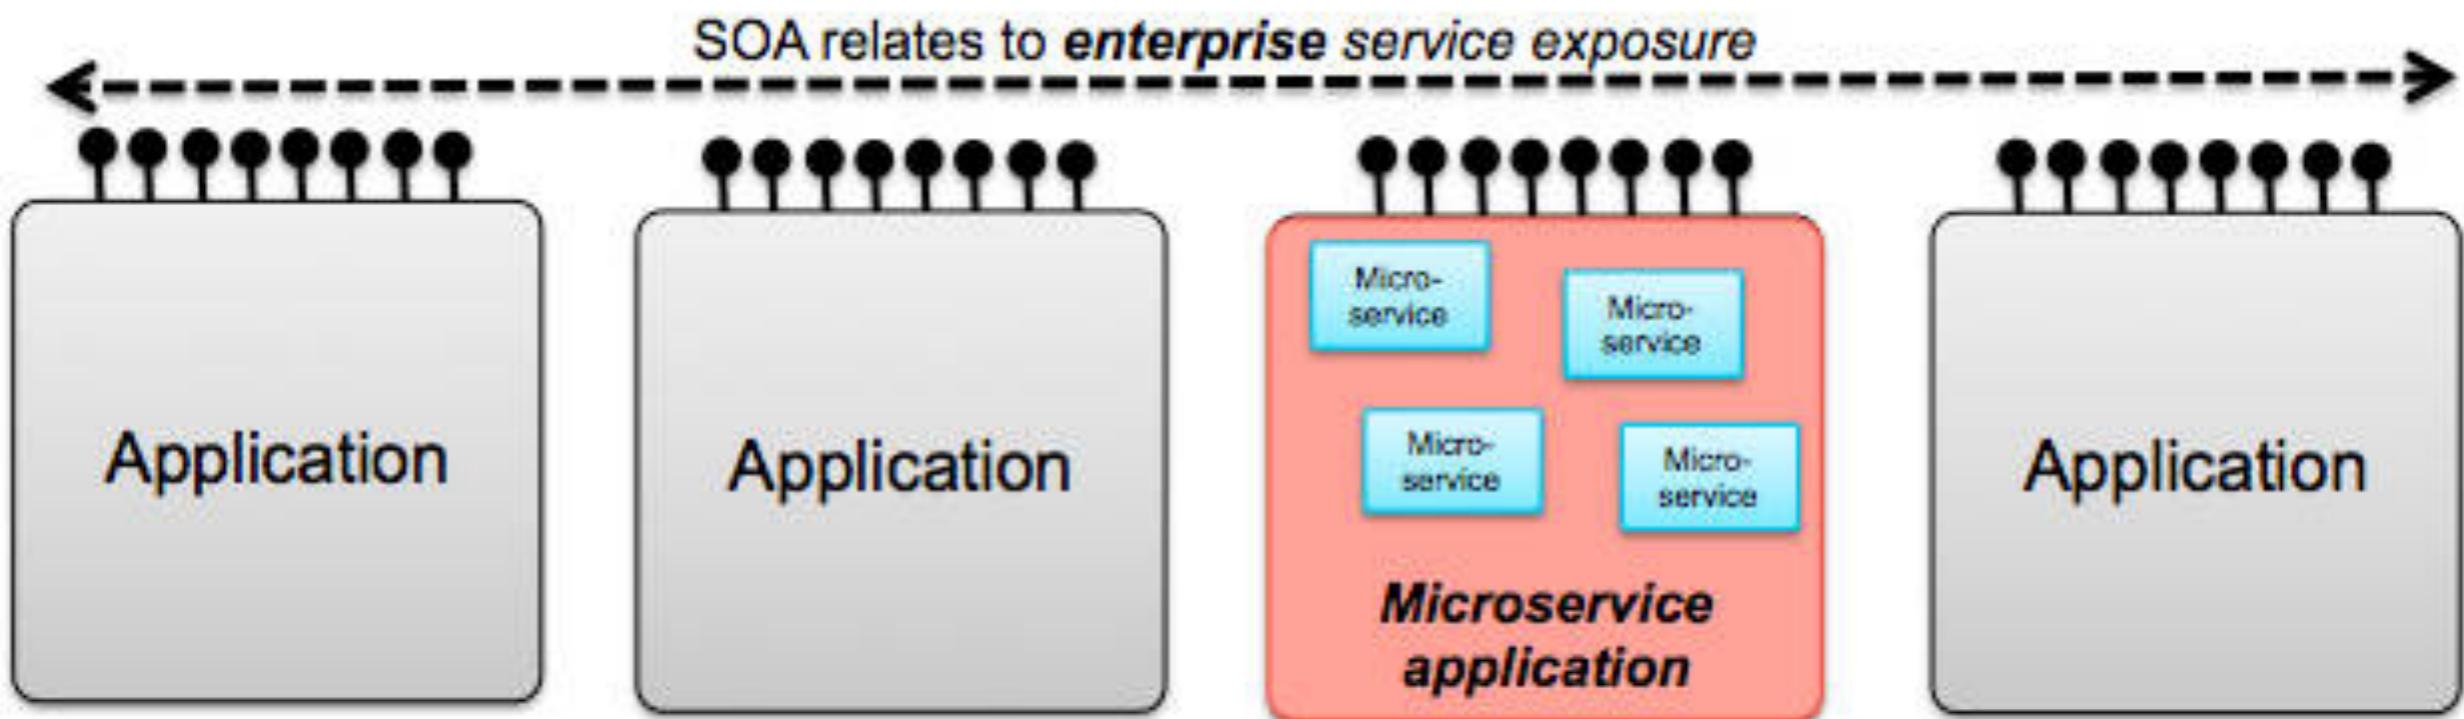
VS



SERVICE
ORIENTED ARCHITECTURE

Service Oriented Architecture





Microservices relate to
application architecture

Day 2

Approach

Performance

Transaction Management

Joins / Reports

Infra Cost

Deployment

Debugging (Finding bugs)

Log Mgmt

Config Mgmt

Authentication

Authorization

Monitoring / Alerting

Isolated Env.

Virtual Machine

App 1

App 2

VM1

VM2

My app 1

Middlewear

Unix Guest Os

My app 2

Middlewear

Win Guest Os

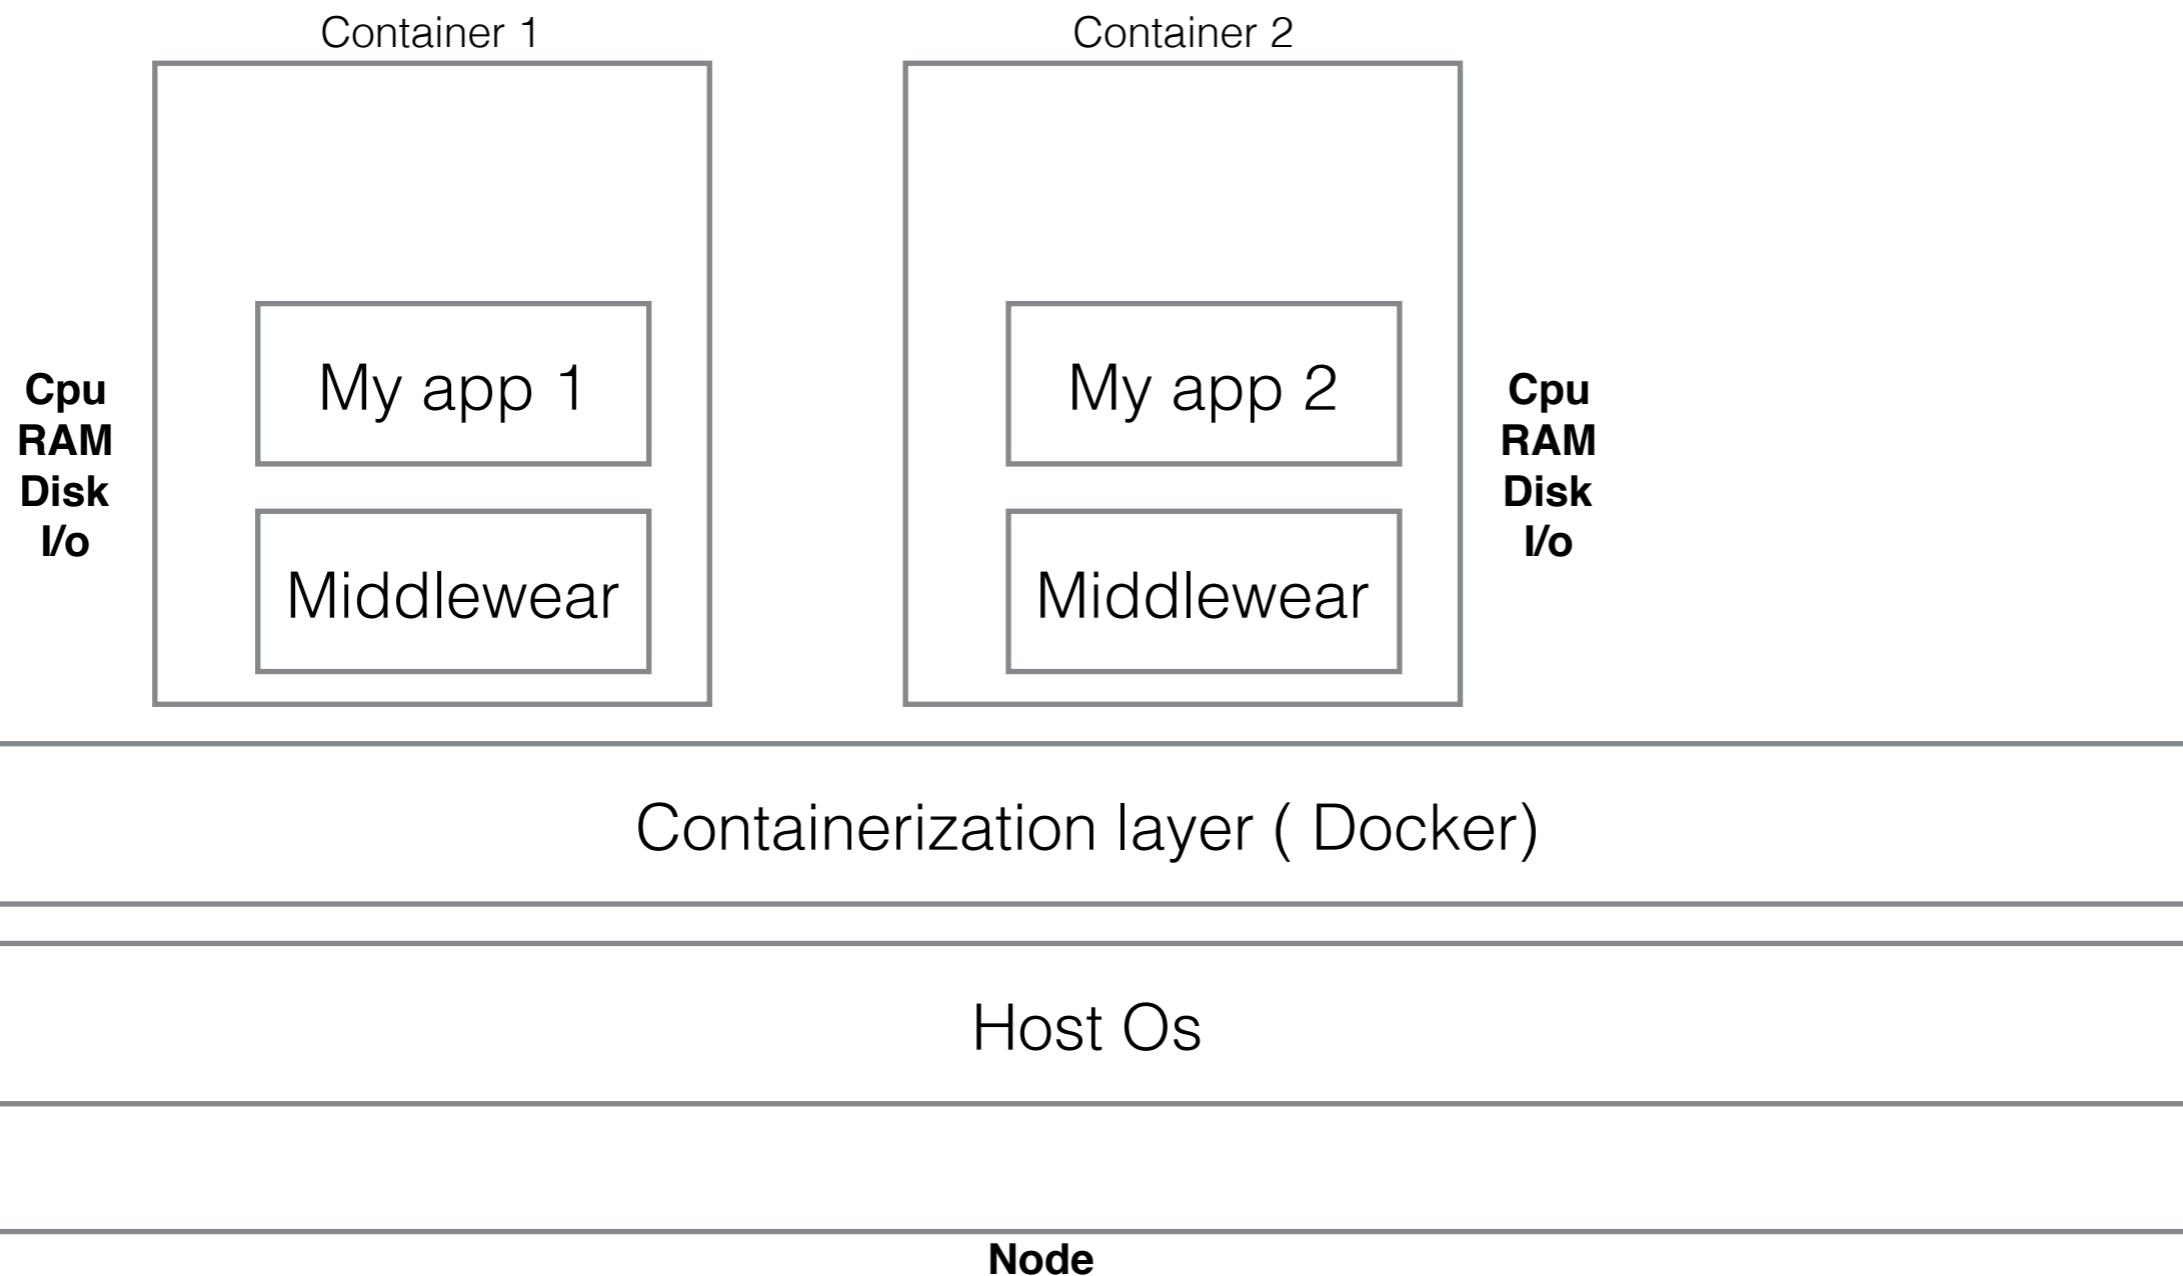
Cpu
RAM
Disk
I/o

Cpu
RAM
Disk
I/o

Virtualisation layer

Host Os

Node



VM1

VM2

Container 1

Container 2

My app 1

Middlewear

Cpu
RAM
Disk
I/o

My app 2

Middlewear

Cpu
RAM
Disk
I/o

Containerization layer (Docker)

Unix Guest Os

Cpu
RAM
Disk
I/o

Win Guest Os

Cpu
RAM
Disk
I/o

Virtualisation layer

Host Os (azure)

Machine

Approach

Performance

Transaction Management

Joins / Reports

Infra Cost

Deployment

Debugging (Finding bugs)

Log Mgmt

Config Mgmt

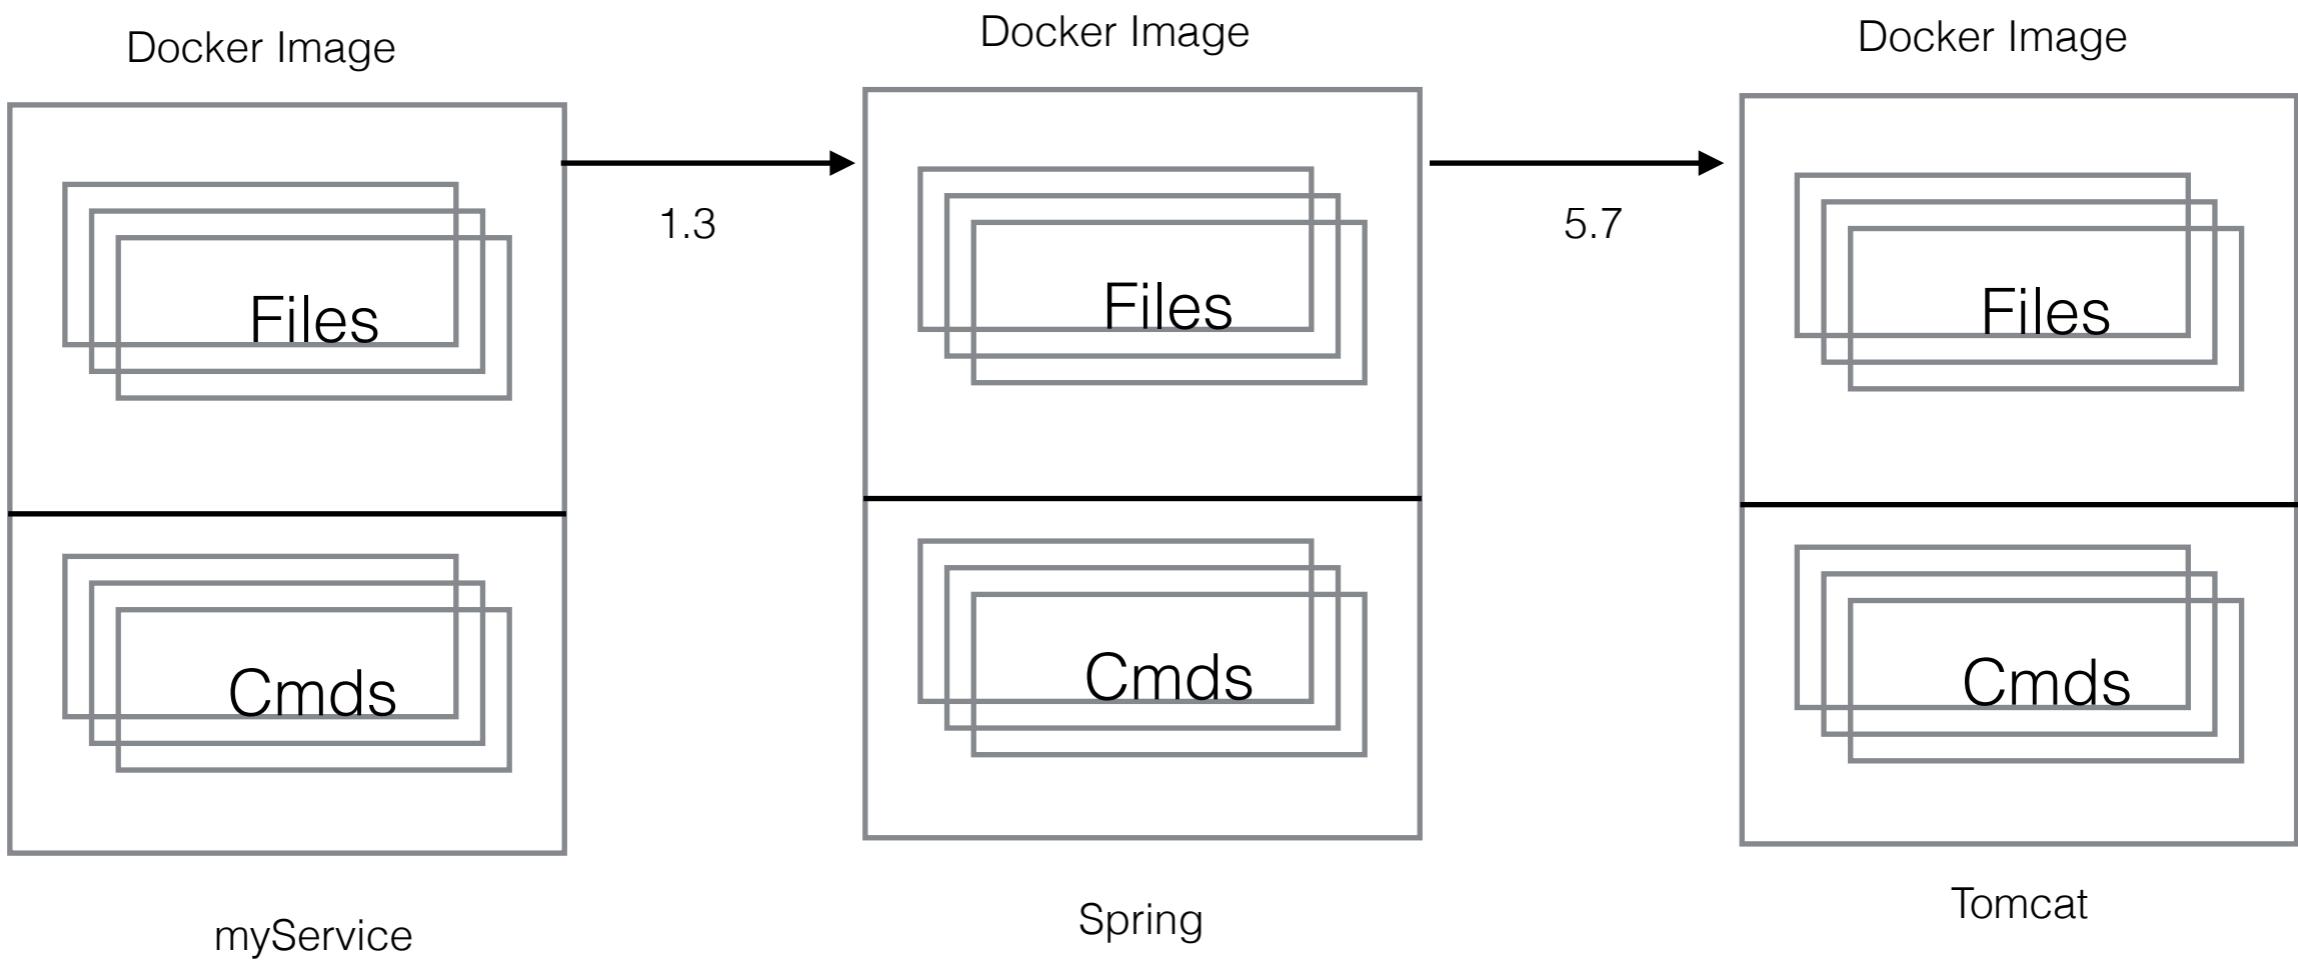
Authentication

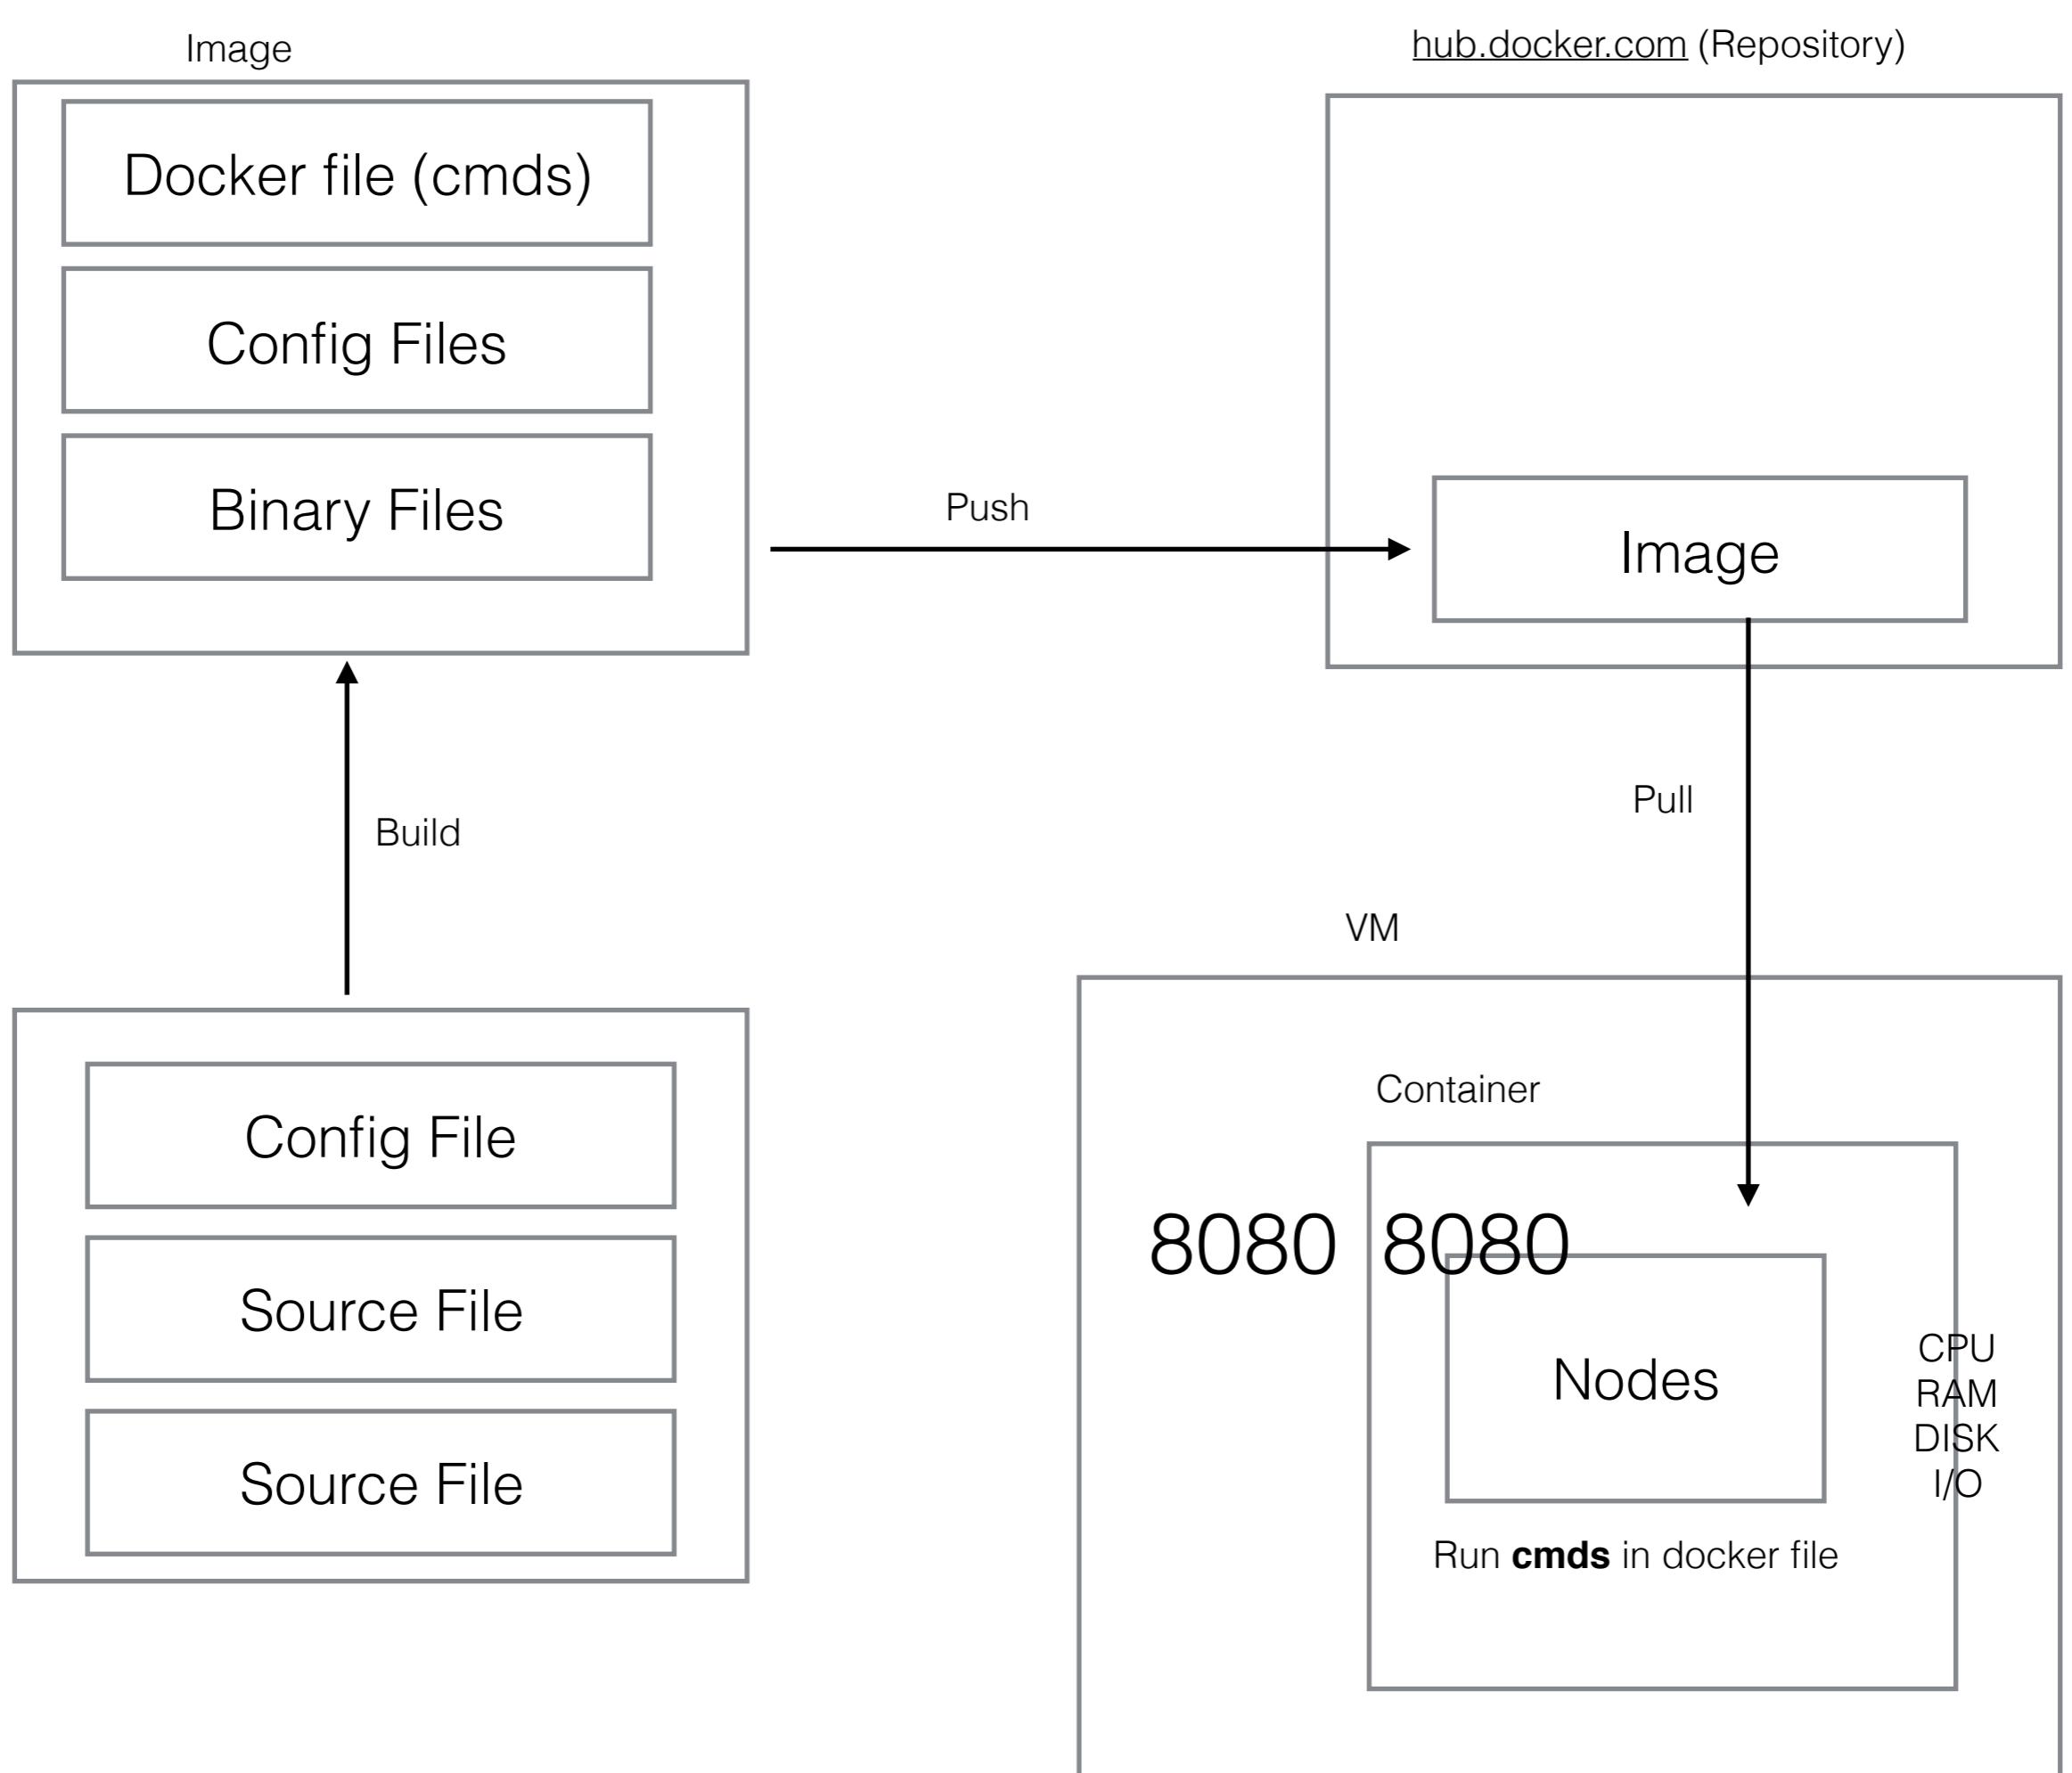
Authorization

Monitoring / Alerting

Container

Reproducible Env.





Approach

Performance

Transaction Management

Joins / Reports

Infra Cost

Deployment

Debugging (Finding bugs)

Log Mgmt

Config Mgmt

Authentication

Authorization

Monitoring / Alerting

Container

Docker Image, Hub

Managing Container Life cycle

VM1

VM2

Container 1

Container 2

Cpu
RAM
Disk
I/o

My app 1

Middlewear

Kubernetes (slave)

Containerization layer (Docker)

Unix Guest Os

Container 1

Container 2

Cpu
RAM
Disk
I/o

My app 1

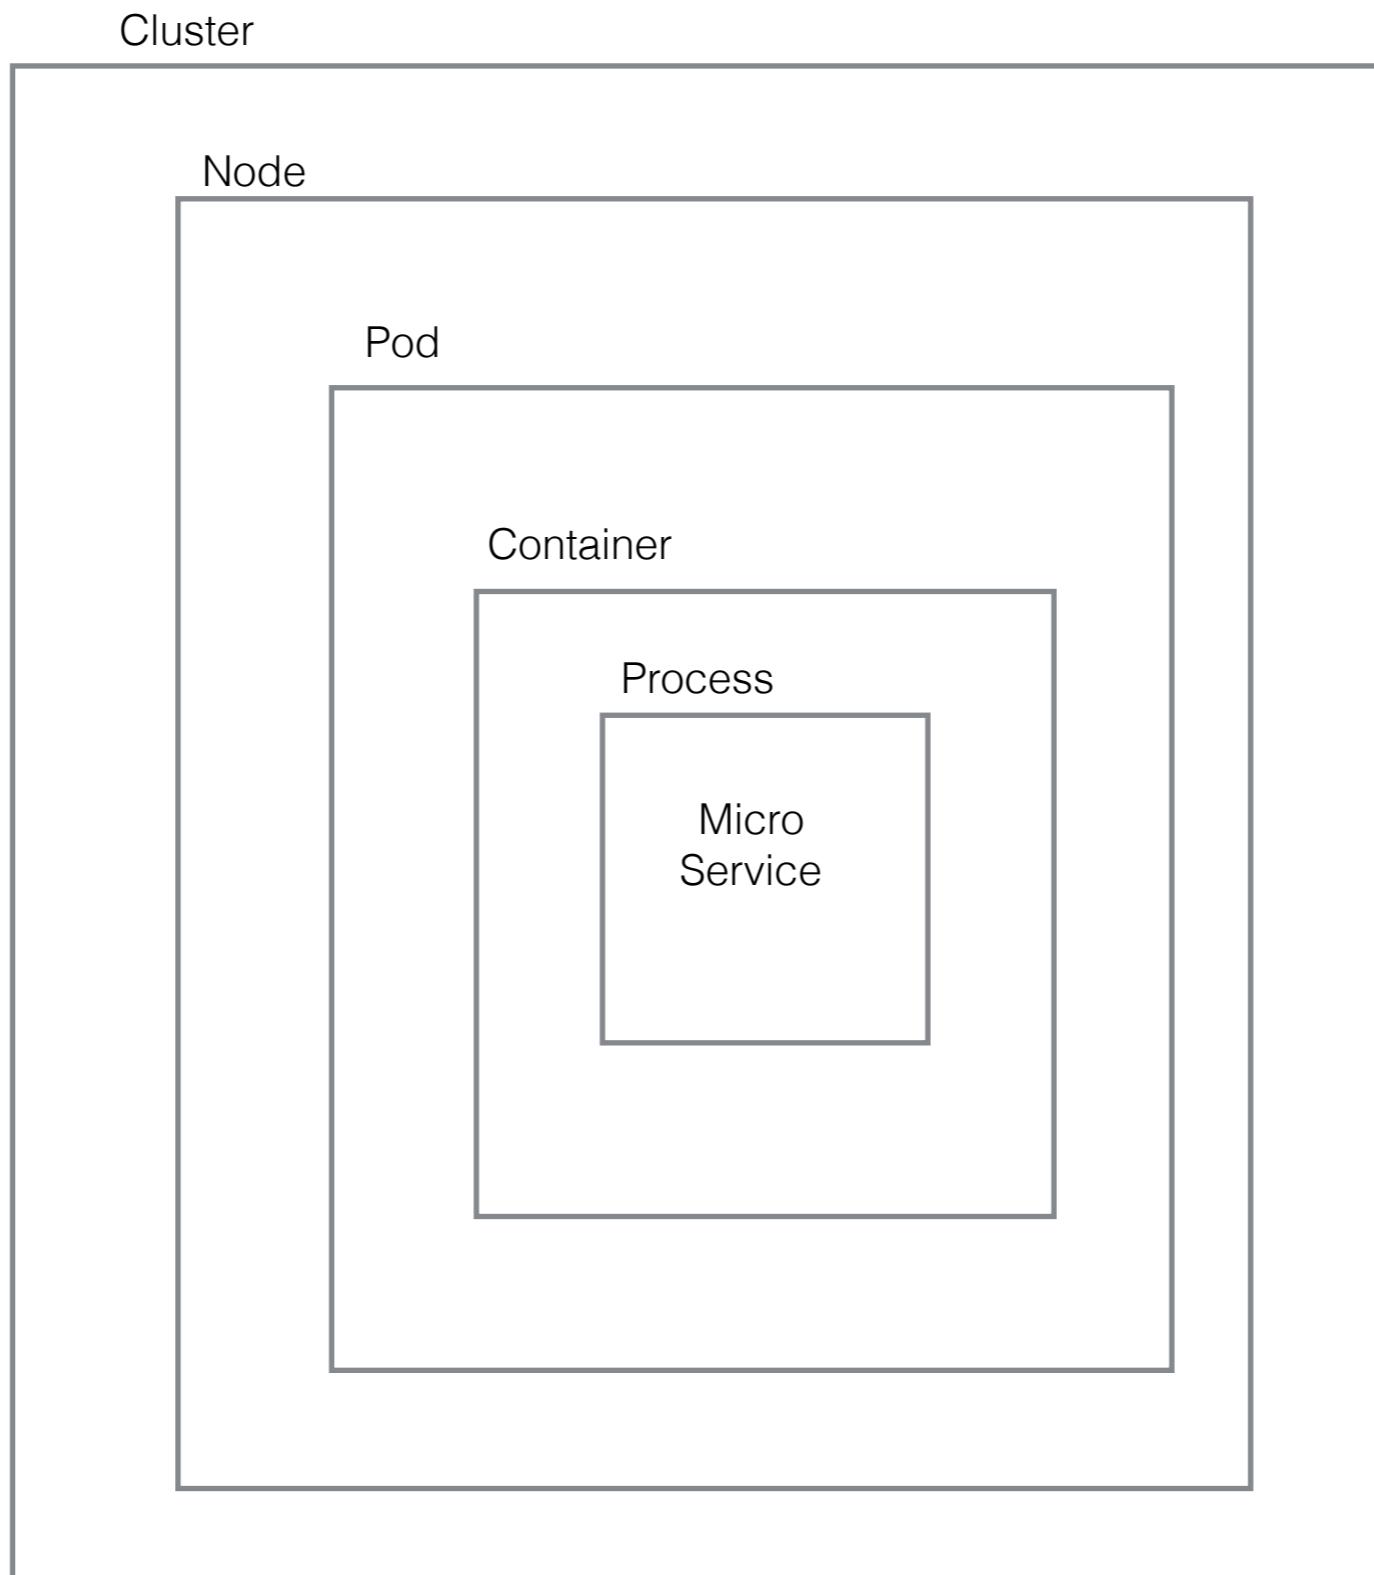
Middlewear

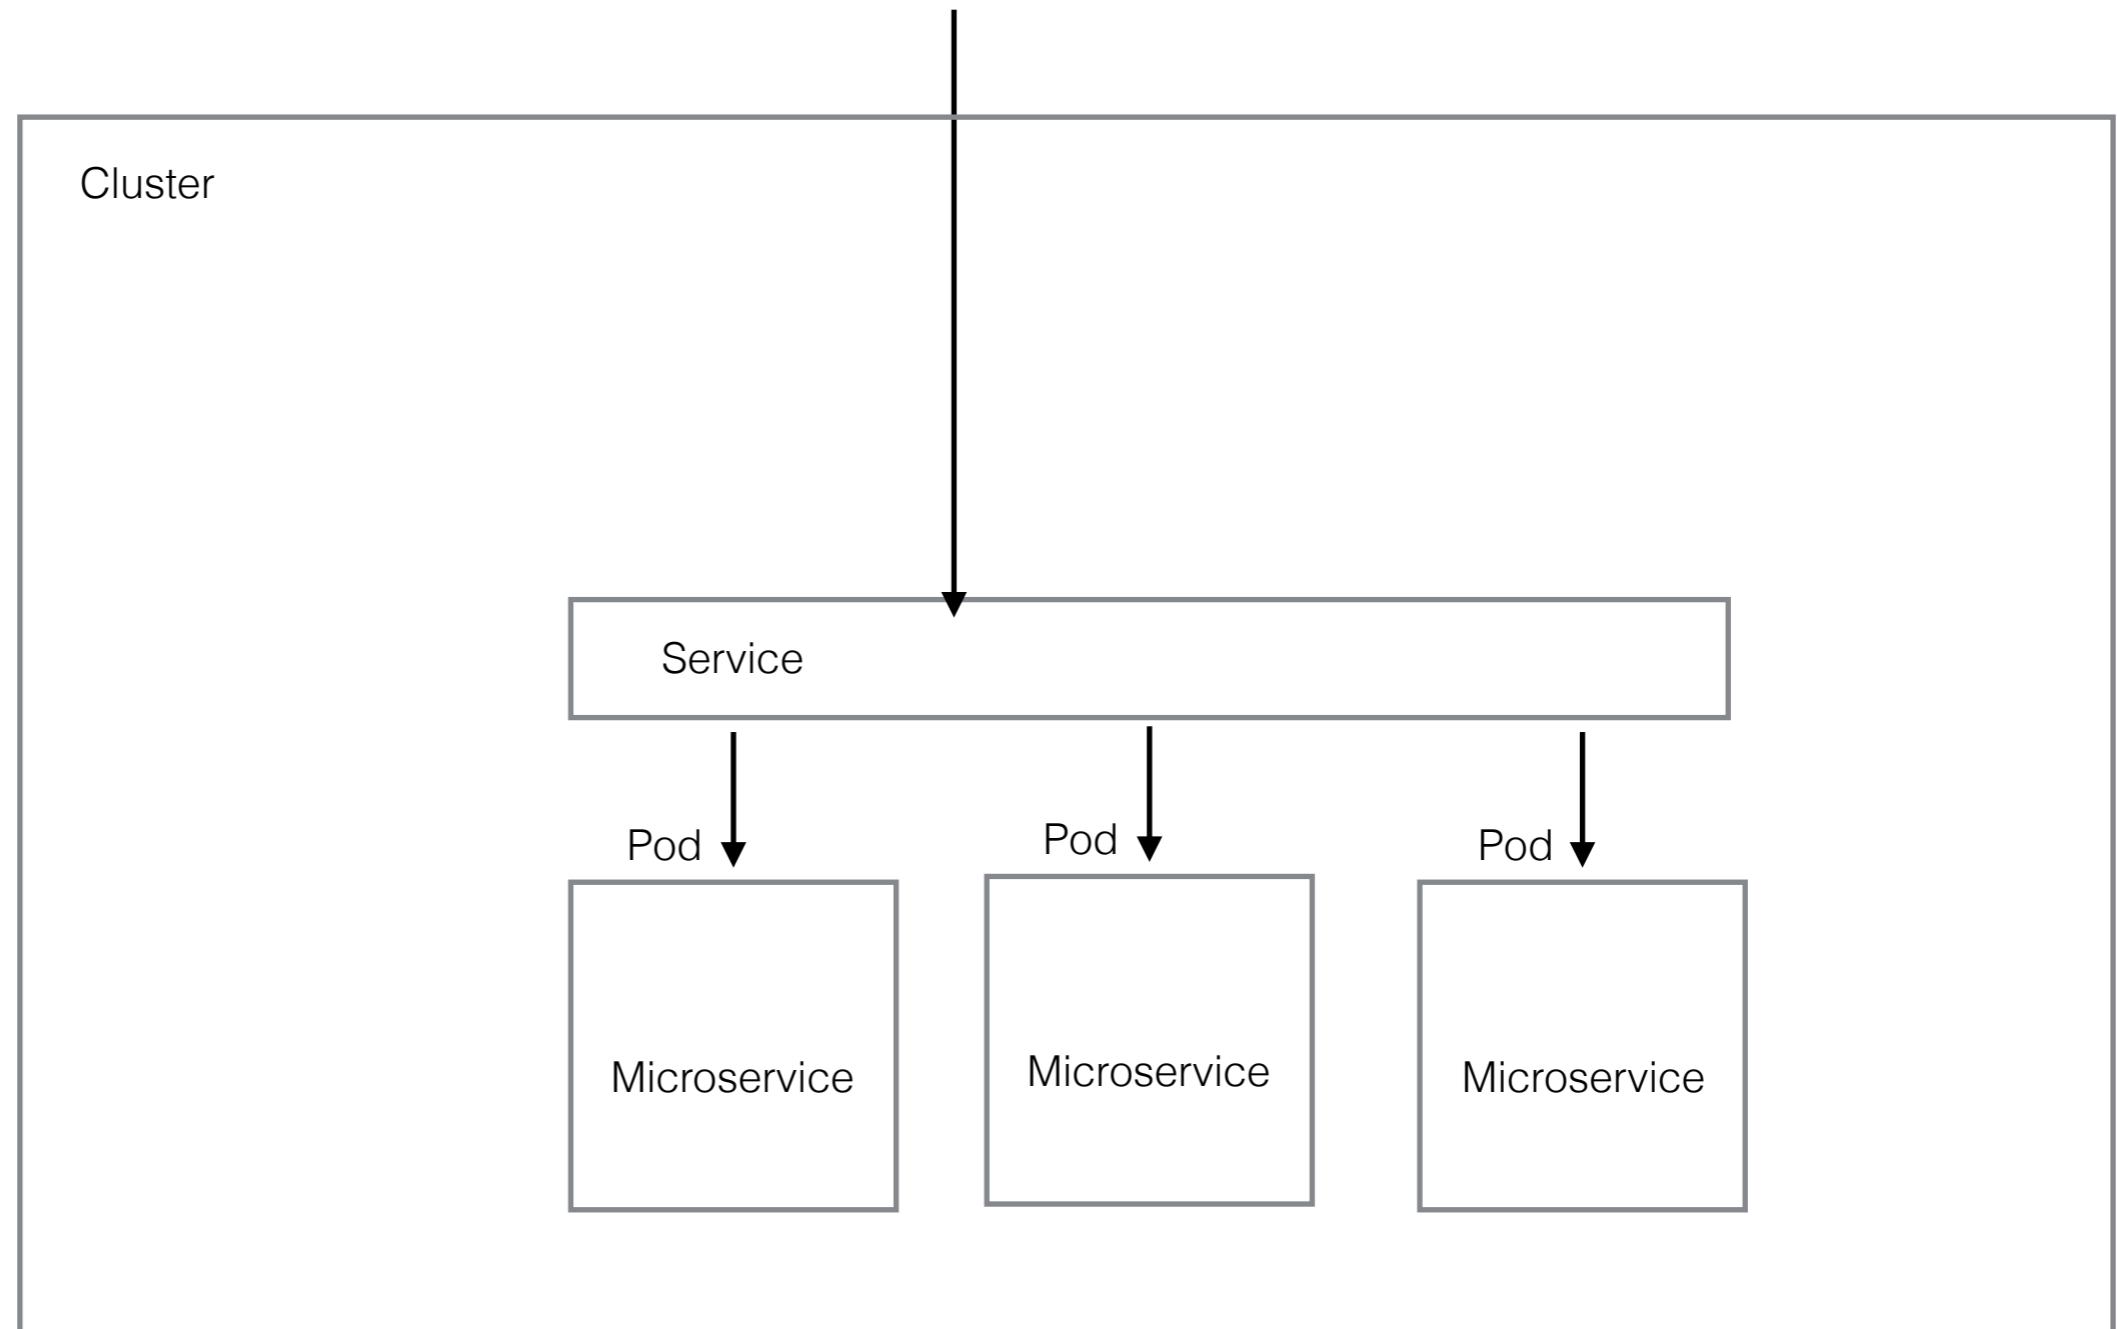
Kubernetes (Slave)

Containerization layer (Docker)

Unix Guest Os

Kubernetes (Master)





Approach

Performance

Transaction Management

Joins / Reports

Infra Cost

Deployment

Debugging (Finding bugs)

Log Mgmt

Config Mgmt

Authentication

Authorization

Monitoring / Alerting

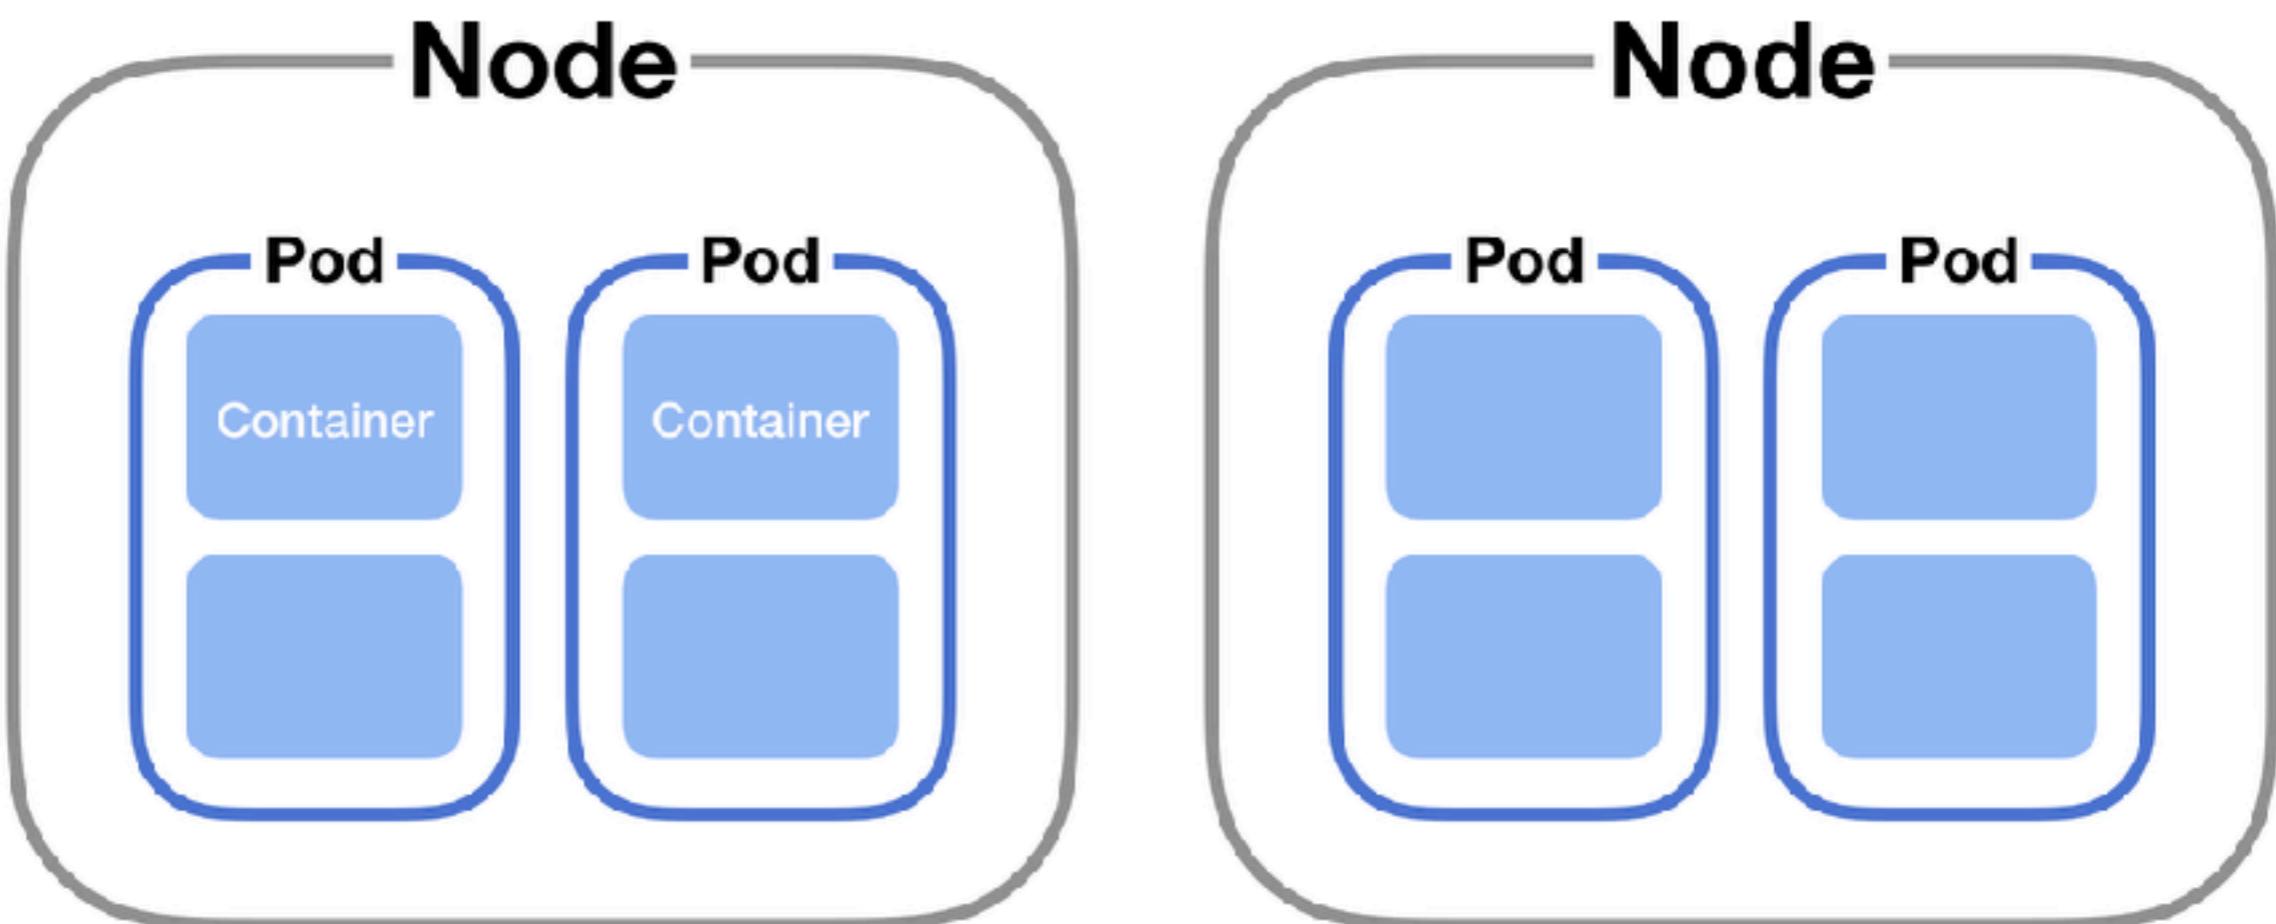
Container, Kubernetes

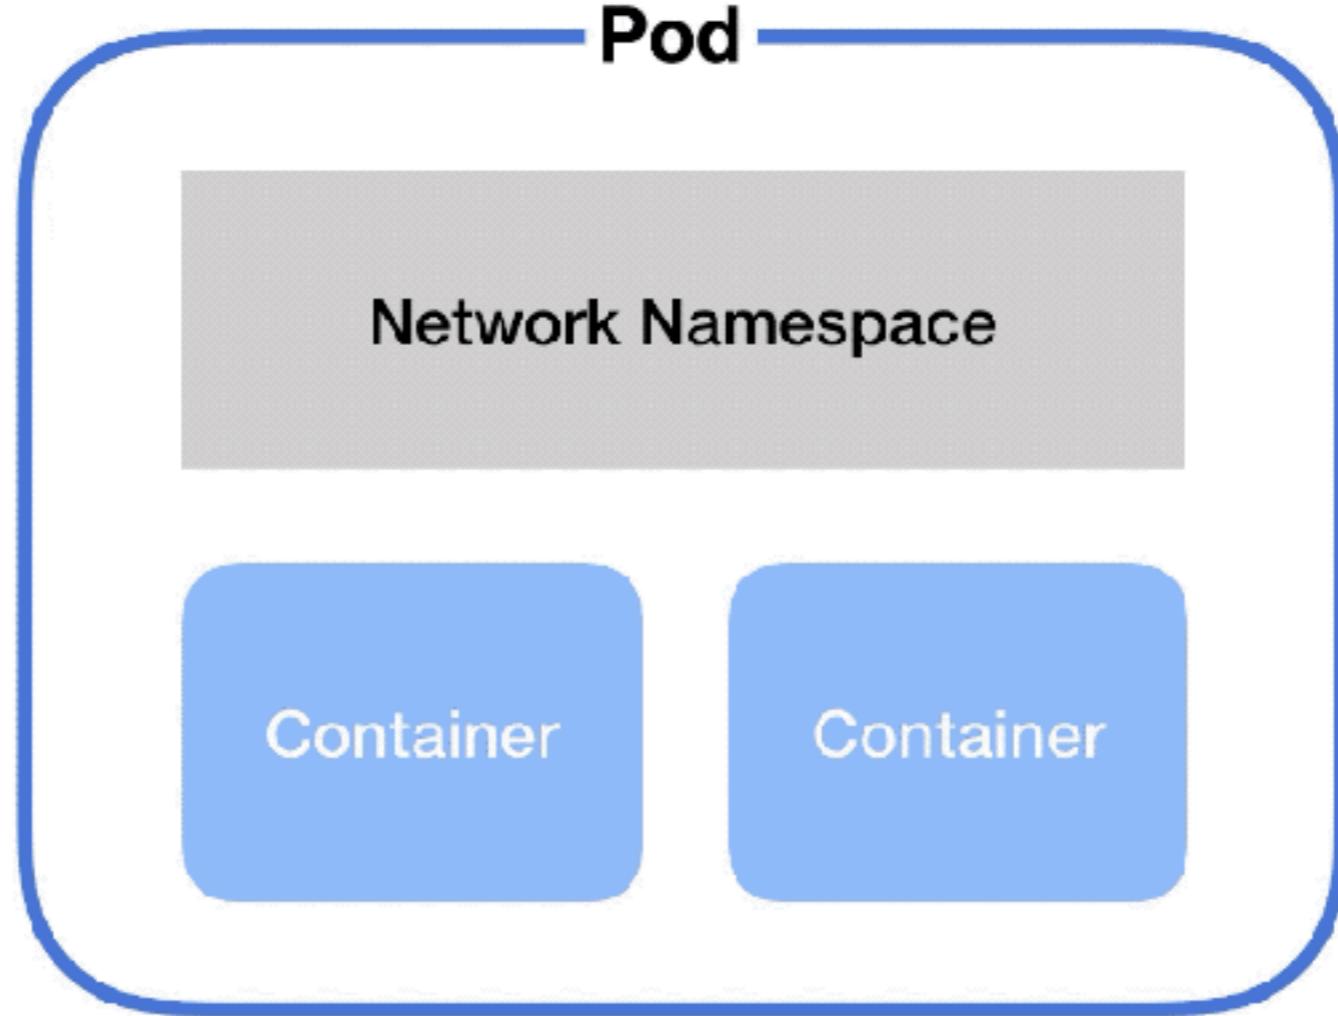
Docker Image, Hub

Kubernetes

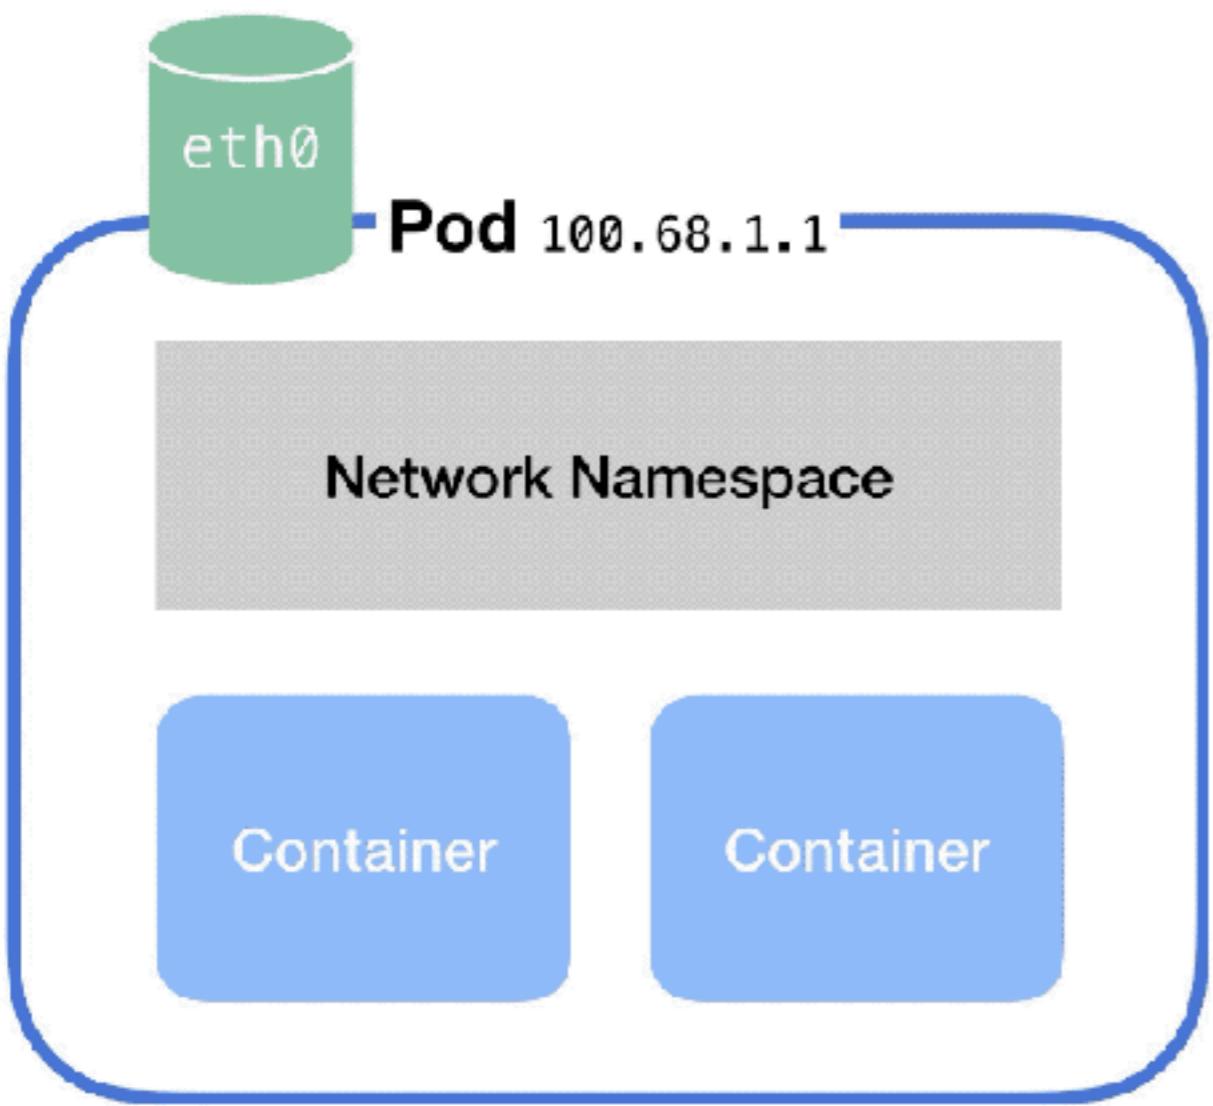
Service Discovery

Cluster

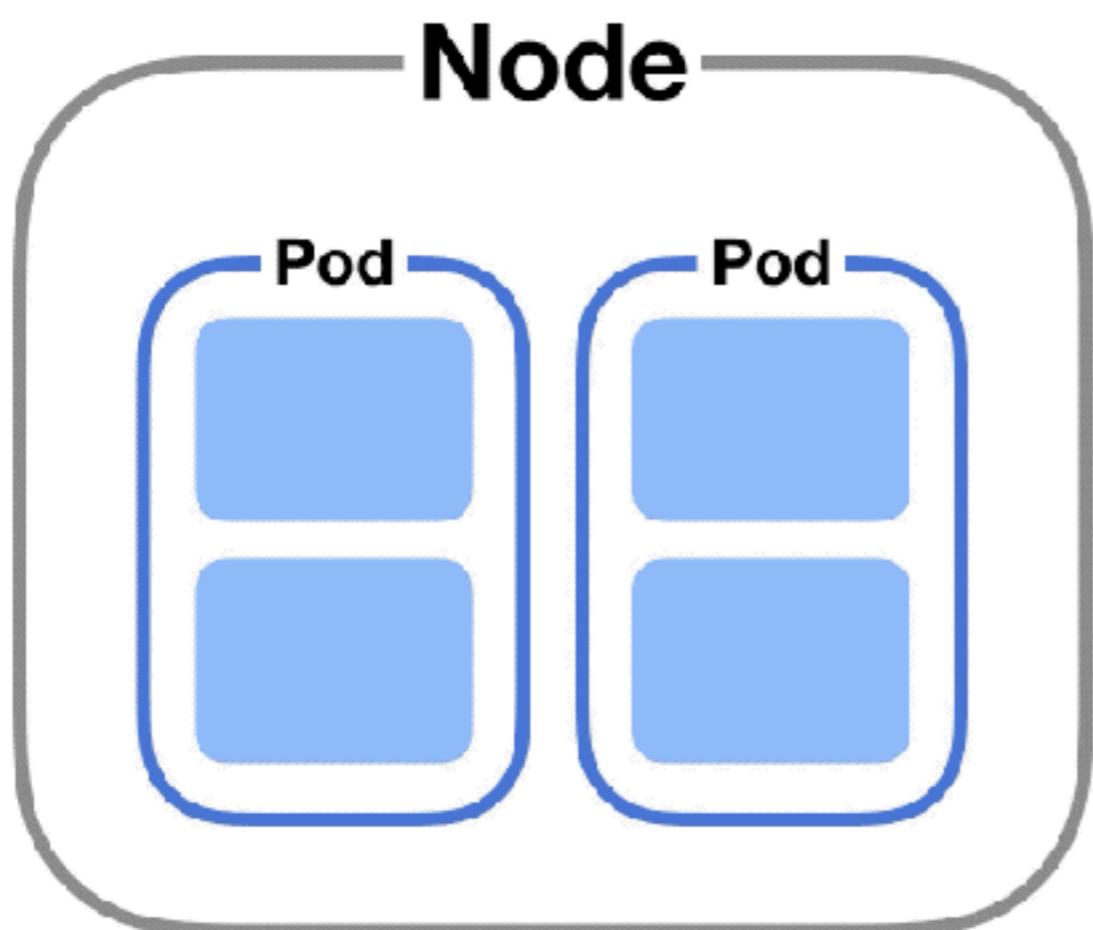


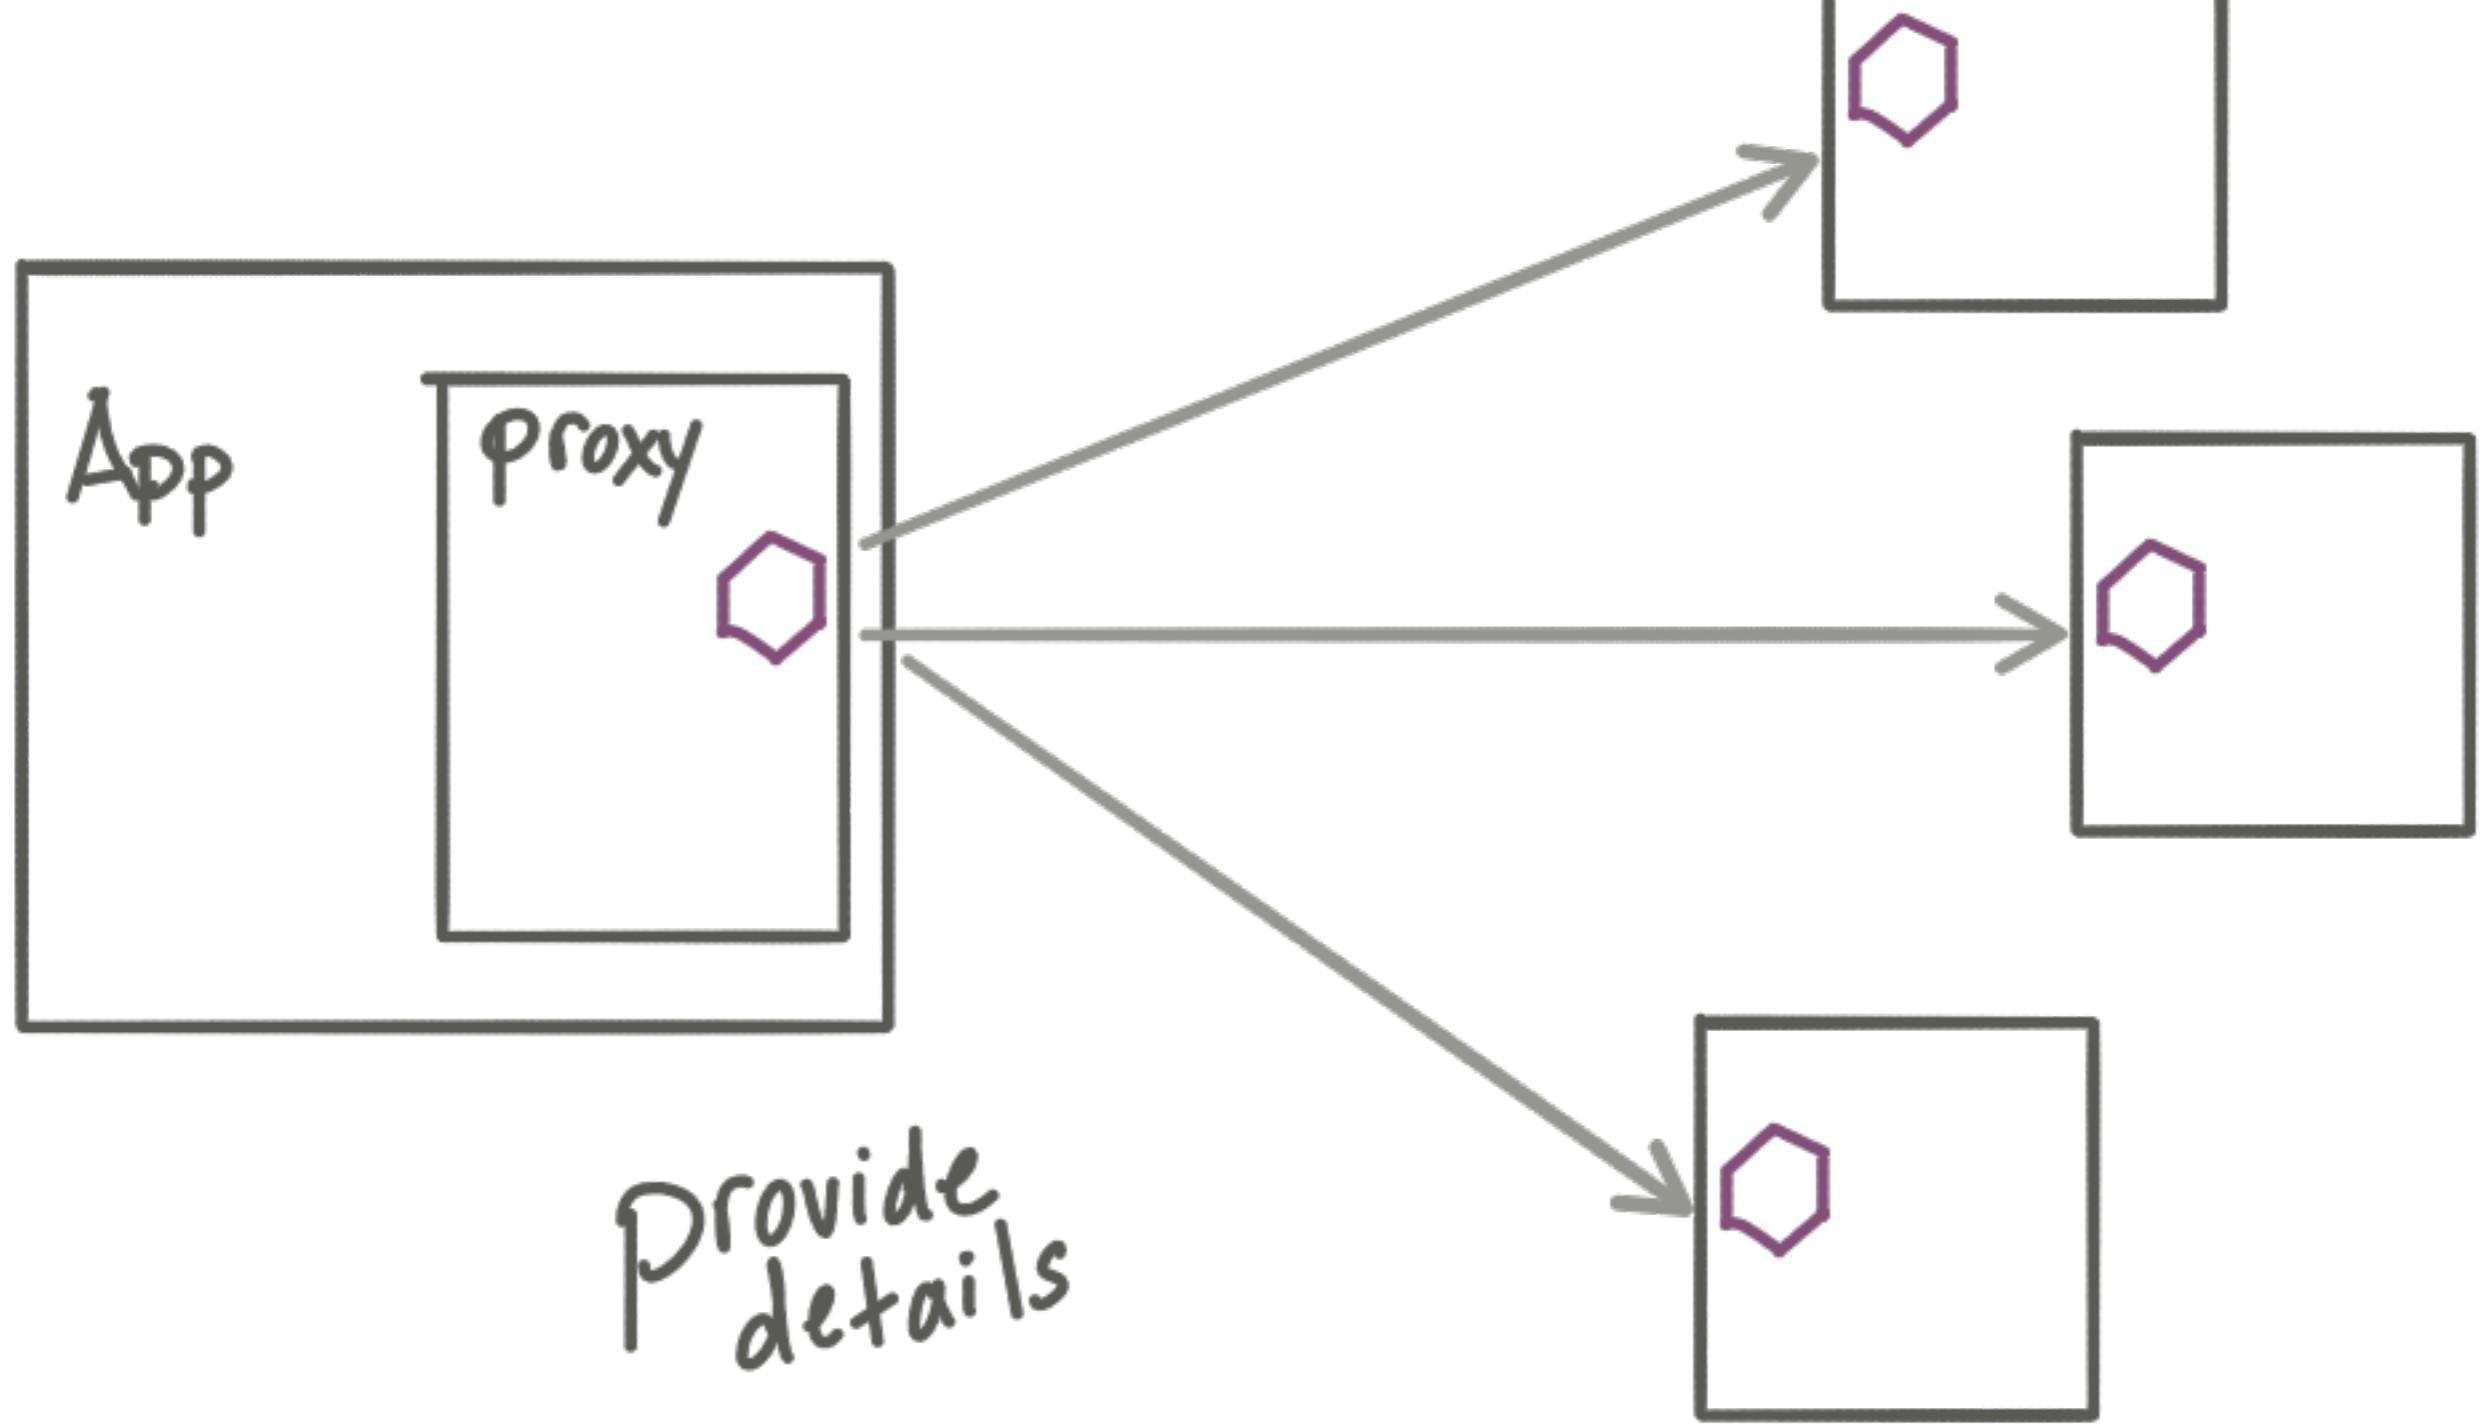


Node



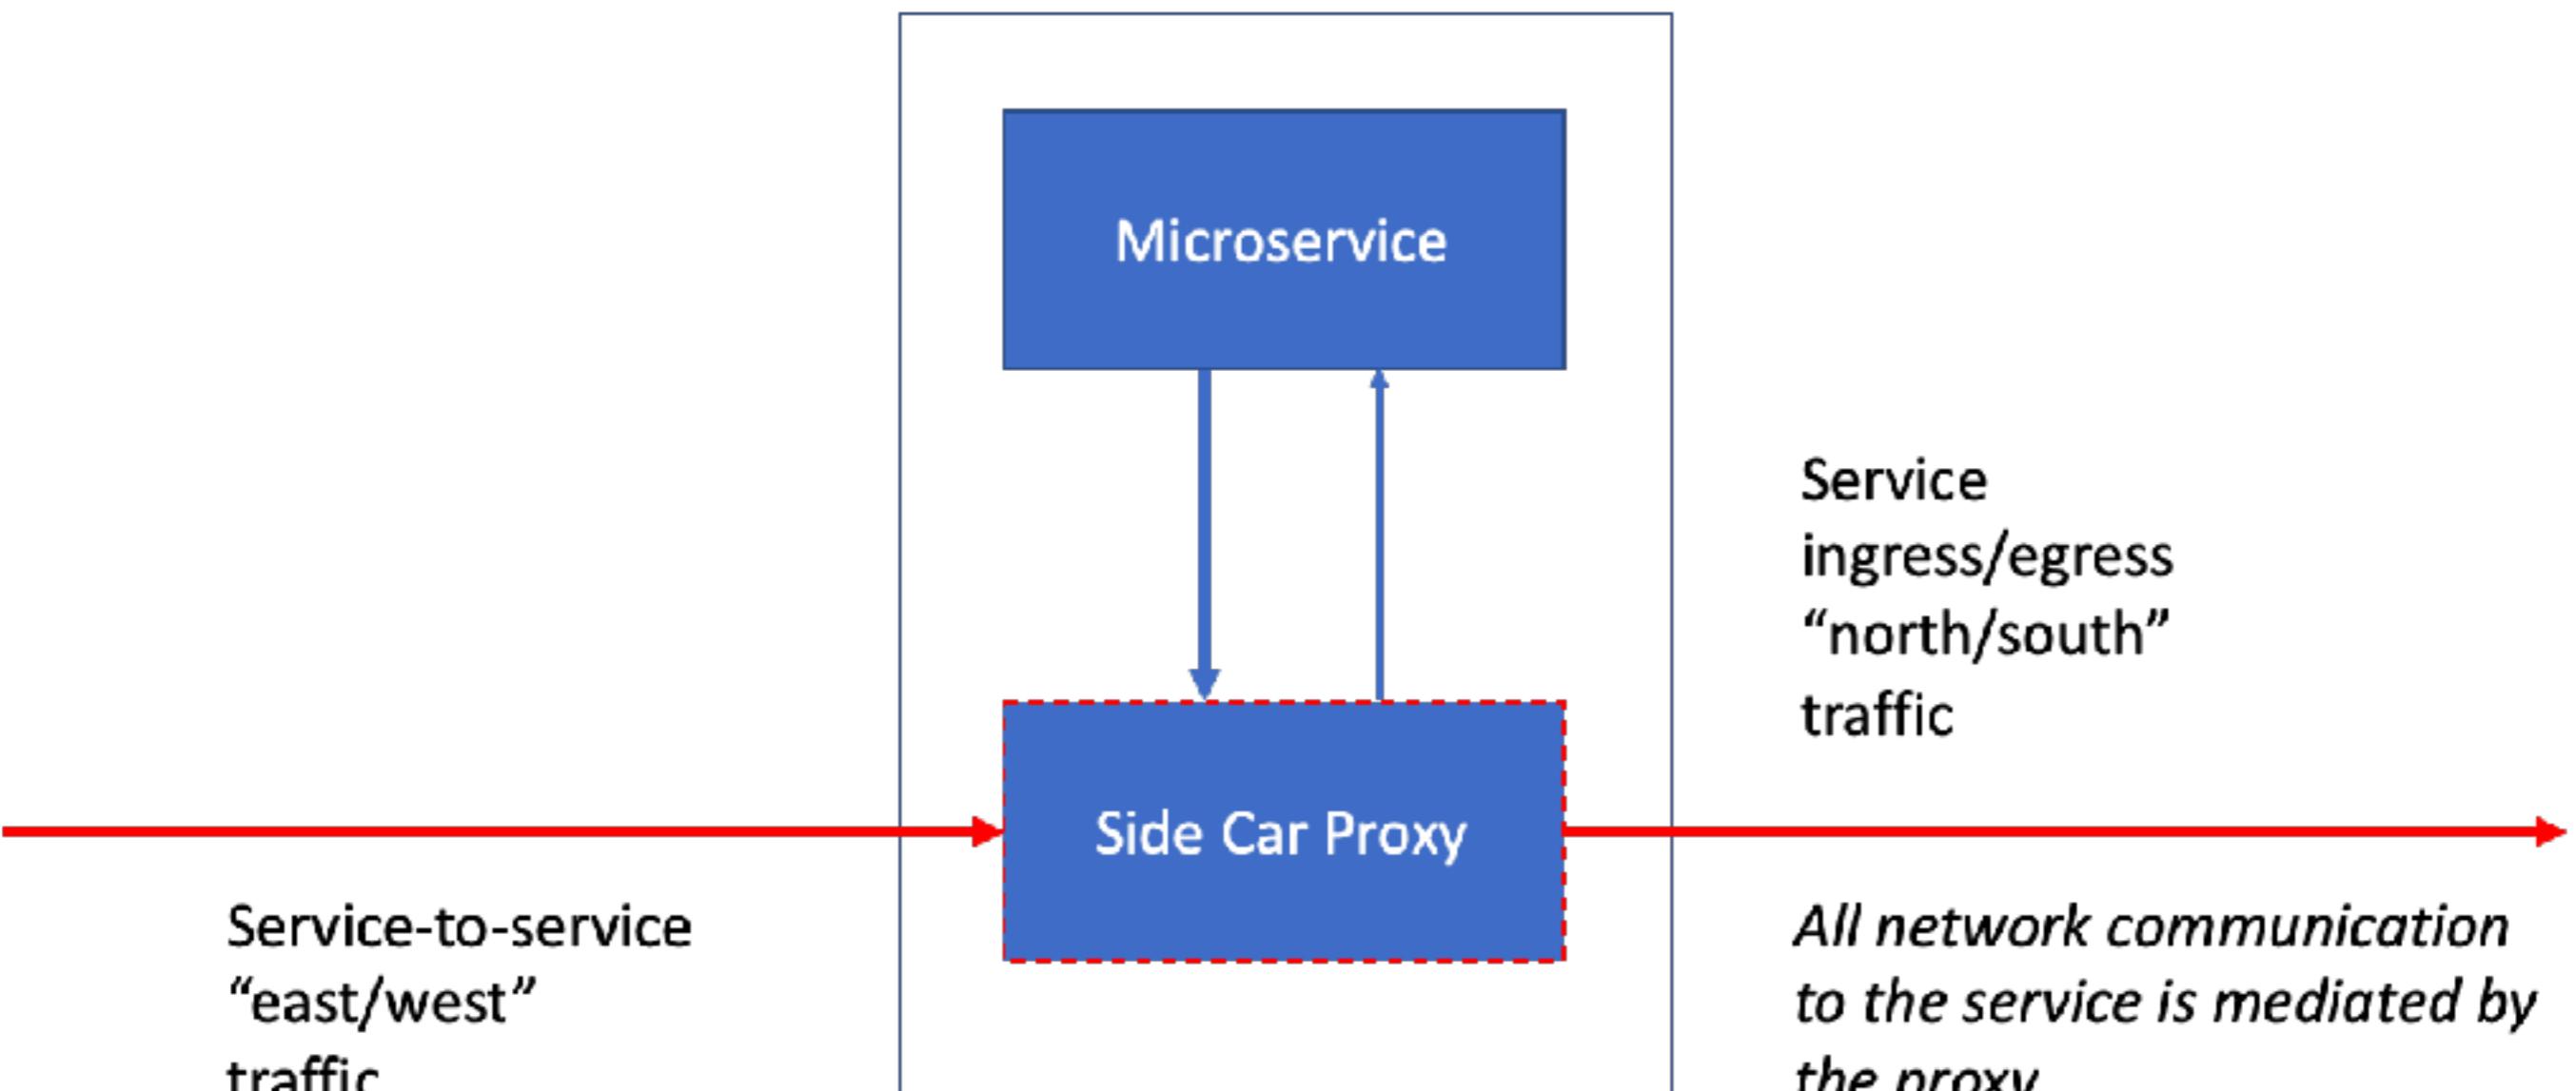
100.68.1.xxx

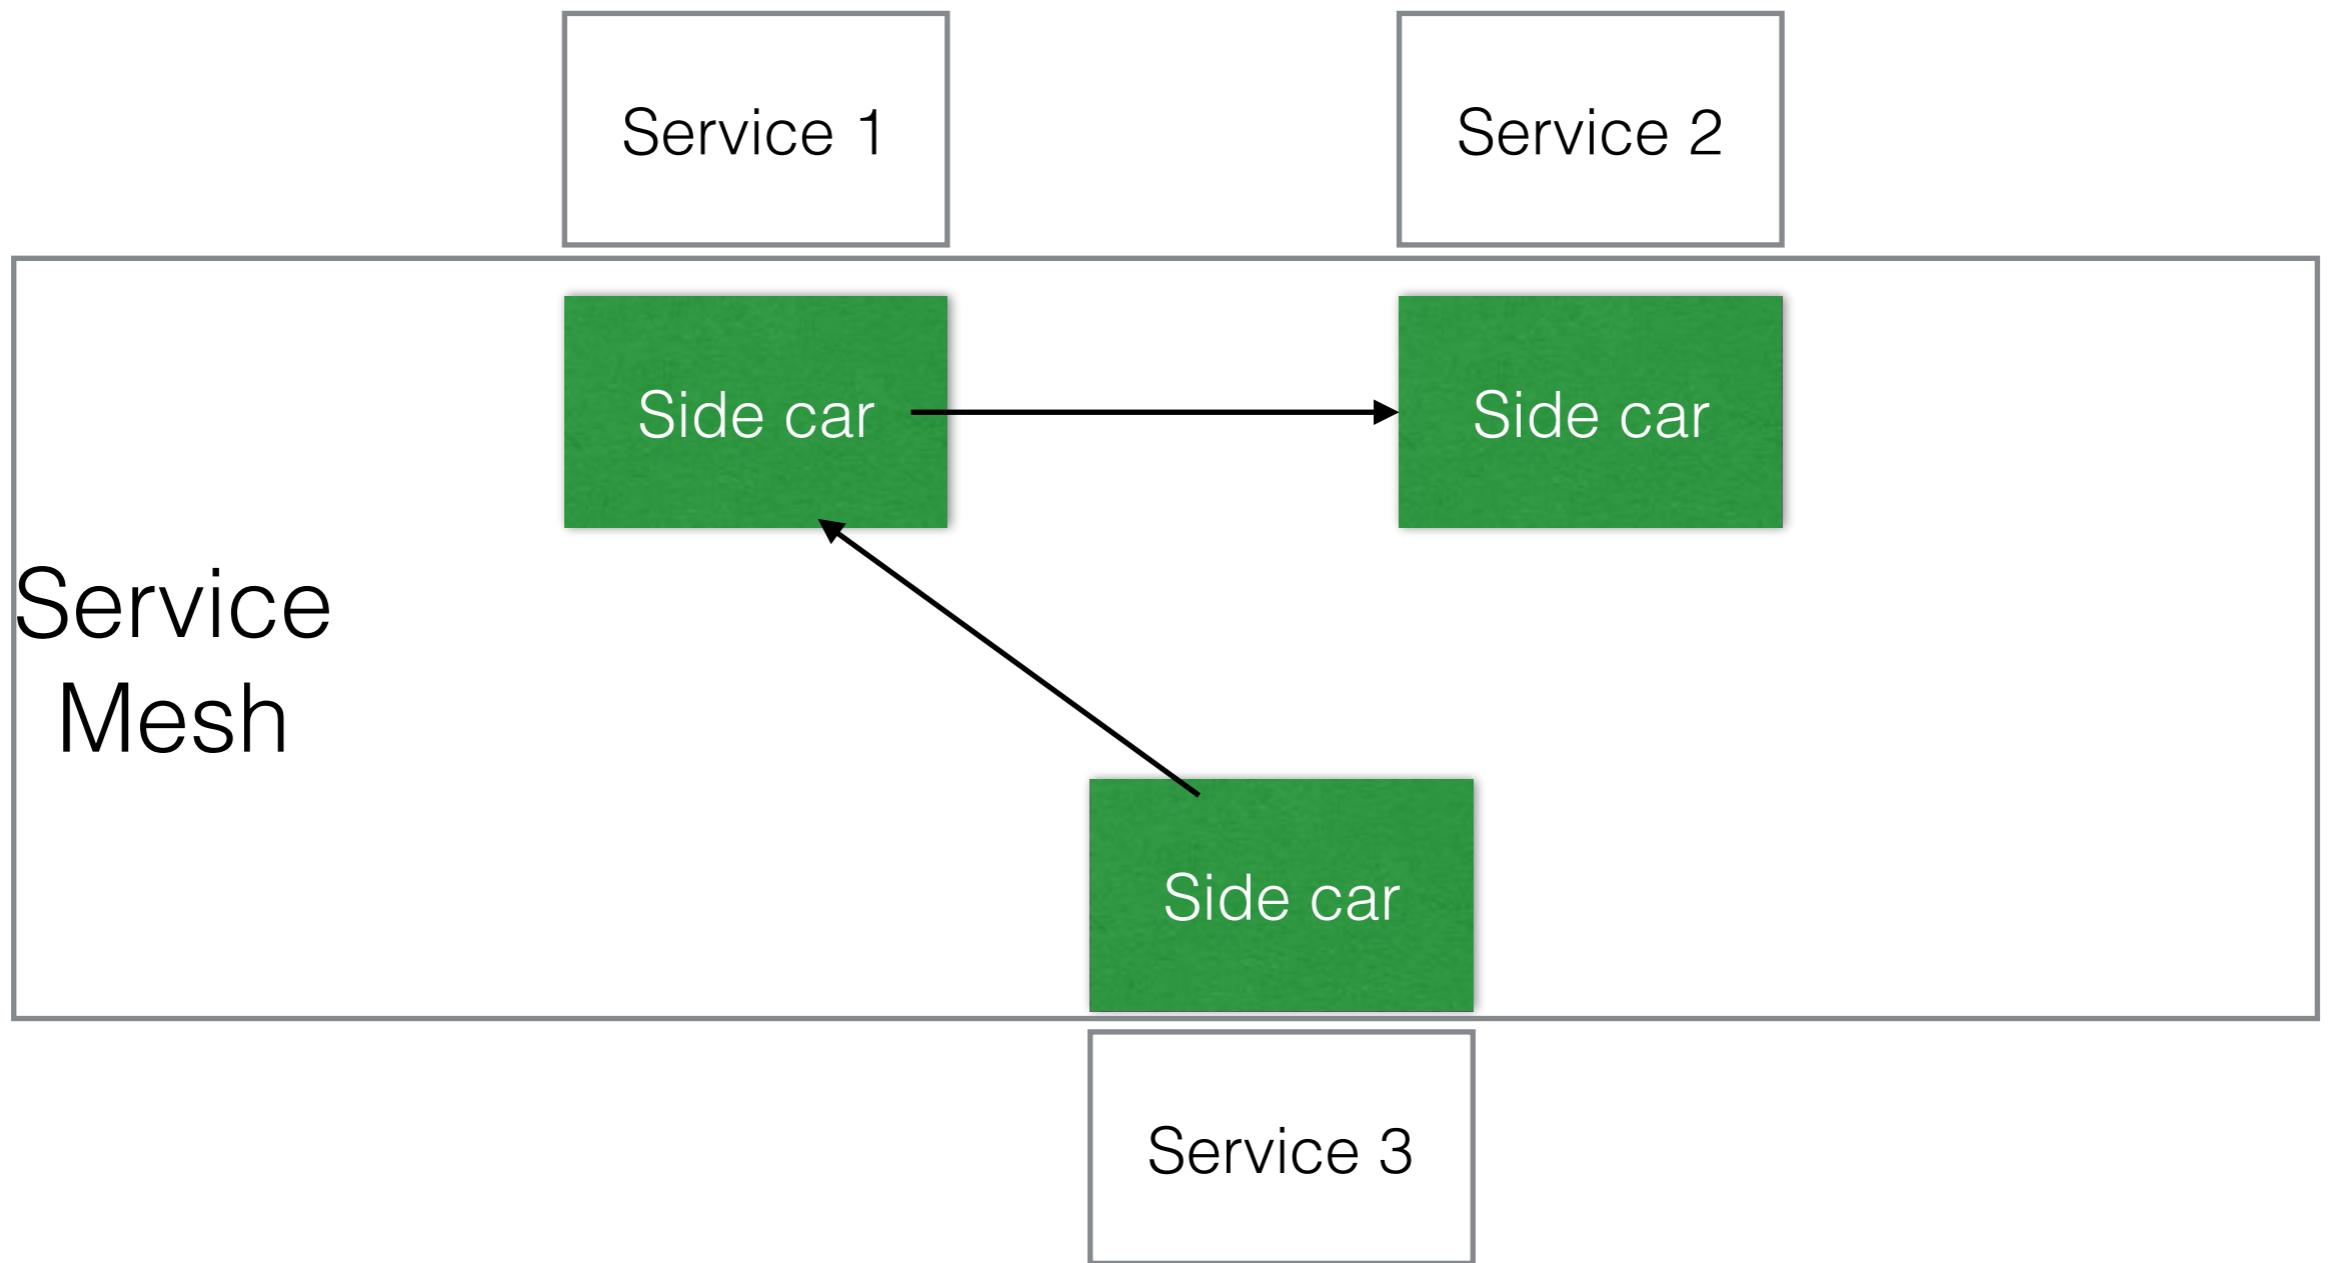


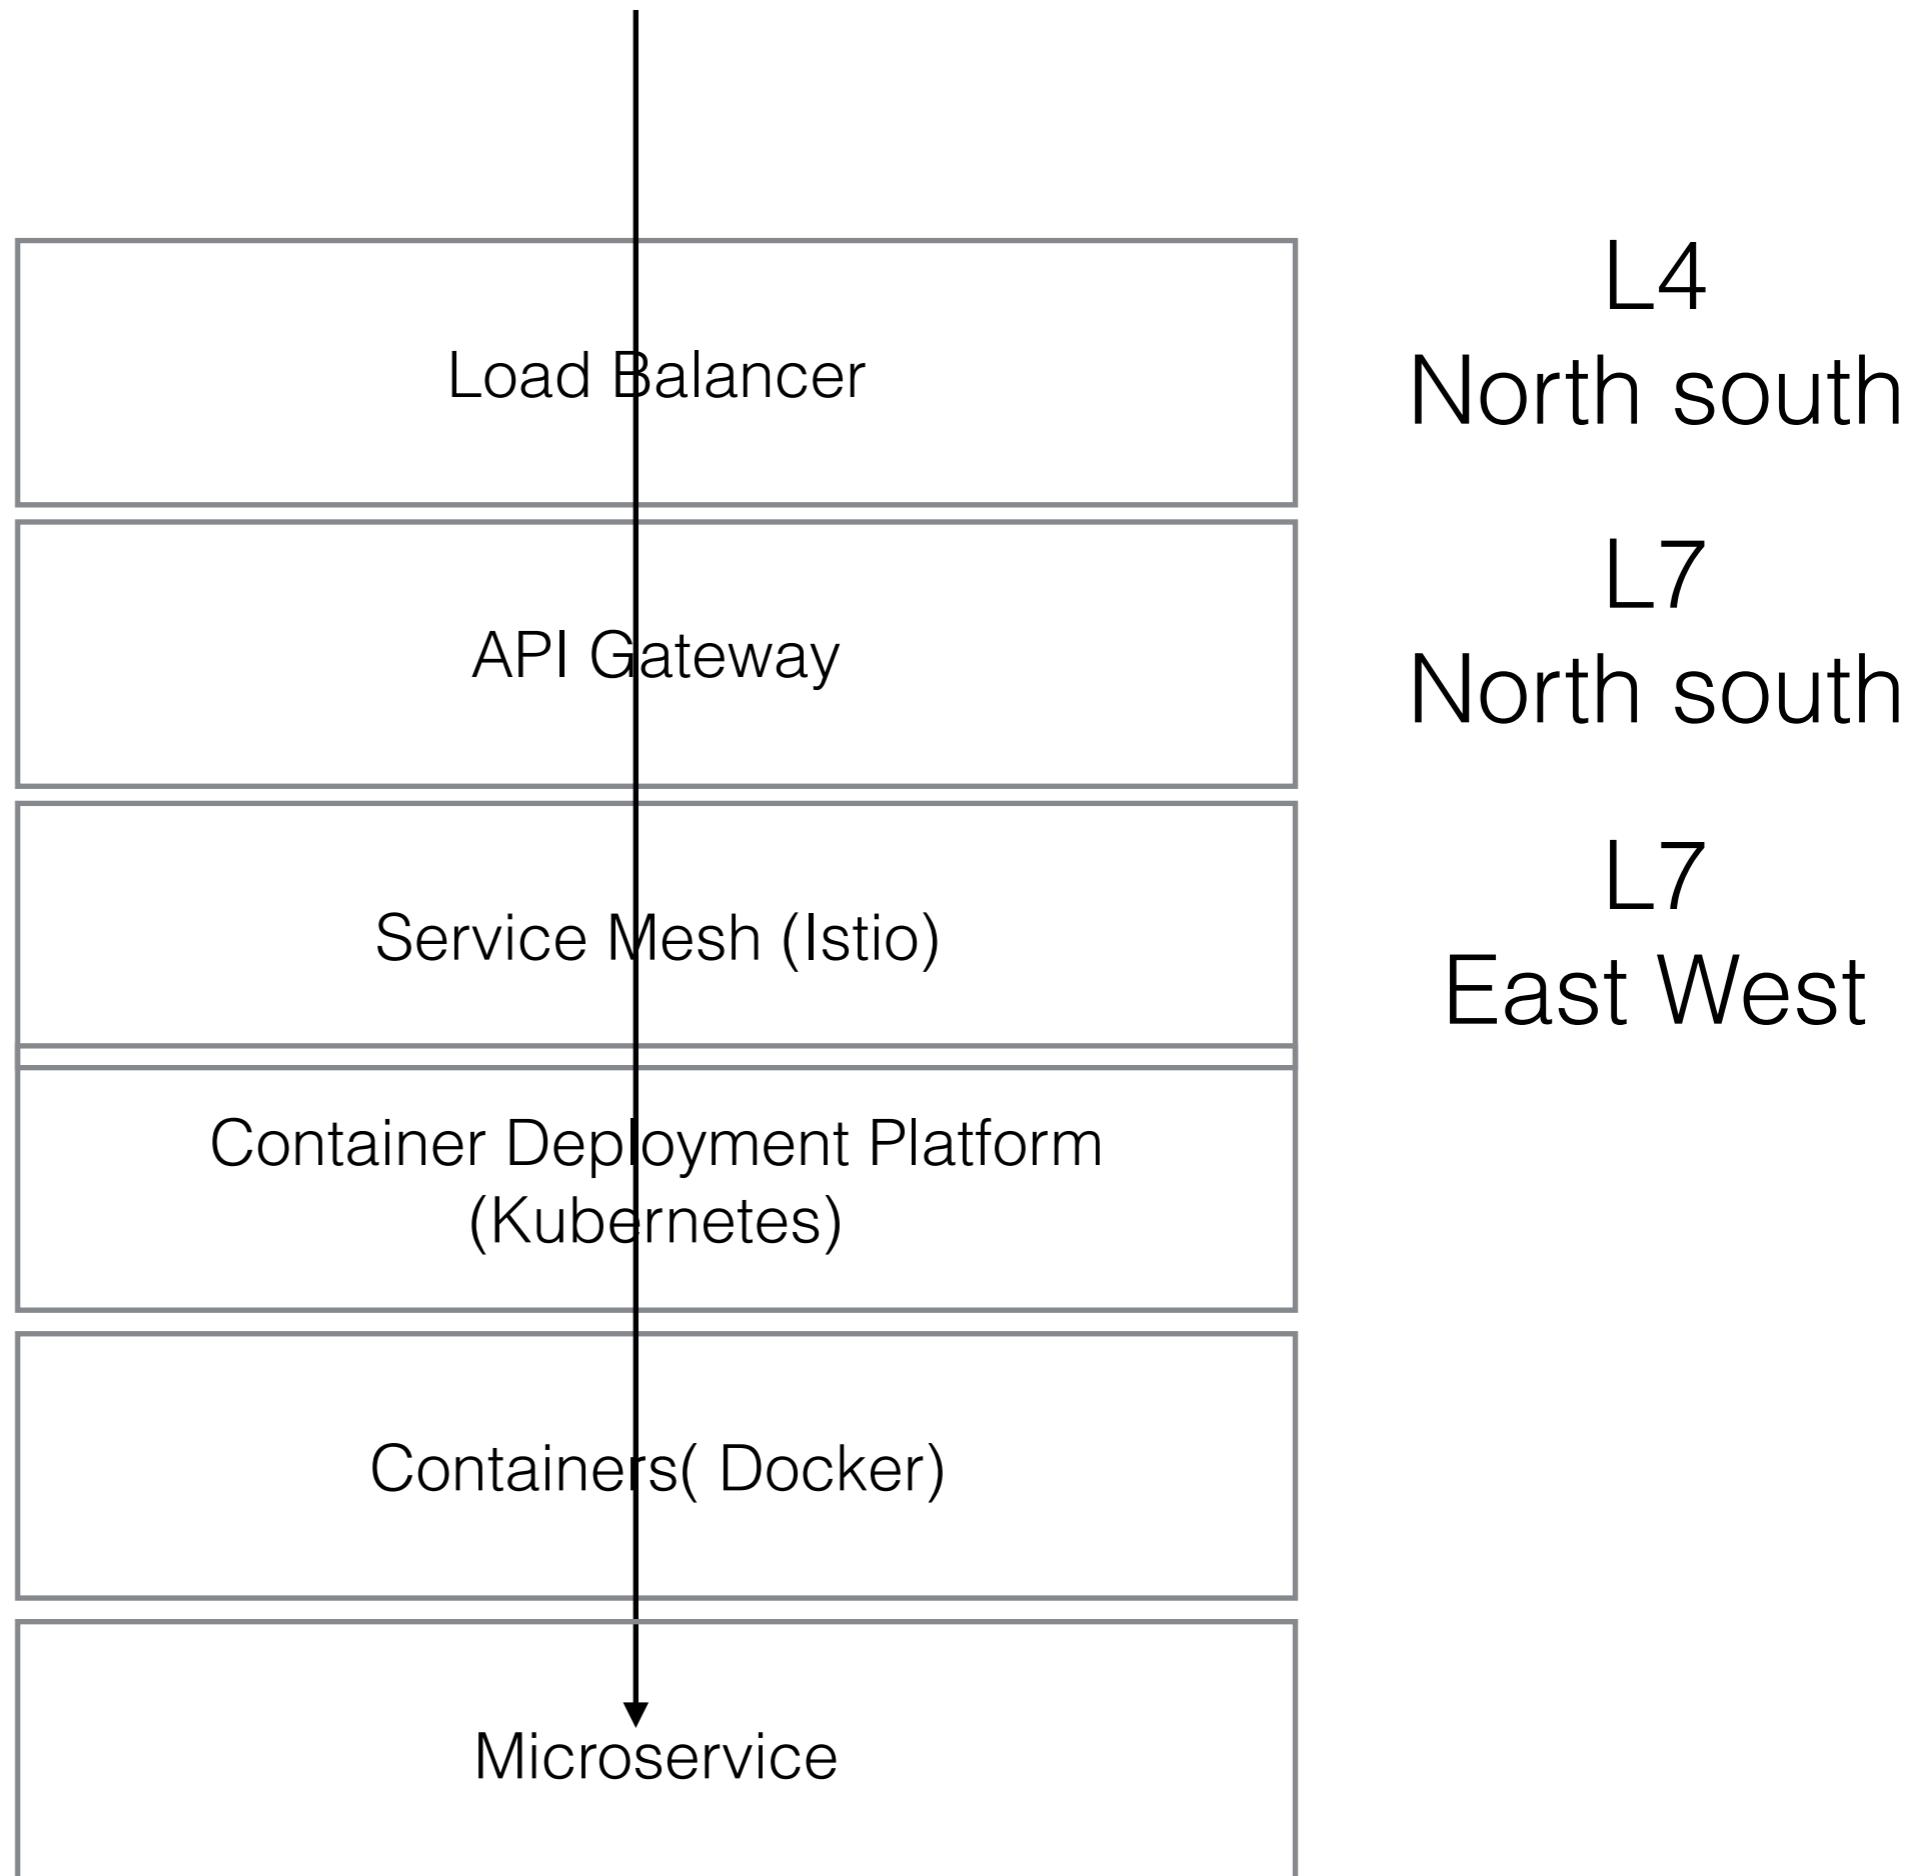


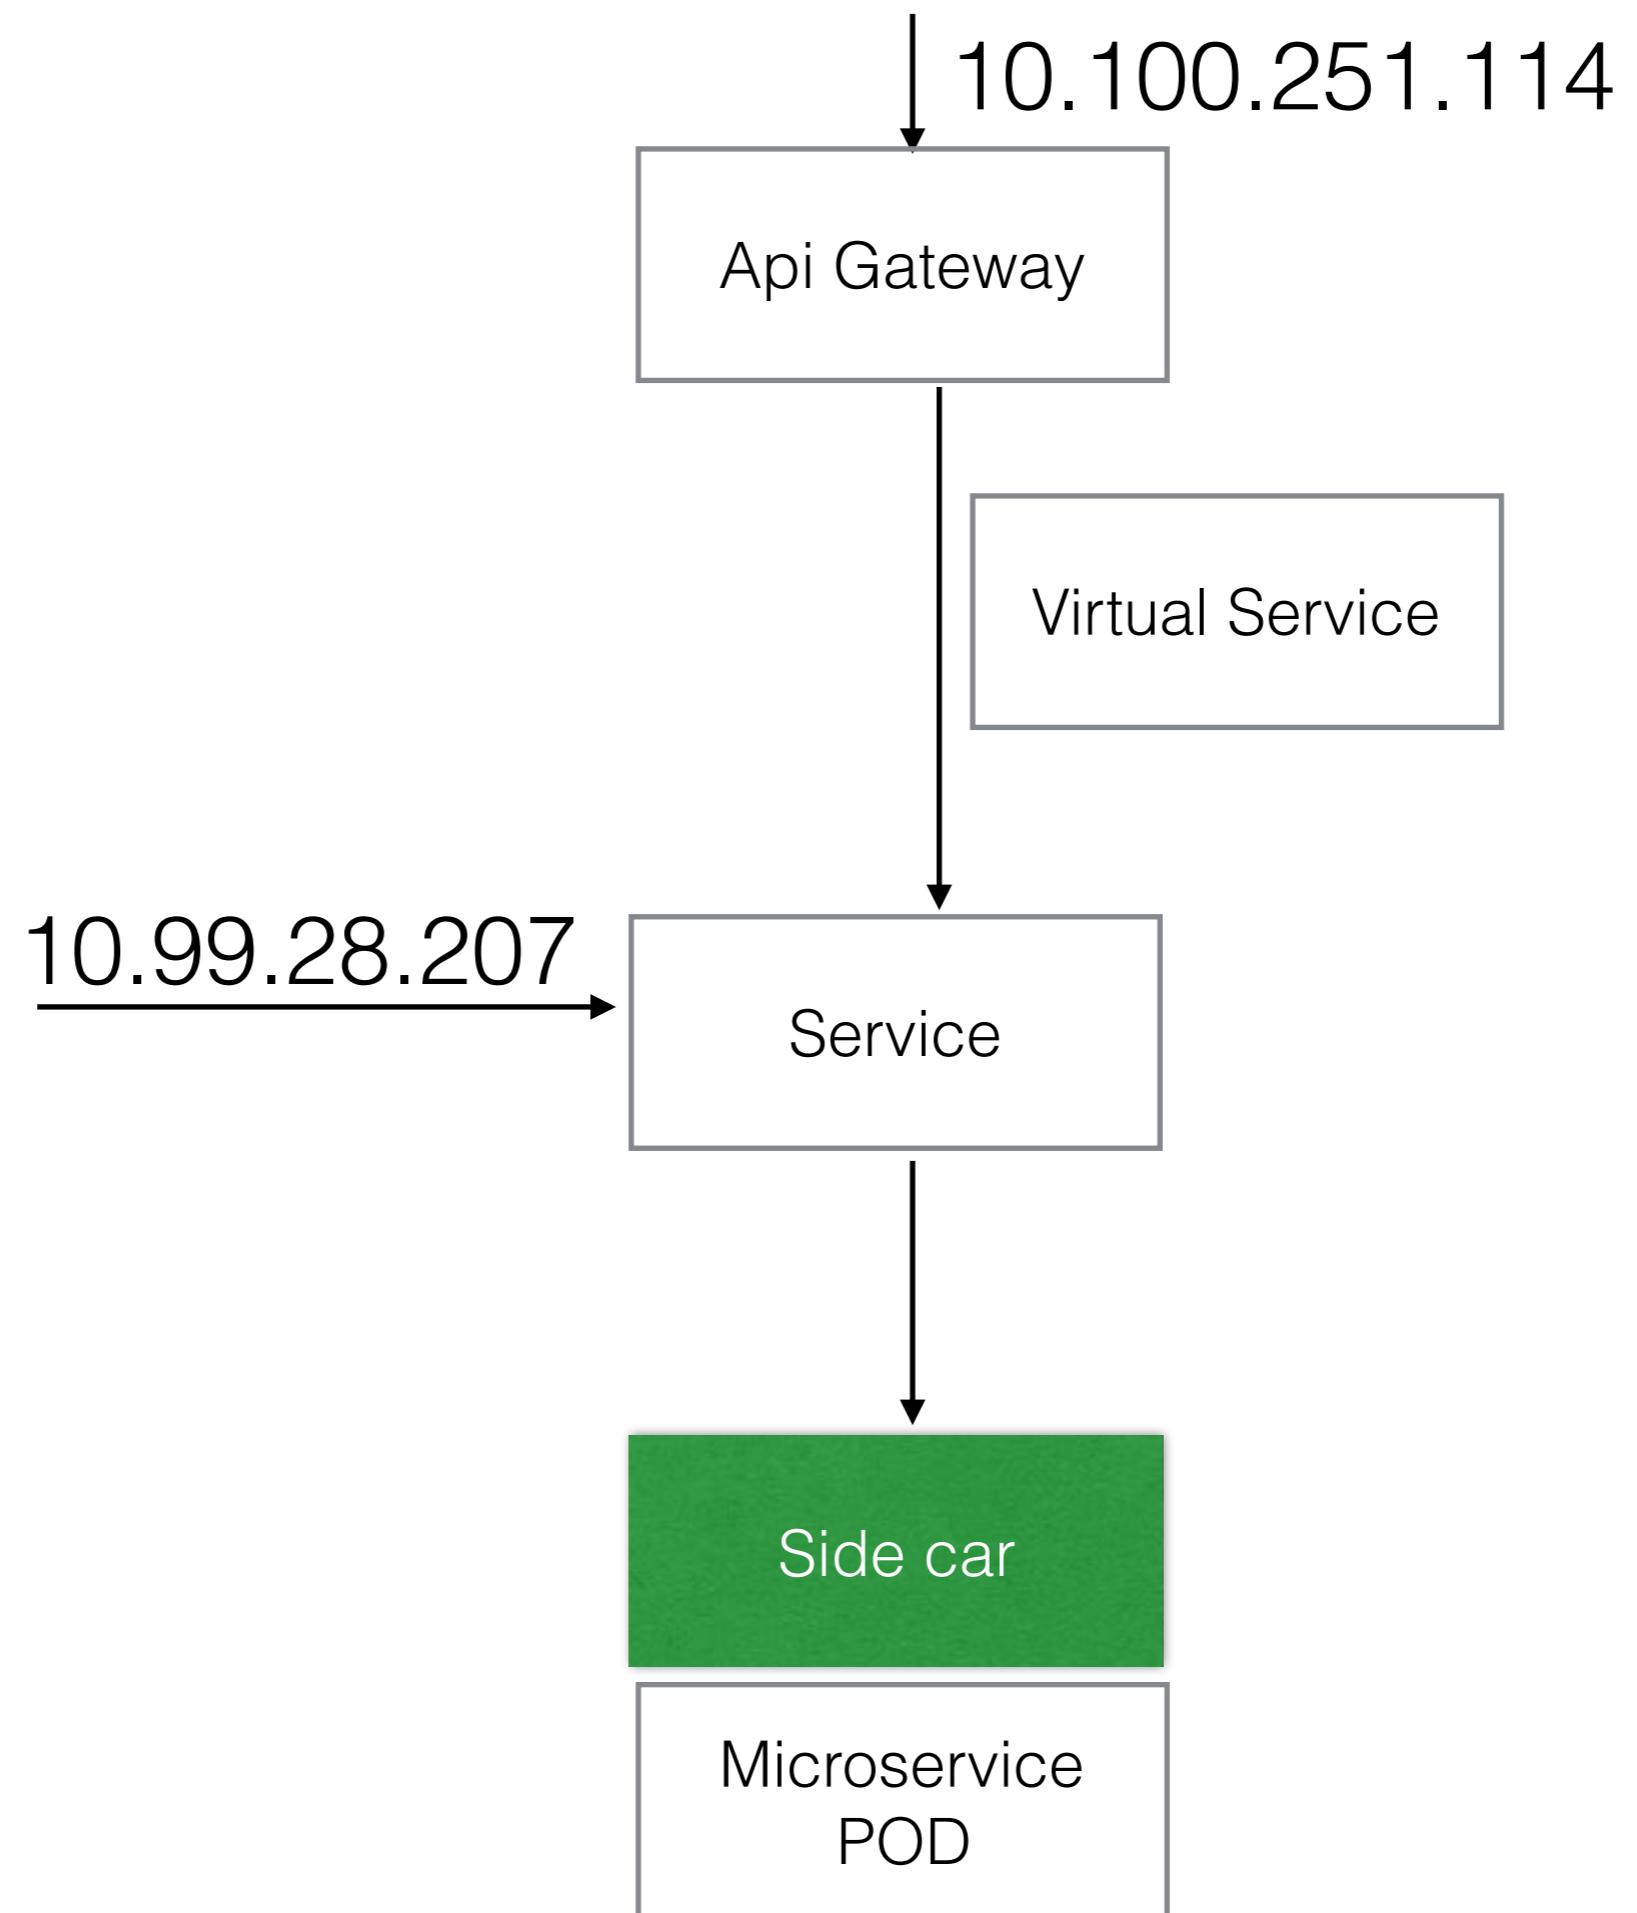
Service Mesh

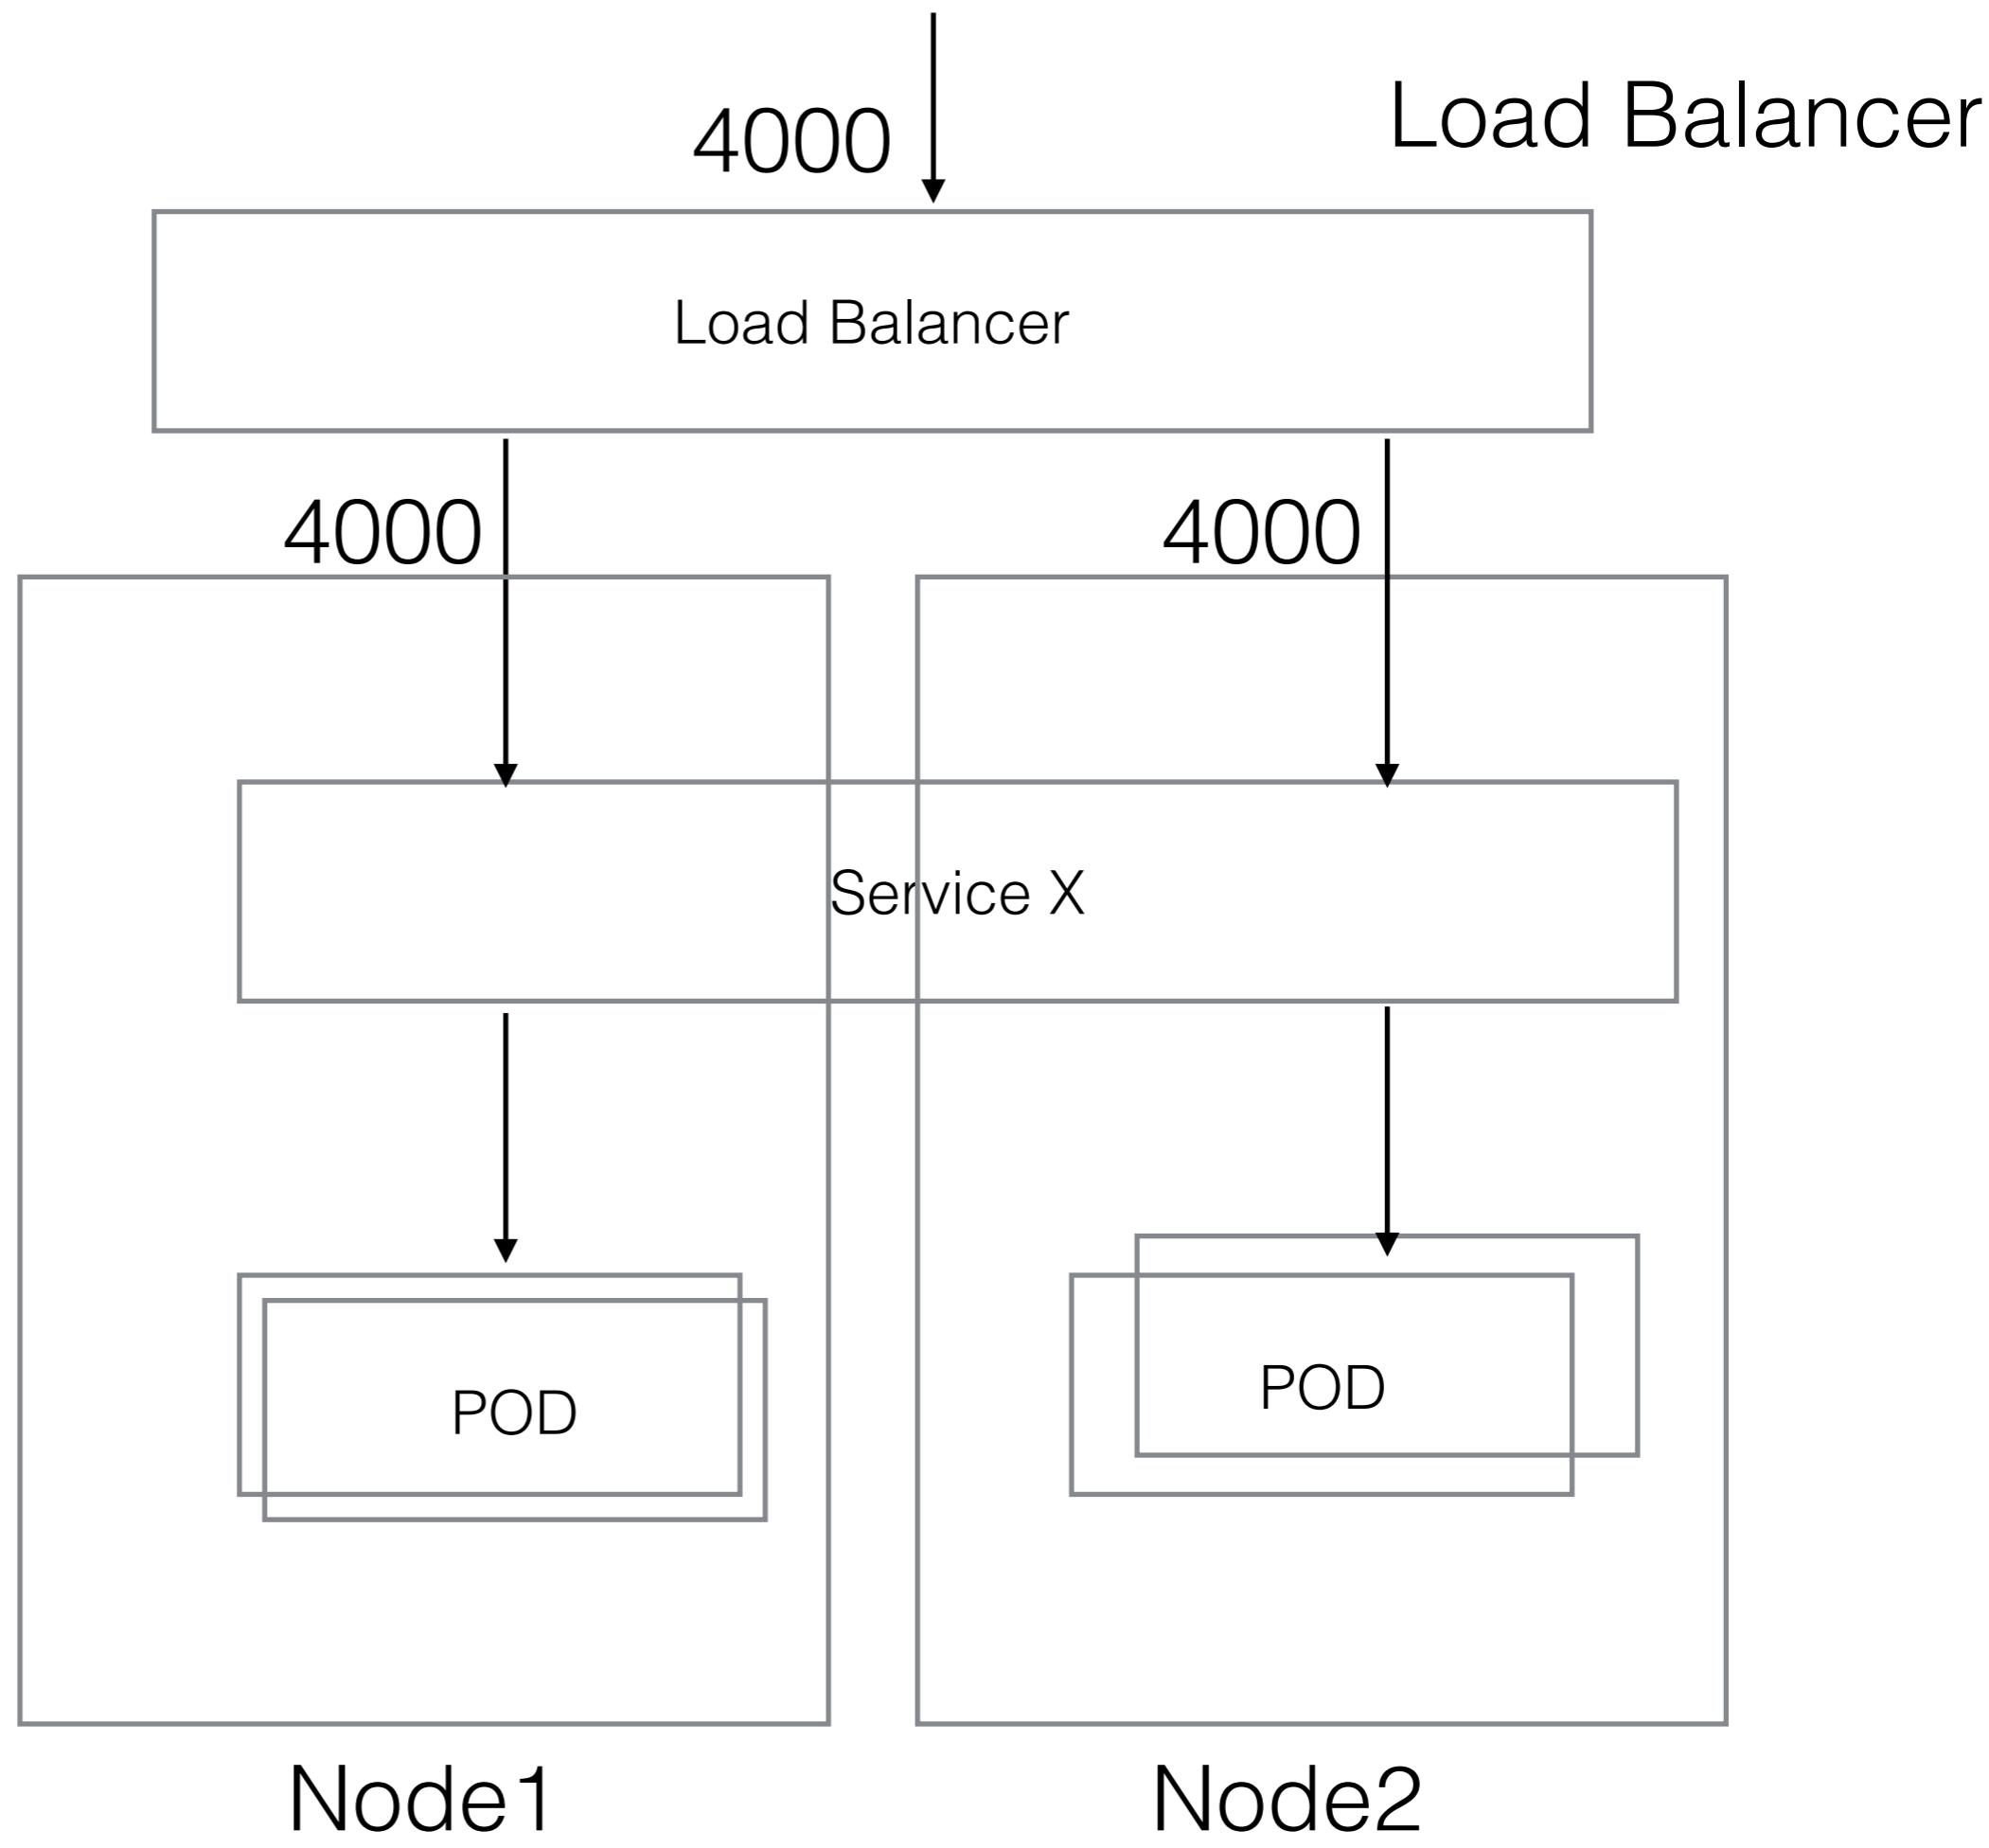
ECS Task/Kubernetes Pod

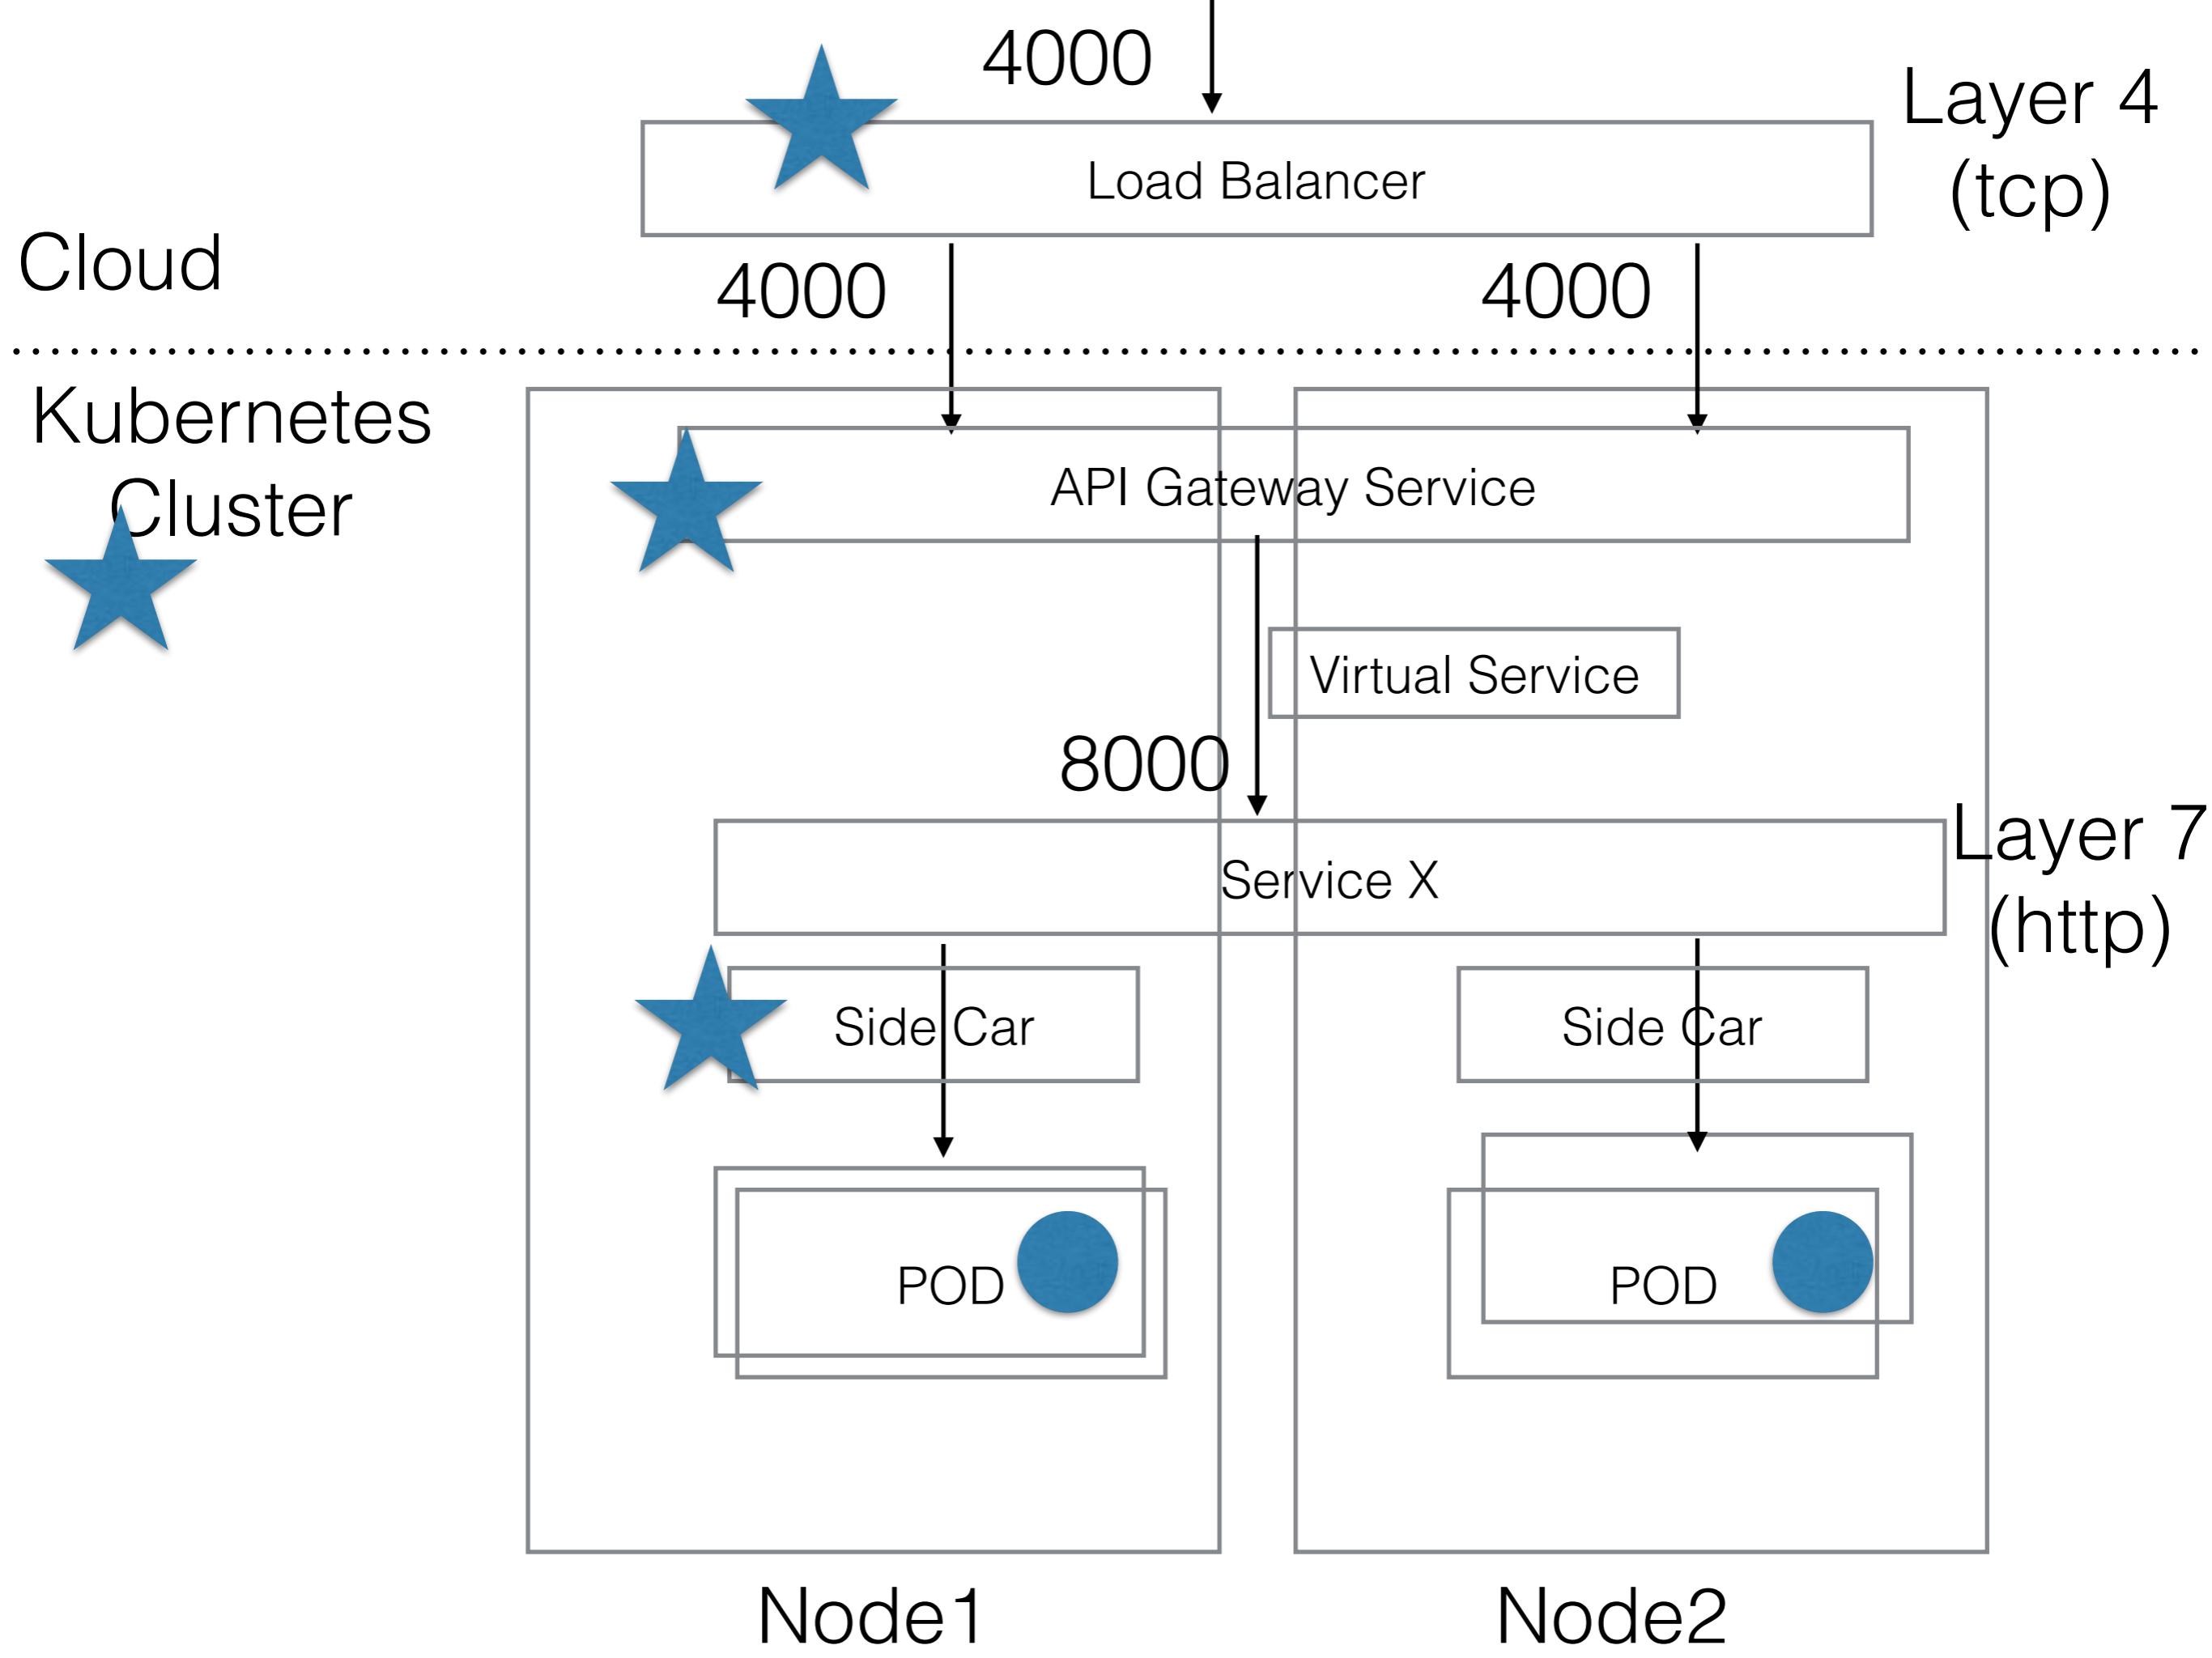




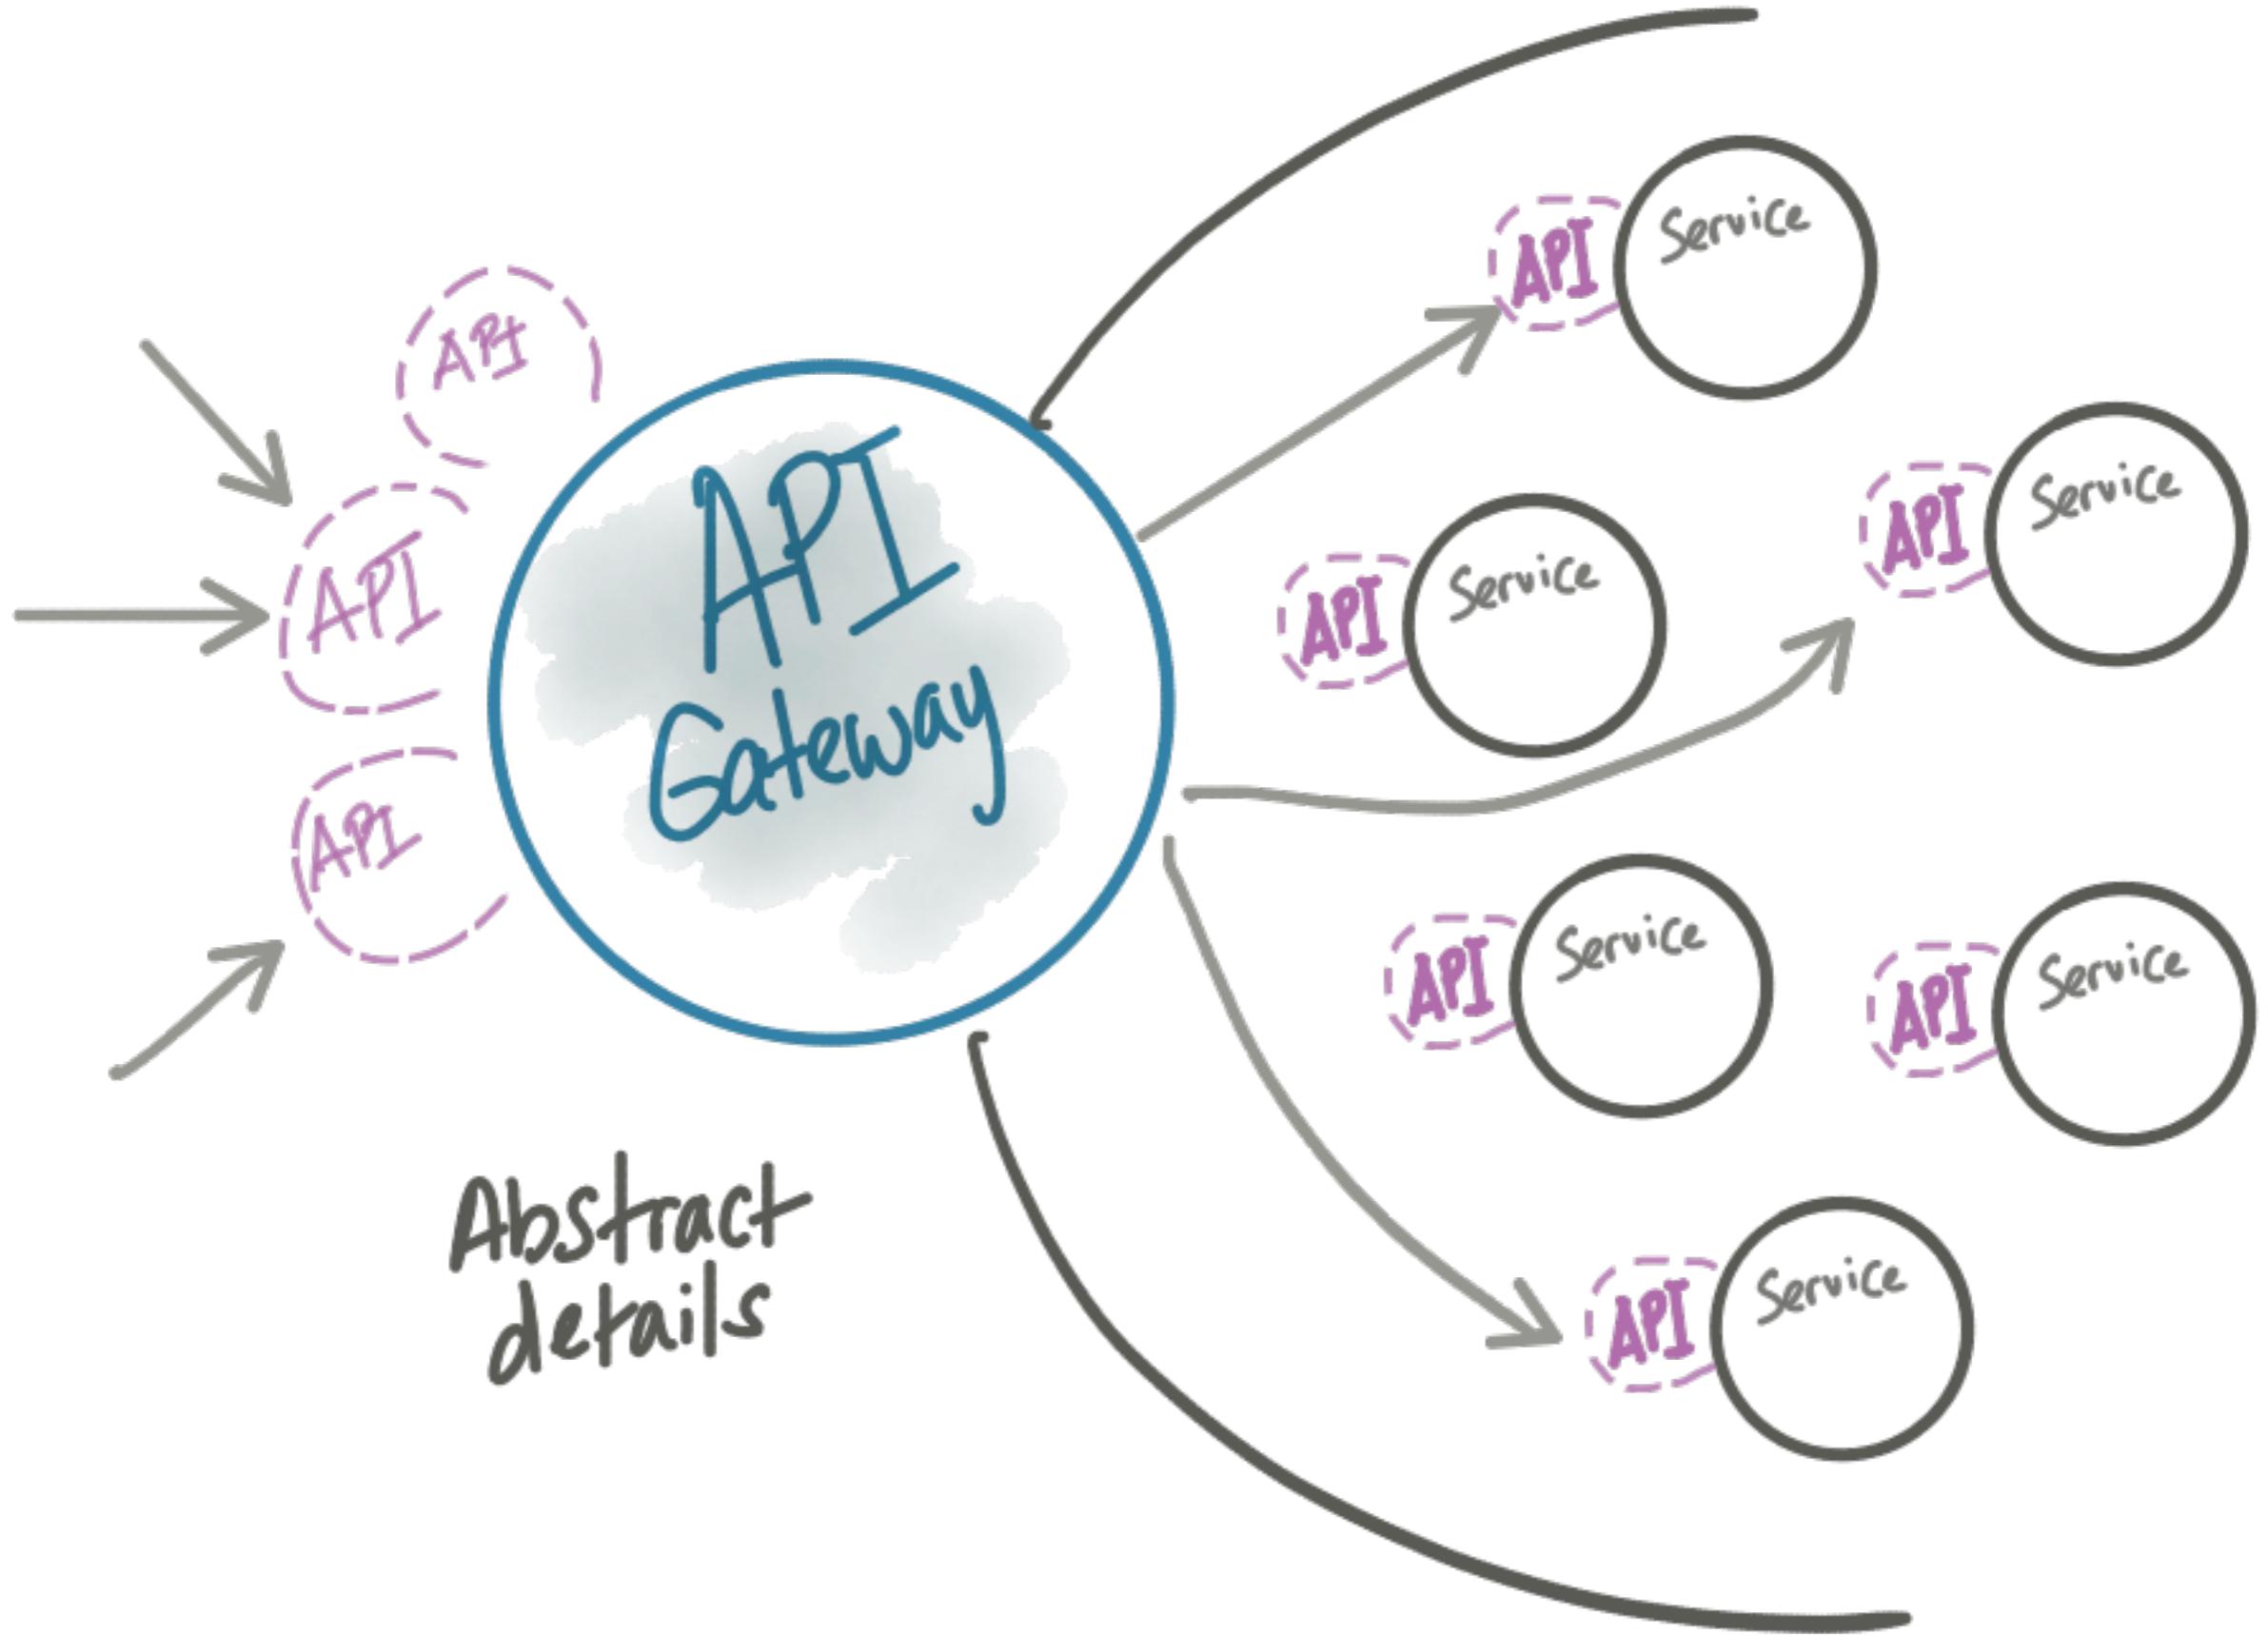


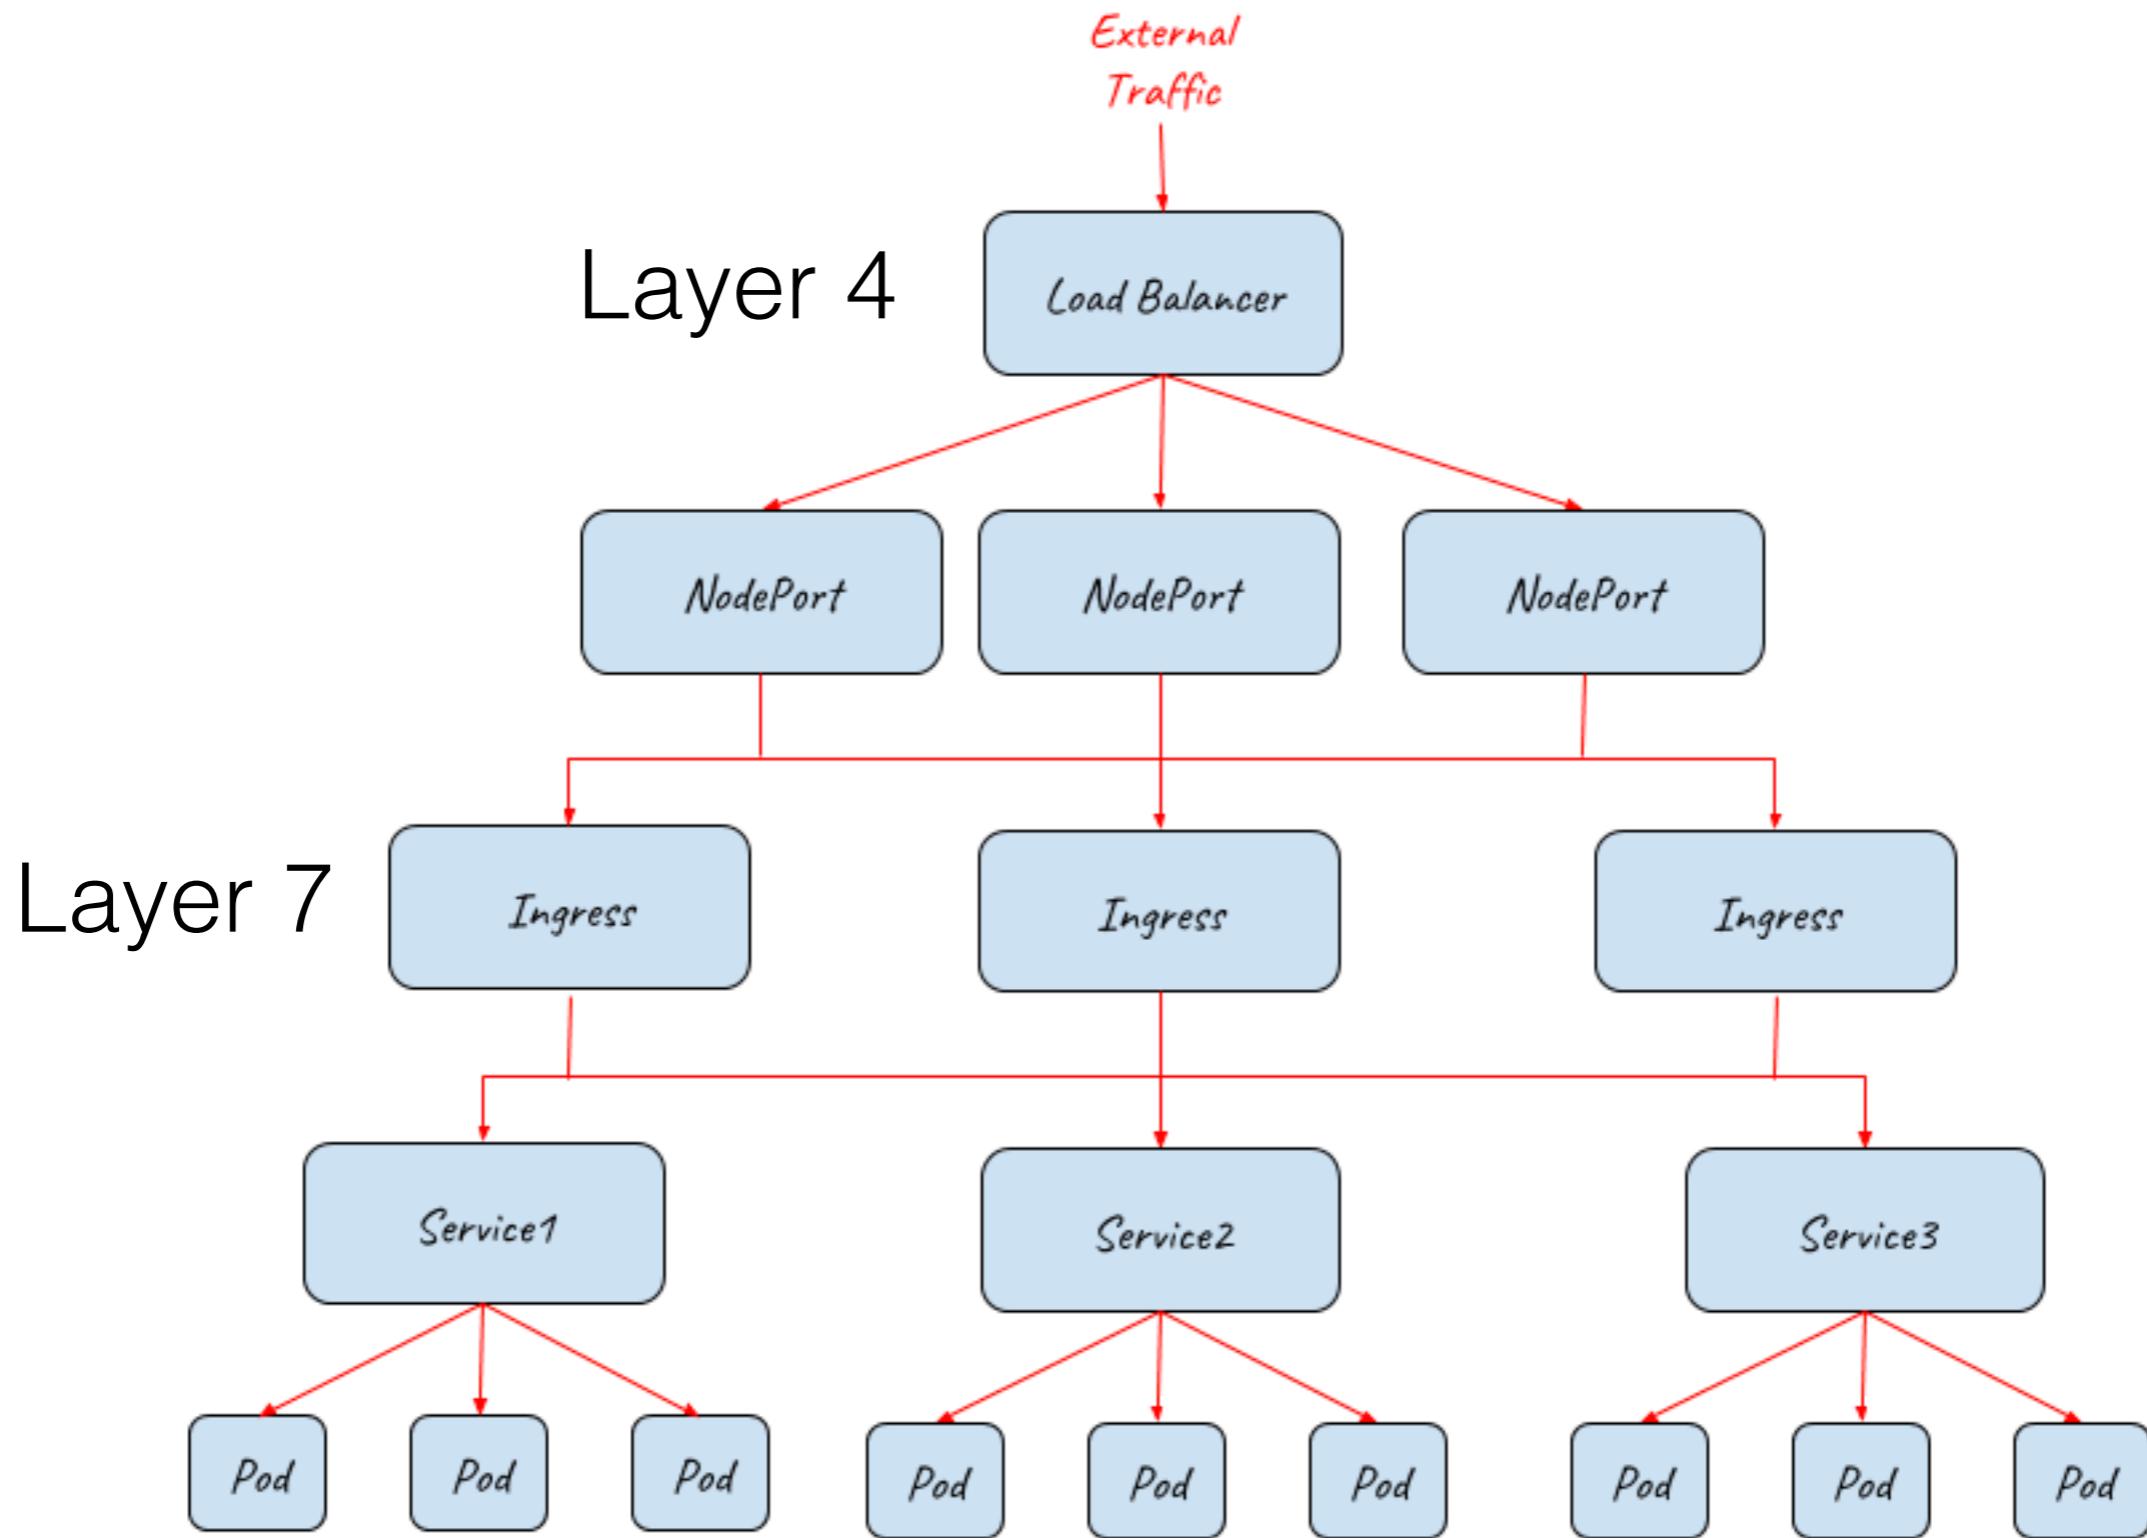




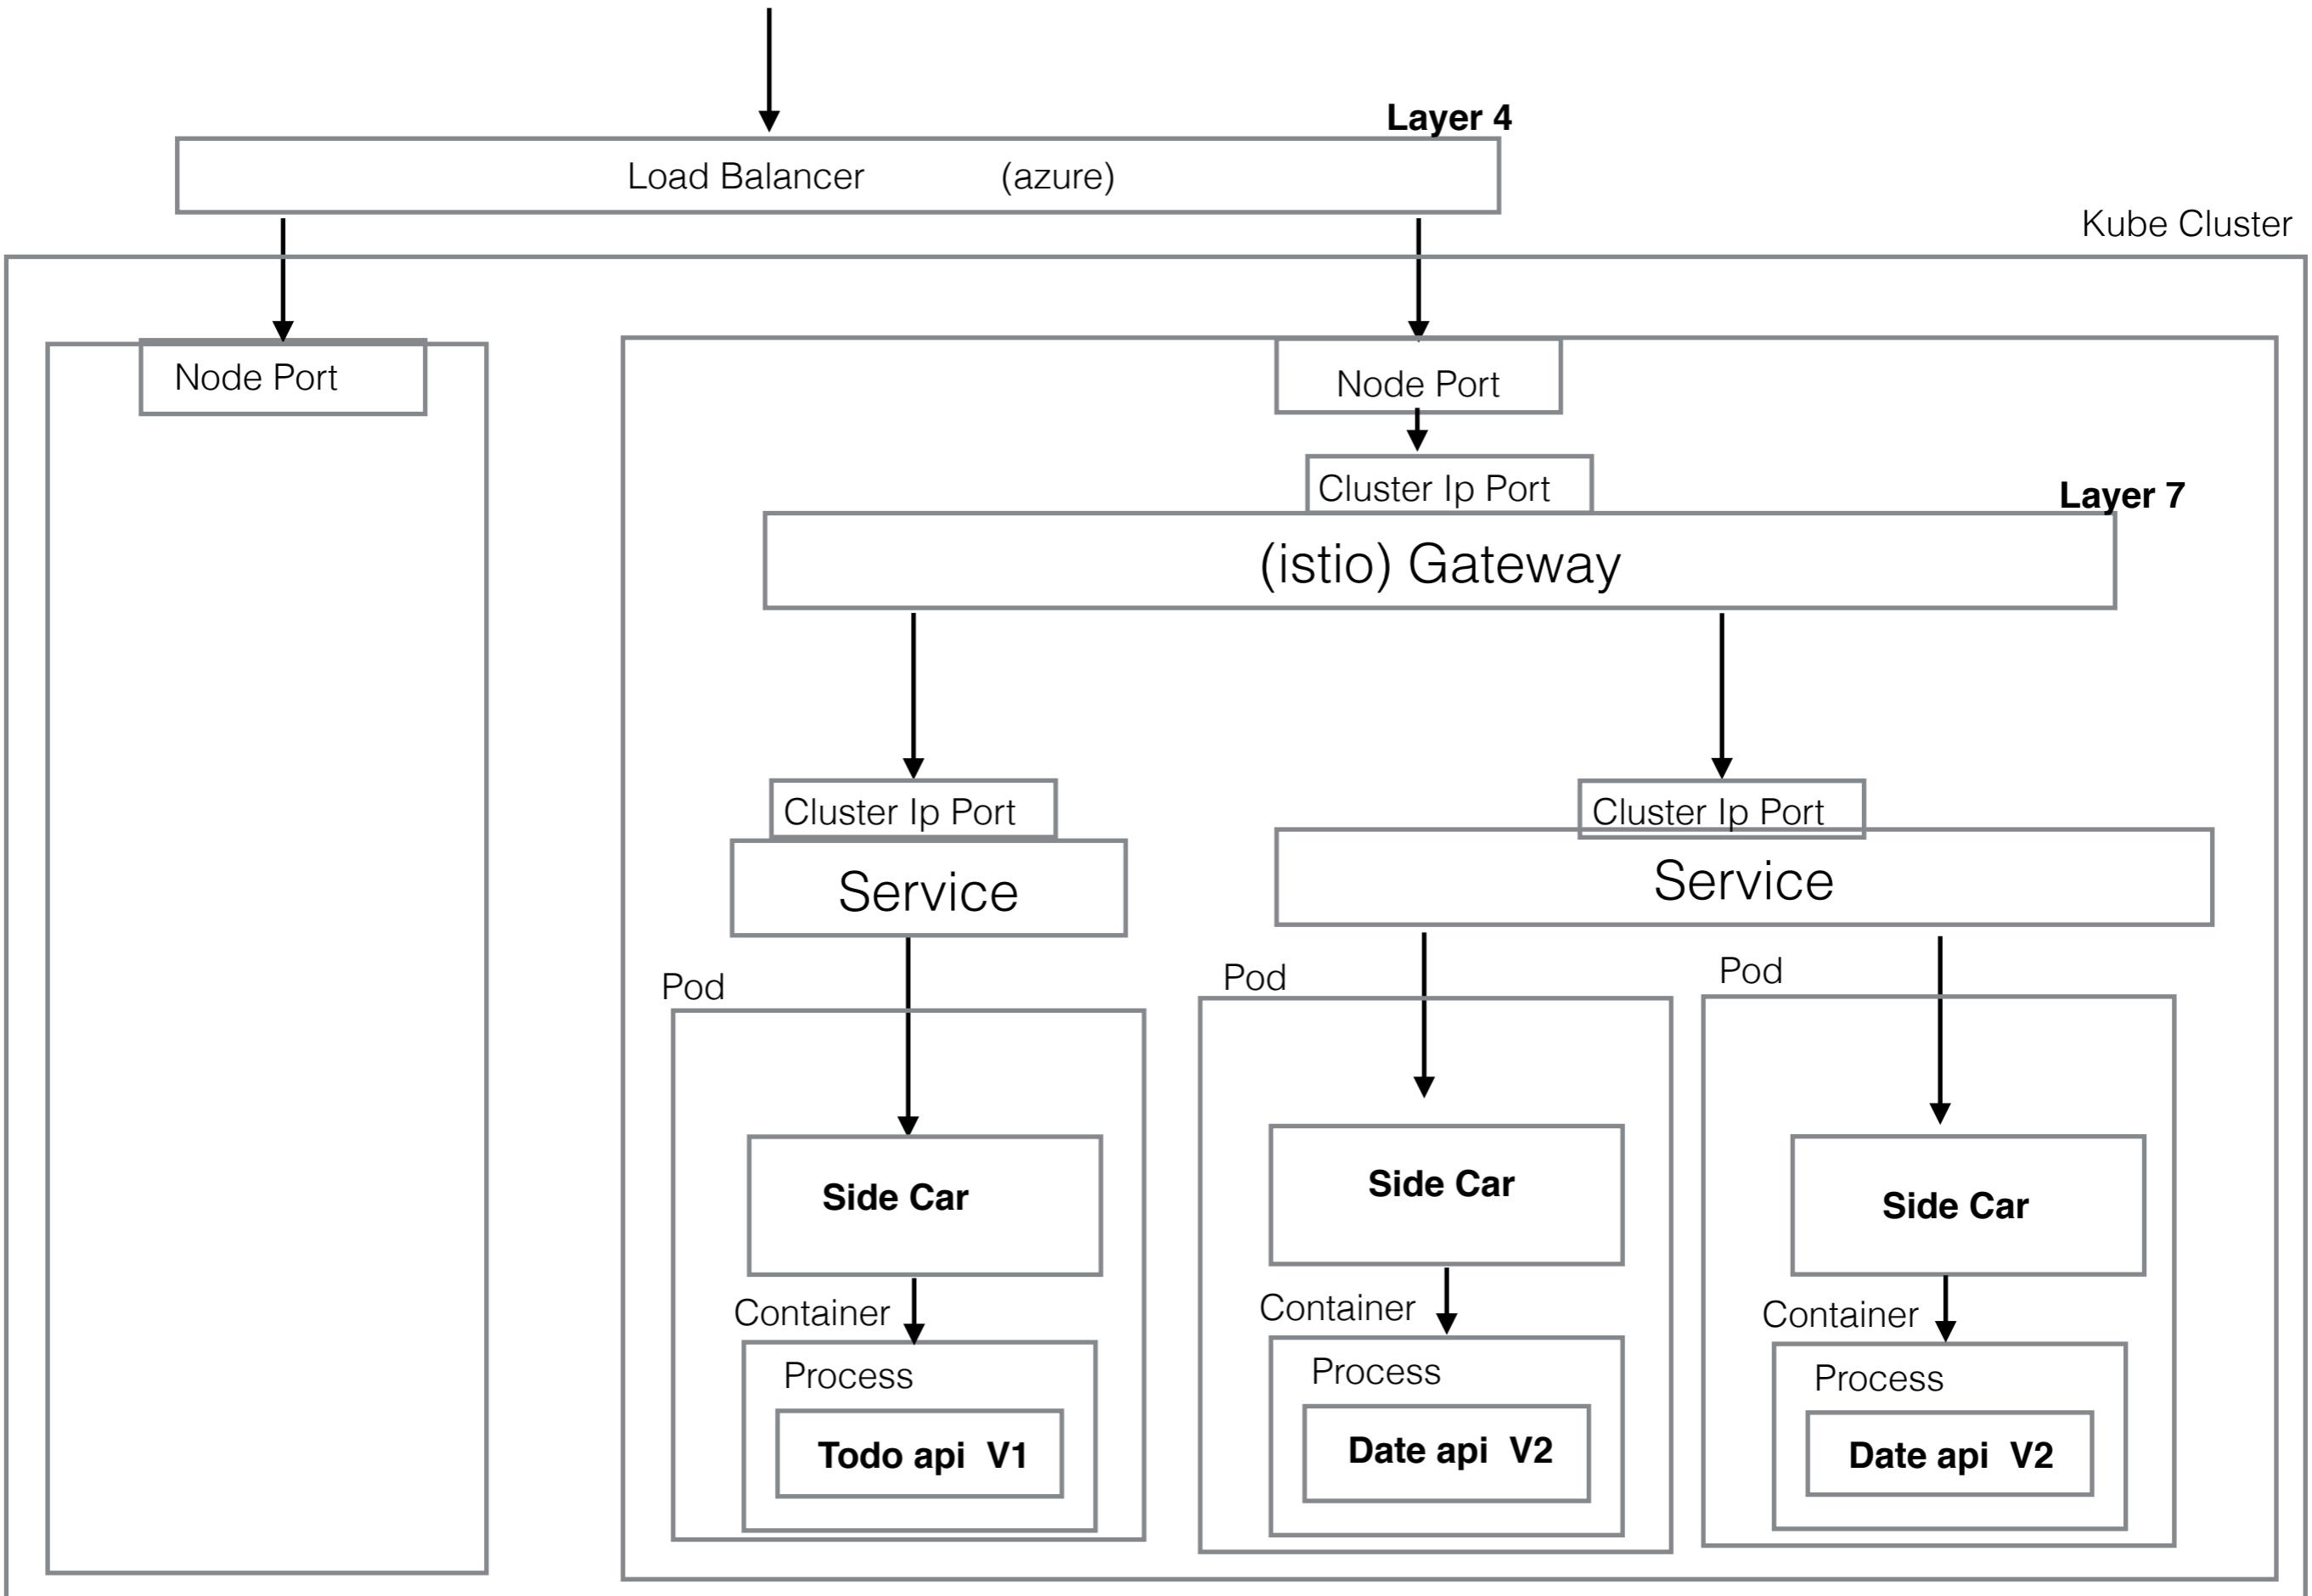


	Approach
Performance	
Transaction Management	
Joins / Reports	
Infra Cost	Container, Kubernetes
Deployment	Docker Image, Hub
Debugging (Finding bugs)	ServiceMesh
Log Mgmt	
Config Mgmt	
Authentication	
Authorization	
Monitoring / Alerting	Kubernetes , ServiceMesh





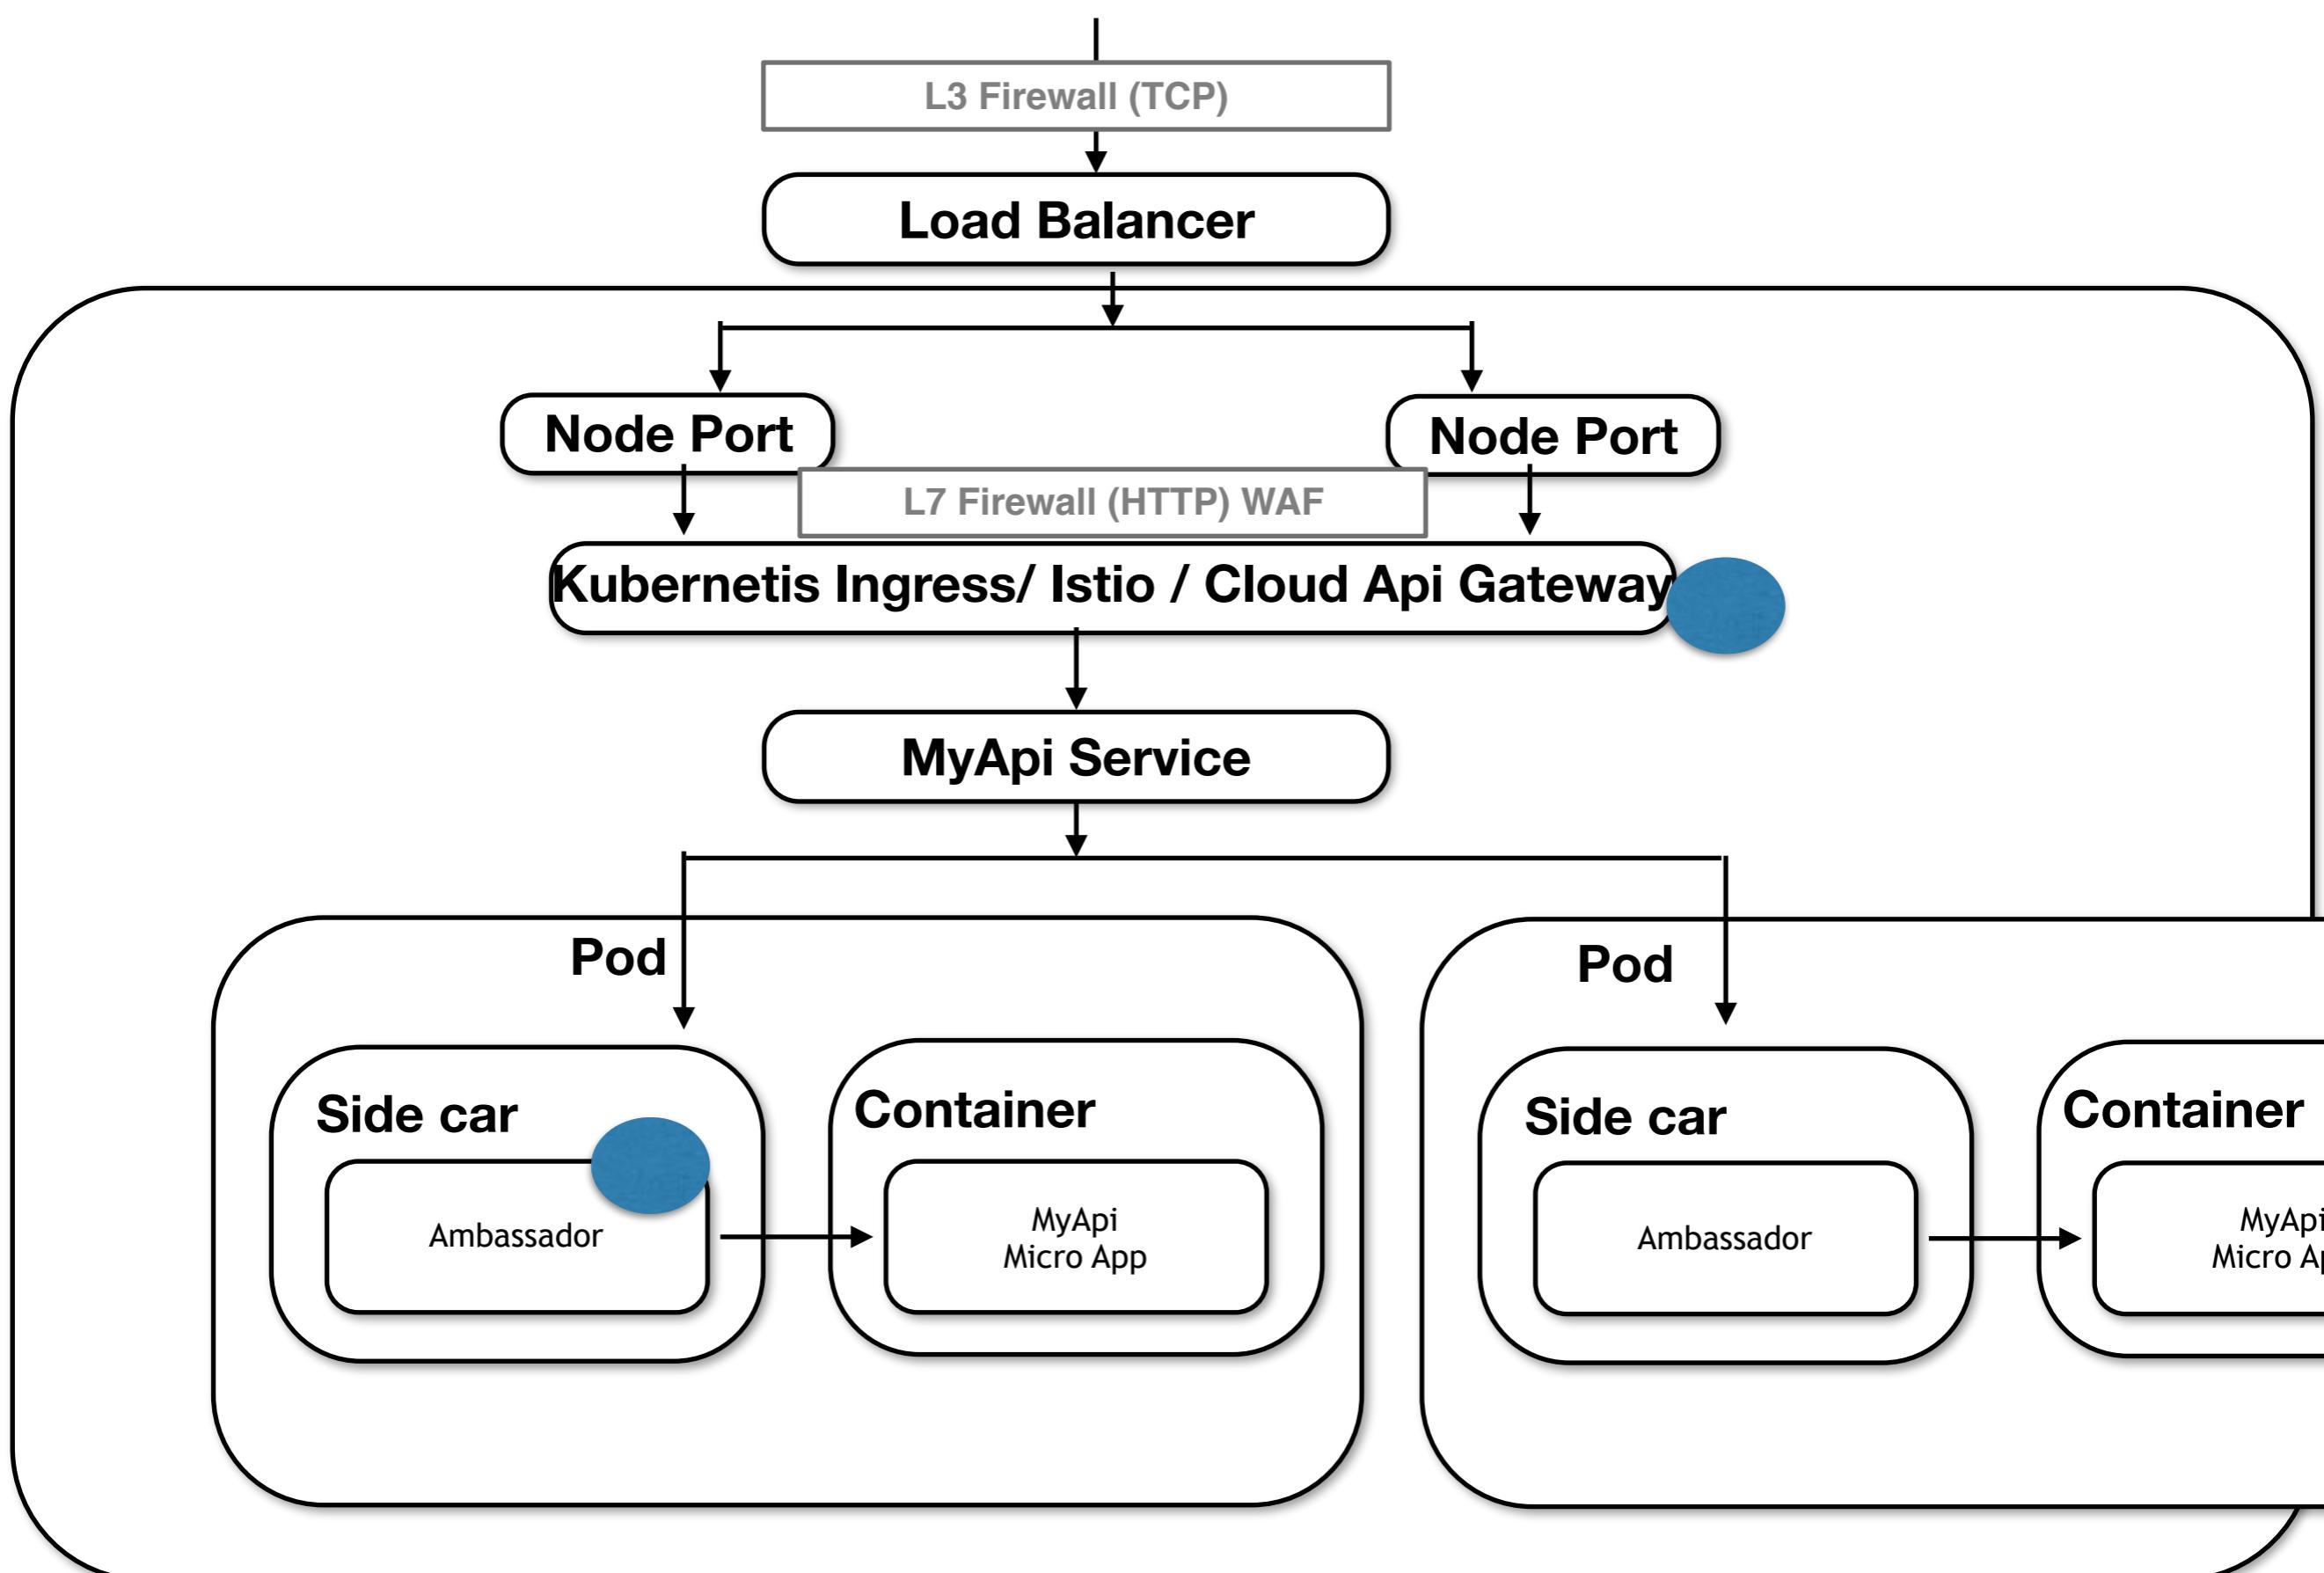
Each of the NodePort, Ingress or Pod layers can be scale out/in accordingly to handle different working loads.

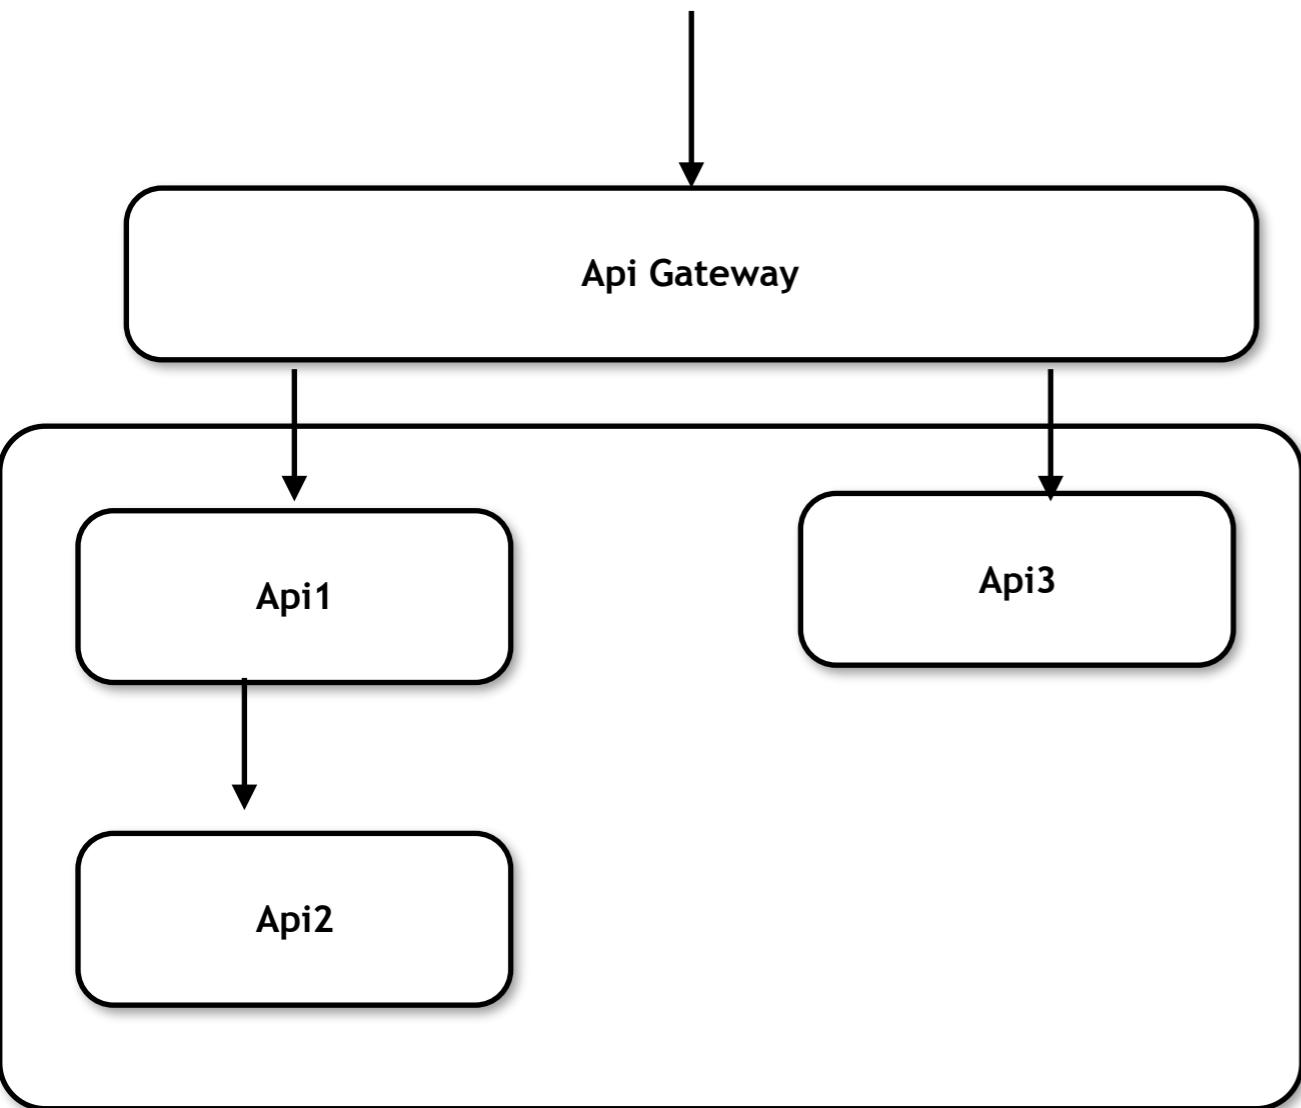


Kubernetes Ingress v/s Istio Gateway v/s API Gateway?

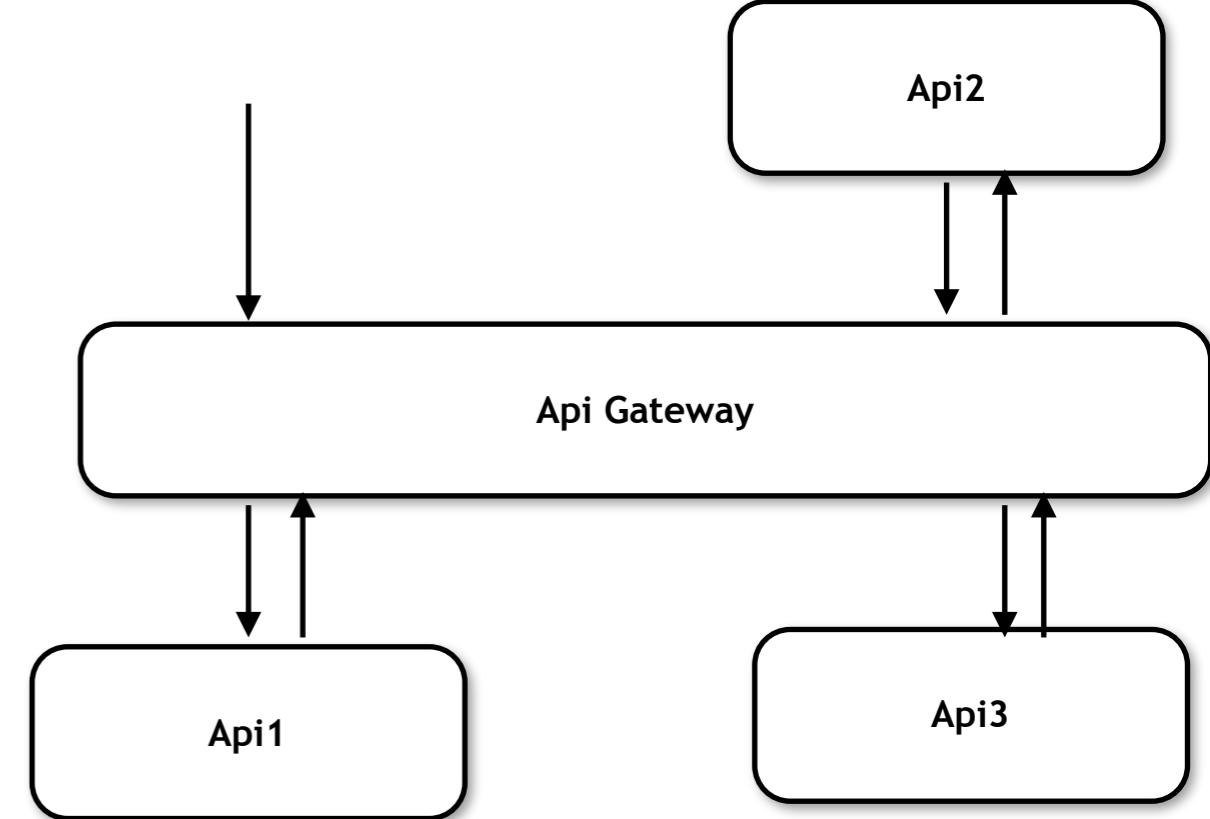
Given that it's difficult to find an ideal out-of-box implementation which can provide both the functions of an application-layer API gateway and an Istio ingress gateway, a practical solution could be using a cascade of an API Gateway and a mesh sidecar proxy as the external traffic entrance.

Deployment Diagram



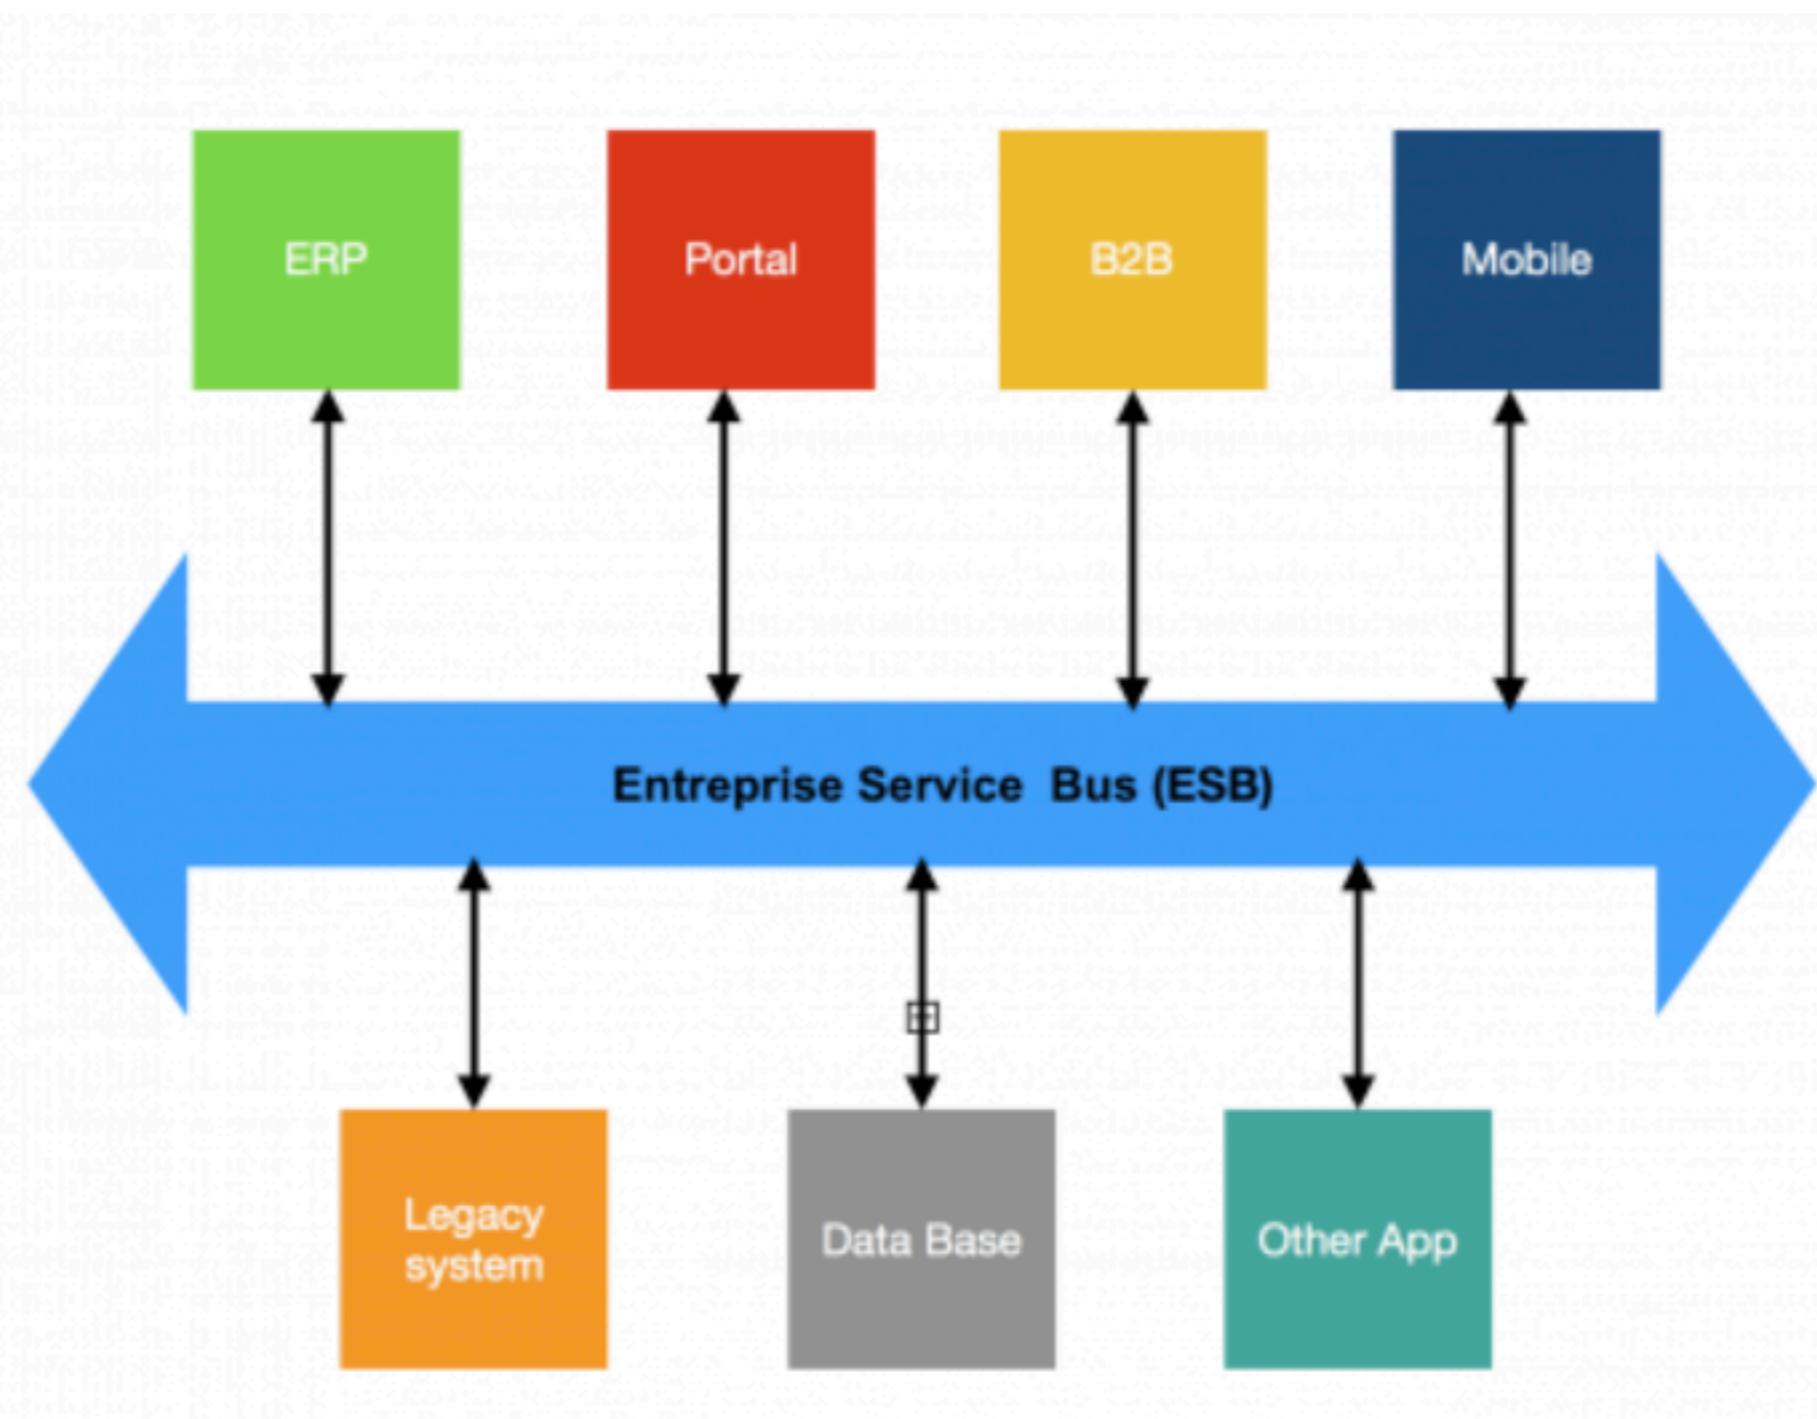


Facade



Bus

API Gateway vs ESB



Api Gateway vs Service Mesh

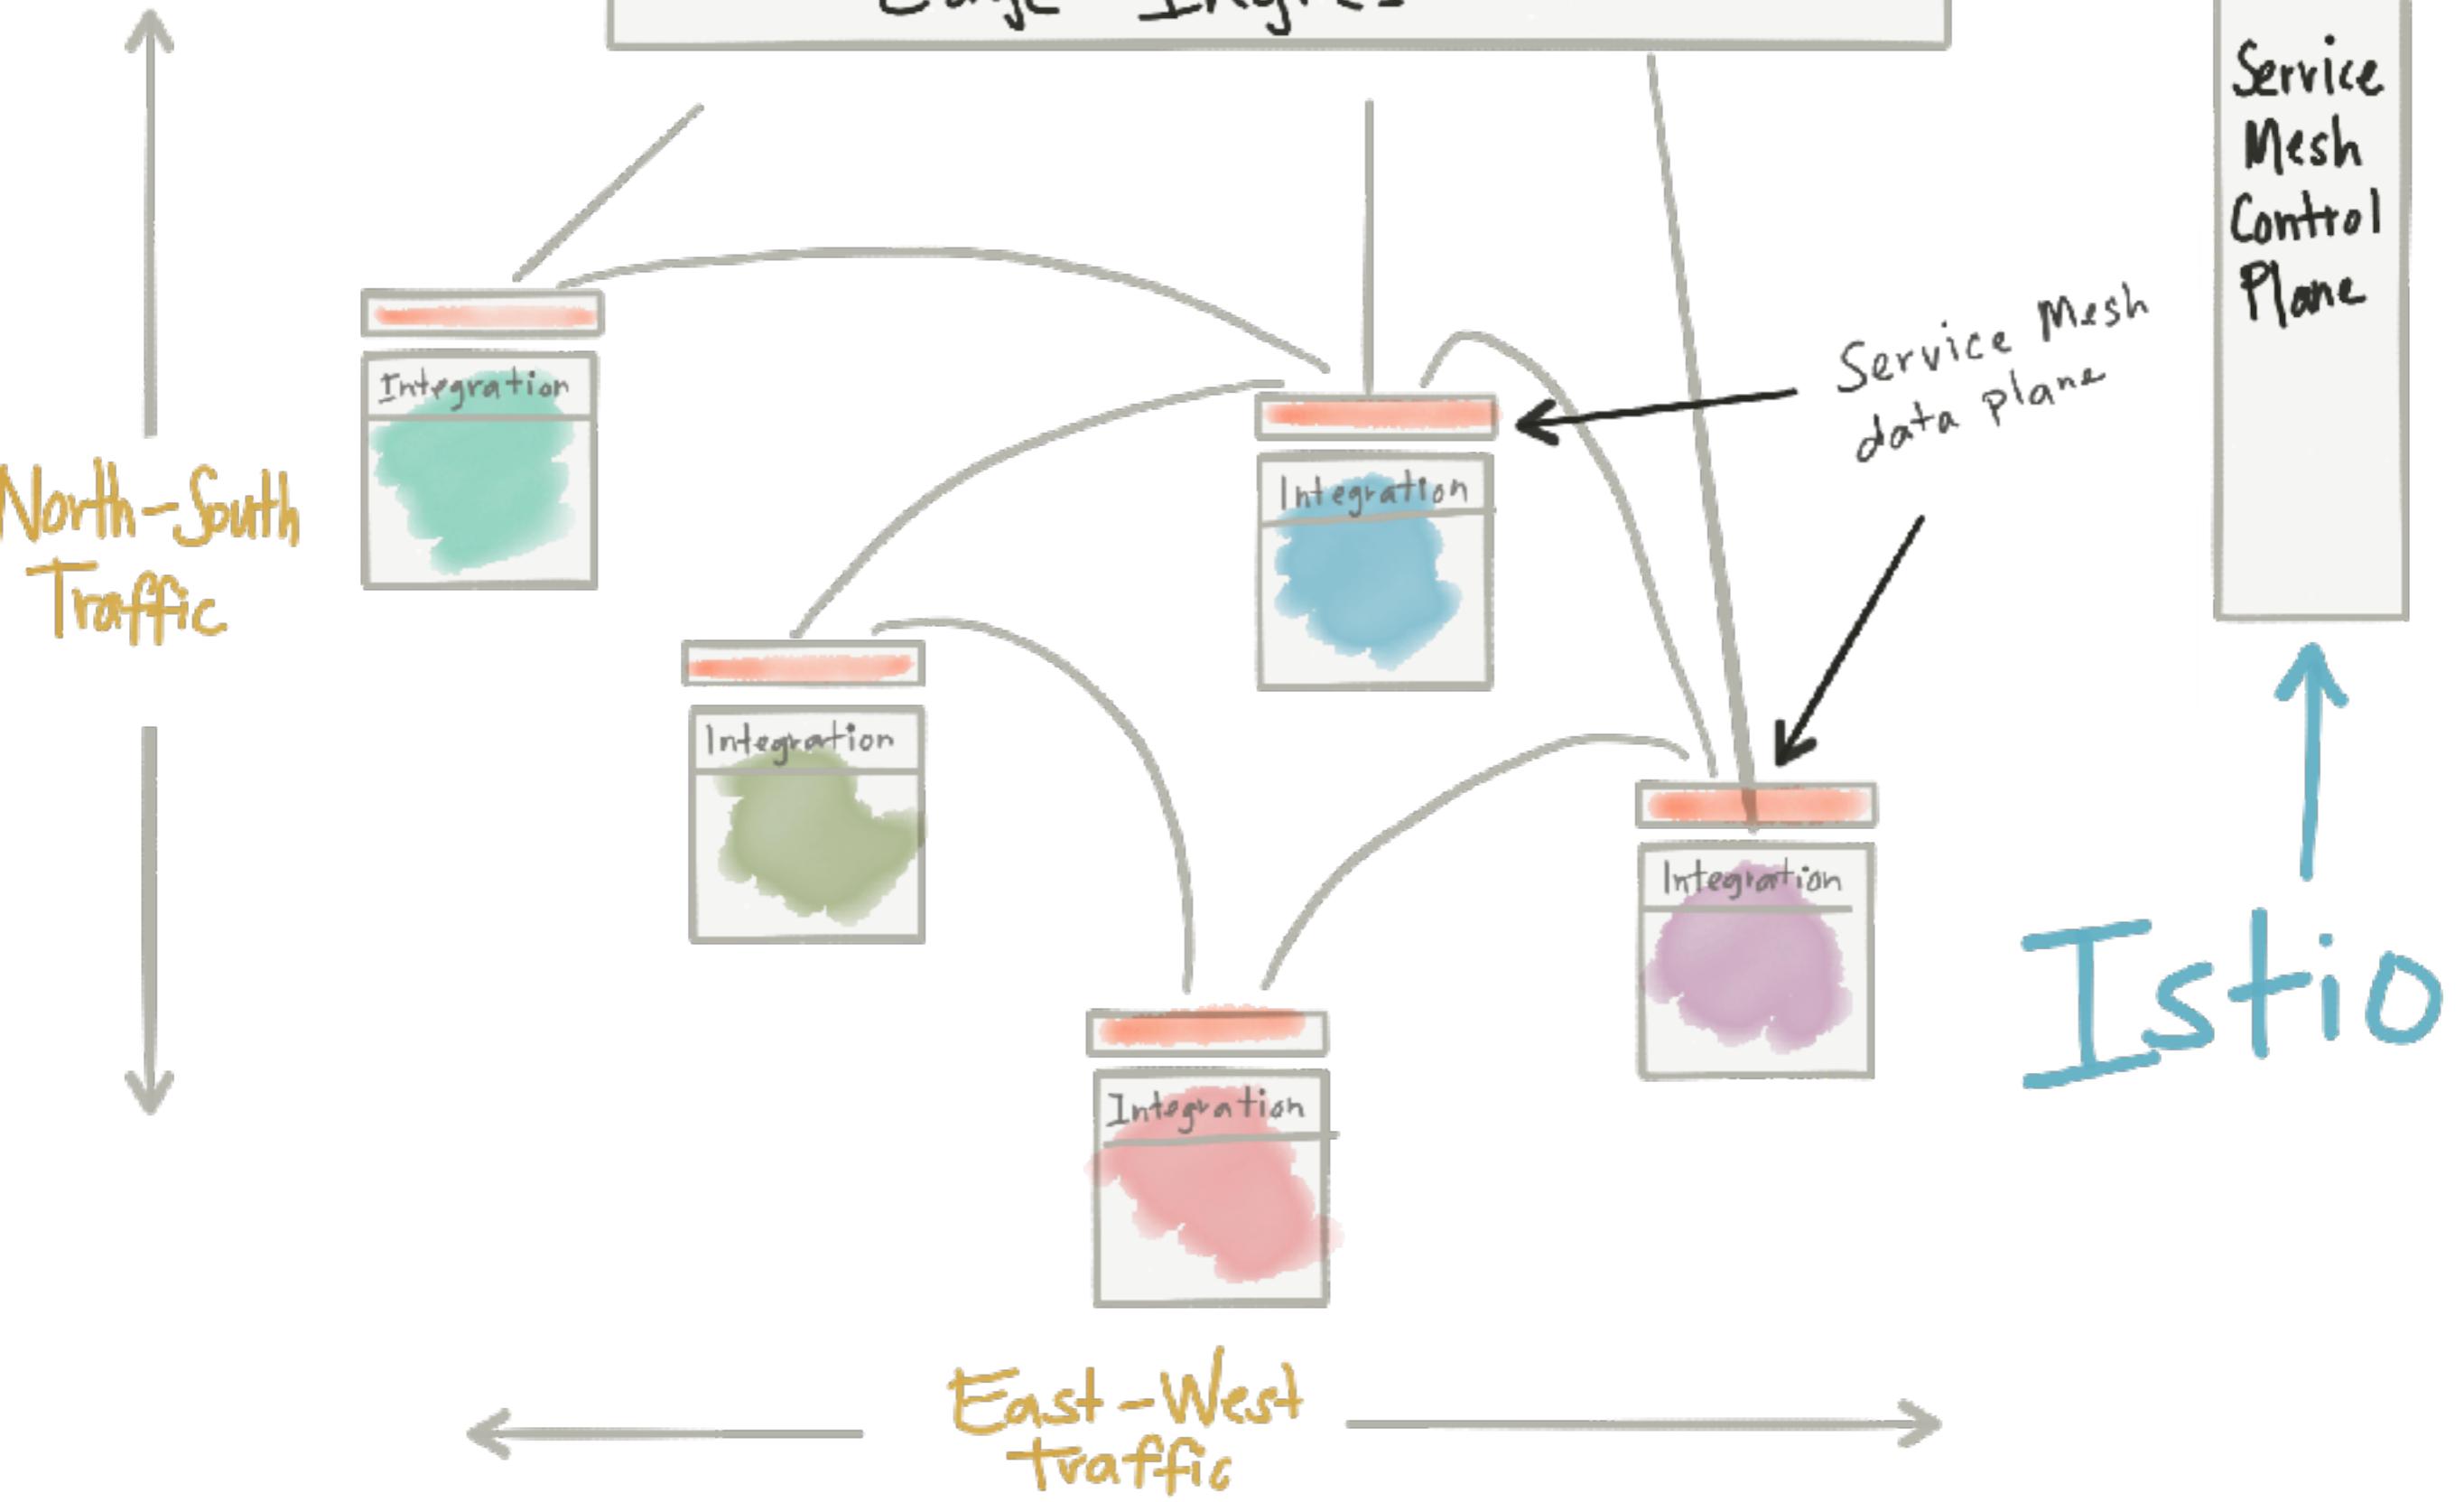
Load Balancer

API Gateway (Kubernetes Ingress, Istio Ingress, aws)

Service Mesh (Istio)

Deployment Platform (Kubernetes)

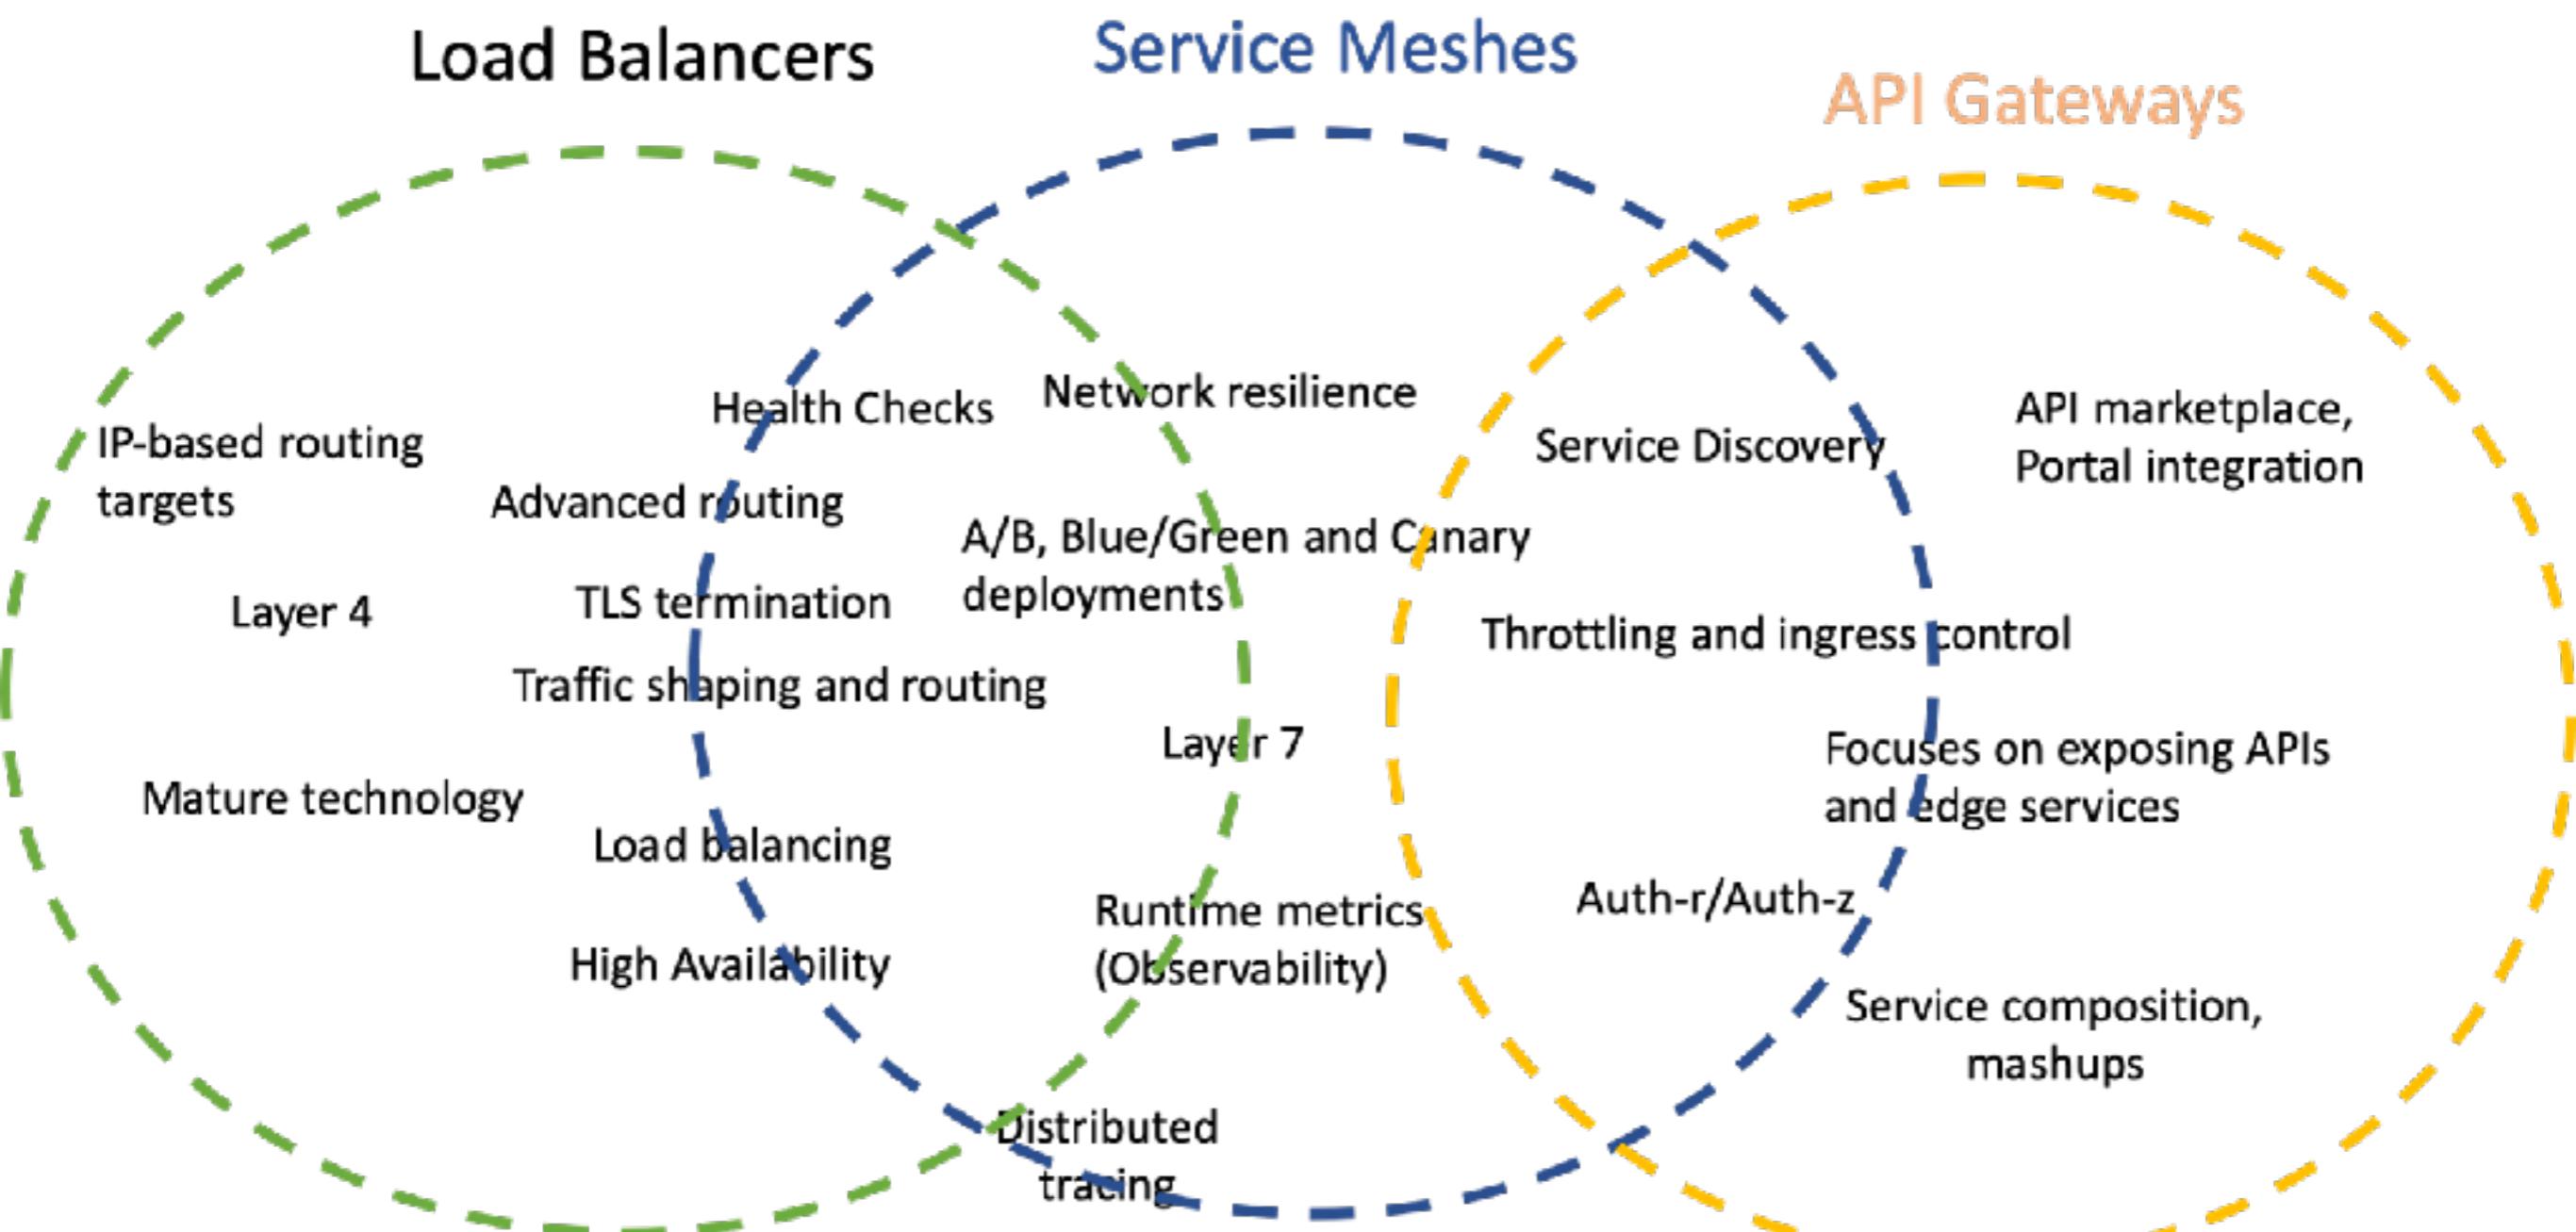
Edge Ingres



API Gateway : Abstraction,
decoupling,
edge routing / security

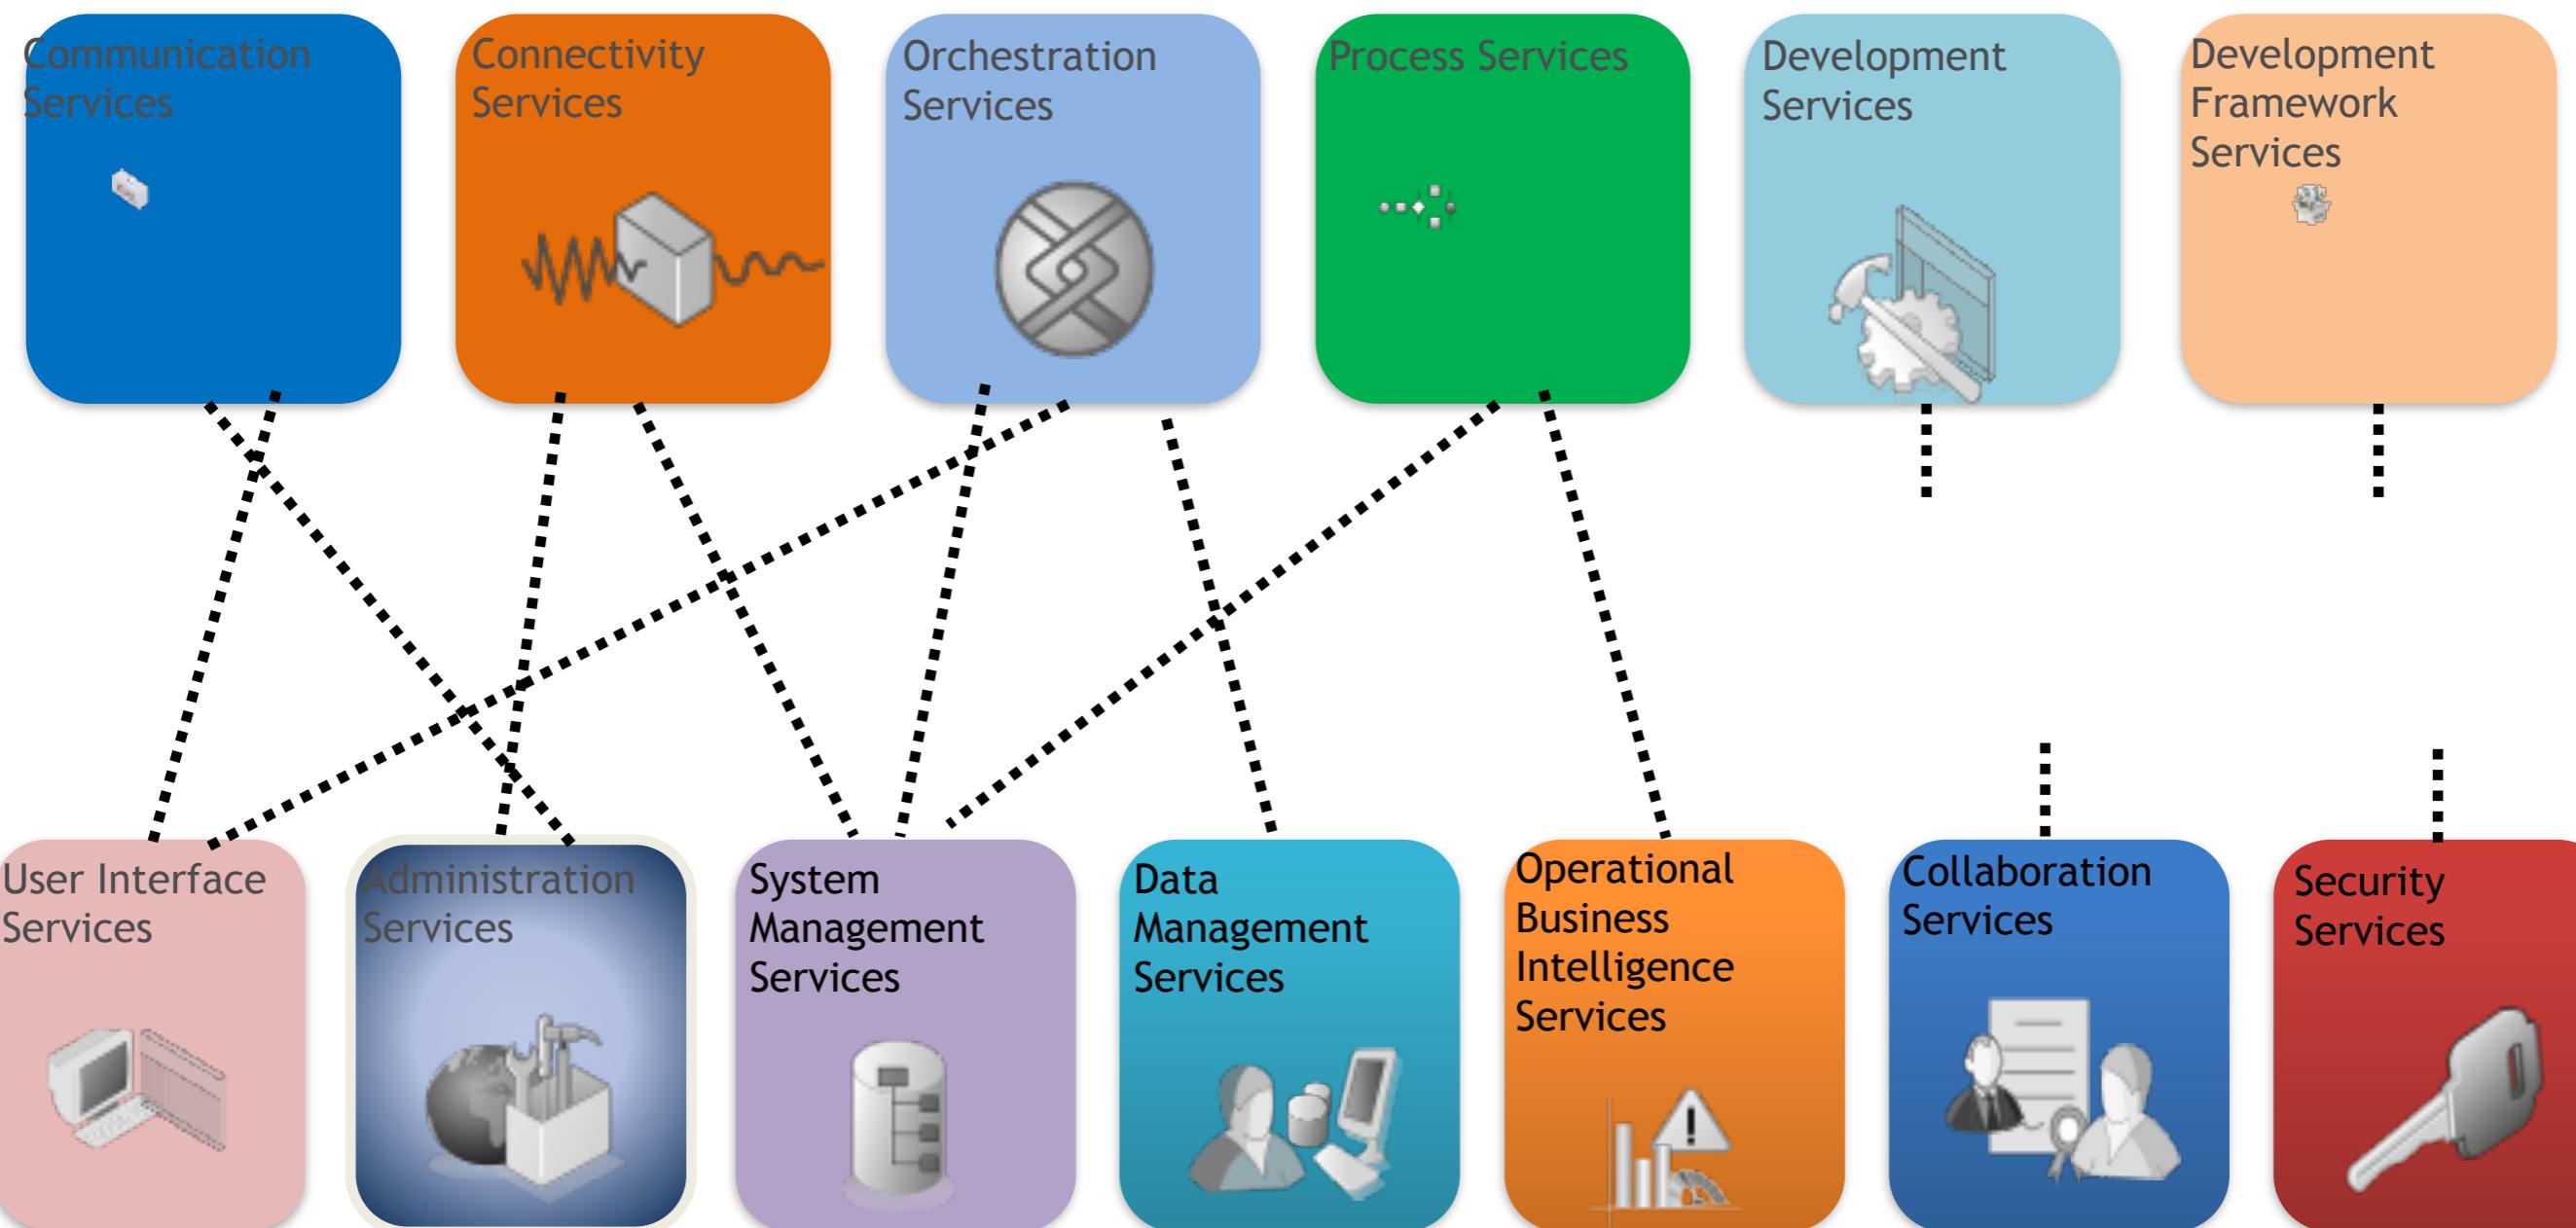
Service Mesh : endpoints, hosts, ports,
metric collection,
traffic routing, security

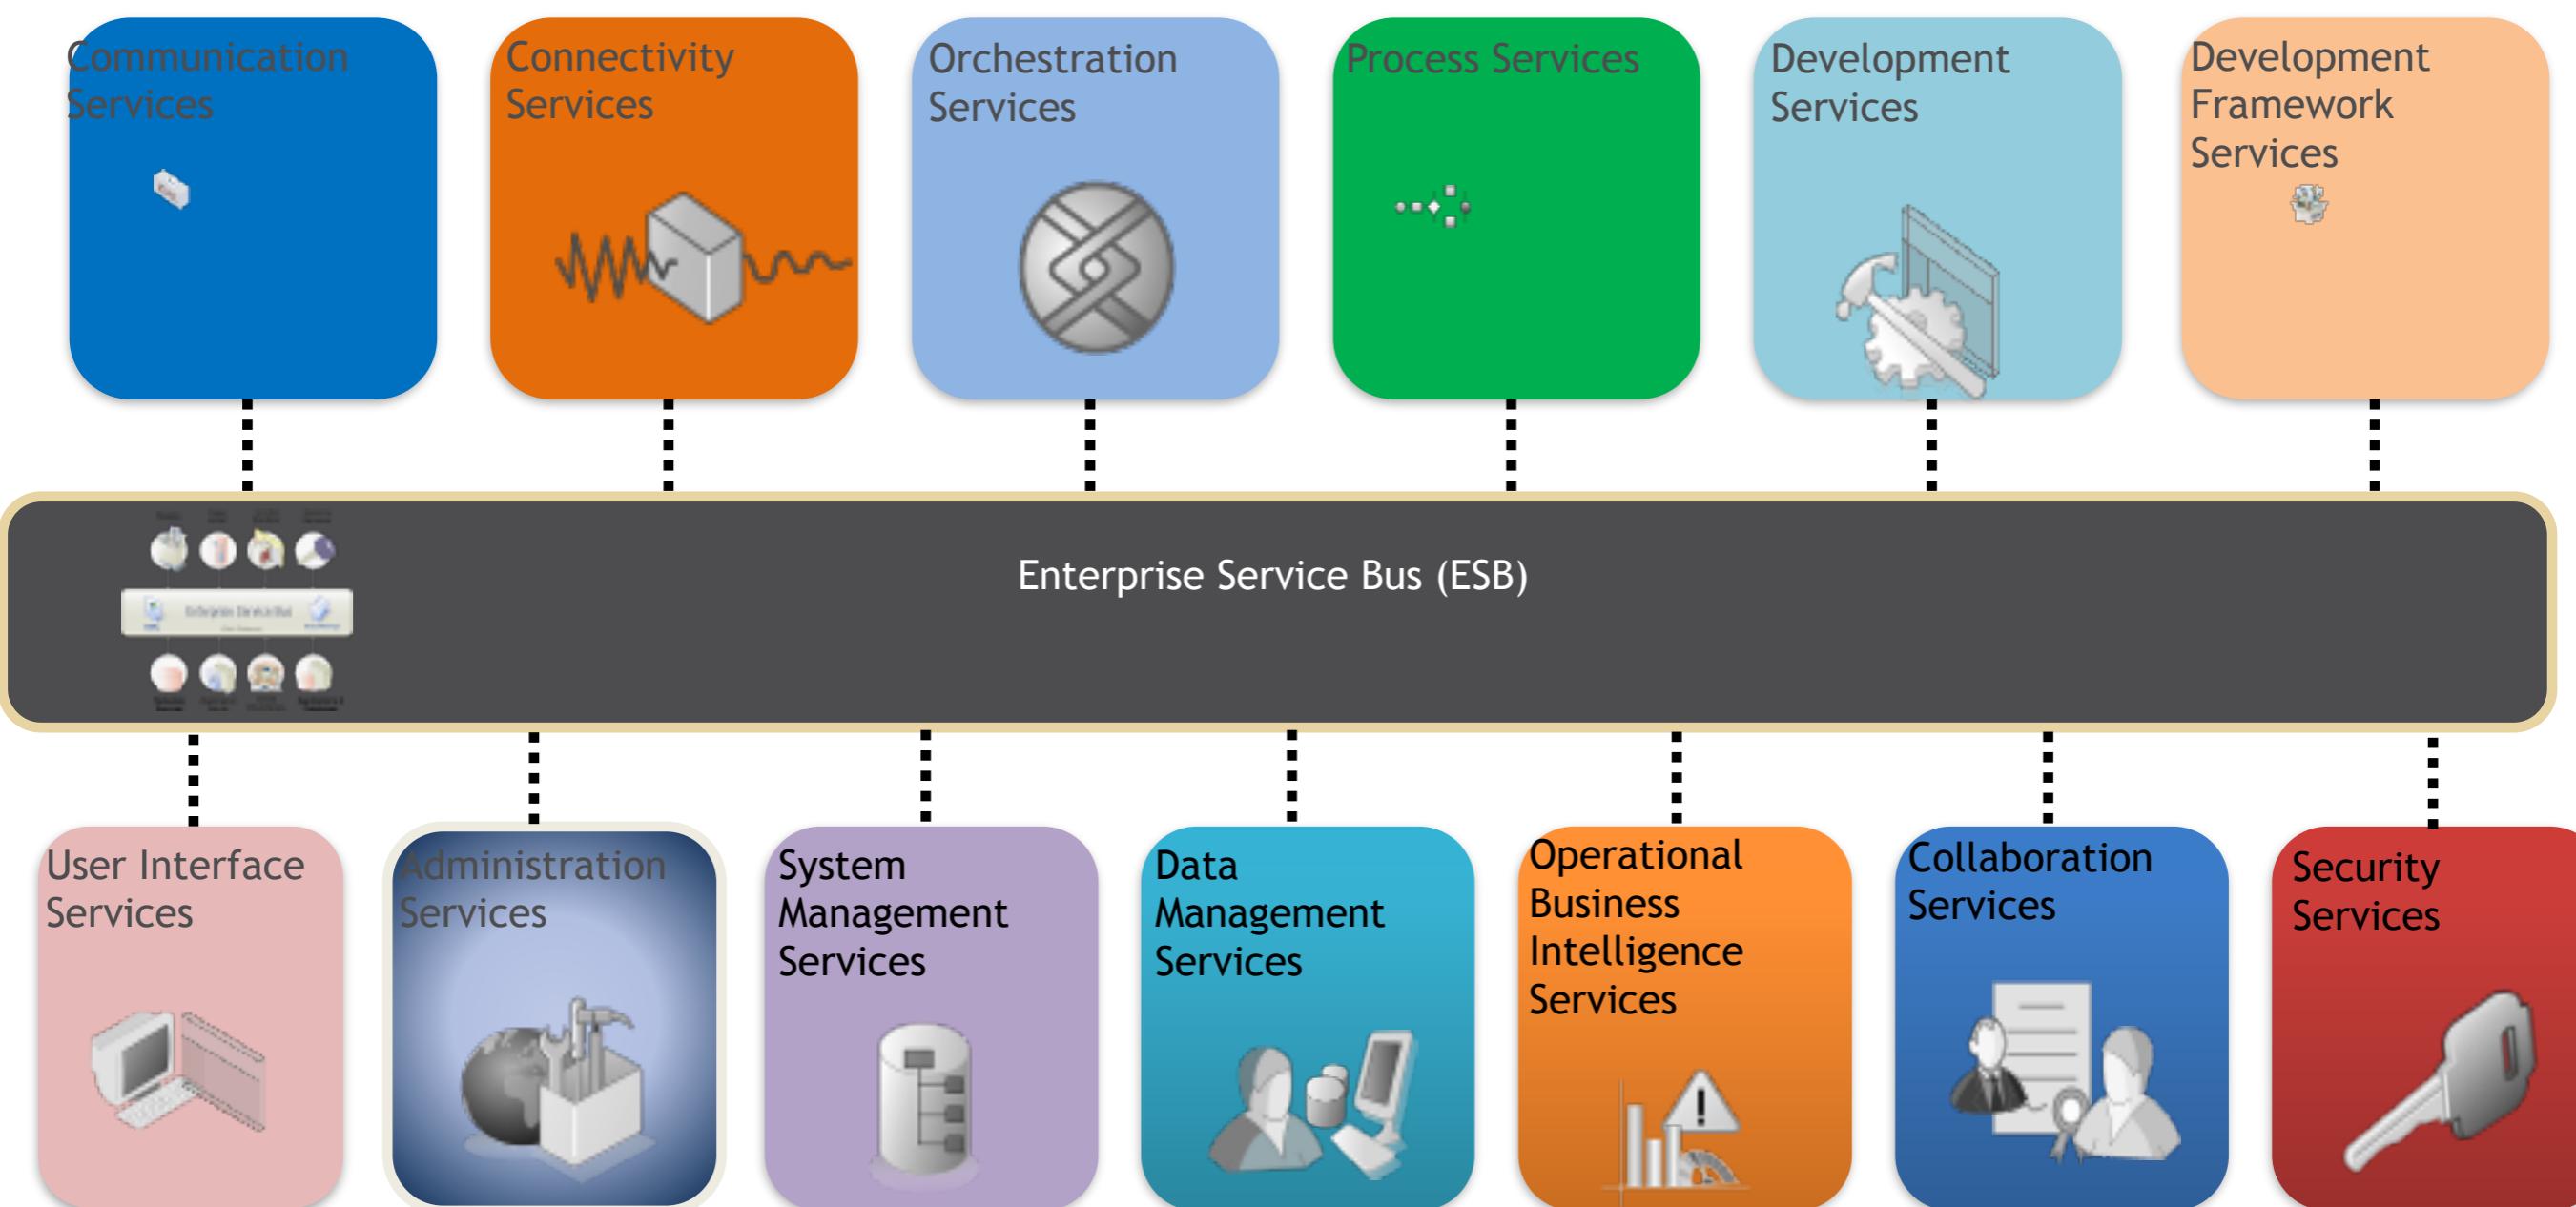
Deployment Platform : Cluster nodes,
container schedule,
resource management



	Approach
Performance	
Transaction Management	
Joins / Reports	
Infra Cost	Container, Kubernetes, API Gateway
Deployment	Docker Image, Hub
Debugging (Finding bugs)	ServiceMesh
Log Mgmt	
Config Mgmt	
Authentication	API Gateway
Authorization	
Monitoring / Alerting	Kubernetes , ServiceMesh

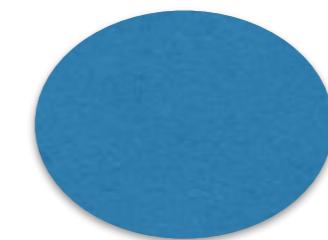
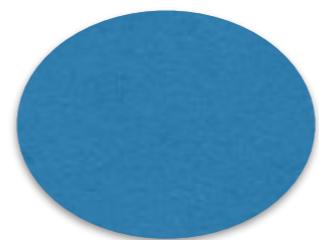
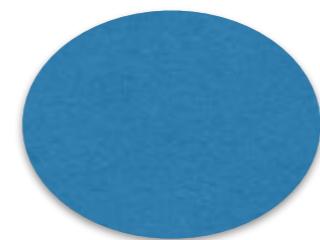
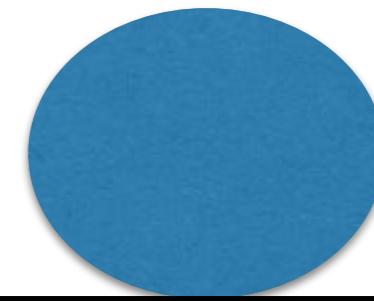
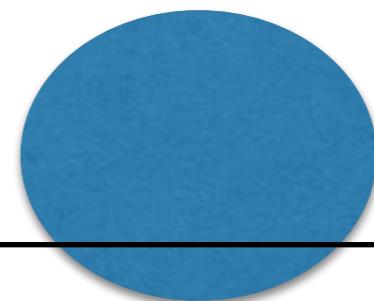
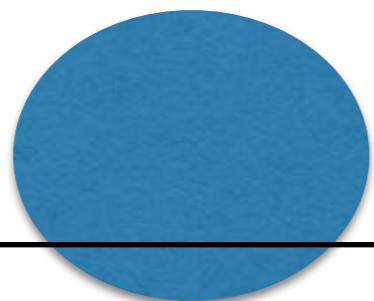
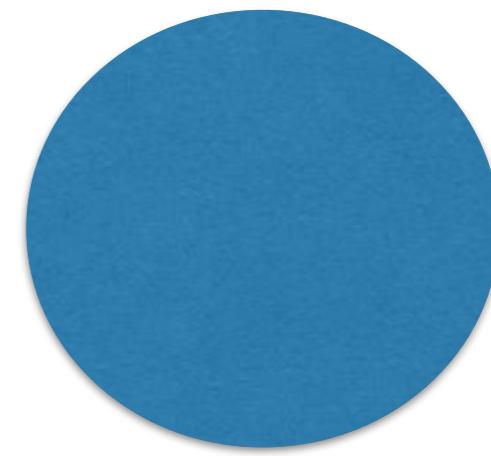
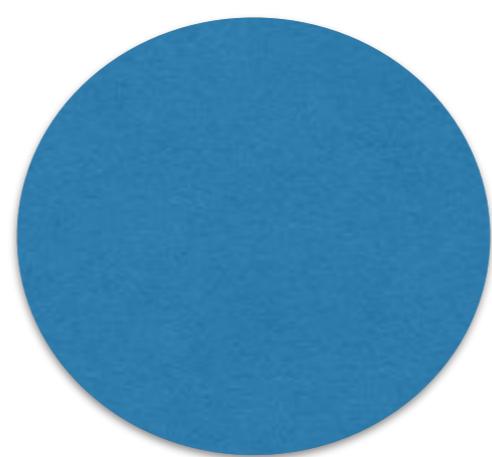
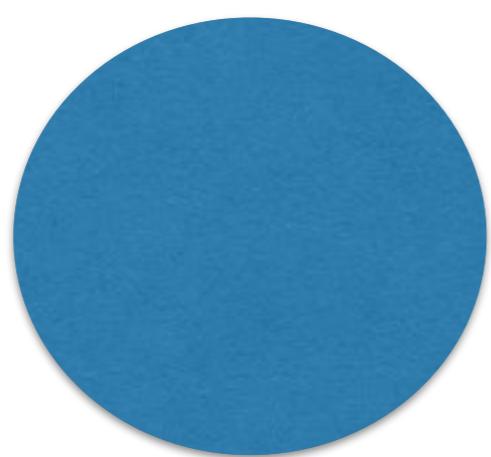
Event Driven Architecture





Case Study

Coarse

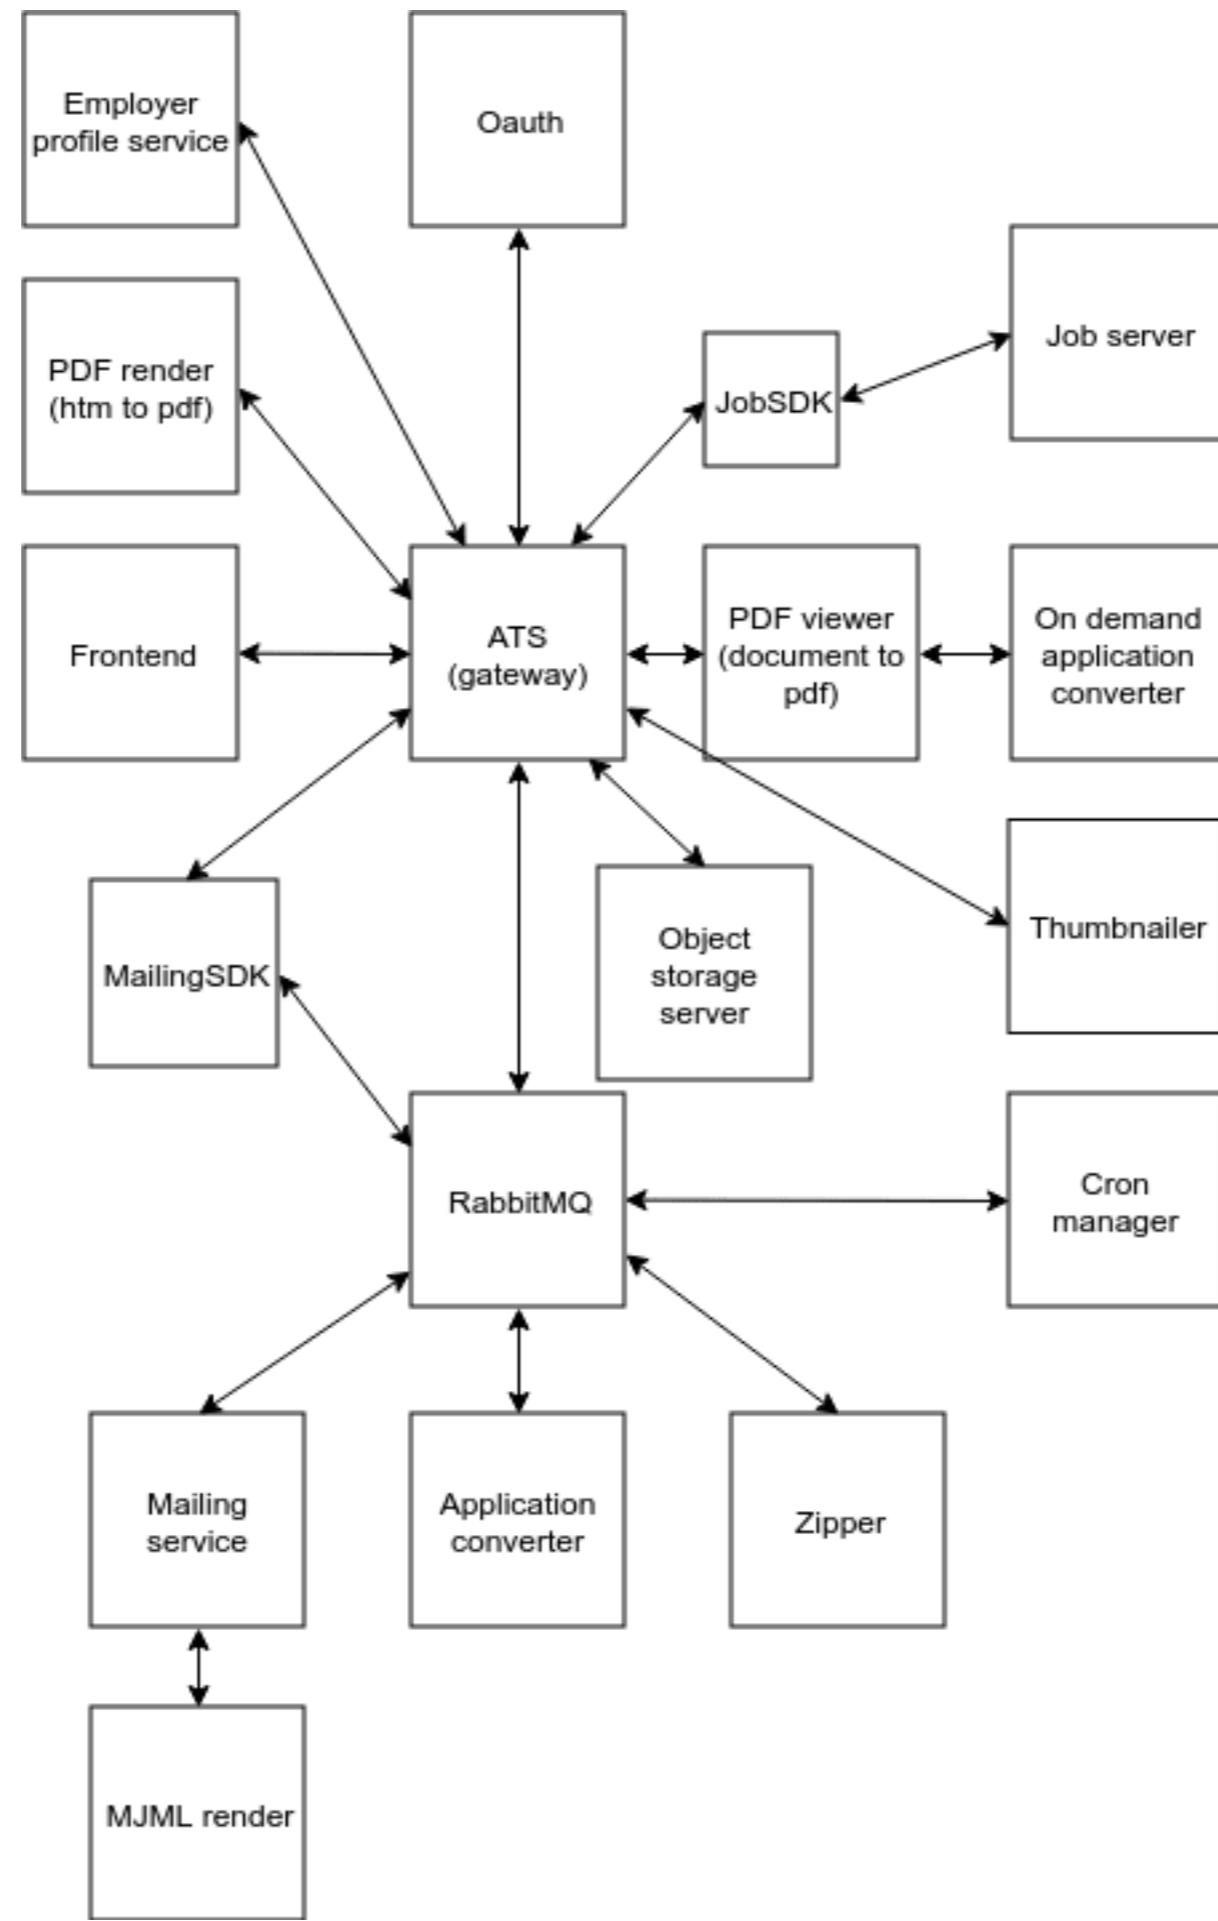


Fine

Application Tracking System (ATS)

1. After adding a new feature, it goes into the manual testing phase. Once it's greenlit, you deploy it to production — but what often happens is you miss a testing scenario, and you end up with lots of bugs in production.
2. This problem is tightly connected to the previous one: If you have an older codebase, its foundation starts to slowly break down, so adding new code on top of it ends up relying on other code. With that kind of codebase, sometimes you end up breaking code that should be totally irrelevant to what you're currently writing.
3. When you have 11 people working on the same codebase day in, day out, without any safeties in place whatsoever, problems are bound to arise.

- One Year
- ~50,000 lines of code
- ~2,600 commits
- a team of 3 developers
- consisting of a senior full stack developer and two medior backend developers
- 18 services
- 200 tests



For example, in a hospital domain, a **Patient** being treated in the outpatients department might have a list of **Referrals**, and methods such as *BookAppointment()*. A **Patient** being treated as an Inpatient however, will have a **Ward** property and methods such as *TransferToTheatre()*. Given this, there are two bounded contexts that patients exist in: Outpatients & Inpatients.

In the insurance domain, the sales team put together a **Policy** that has a degree of risk associated to it and therefore cost. But if it reaches the claims department, that information is meaningless to them. They only need to verify whether the policy is valid for the claim. So there are two contexts here: Sales & Claims

the buyer entity might have most of a person's attributes that are defined in the user entity in the profile or identity microservice, including the identity. But the buyer entity in the ordering microservice might have fewer attributes, because only certain buyer data is related to the order process. The context of each microservice or Bounded Context impacts its domain model.

Design Patterns for Microservices

Decomposition Patterns

Integration Patterns

Database Patterns

Observability Patterns

Cross-Cutting Concern Patterns

Decompose by Business Capability

Decompose by Subdomain

Decompose by Transactions

Strangler Pattern

Bulkhead Pattern

Sidecar Pattern

API Gateway Pattern

Aggregator Pattern

Proxy Pattern

Gateway Routing Pattern

Chained Microservice Pattern

Branch Pattern

Client-Side UI Composition Pattern

Database per Service

Shared Database per Service

CQRS

Event Sourcing

Saga Pattern

Log Aggregation

Performance Metrics

Distributed Tracing

Health Check

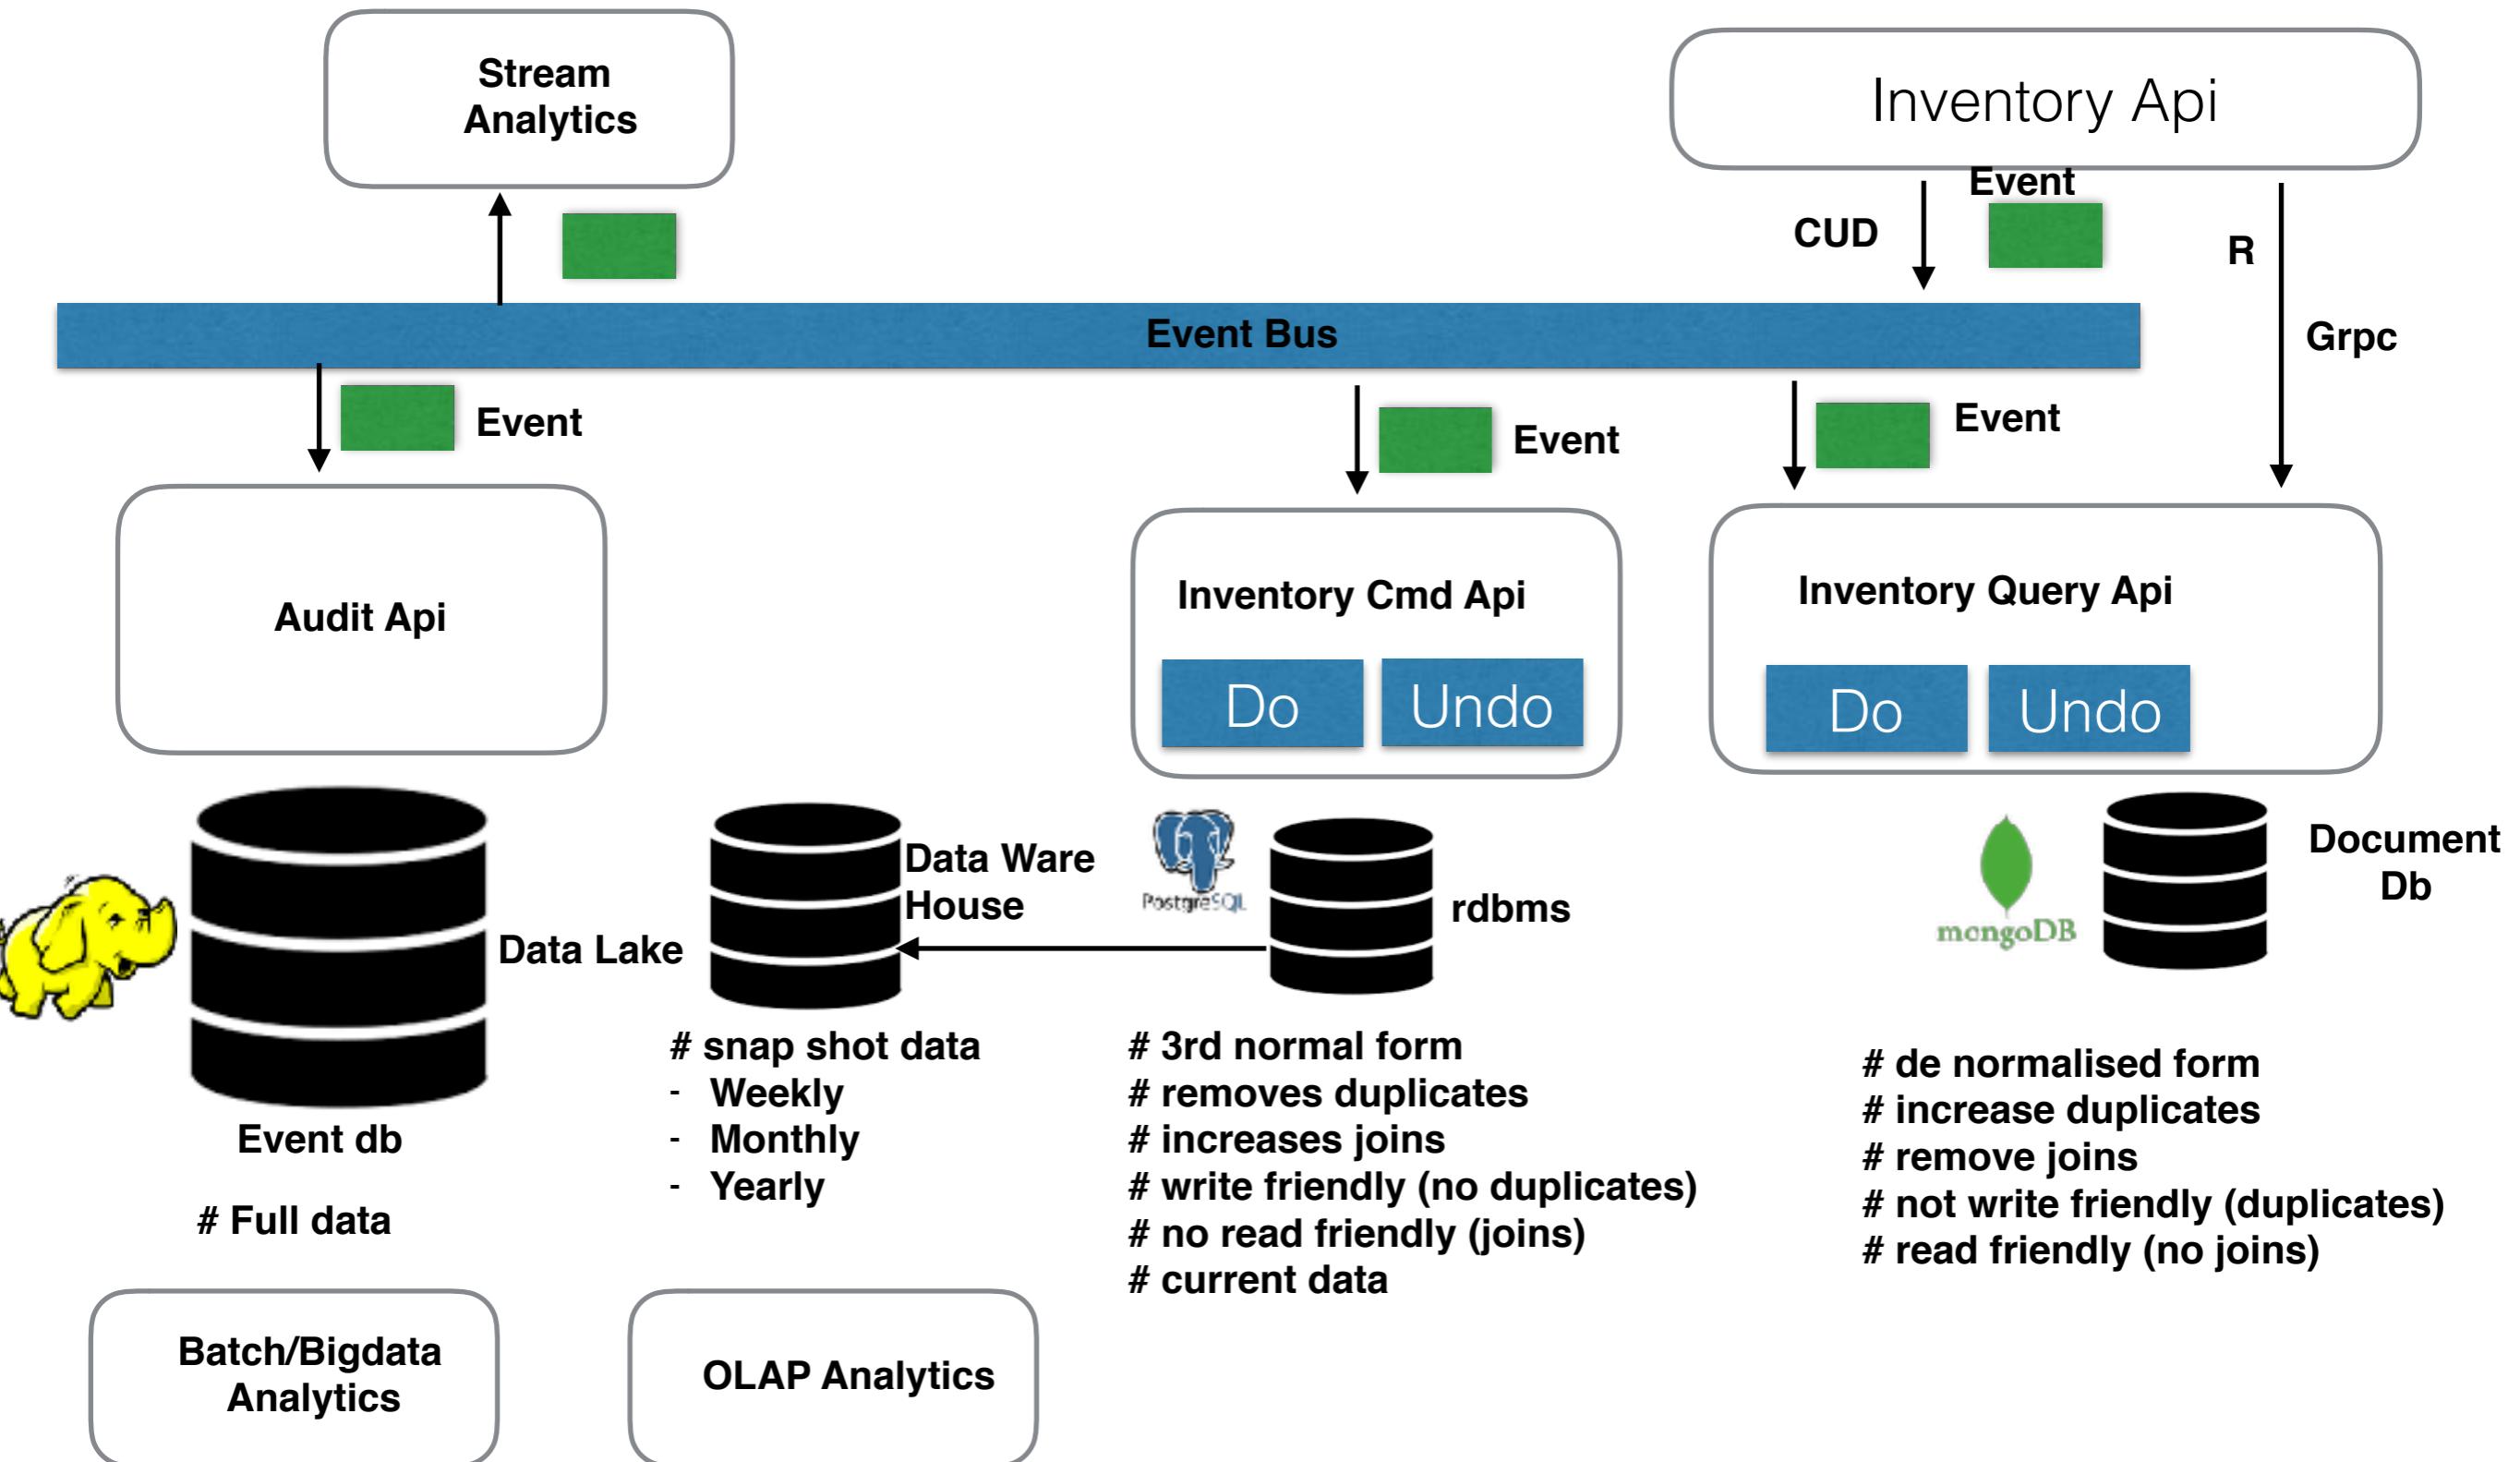
External Configuration

Service Discovery Pattern

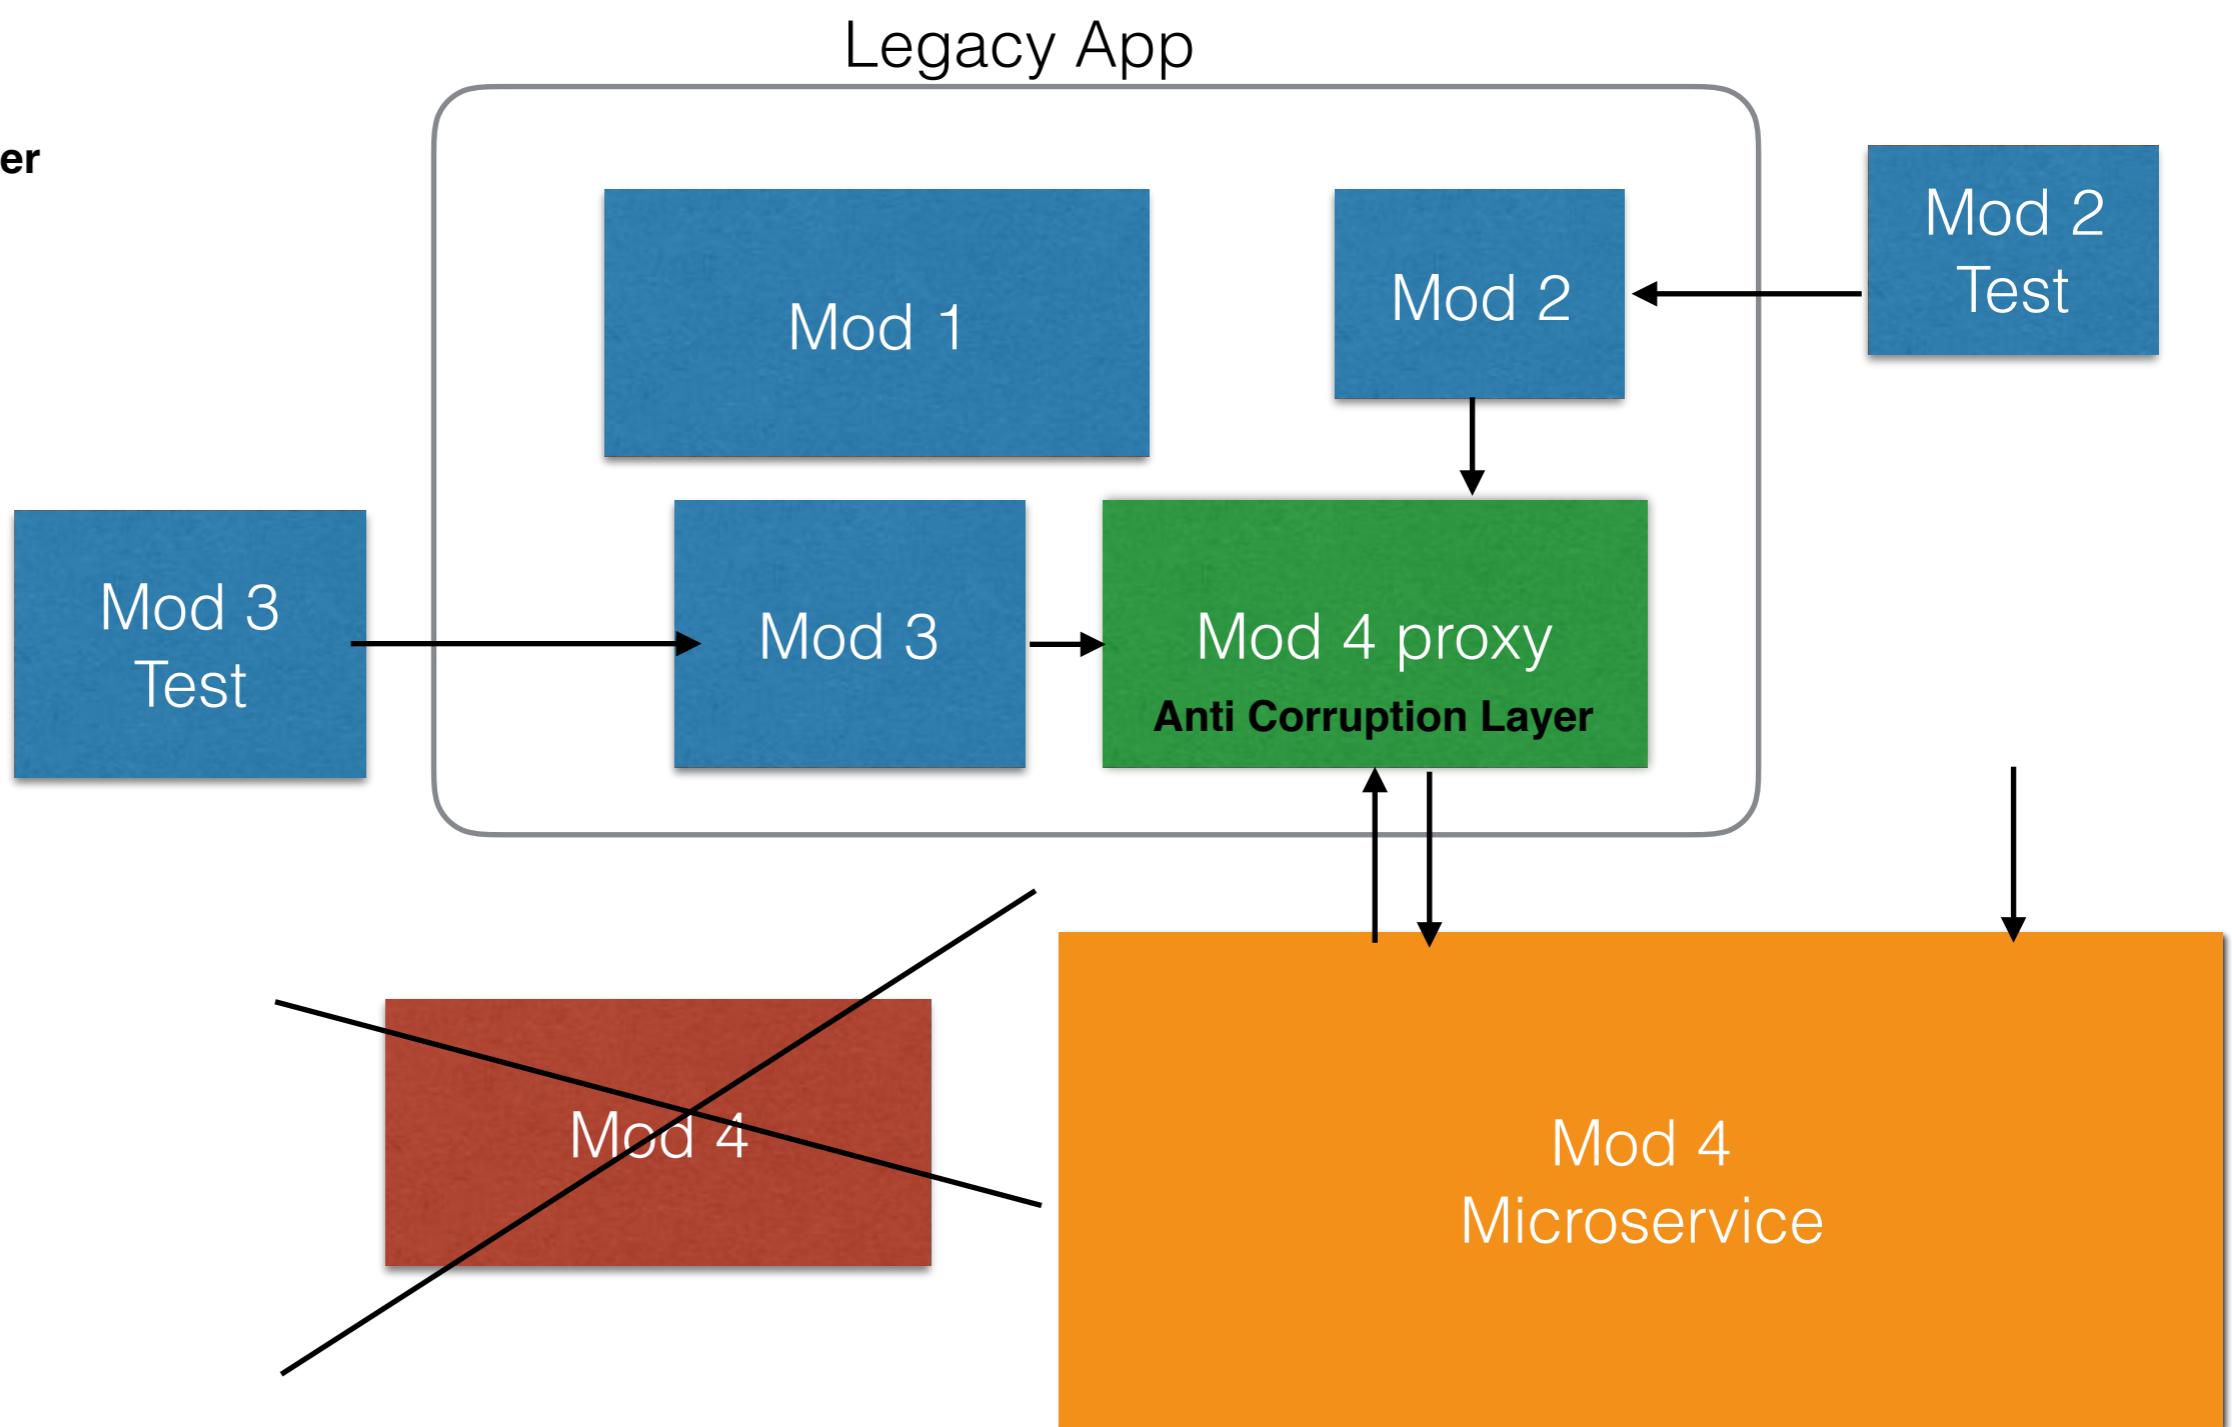
Circuit Breaker Pattern

Blue-Green Deployment Pattern

```
# Event Bus  
# EDA  
# compensatale transaction  
# SAGA  
# Event Sourcing  
# CQRS
```



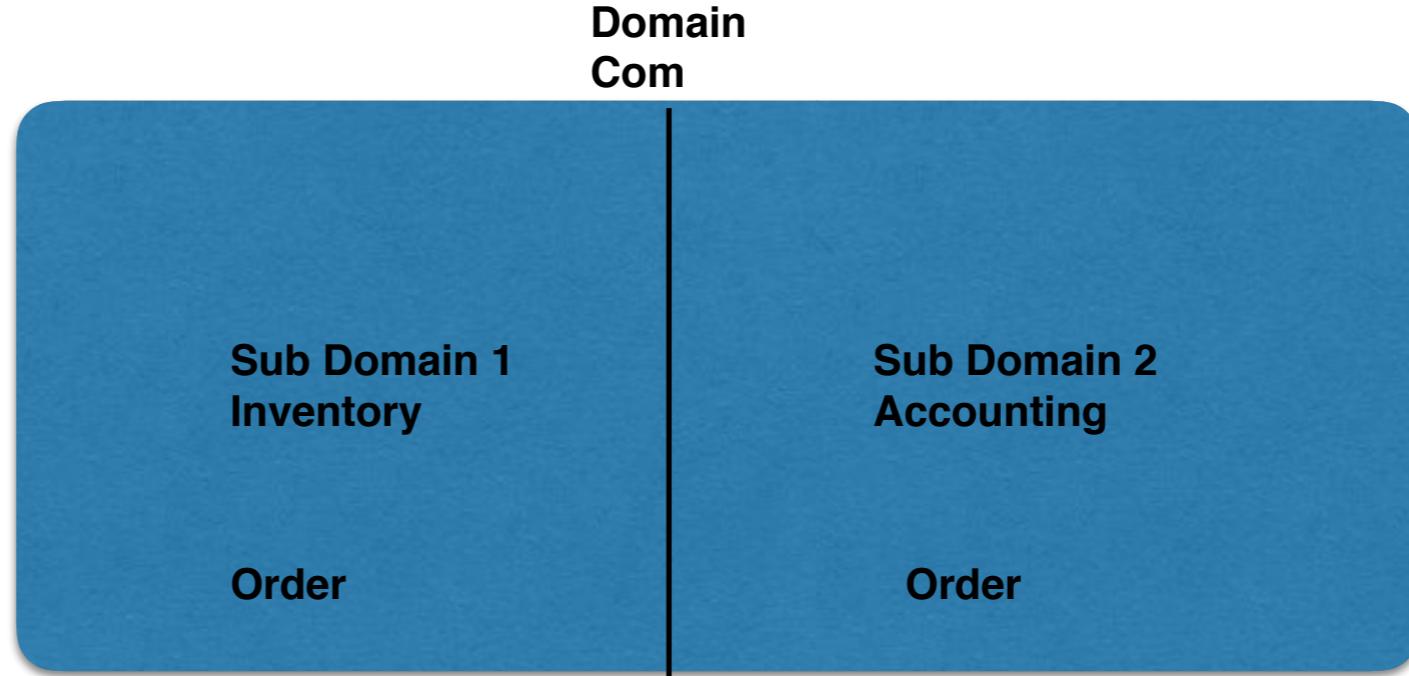
unit test
Strangler Pattern
Anti Corruption Layer



*Transaction
Boundary



Bounded
Context (DDD)



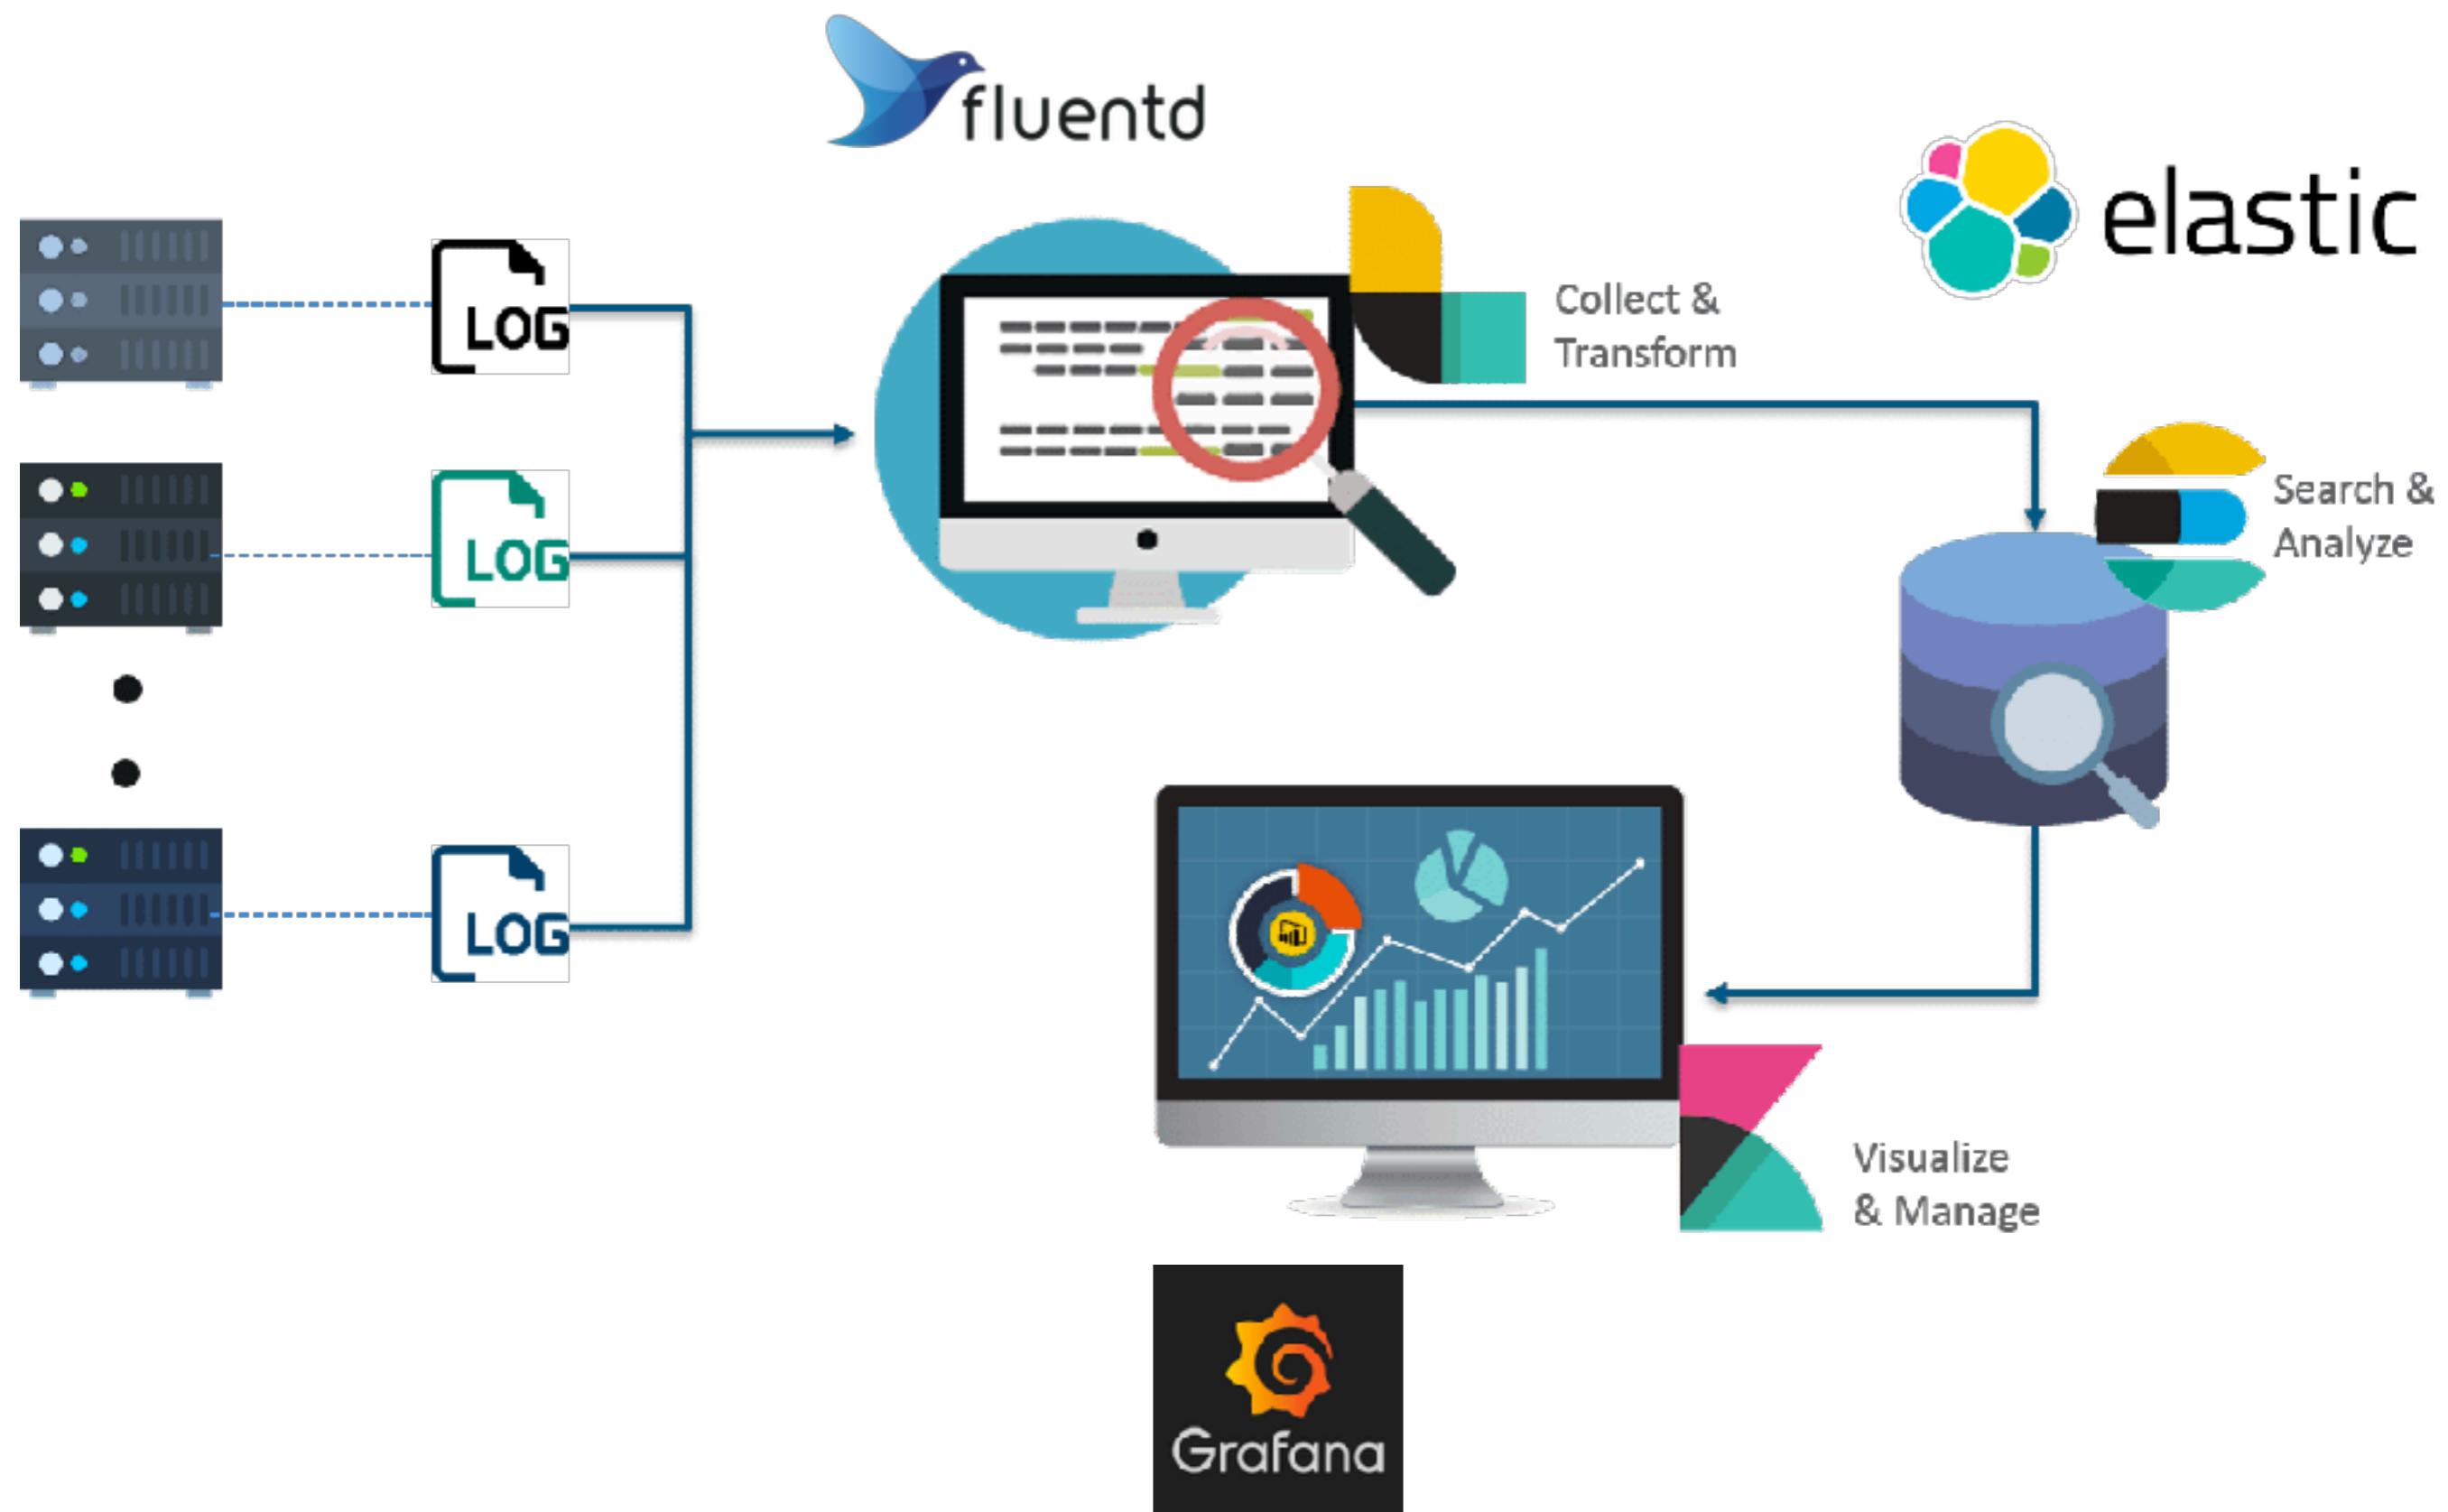
Aggregate
Root (DDD)

**Invoice -> Line Item
-> Address**

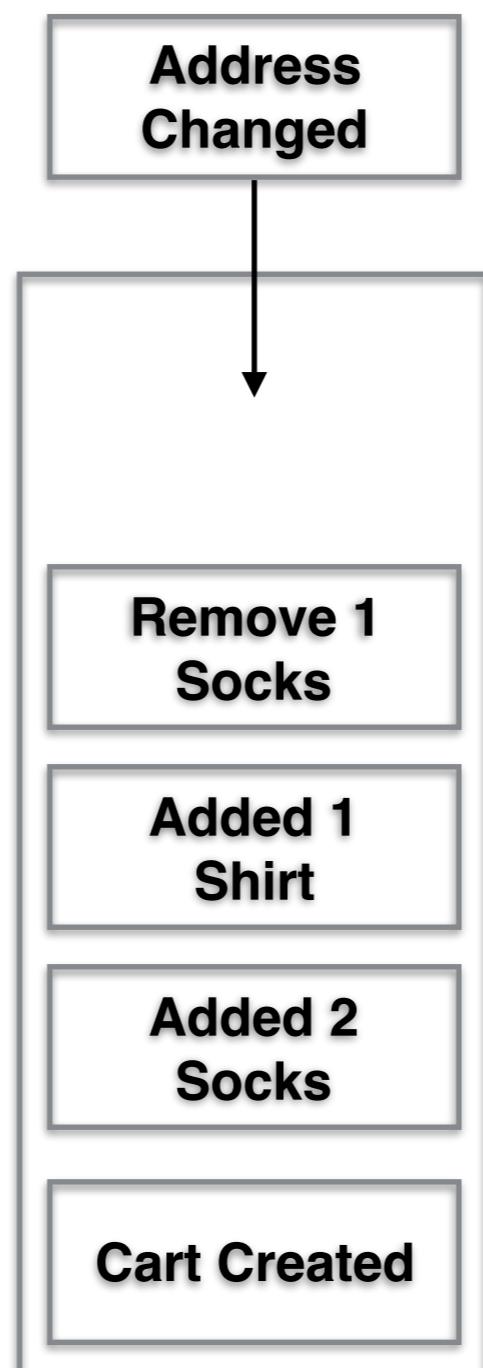
Feature

Patterns

Log Aggregation

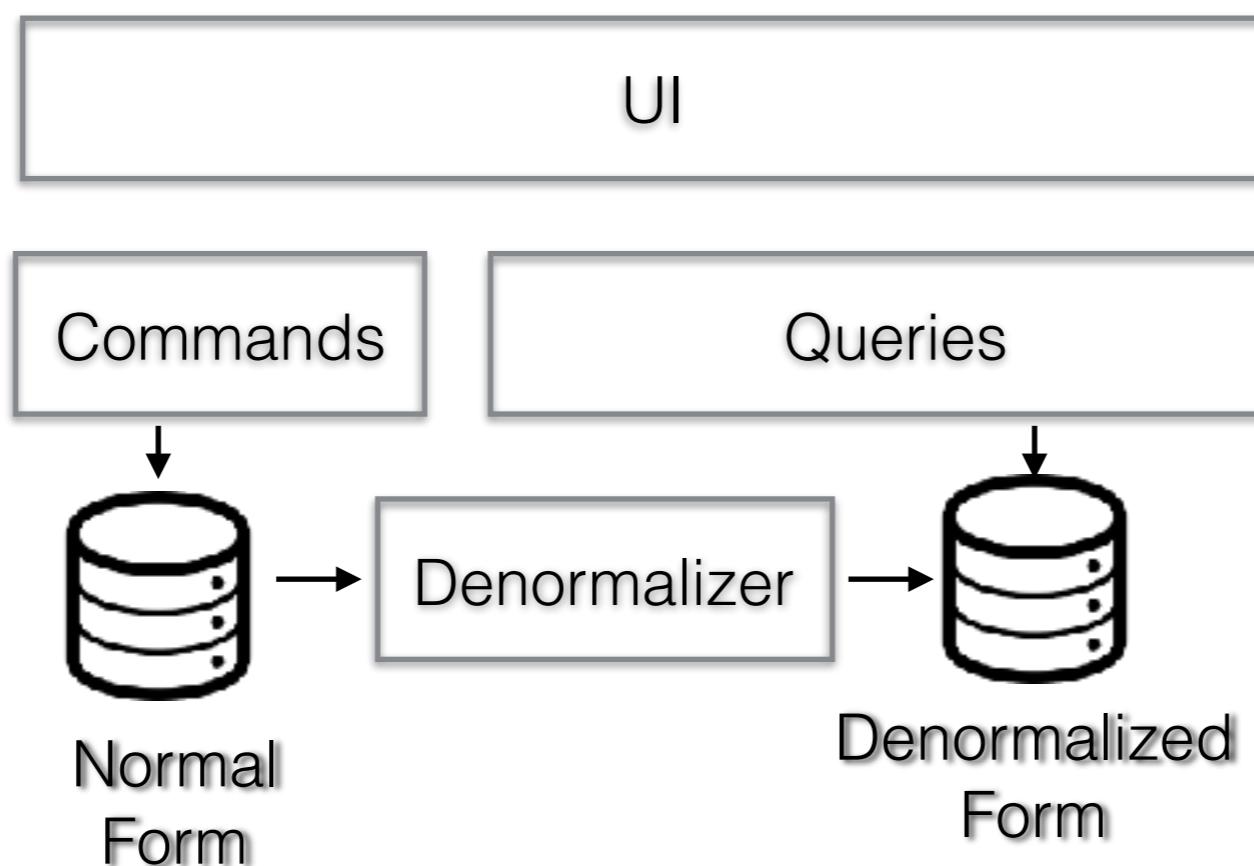


Event Source

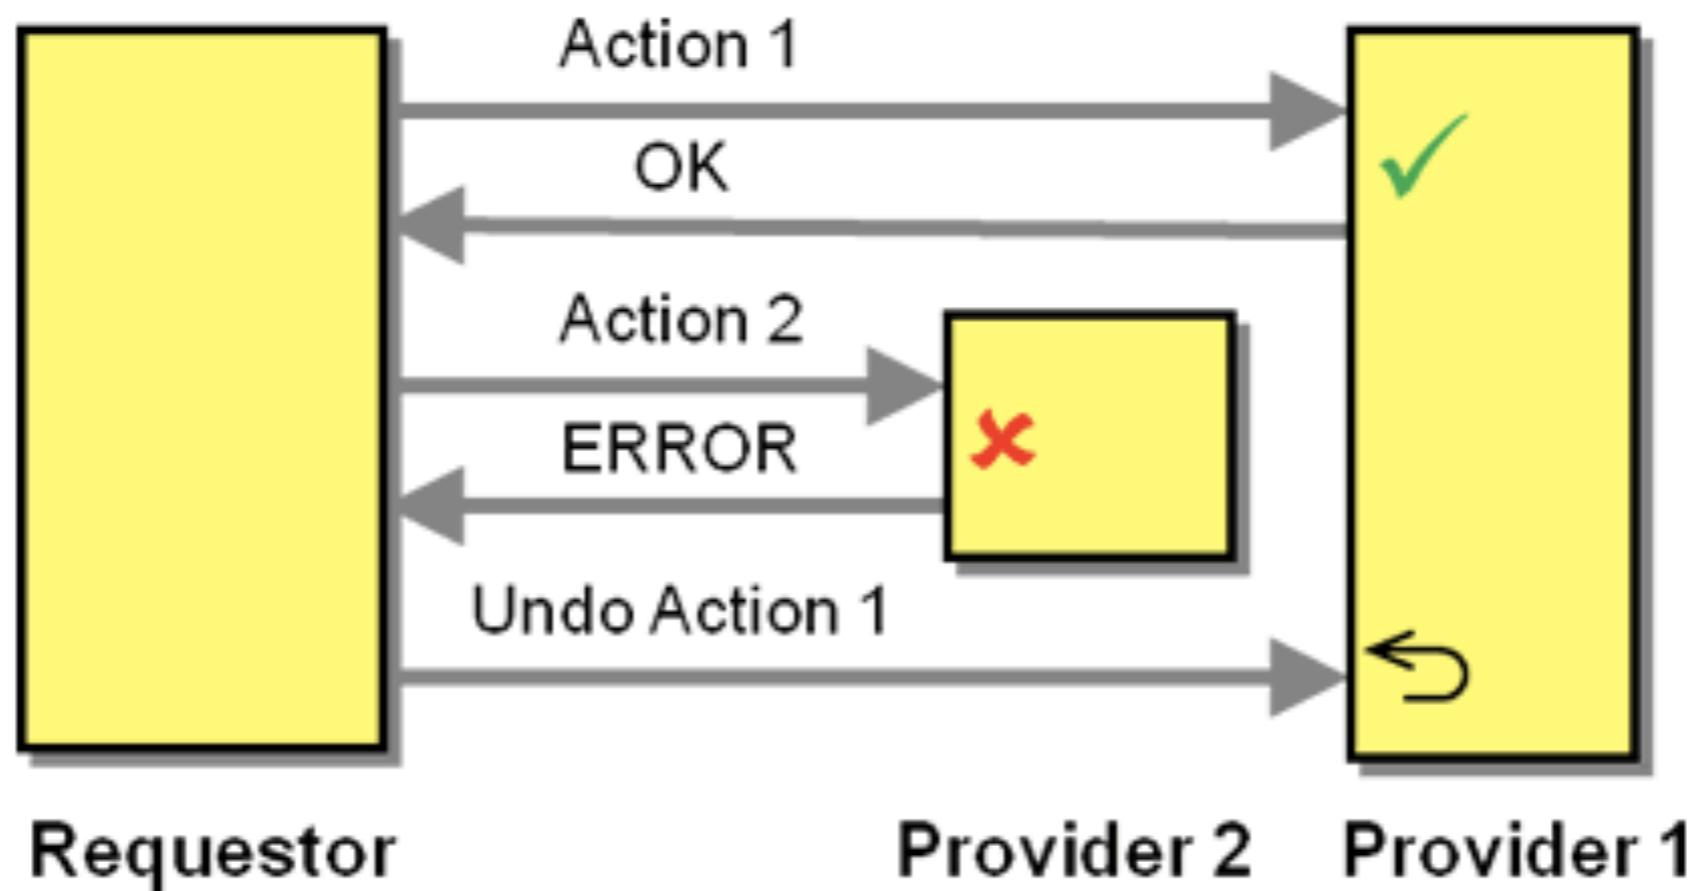


Event Store

CQRS

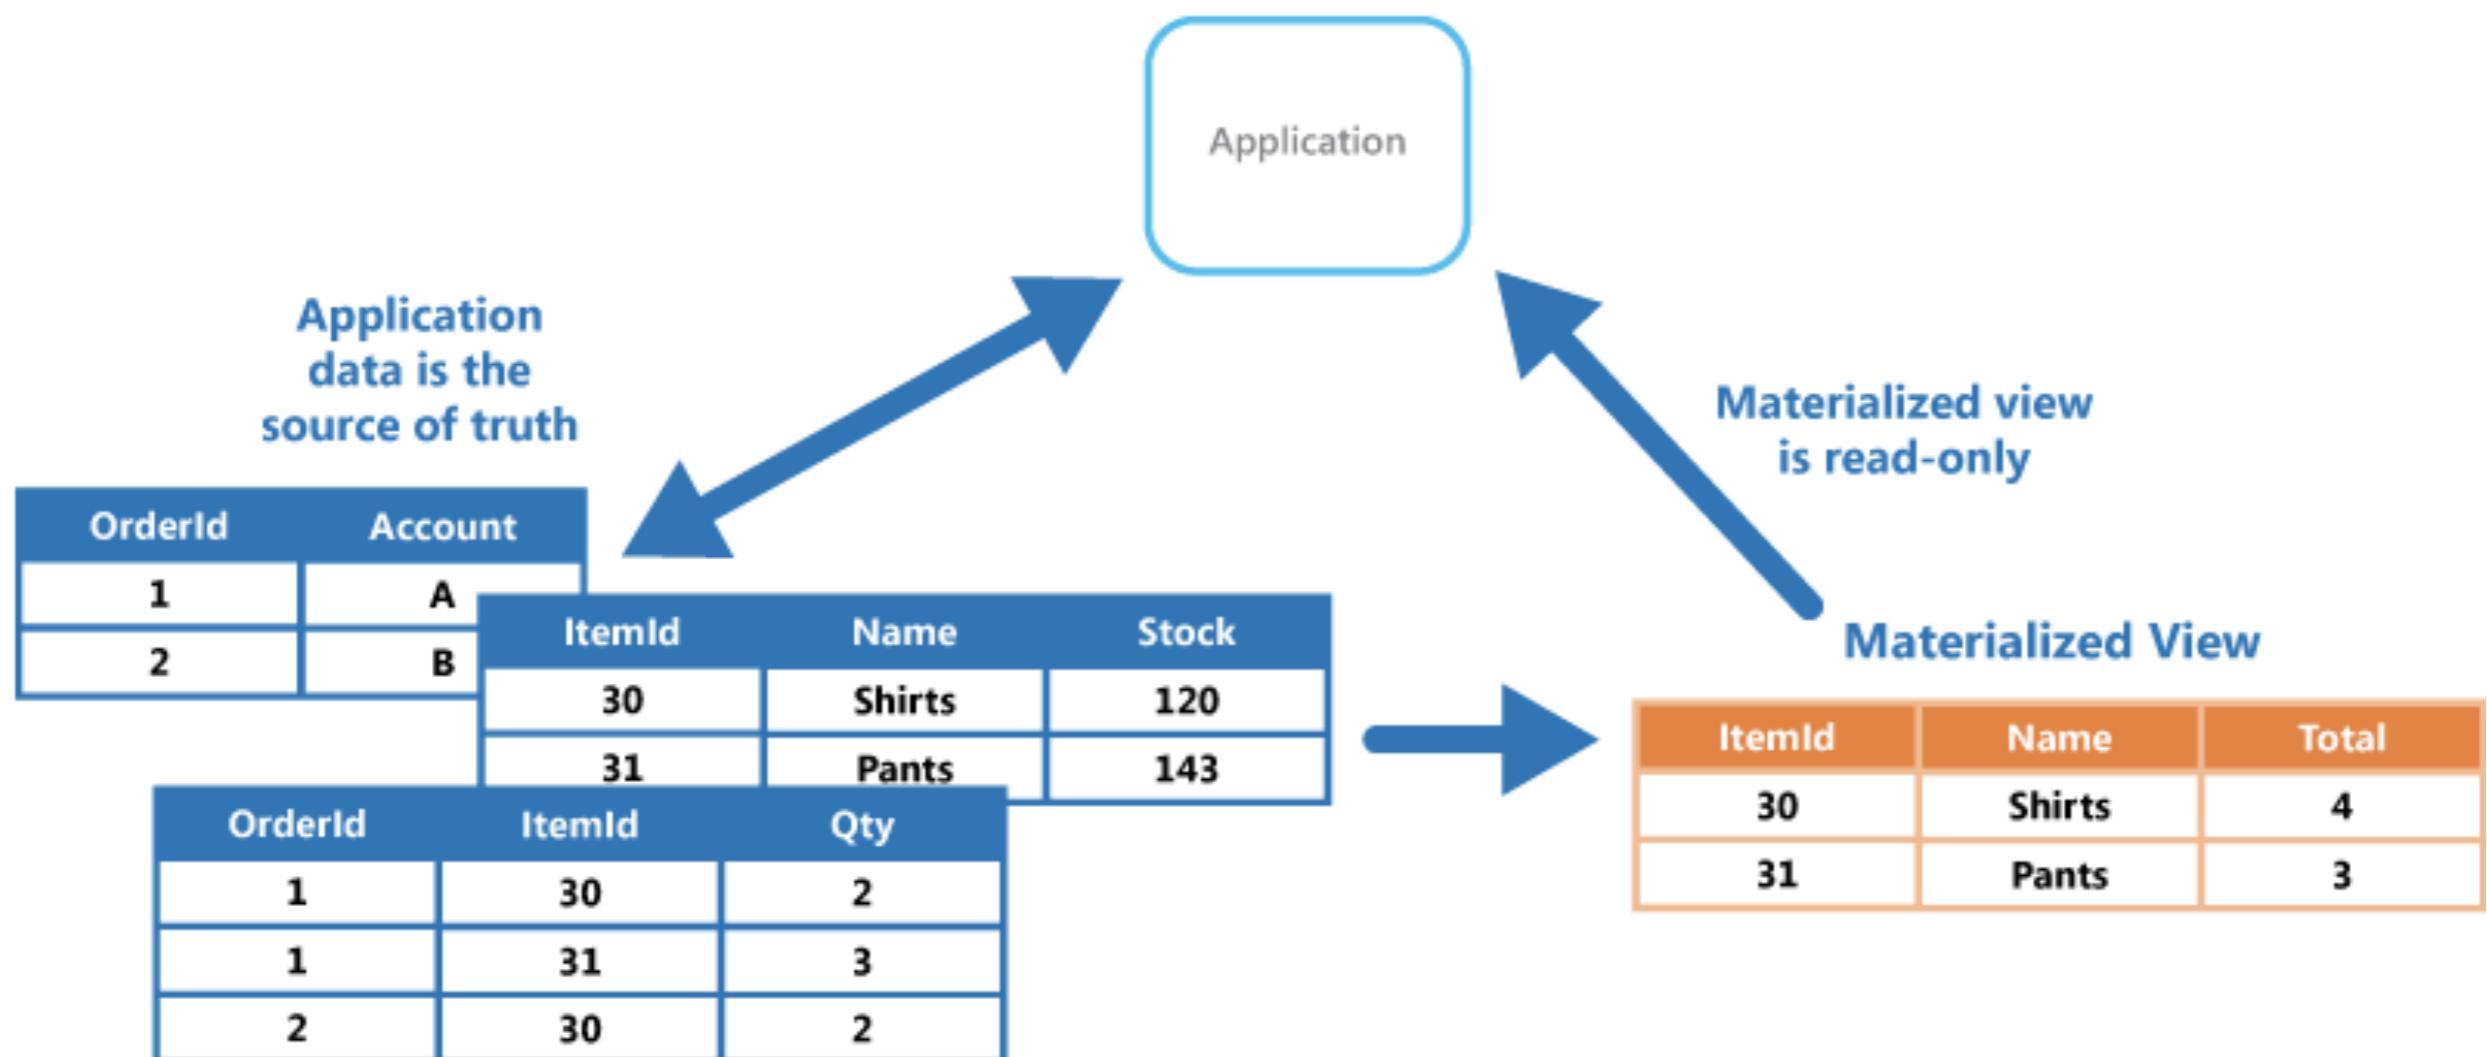


Saga



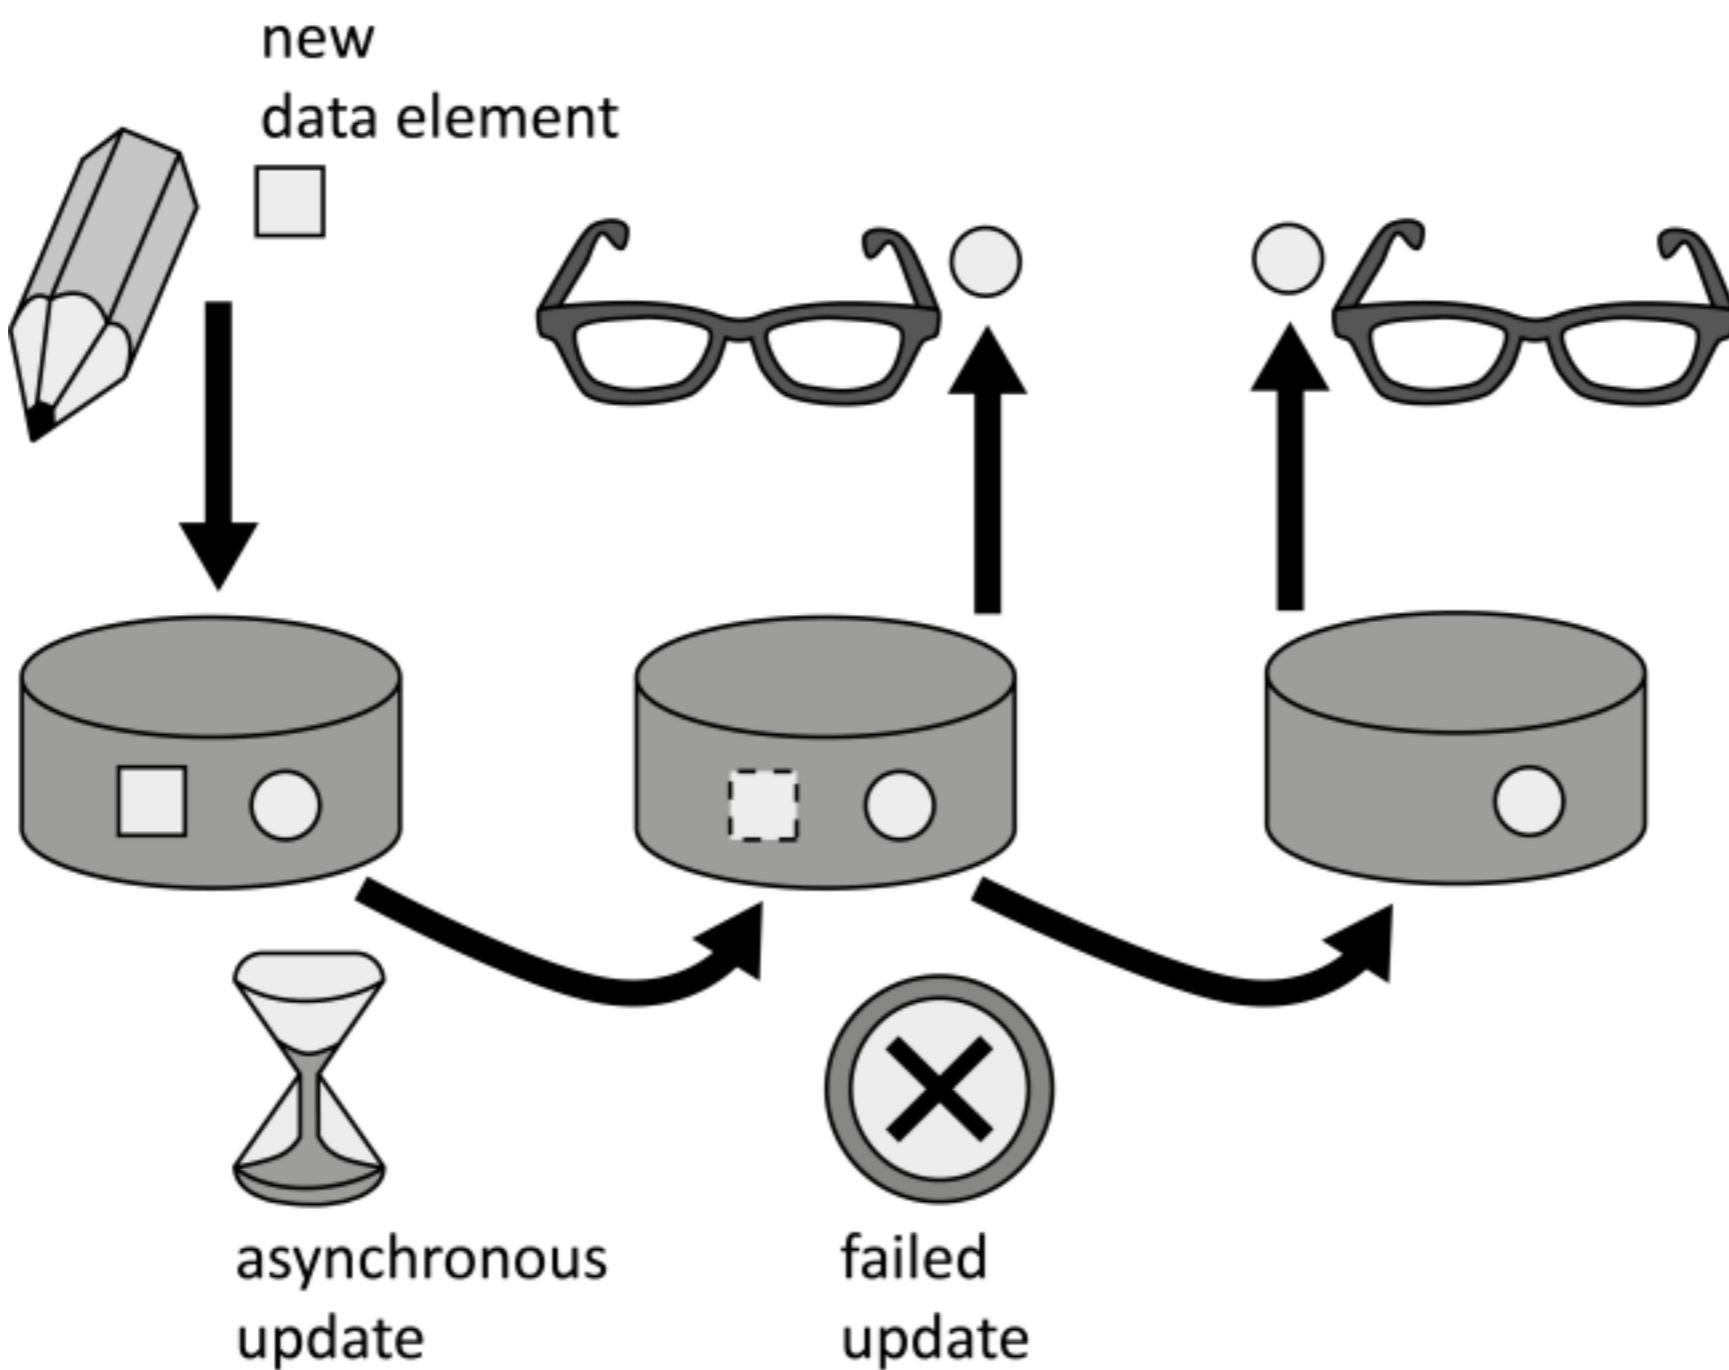
Saga Pattern is a direct solution to implementing distributed transactions in a microservices architecture.

Materialized View

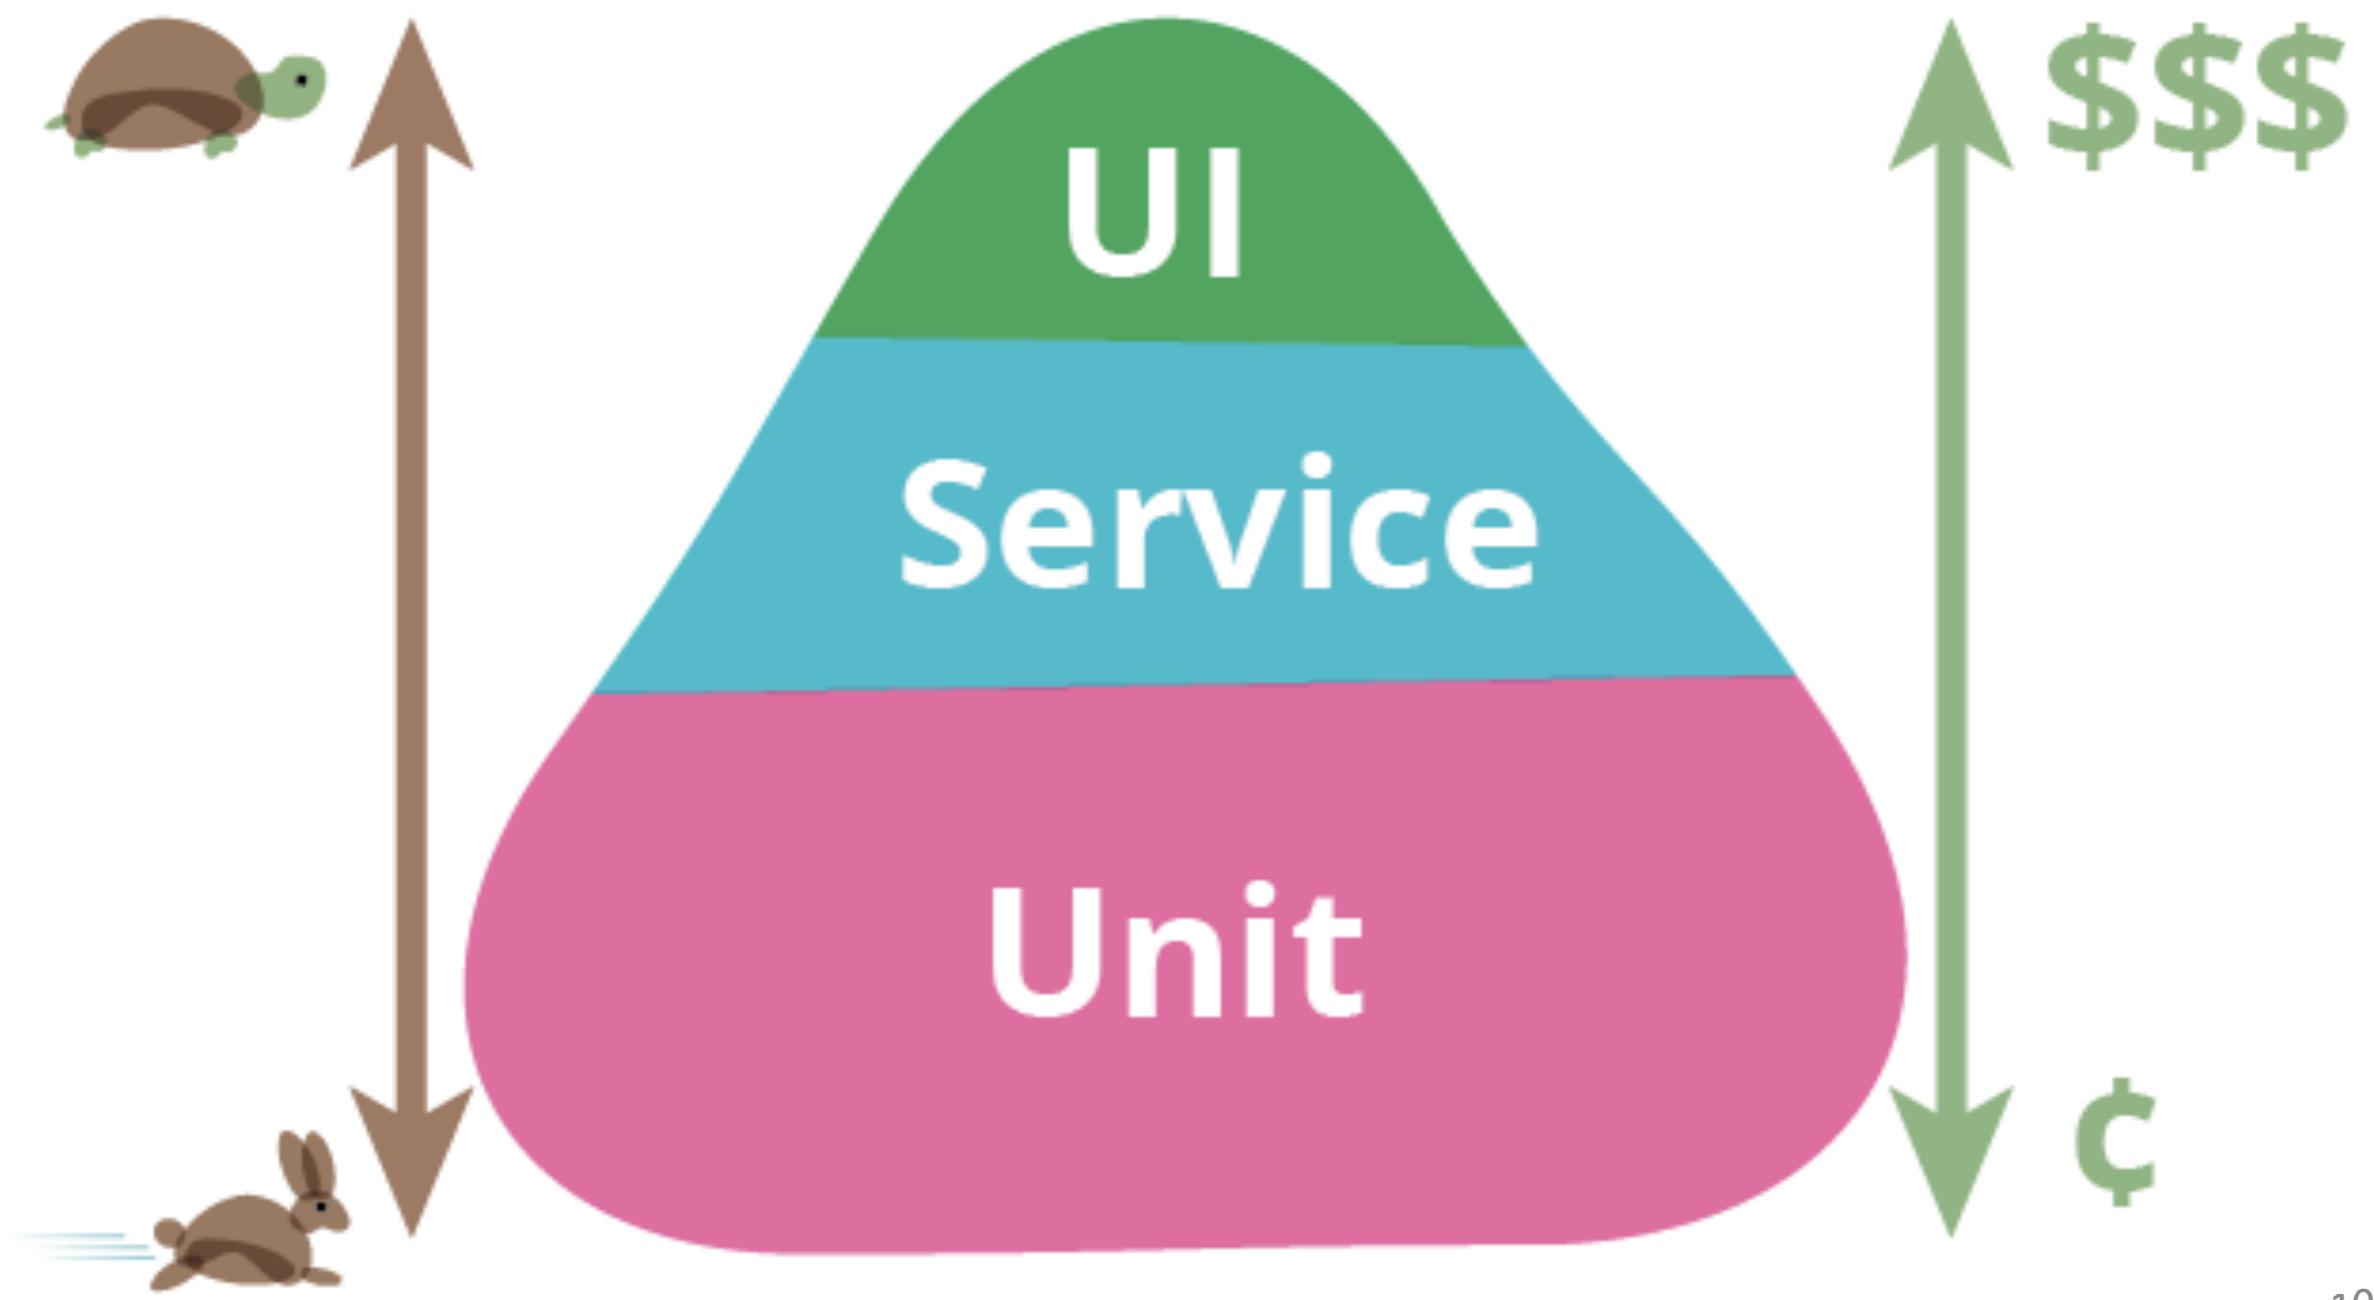


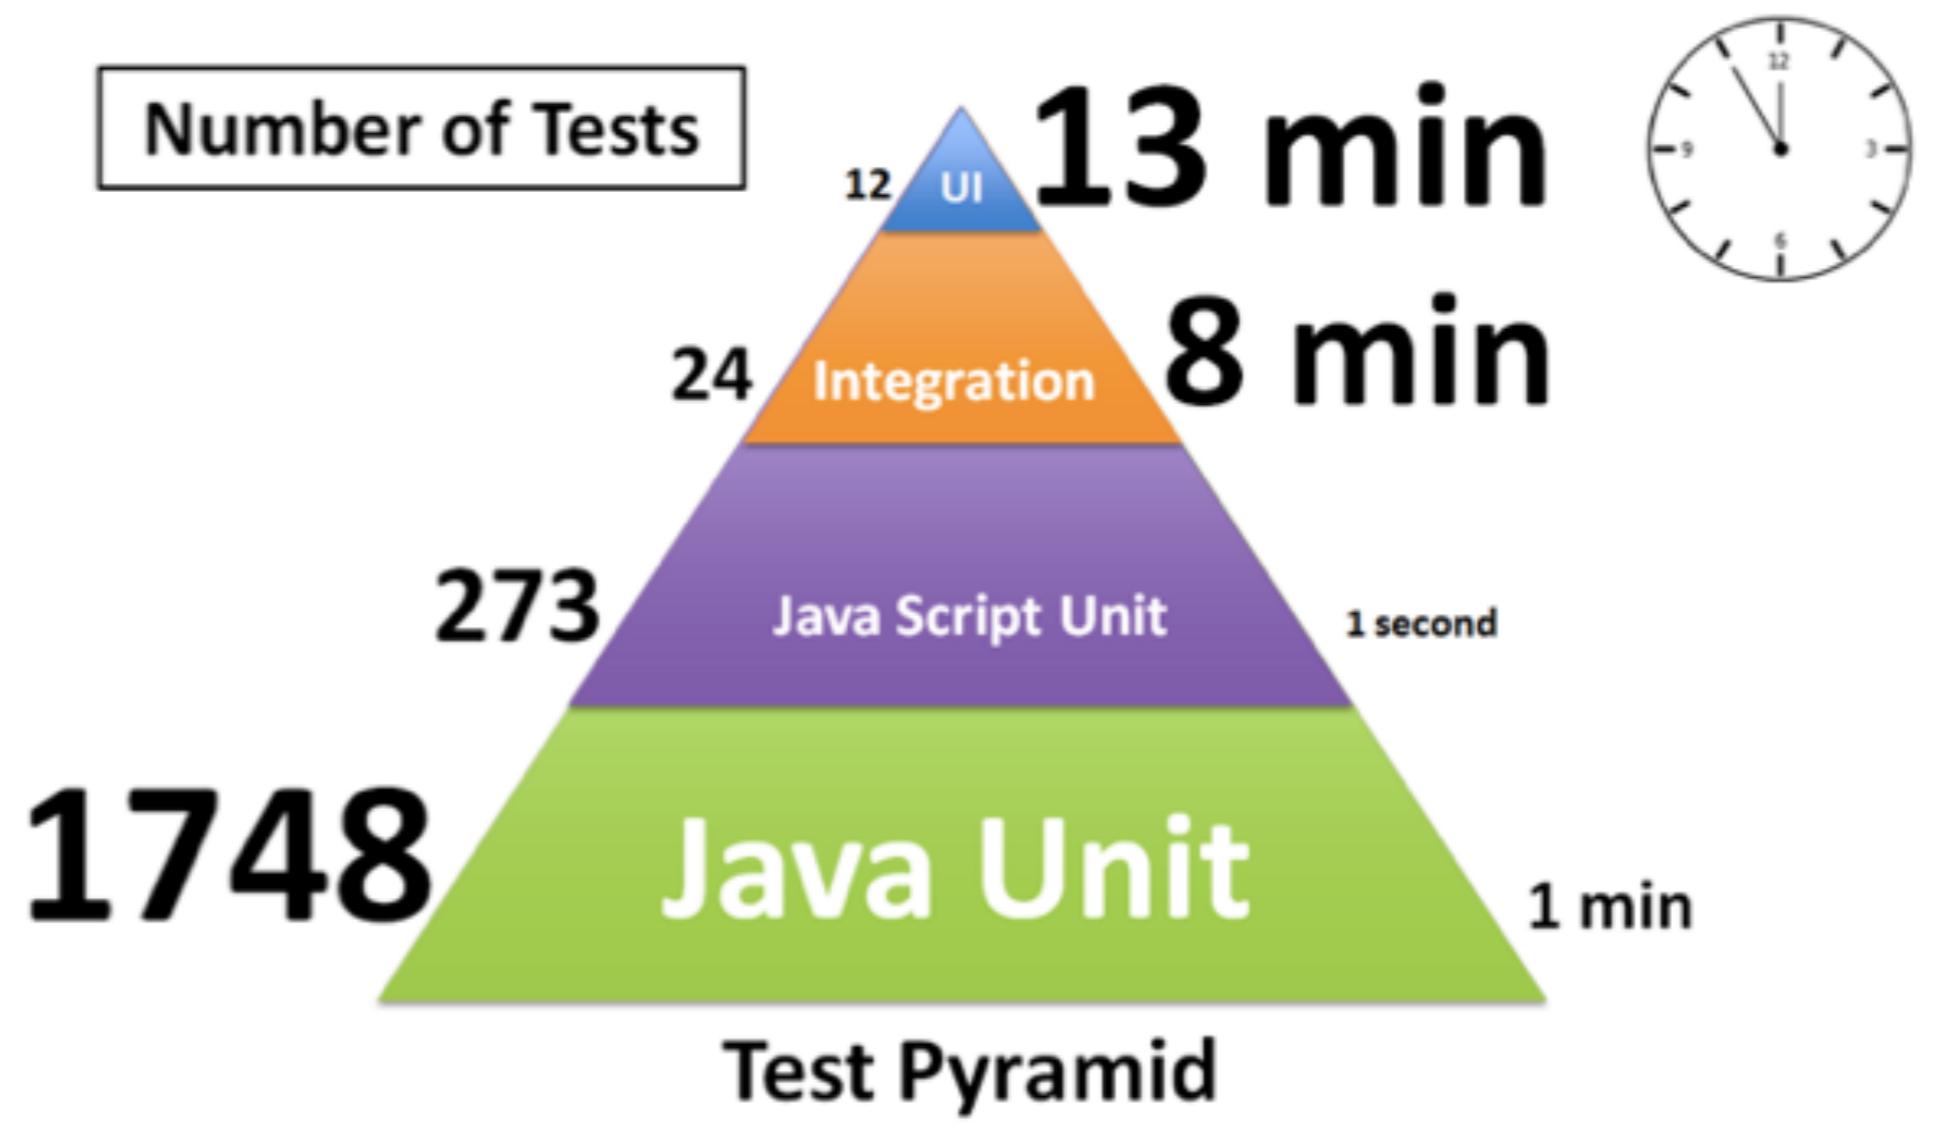
source:msdn

BASE



Test Pyramid





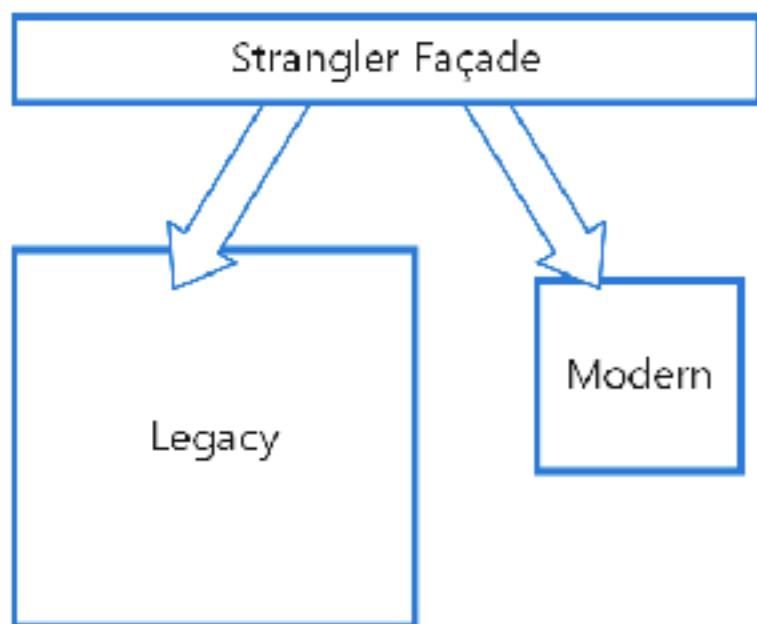
Health Endpoint Monitoring



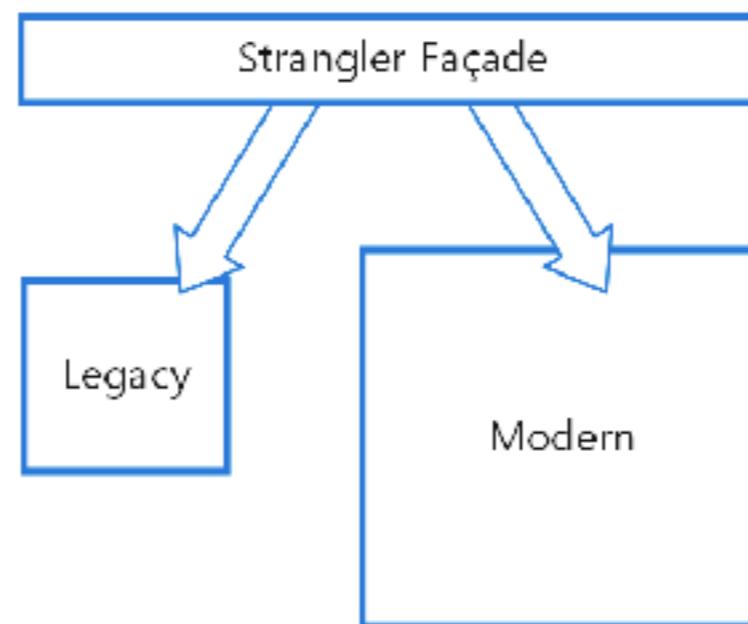
- Implement health monitoring by sending requests to an endpoint on the application. The application should perform the necessary checks, and return an indication of its status.

Strangler pattern

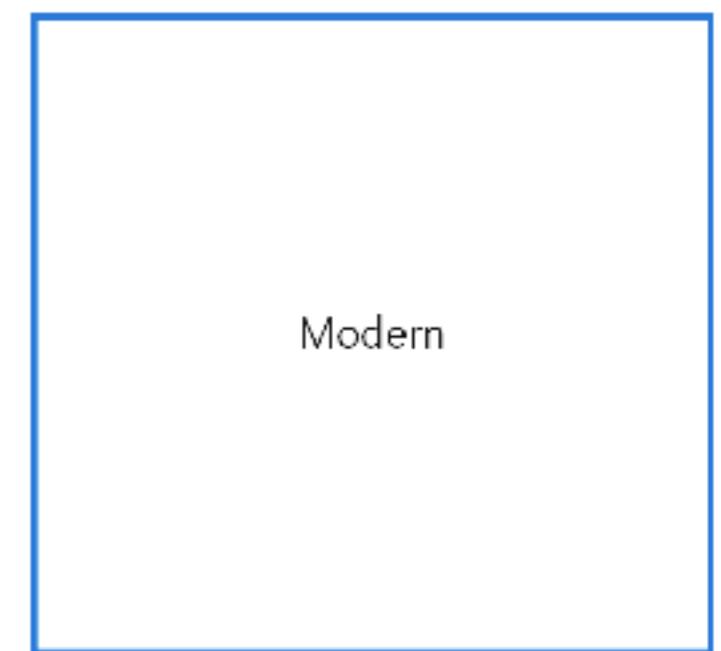
Early migration



Later migration

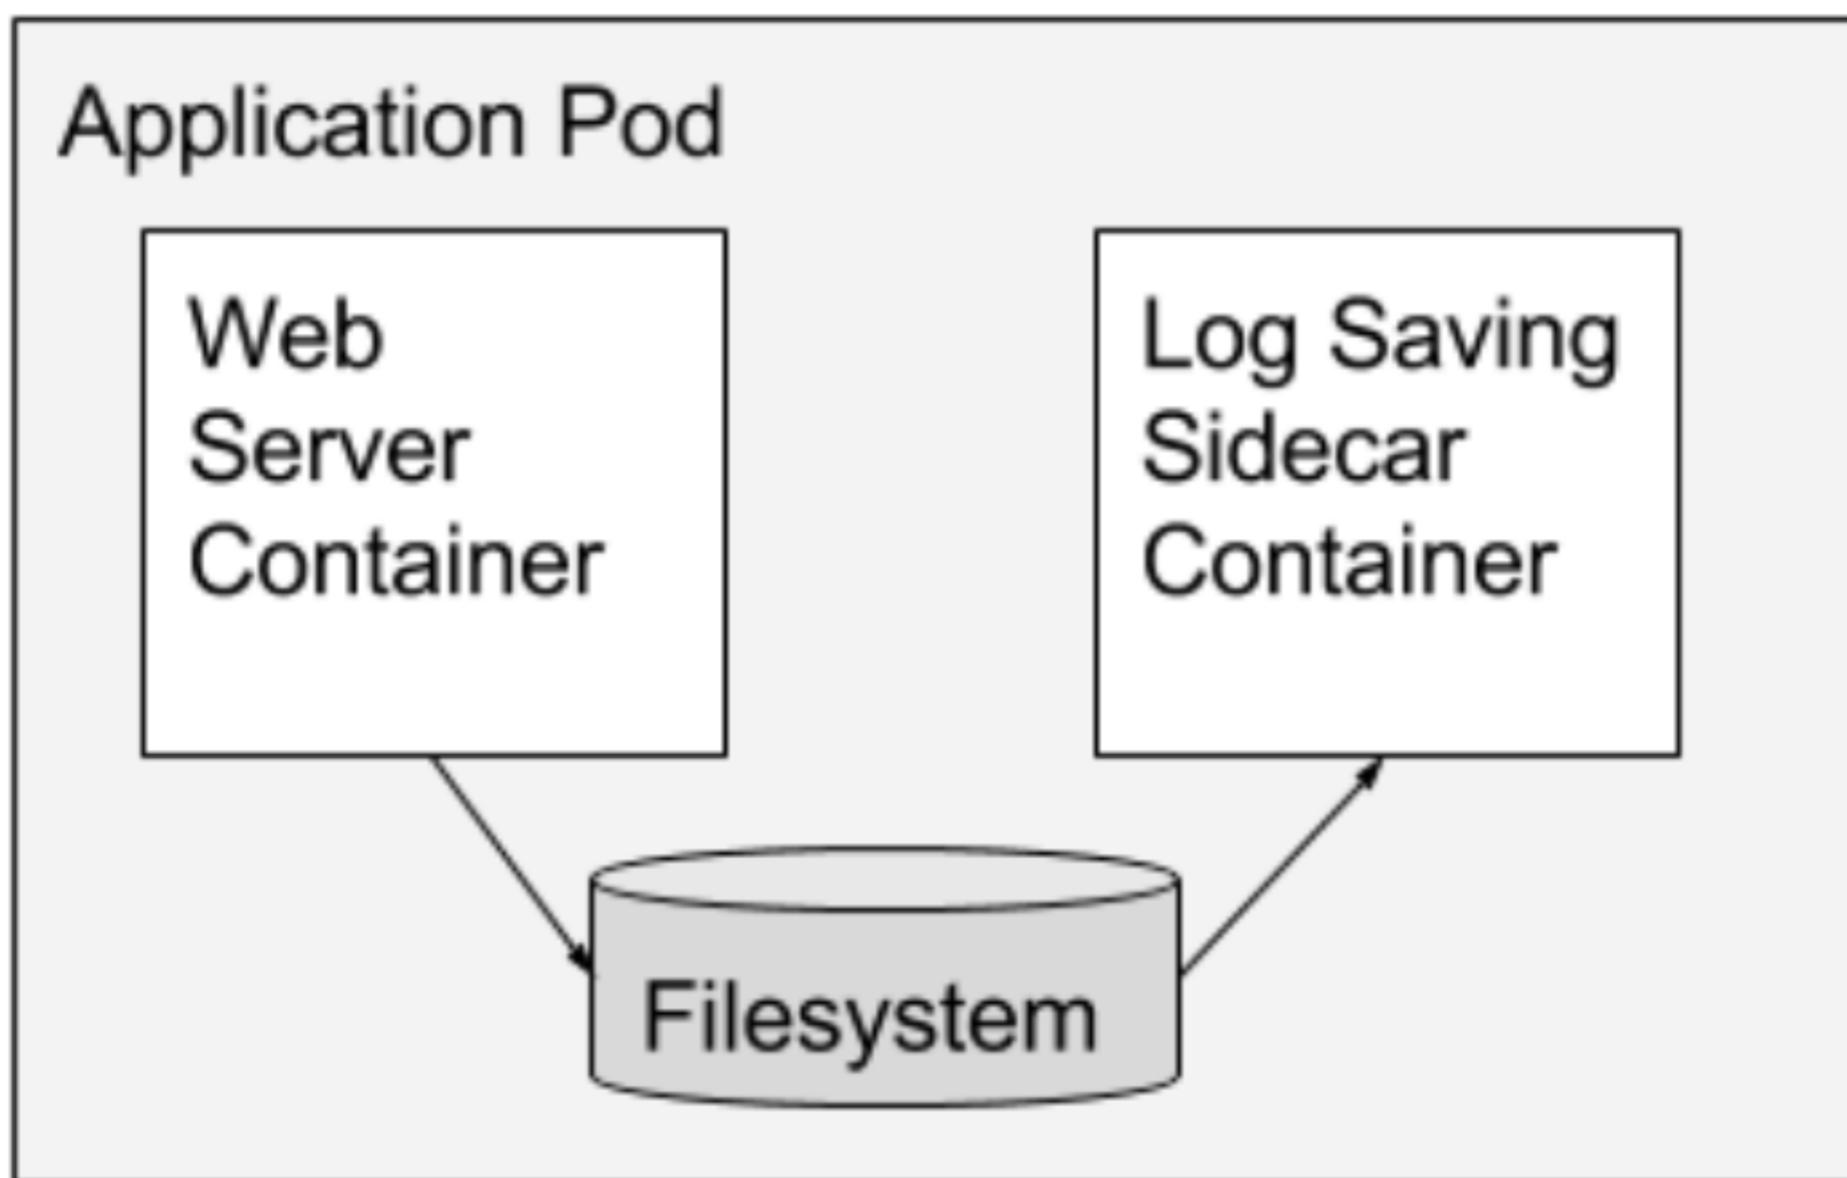


Migration complete



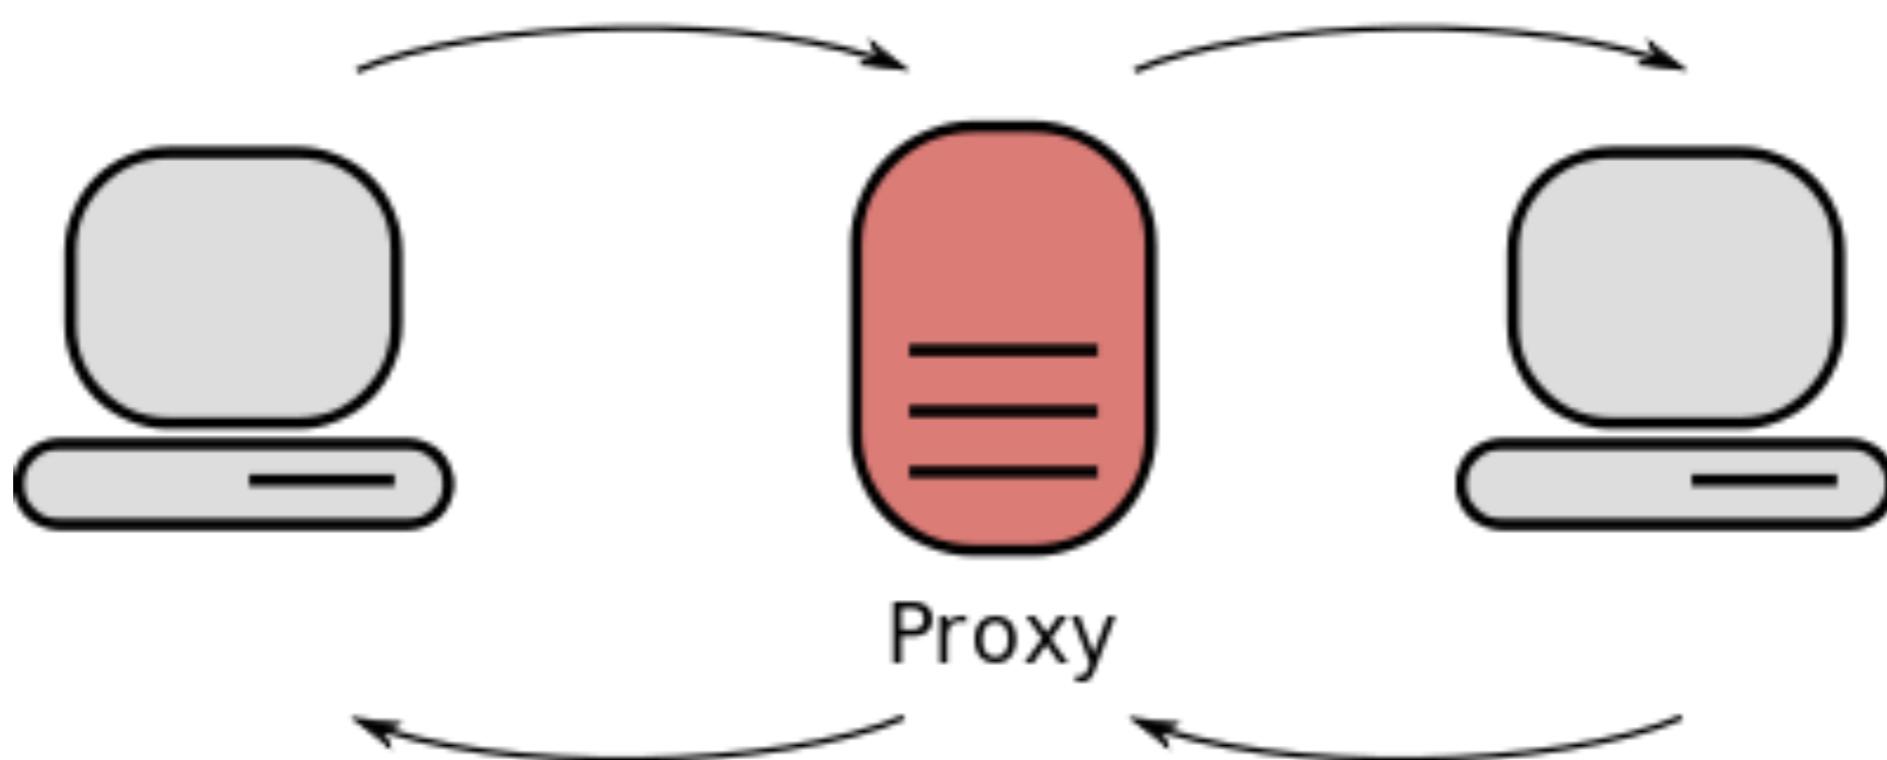
- Create a façade that intercepts requests going to the backend legacy system. Over time, as features are migrated to the new system, the legacy system is eventually "strangled" and is no longer necessary.

Sidecar



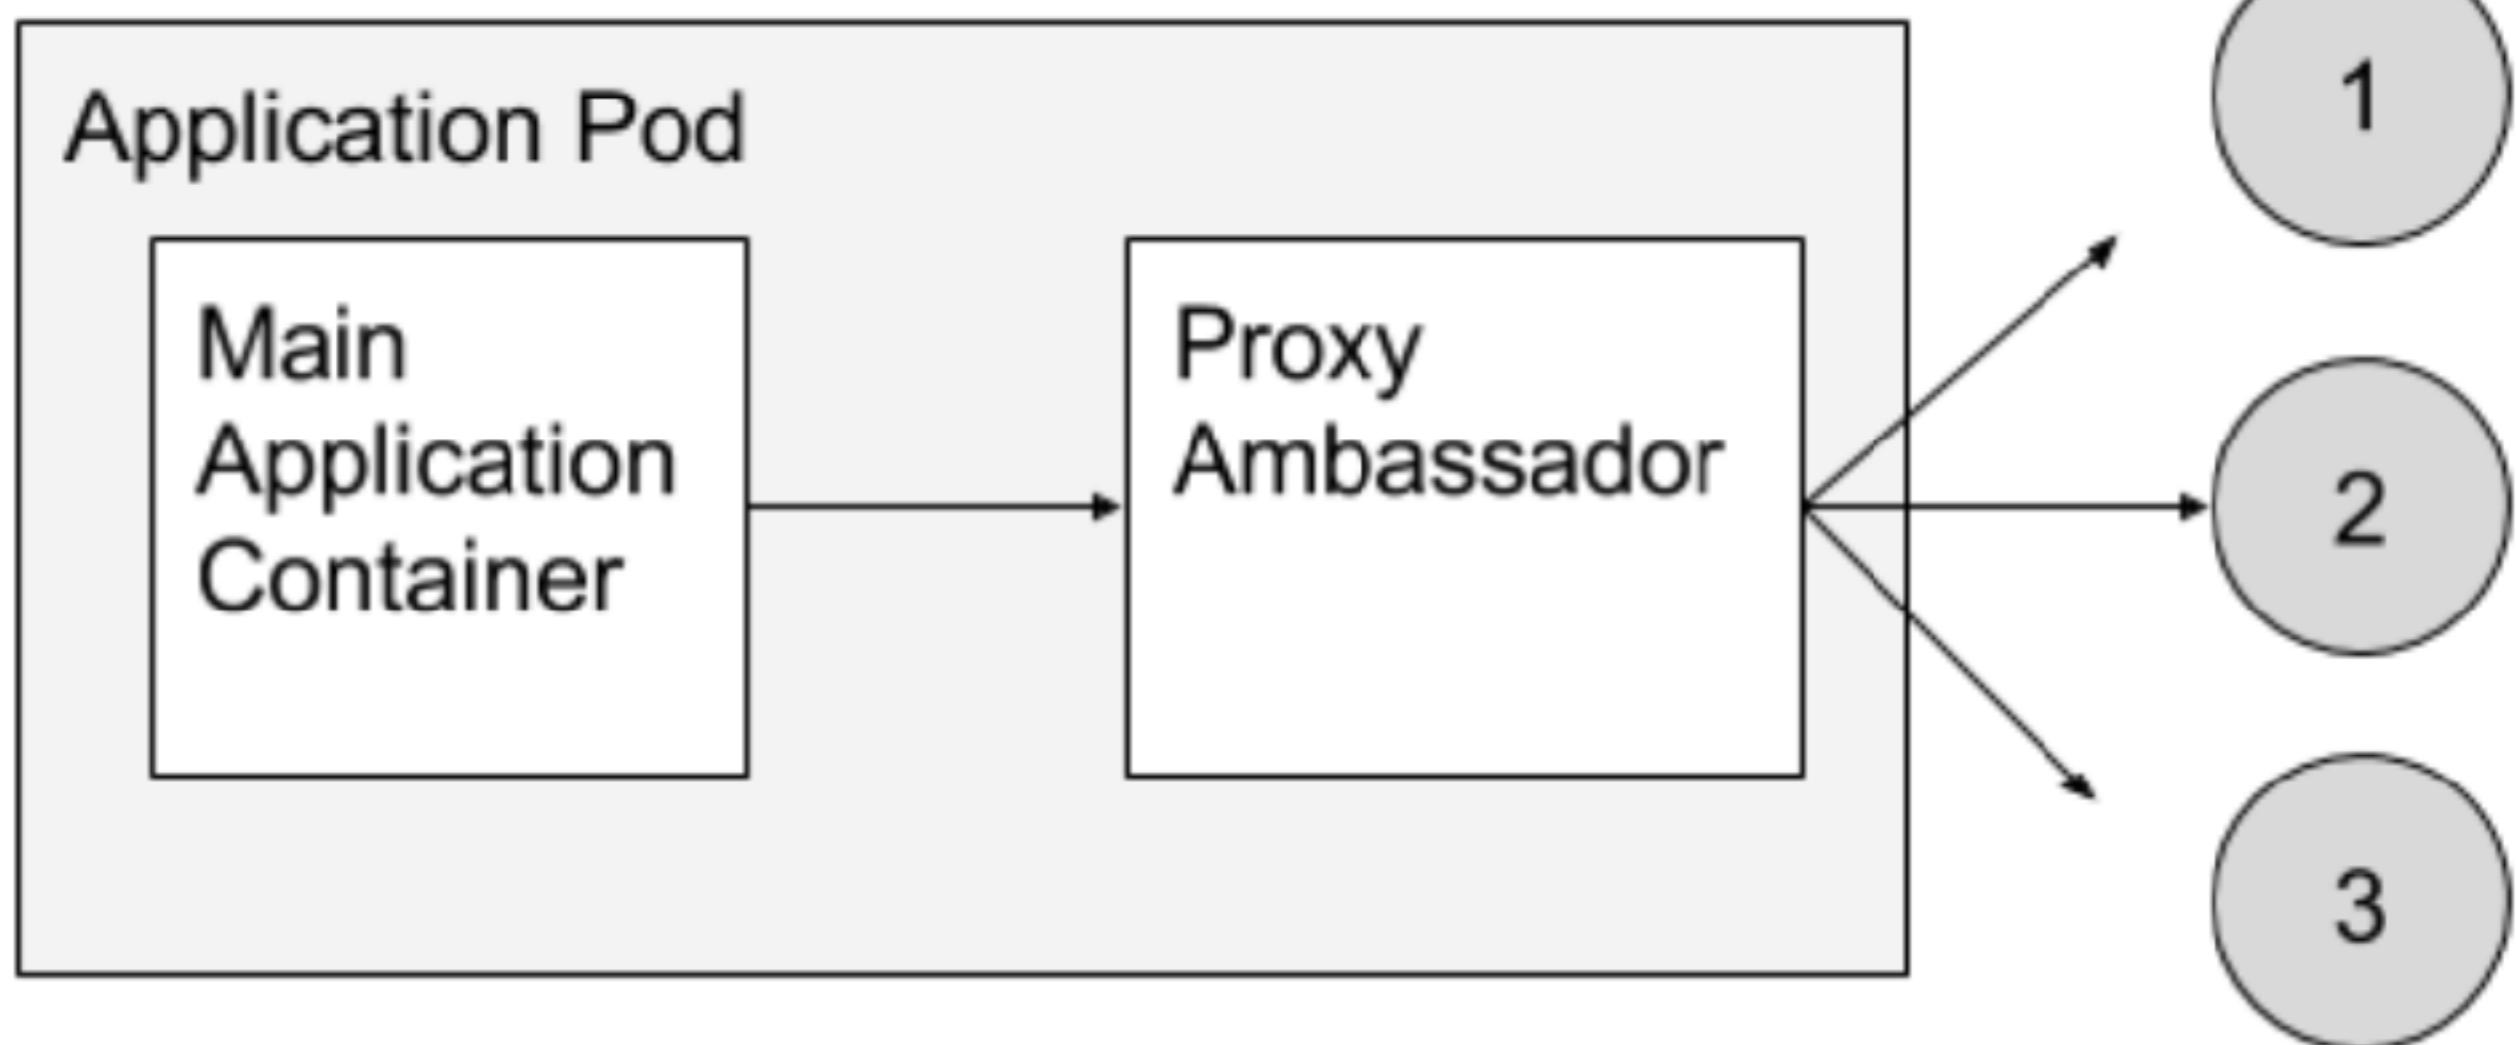
- In a sidecar pattern, the functionality of the main container is extended or enhanced by a sidecar container without strong coupling between two.
- eg. main container is a web server which is paired with a log saver sidecar container that collects the web server's logs from local disk and streams them to centralized log collector.

Proxy pattern



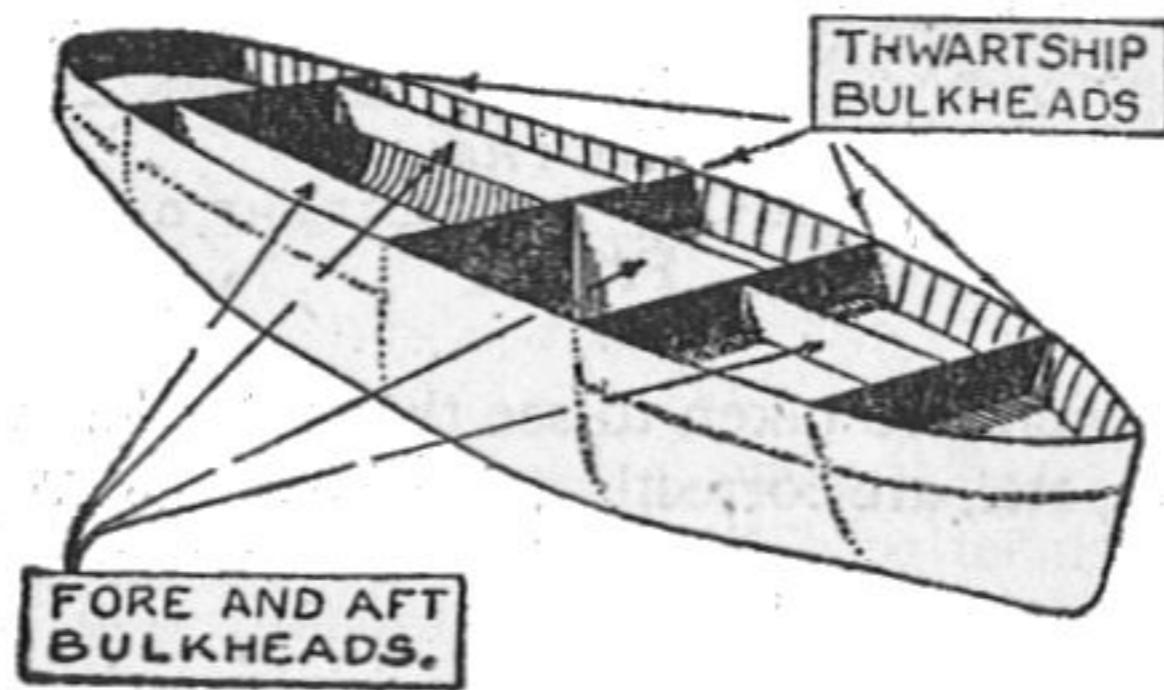
- Proxy pretends being an object when in fact it's not.
There are several types of proxies:
- Remote proxy ("Ambassador") - instead of communicating between server's and client's classes we create a proxy on client side and server side, and only the proxies communicate
- Virtual proxy - creates expensive object on first demand
- Protection proxy - to control access
- Smart proxy - enrich an object

Ambassador pattern



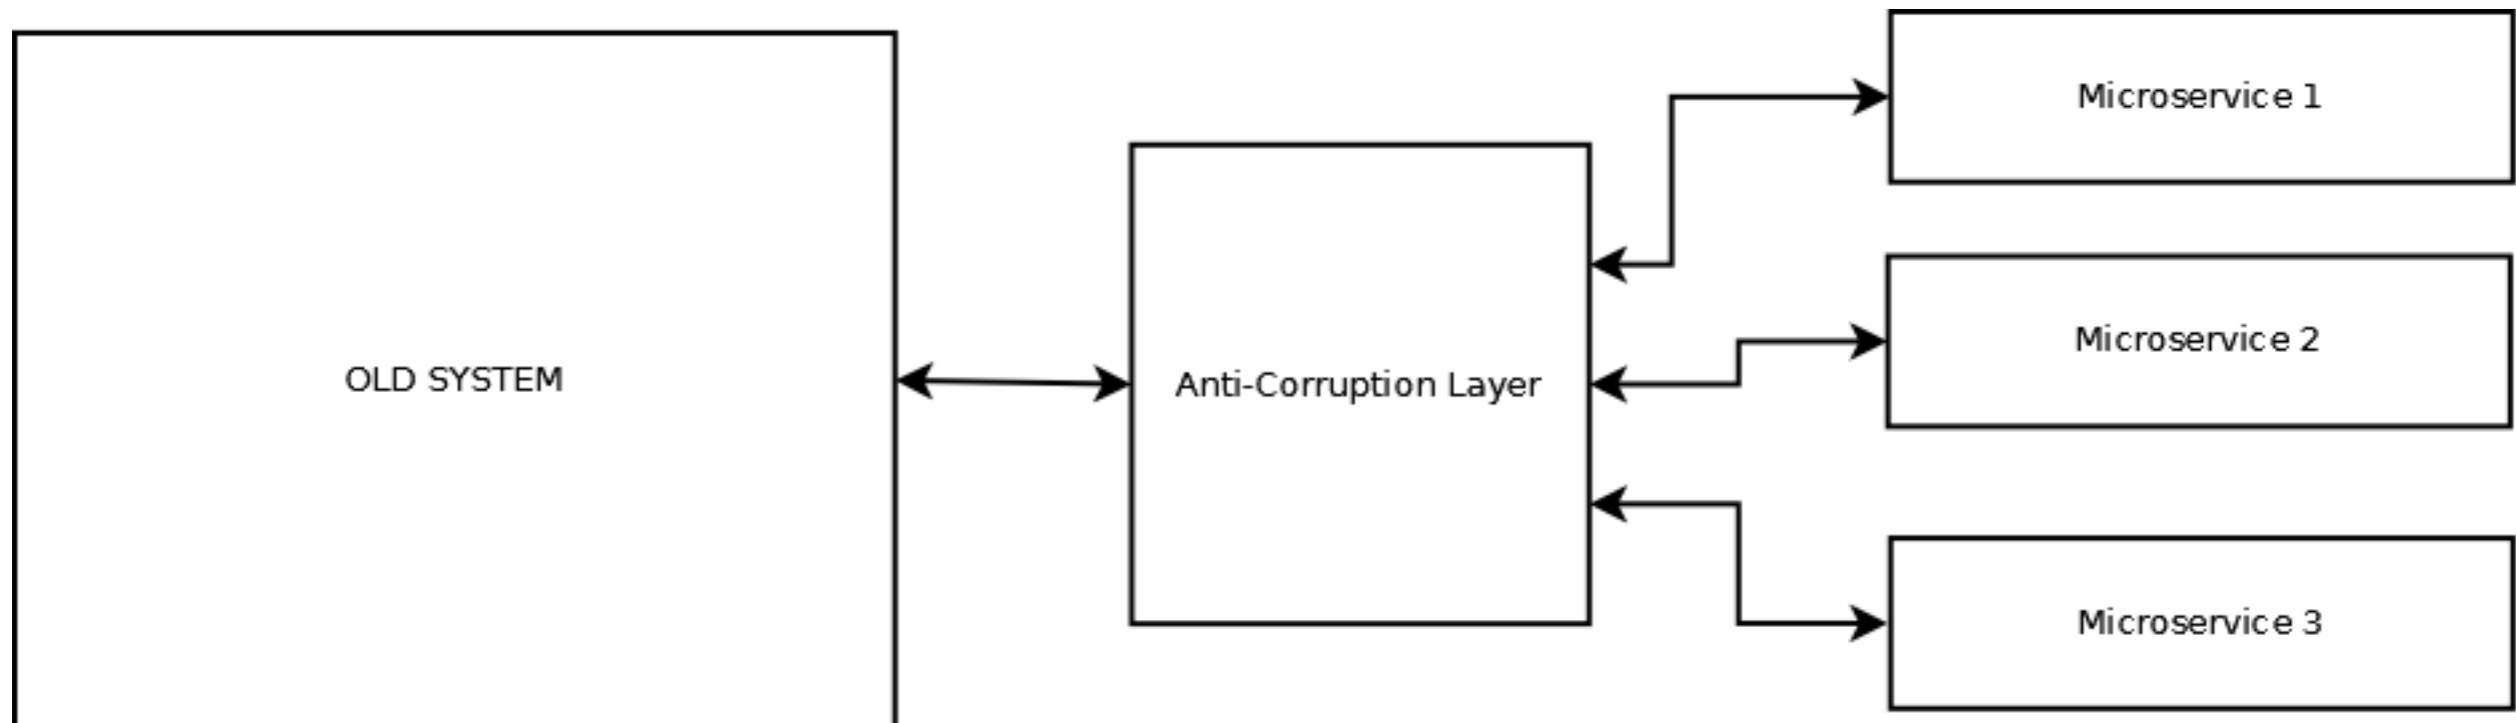
- an Ambassador container act as a proxy between two different types of main containers. Typical use case involves proxy communication related to load balancing and/or sharding to hide the complexity from the application.

bulkhead



- In general, the goal of the bulkhead pattern is to avoid faults in one part of a system to take the entire system down. The term comes from ships where a ship is divided in separate watertight compartments to avoid a single hull breach to flood the entire ship; it will only flood one bulkhead.cir
- Partition service instances into different groups, based on consumer load and availability requirements.

Anti-Corruption Layer

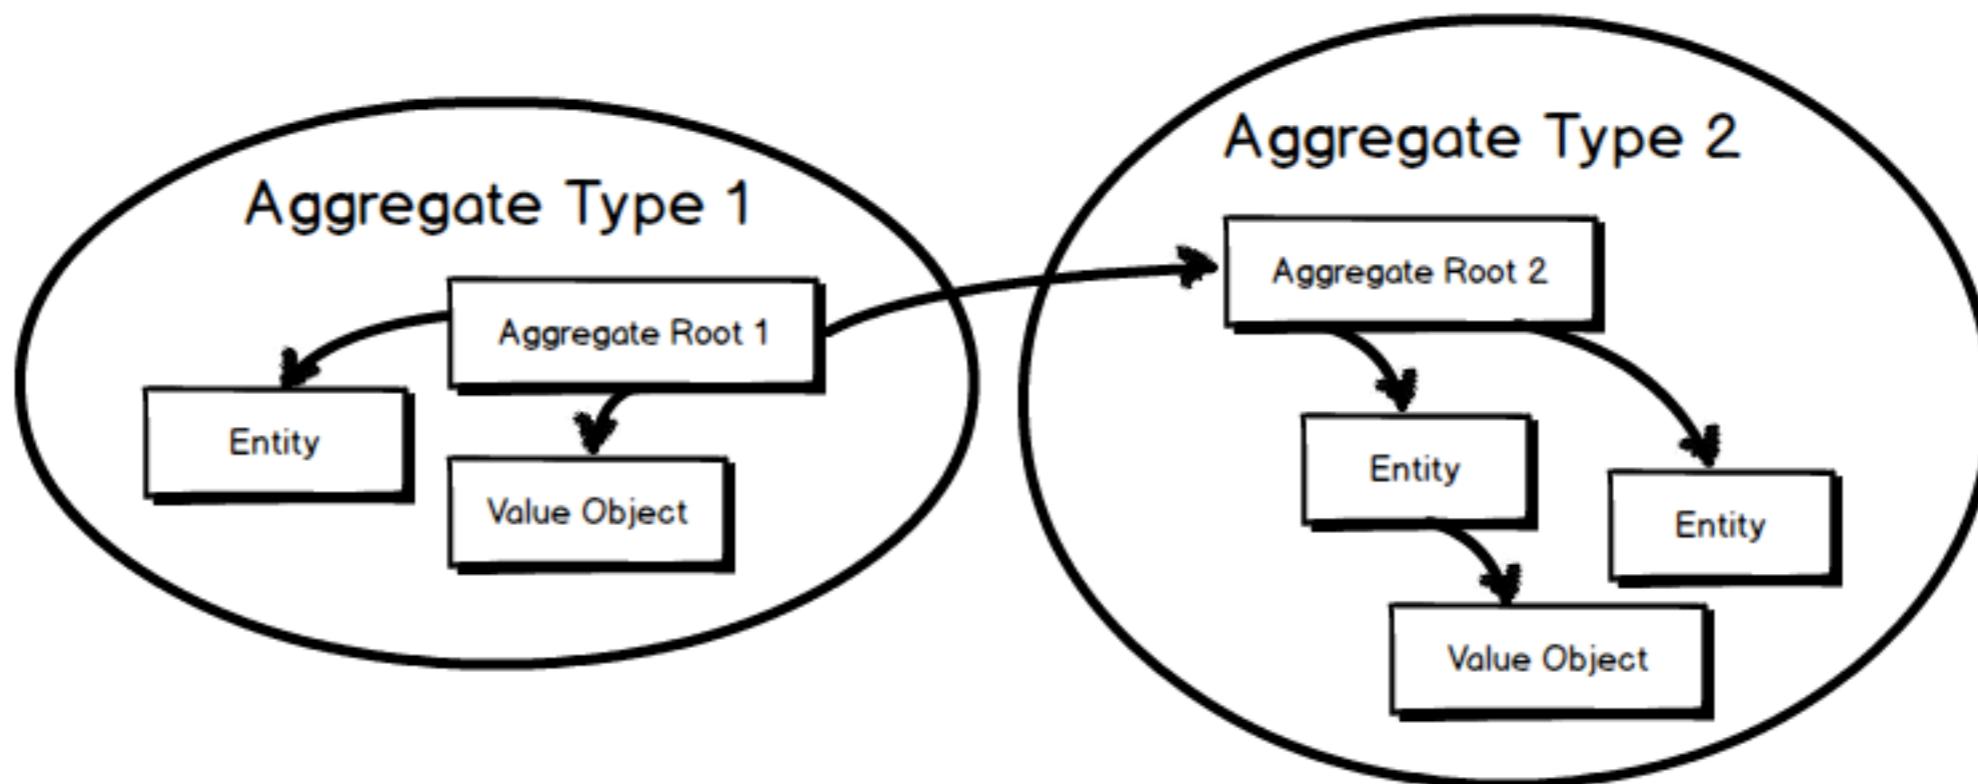


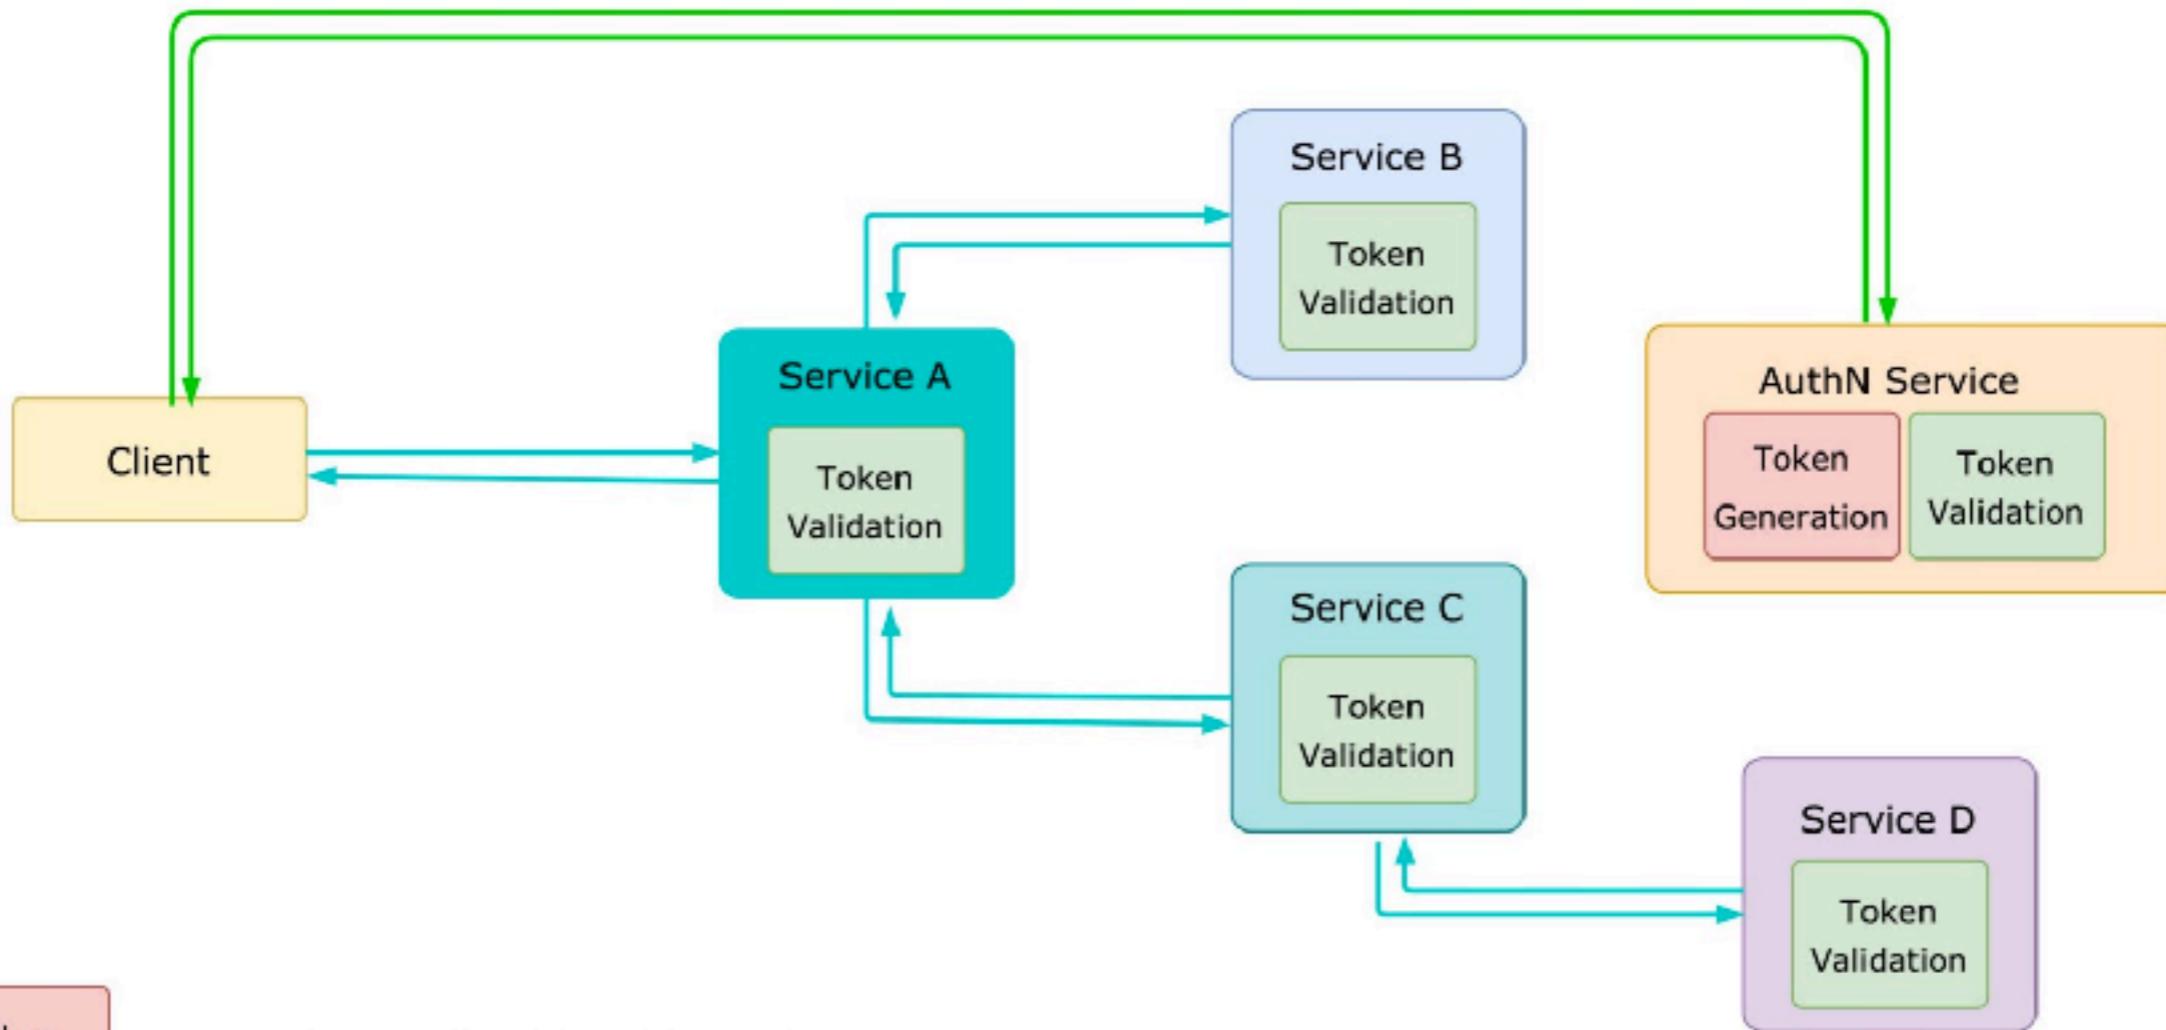
- When you create a new microservice, you notice that some business or technical assumptions or flow should be remodeled. instead of a Big Rewrite ACL maps one domain onto another so that services that use second domain do not have to be "corrupted" by concepts from the first.
- Such modifications may result in changes to the events that a bounded context publishes.

Bounded Context



Aggregate Root





Token
Validation

Token Generation Using Private Key

Token
Validation

Token Validation Using Public Key

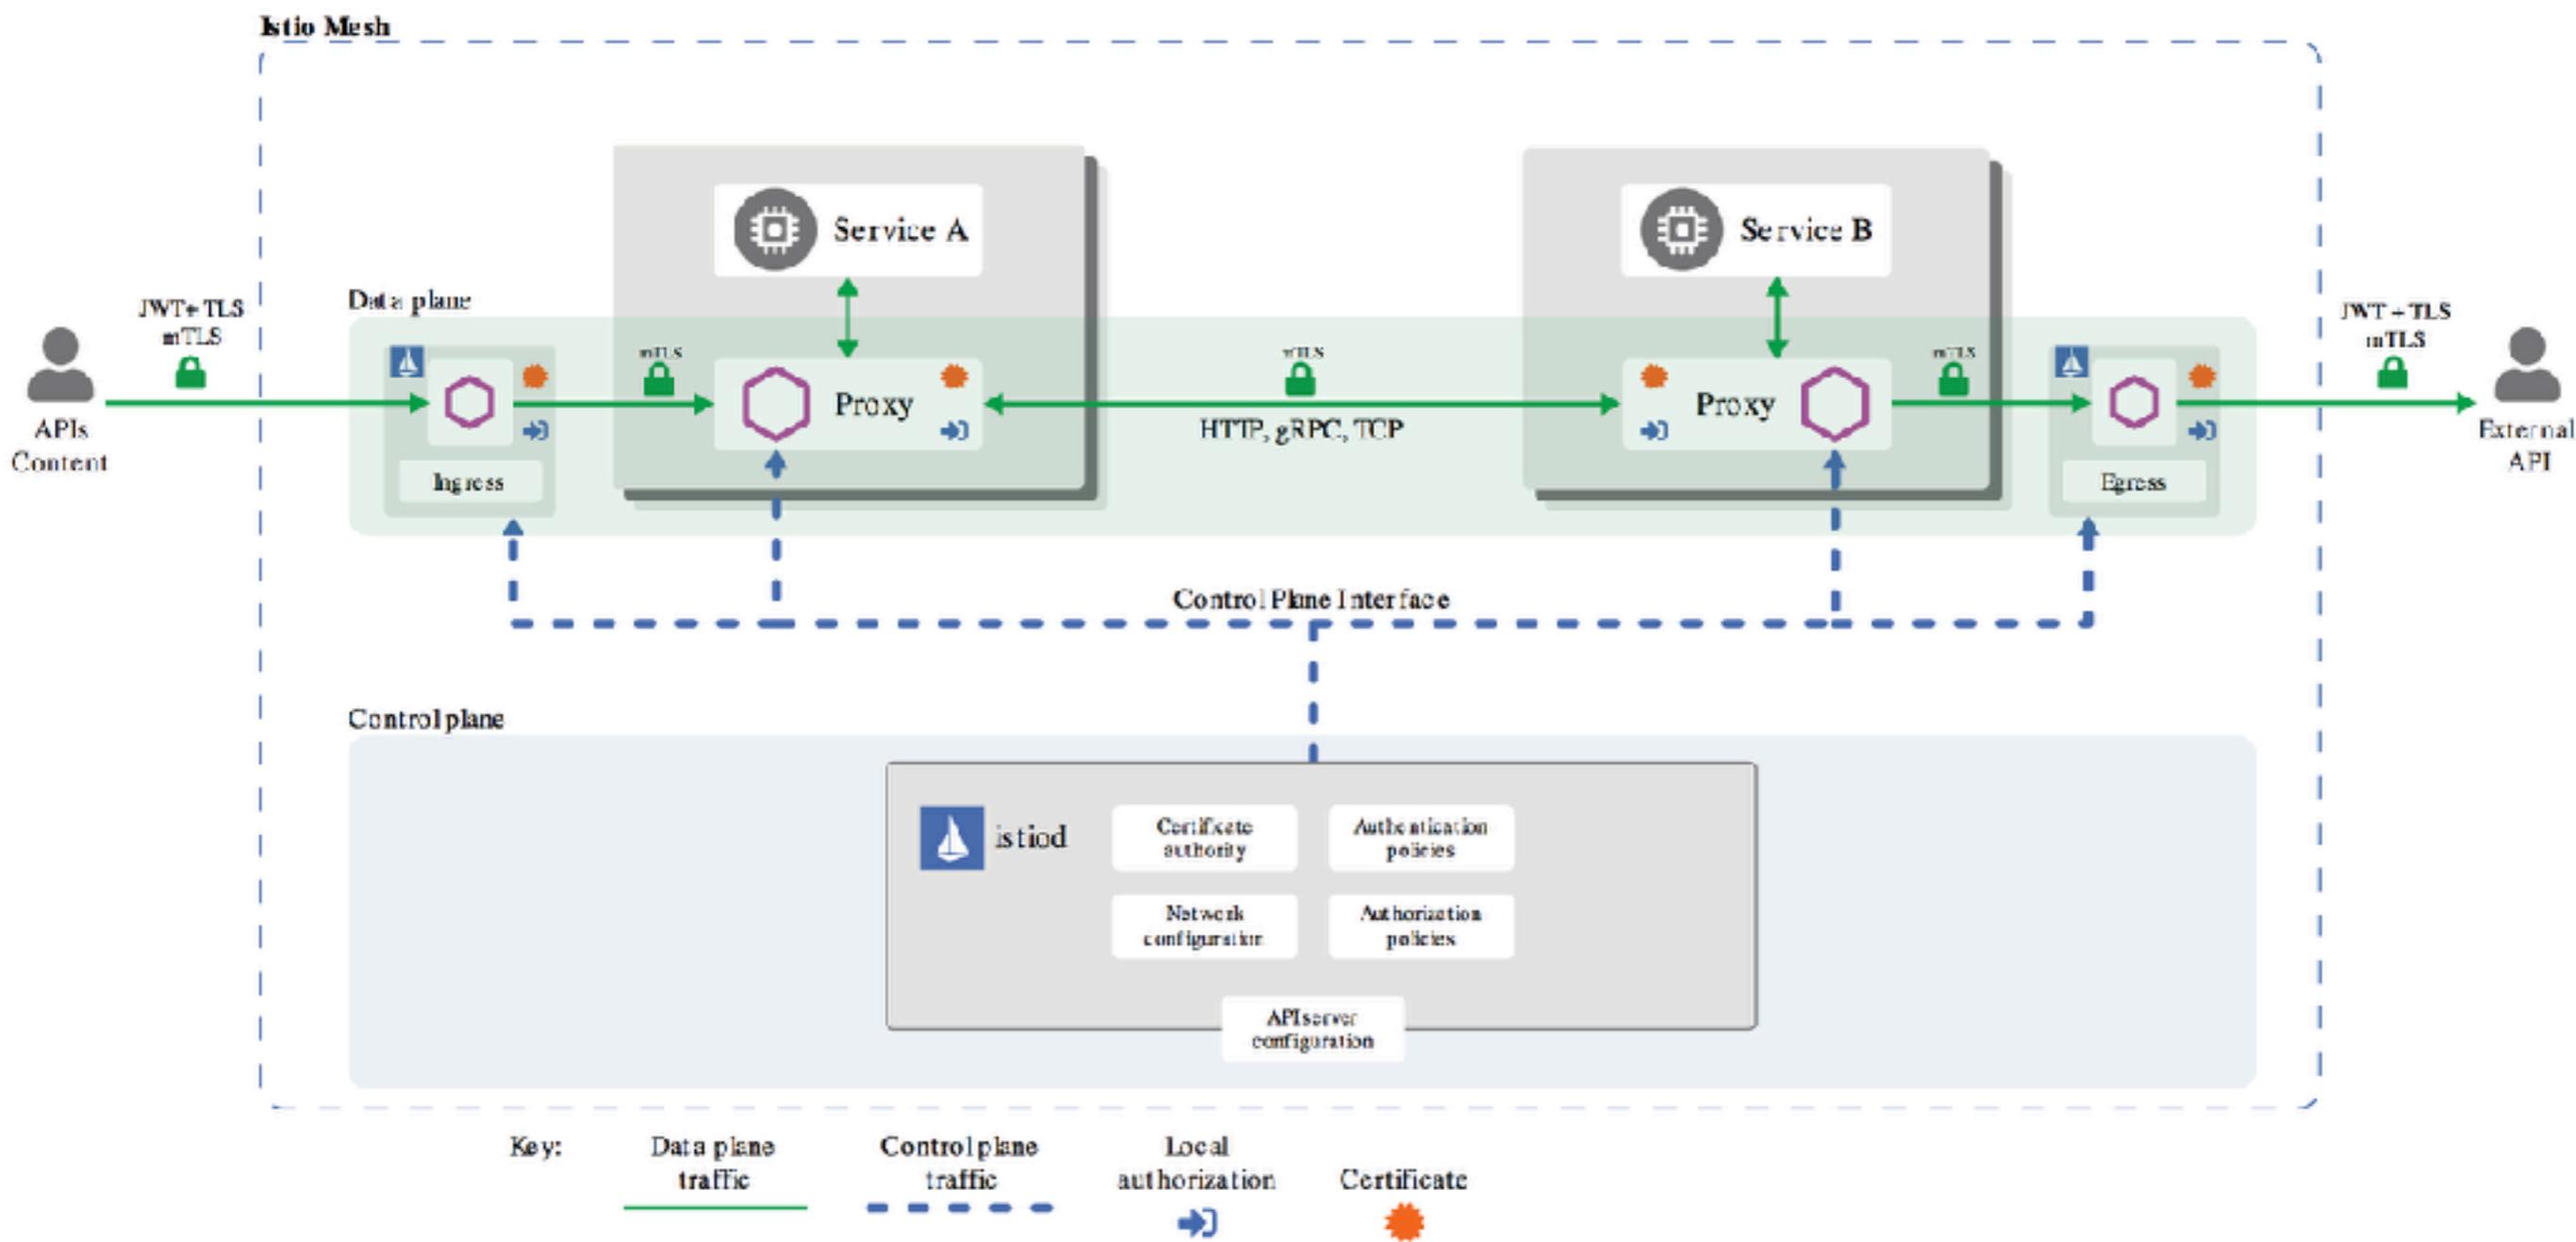
Any service Can Validate Using Public Key

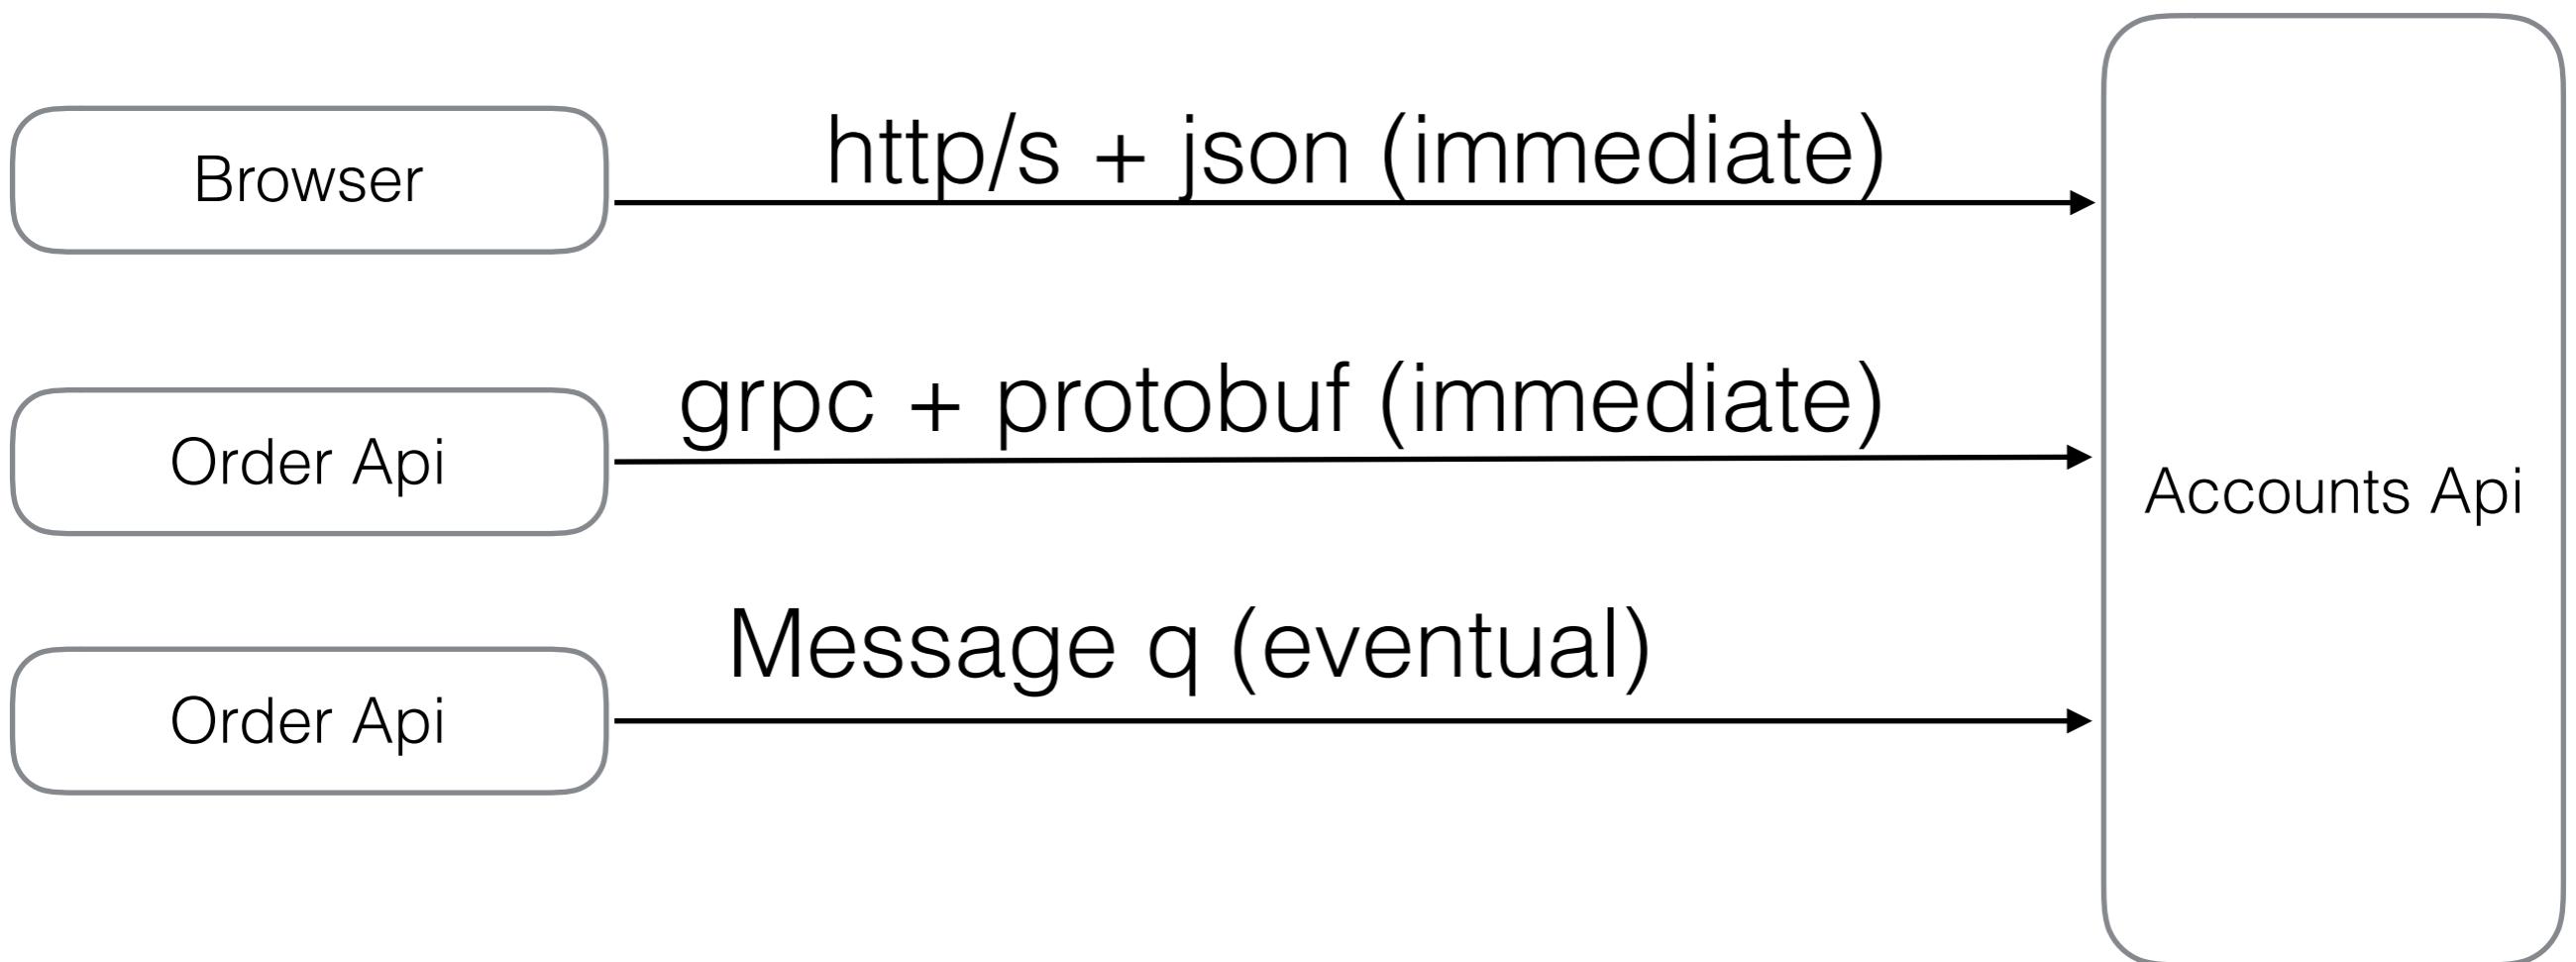
Private Key in One Secure Place, Really Insecure

→ Traffic for Authentication & Generating JWT Token

→ Traffic for Validating JWT Token

→ Traffic for Serving Client Request



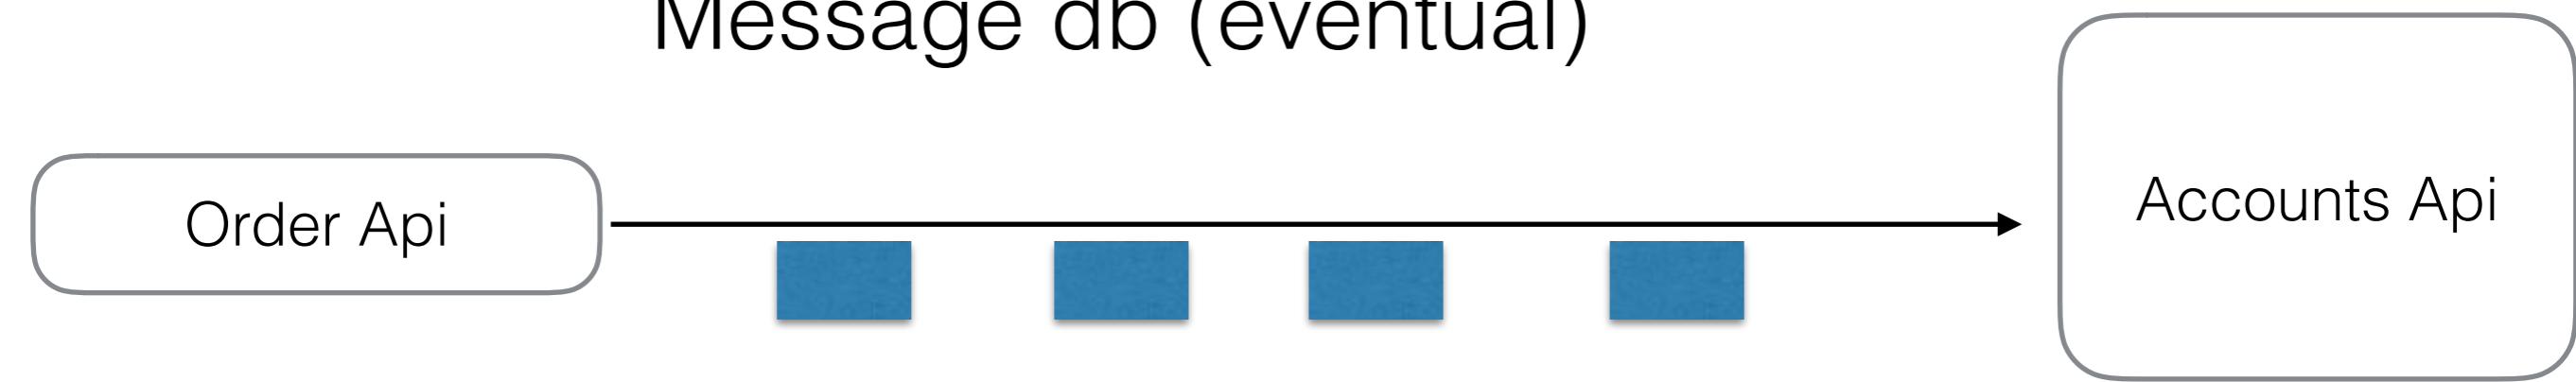


	http	Grpc	Message Q
Pay load	JSON	Protobuf	Any
Immediate	Y	Y	Eventual
Reliable	N	N	Y (retry)
Type	Connected	Connected	Disconnected
Scale	- (chatty)	+ (chunky)	+++ (load leveling)
Protocol	Request-response	Request-response	One way

Message Q

- One way protocol (result and error)
- UnOrdered Delivery
- Duplicate delivery (Idempotent)
- Retention Period

Message db (eventual)



`msg = getmessage()`

...

...

`msg.ack()`

Create(order 101)



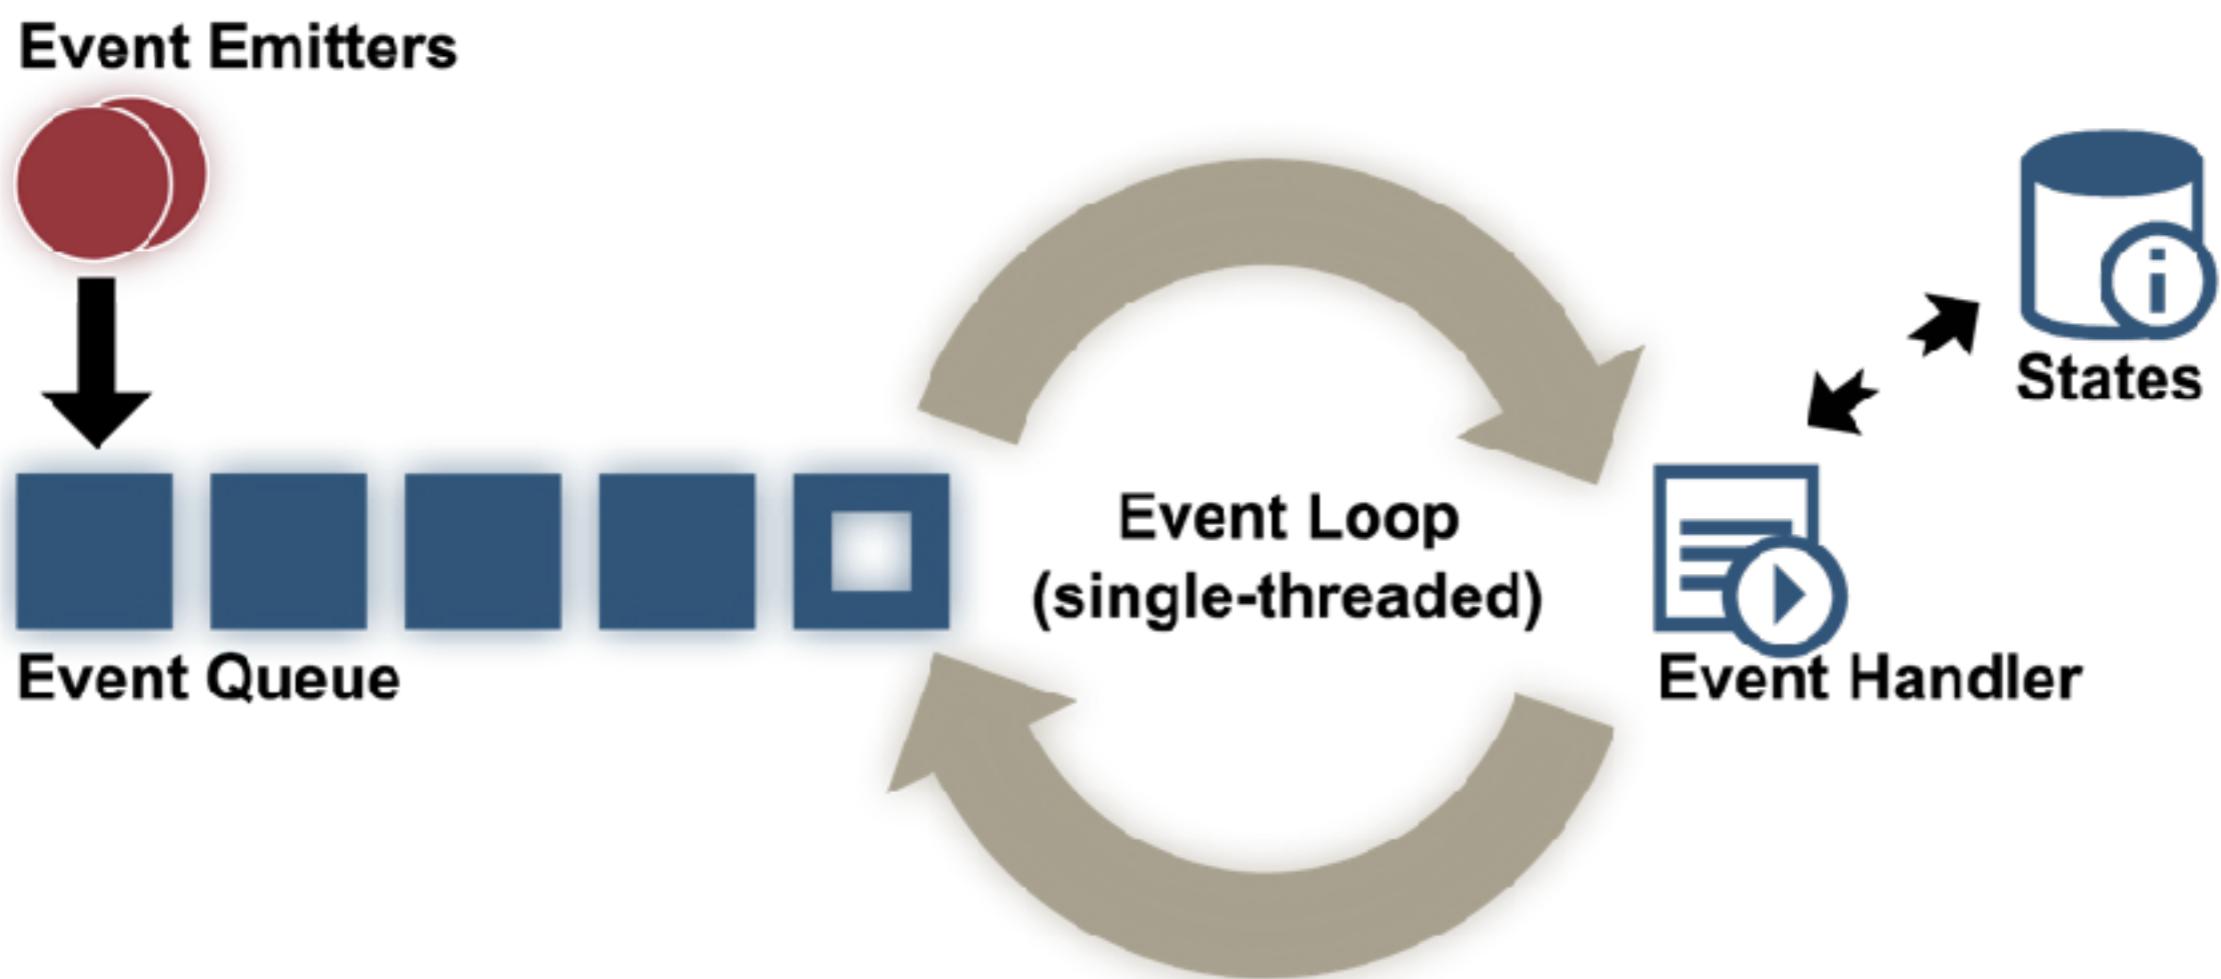
Create(order 101)



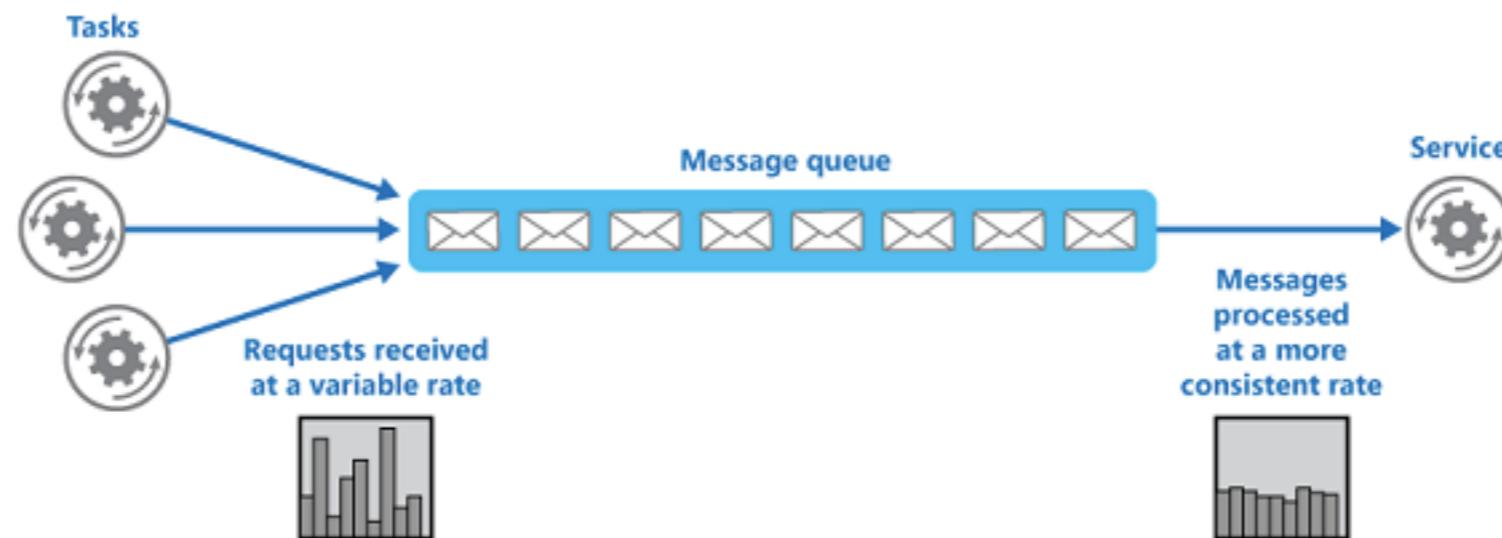
Accounts Api



Event Driven Architecture

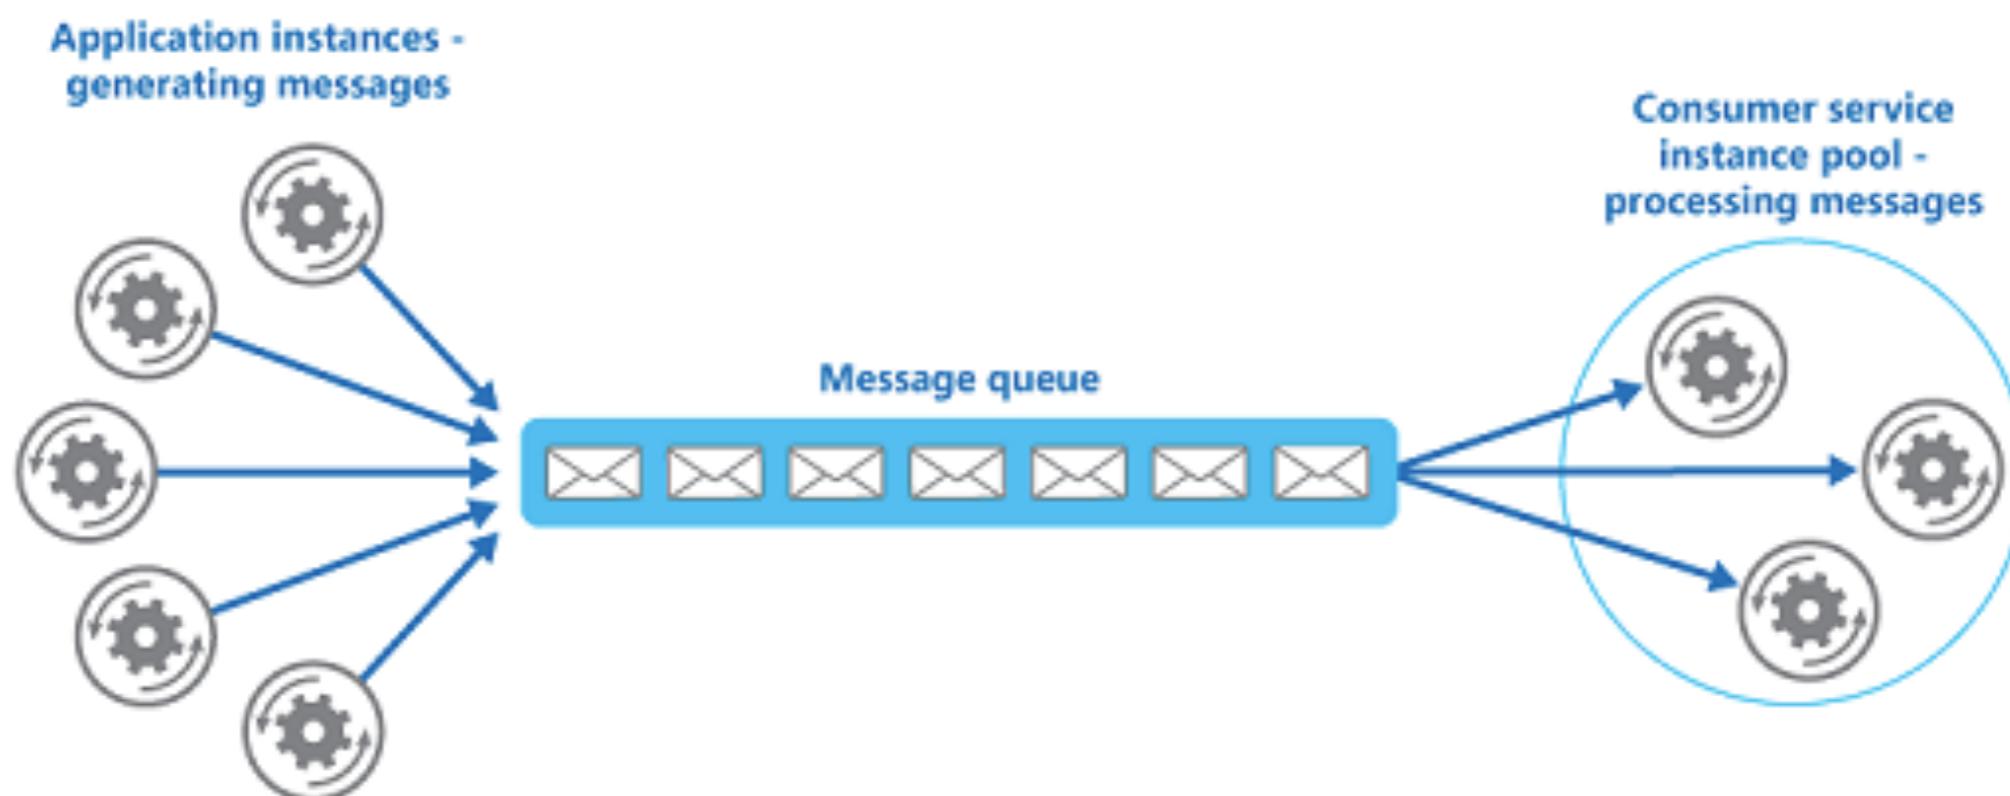


Queue-based Load Levelling



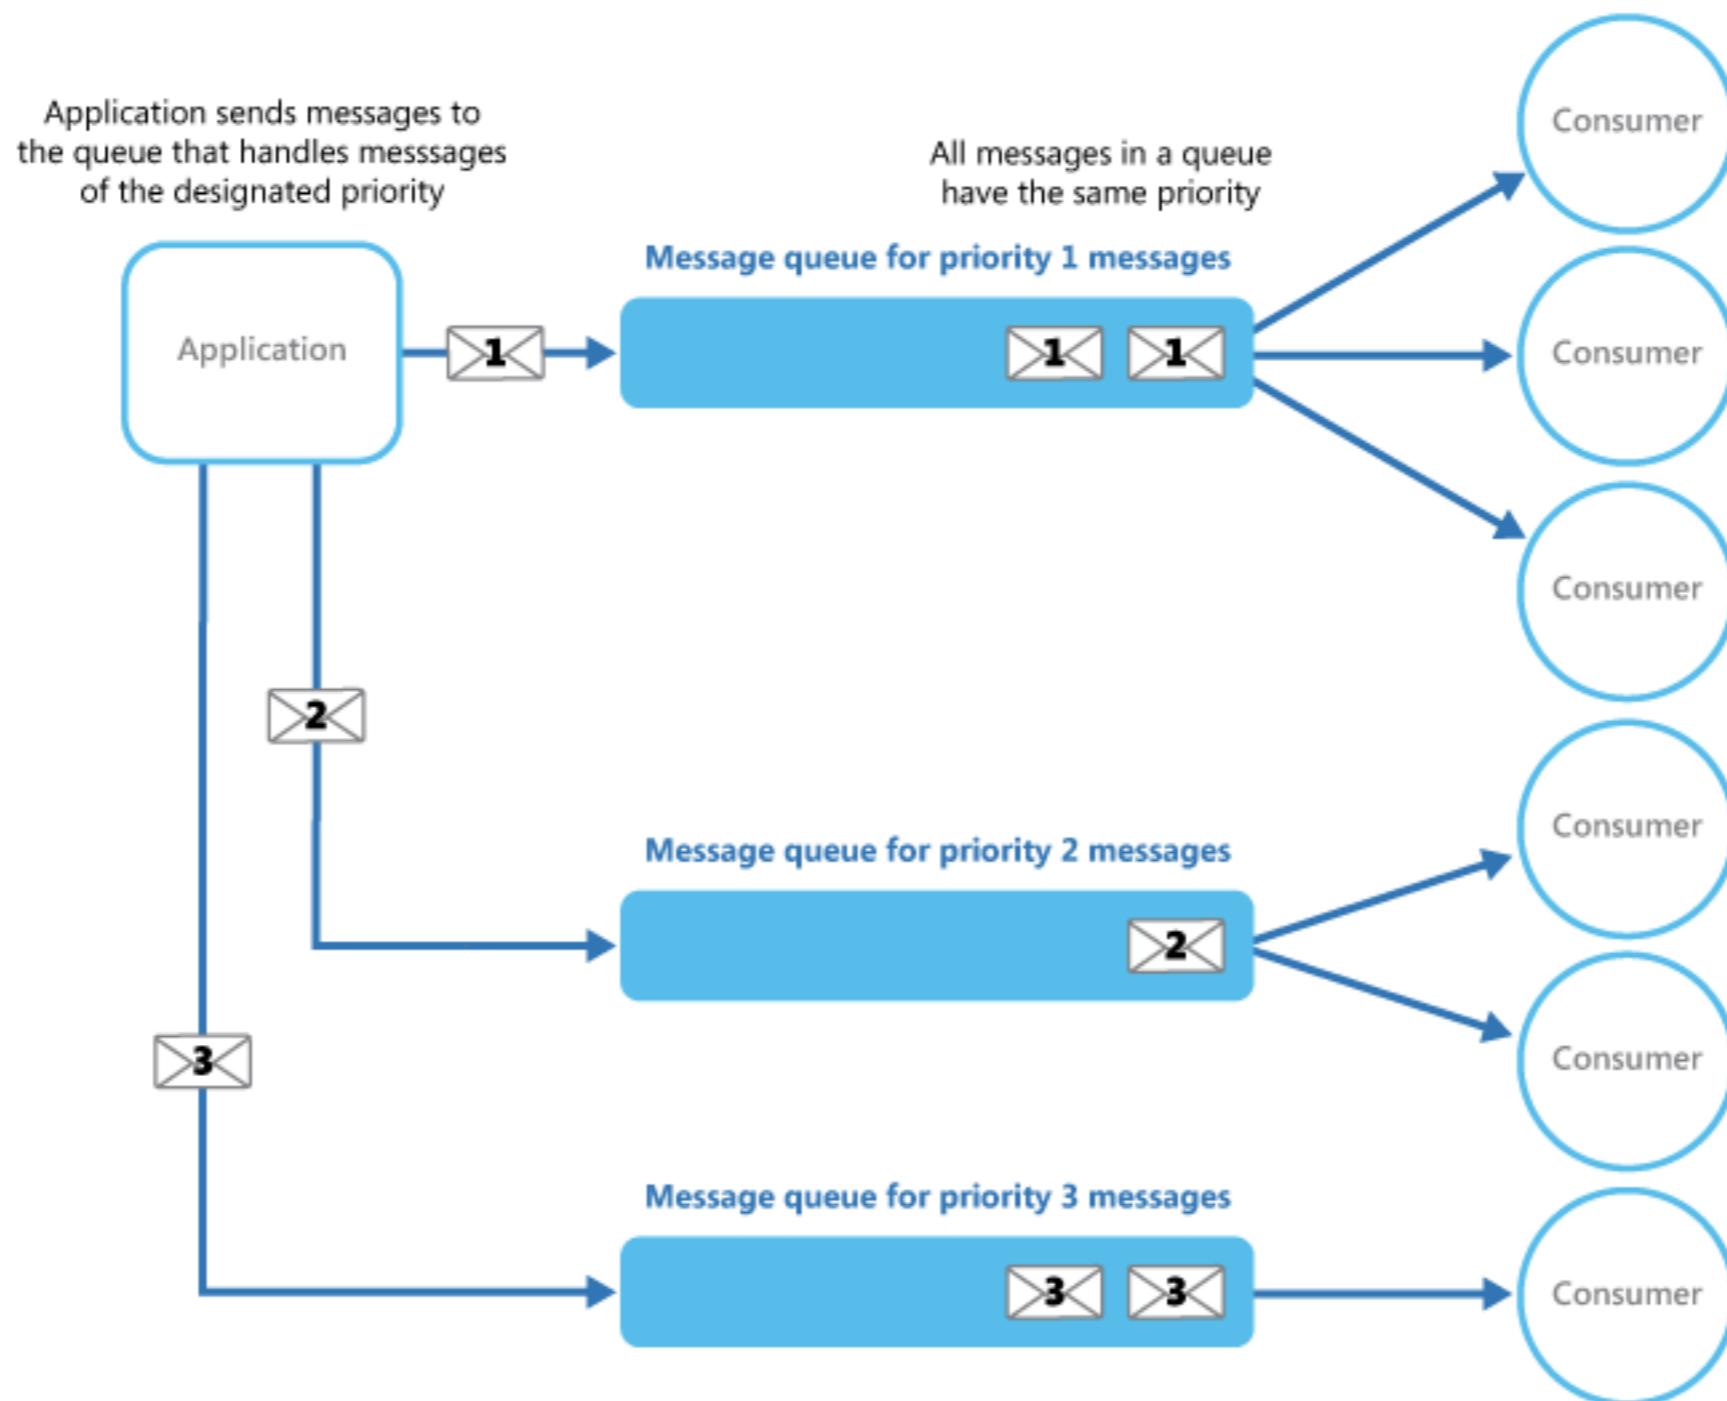
source:msdn

Competing Consumers



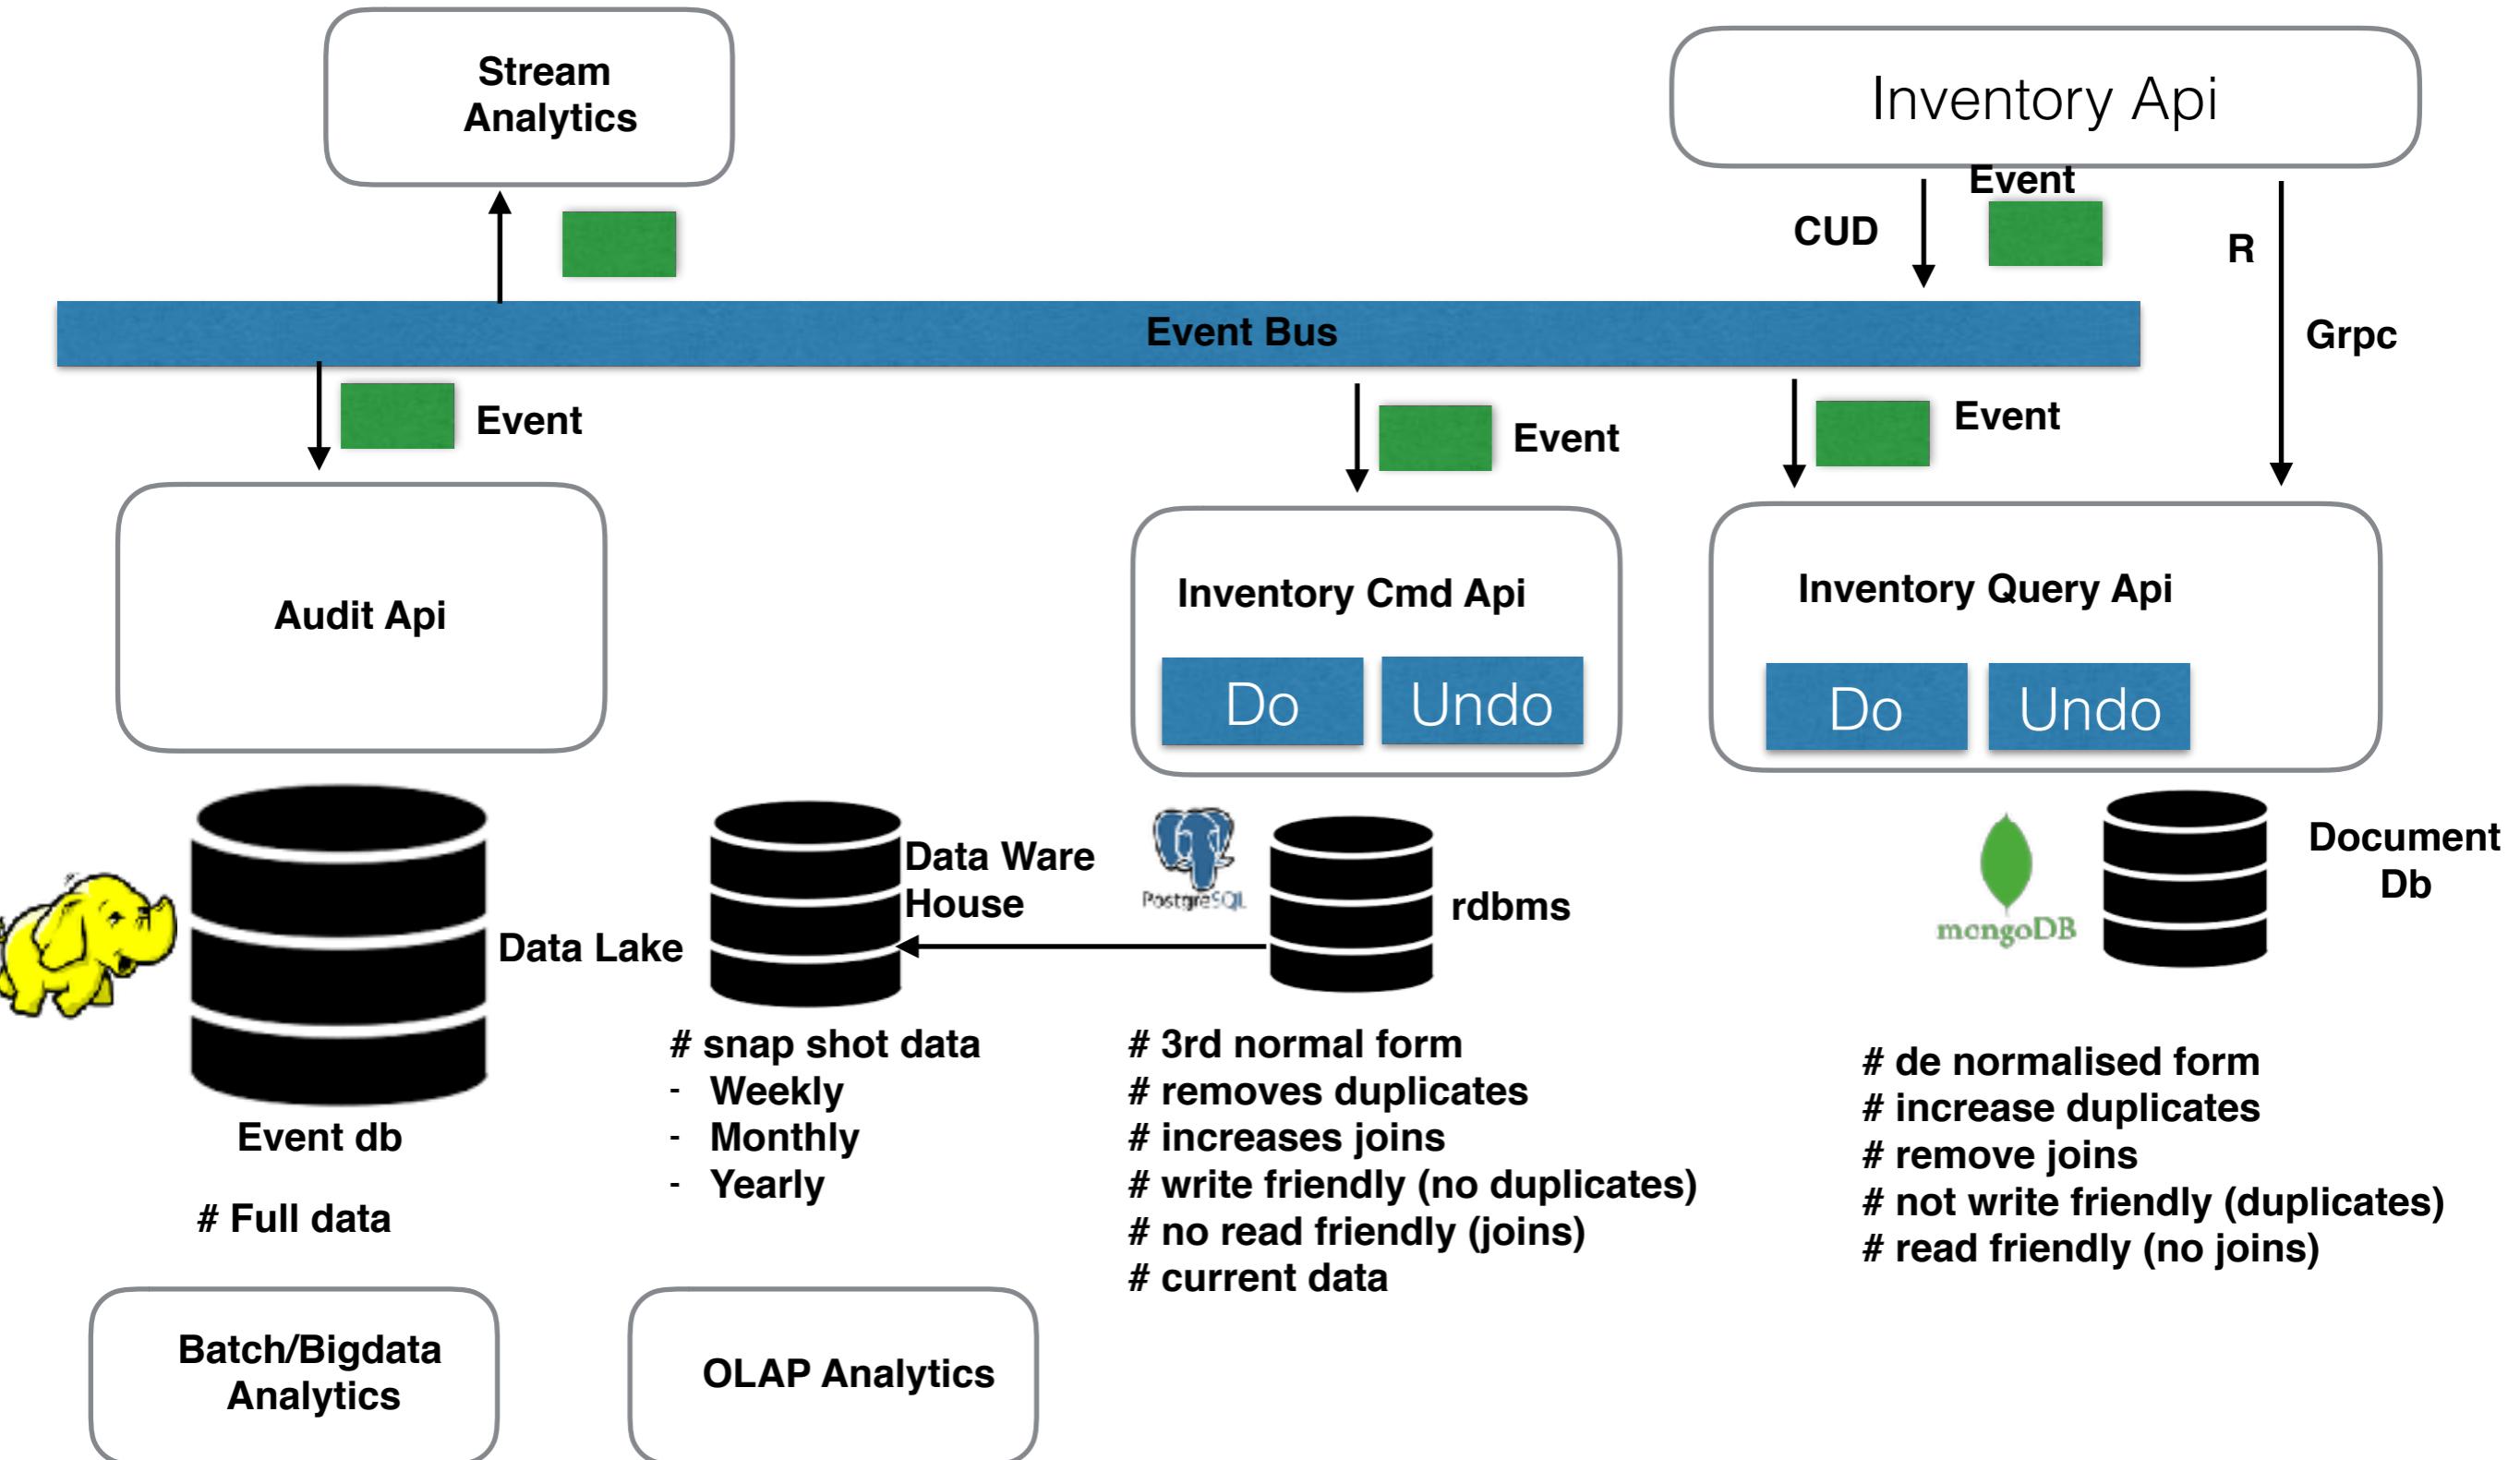
source:msdn

Priority Queue



source:msdn

```
# Event Bus  
# EDA  
# compensatale transaction  
# SAGA  
# Event Sourcing  
# CQRS
```



TABLE_BOOK

Book ID	Genre ID	Price
1	1	25.99
2	2	14.99
3	1	10.00
4	3	12.99
5	2	17.99

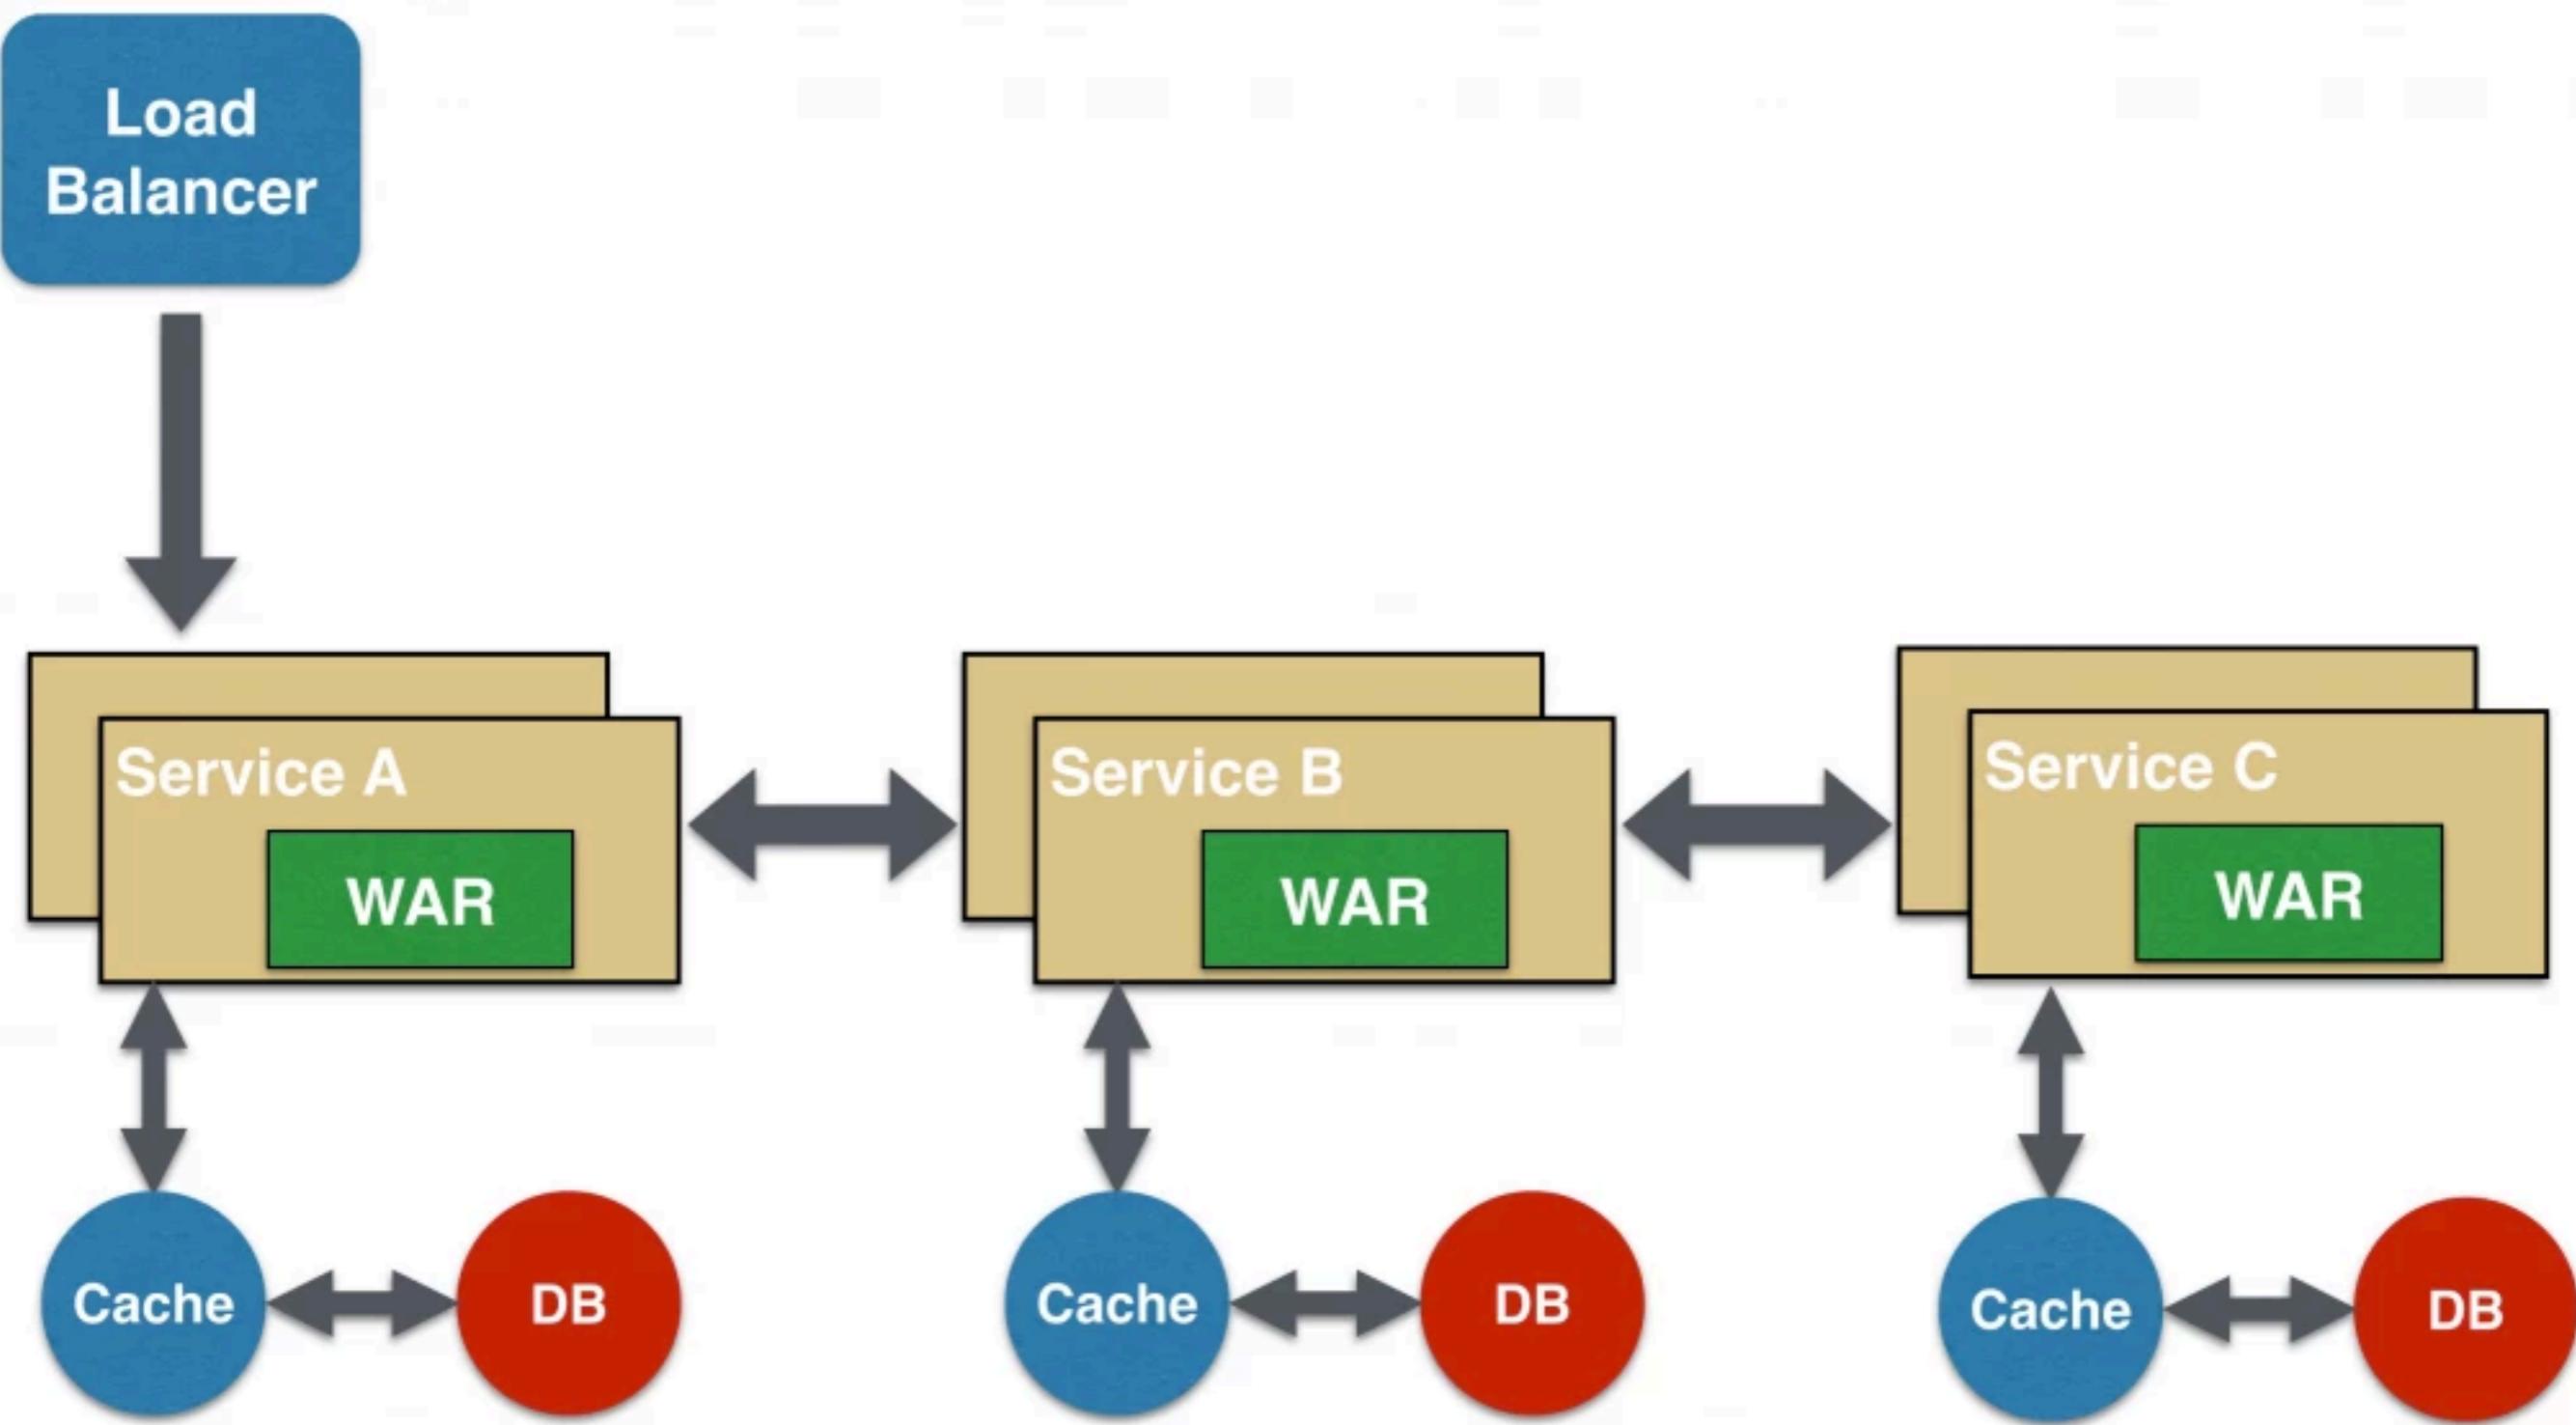
TABLE_GENRE

Genre ID	Genre Type
1	Gardening
2	Sports
3	Travel

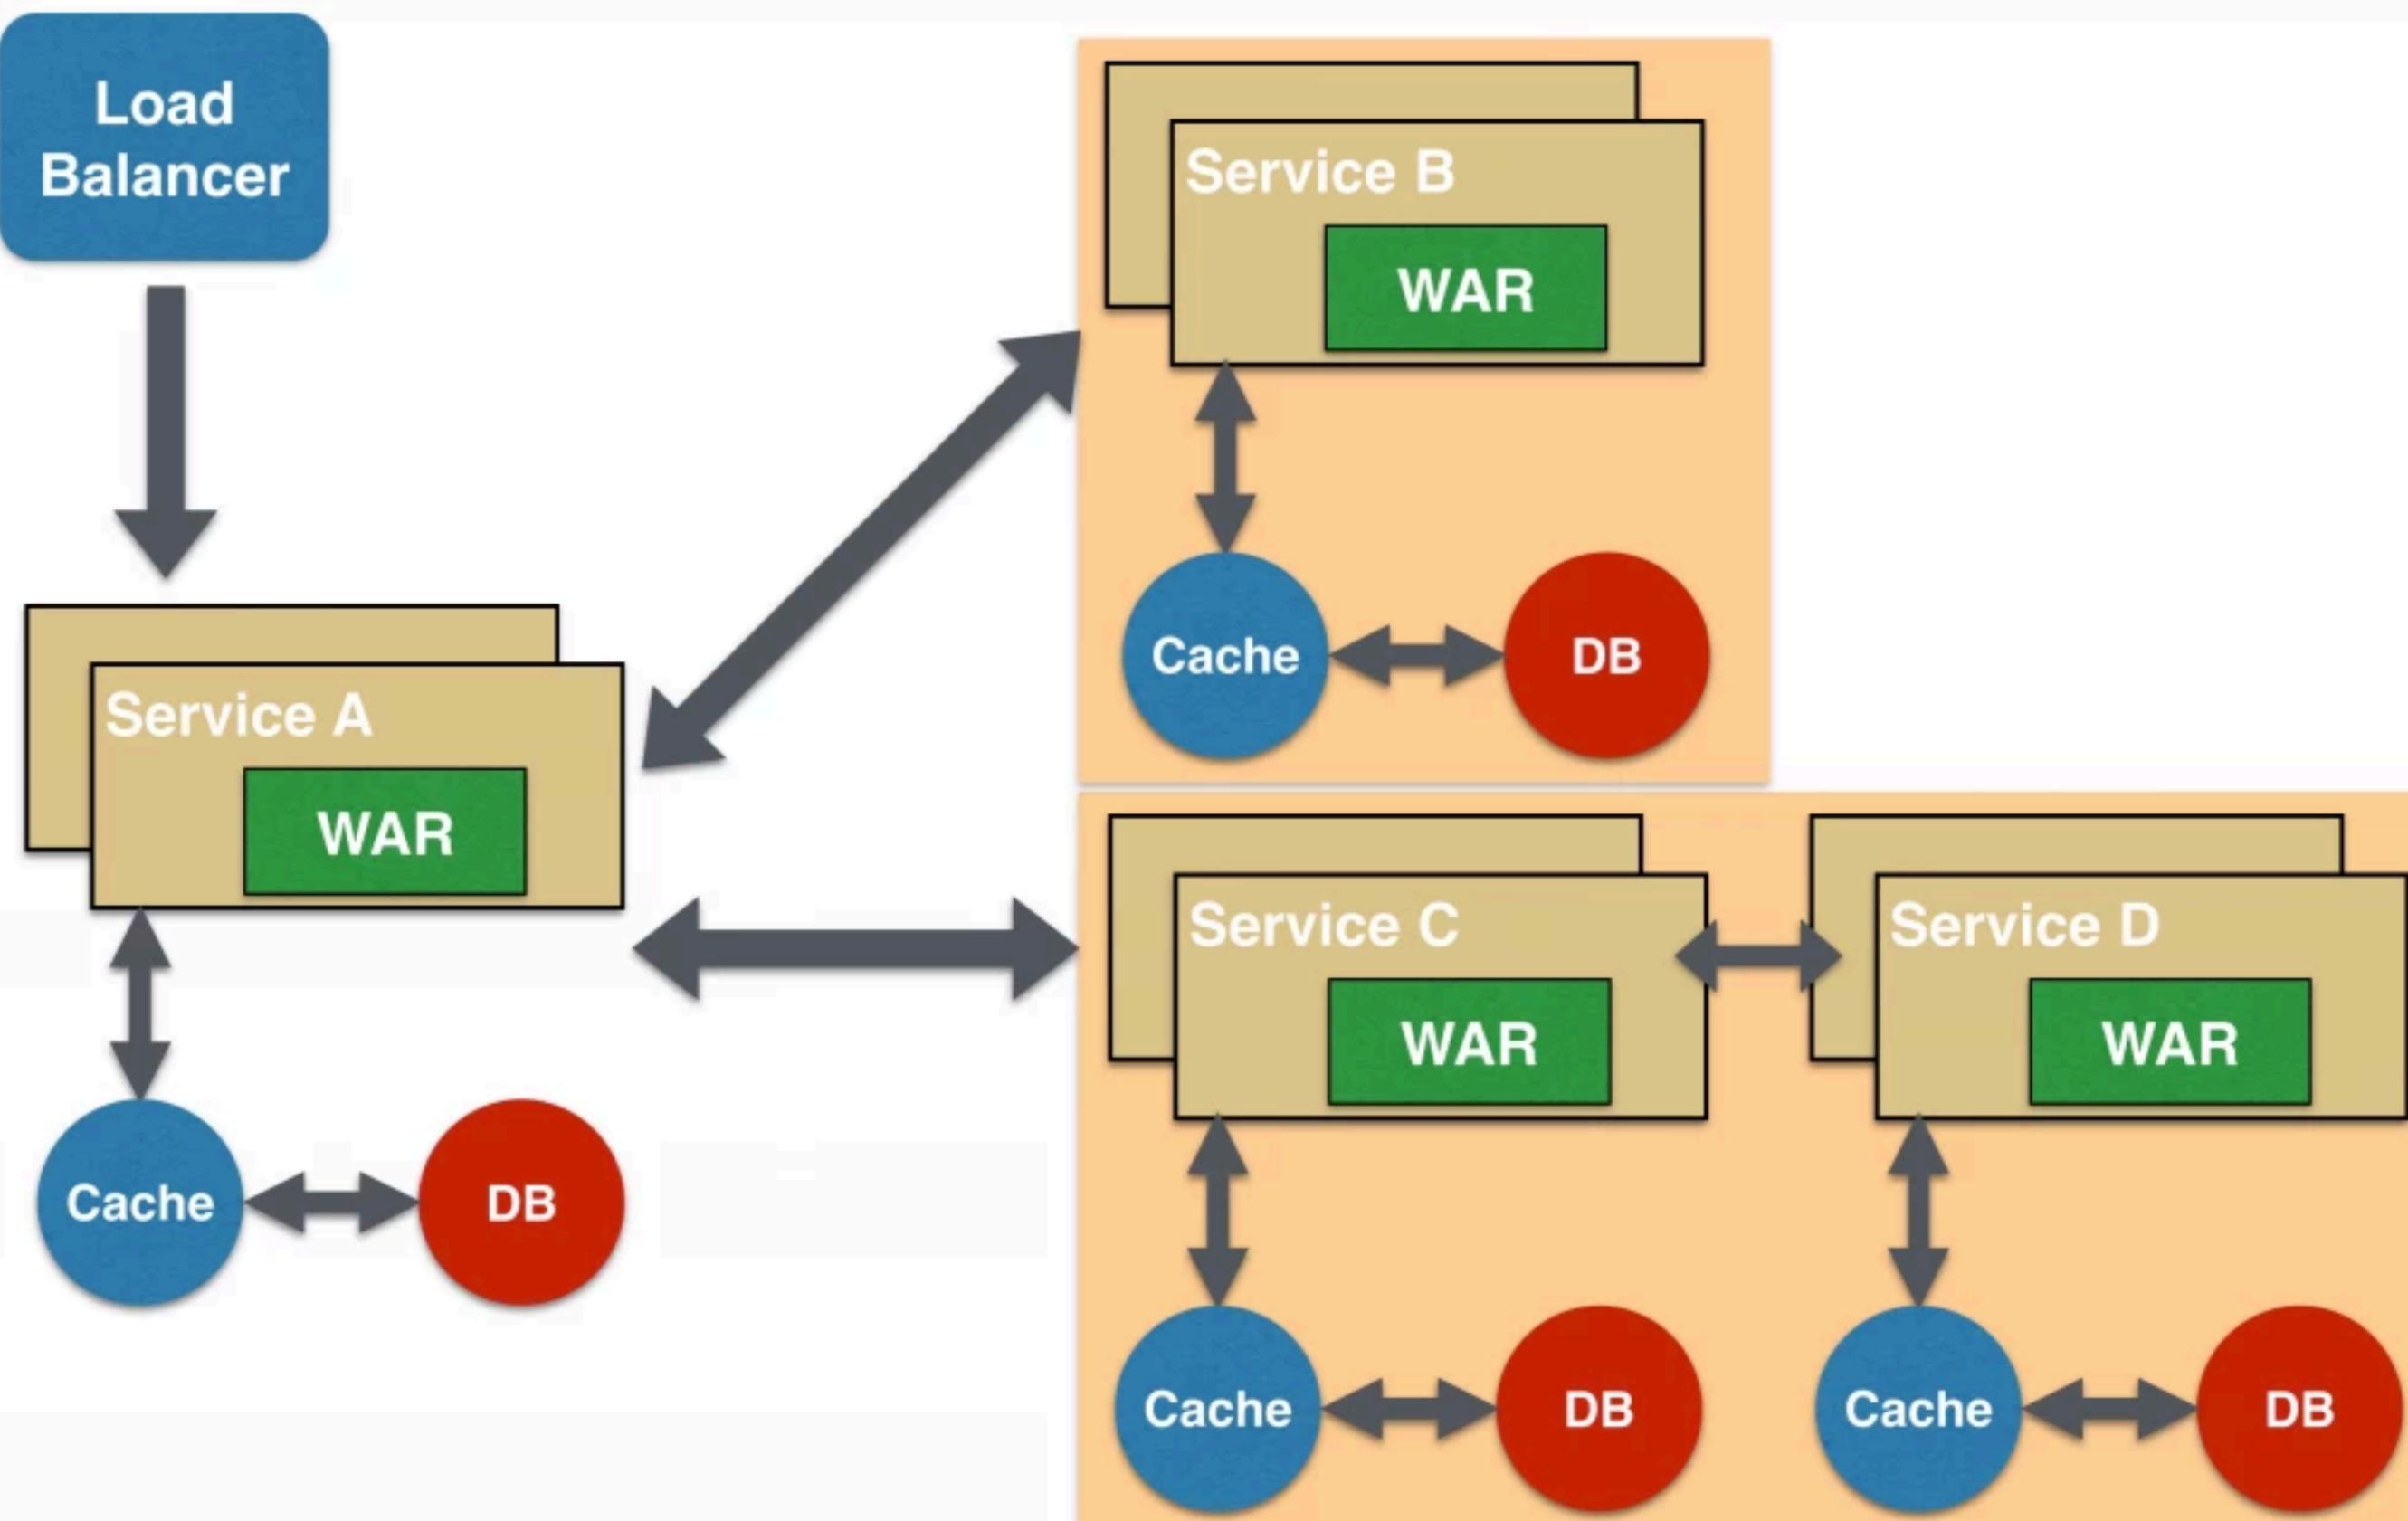
Weather

city	state	high	low
Phoenix	Arizona	105	90
Tucson	Arizona	101	92
Flagstaff	Arizona	88	69
San Diego	California	77	60
Albuquerque	New Mexico	80	72

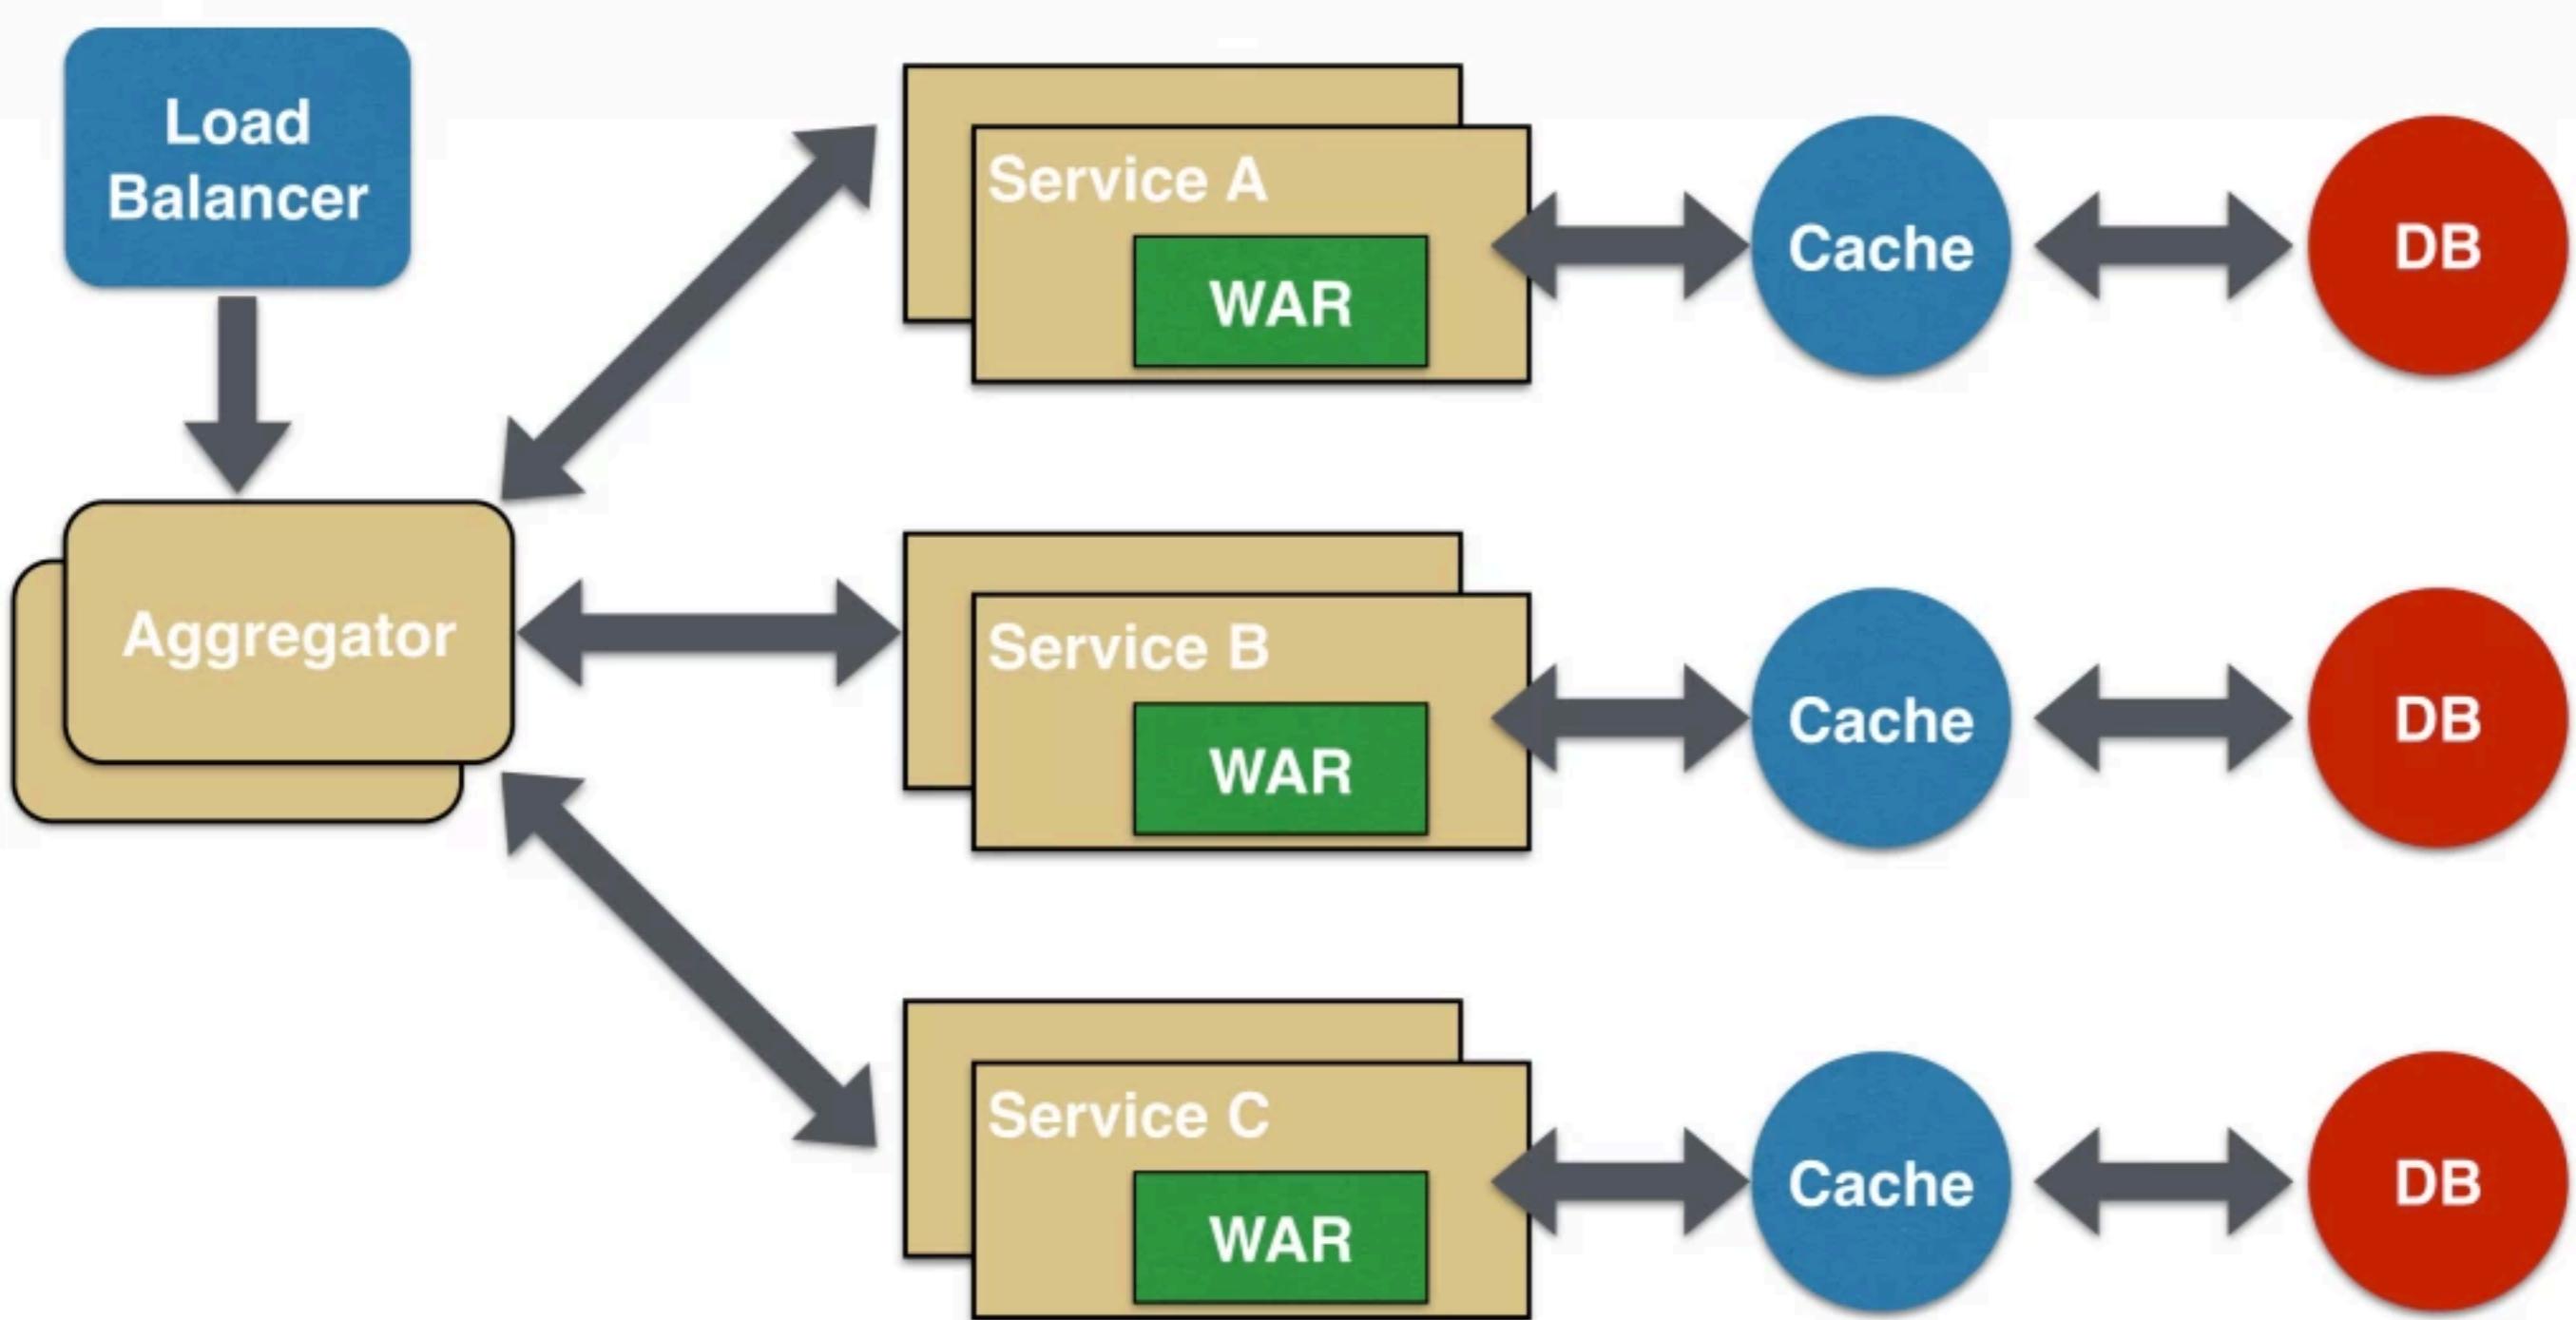
Chained Microservice

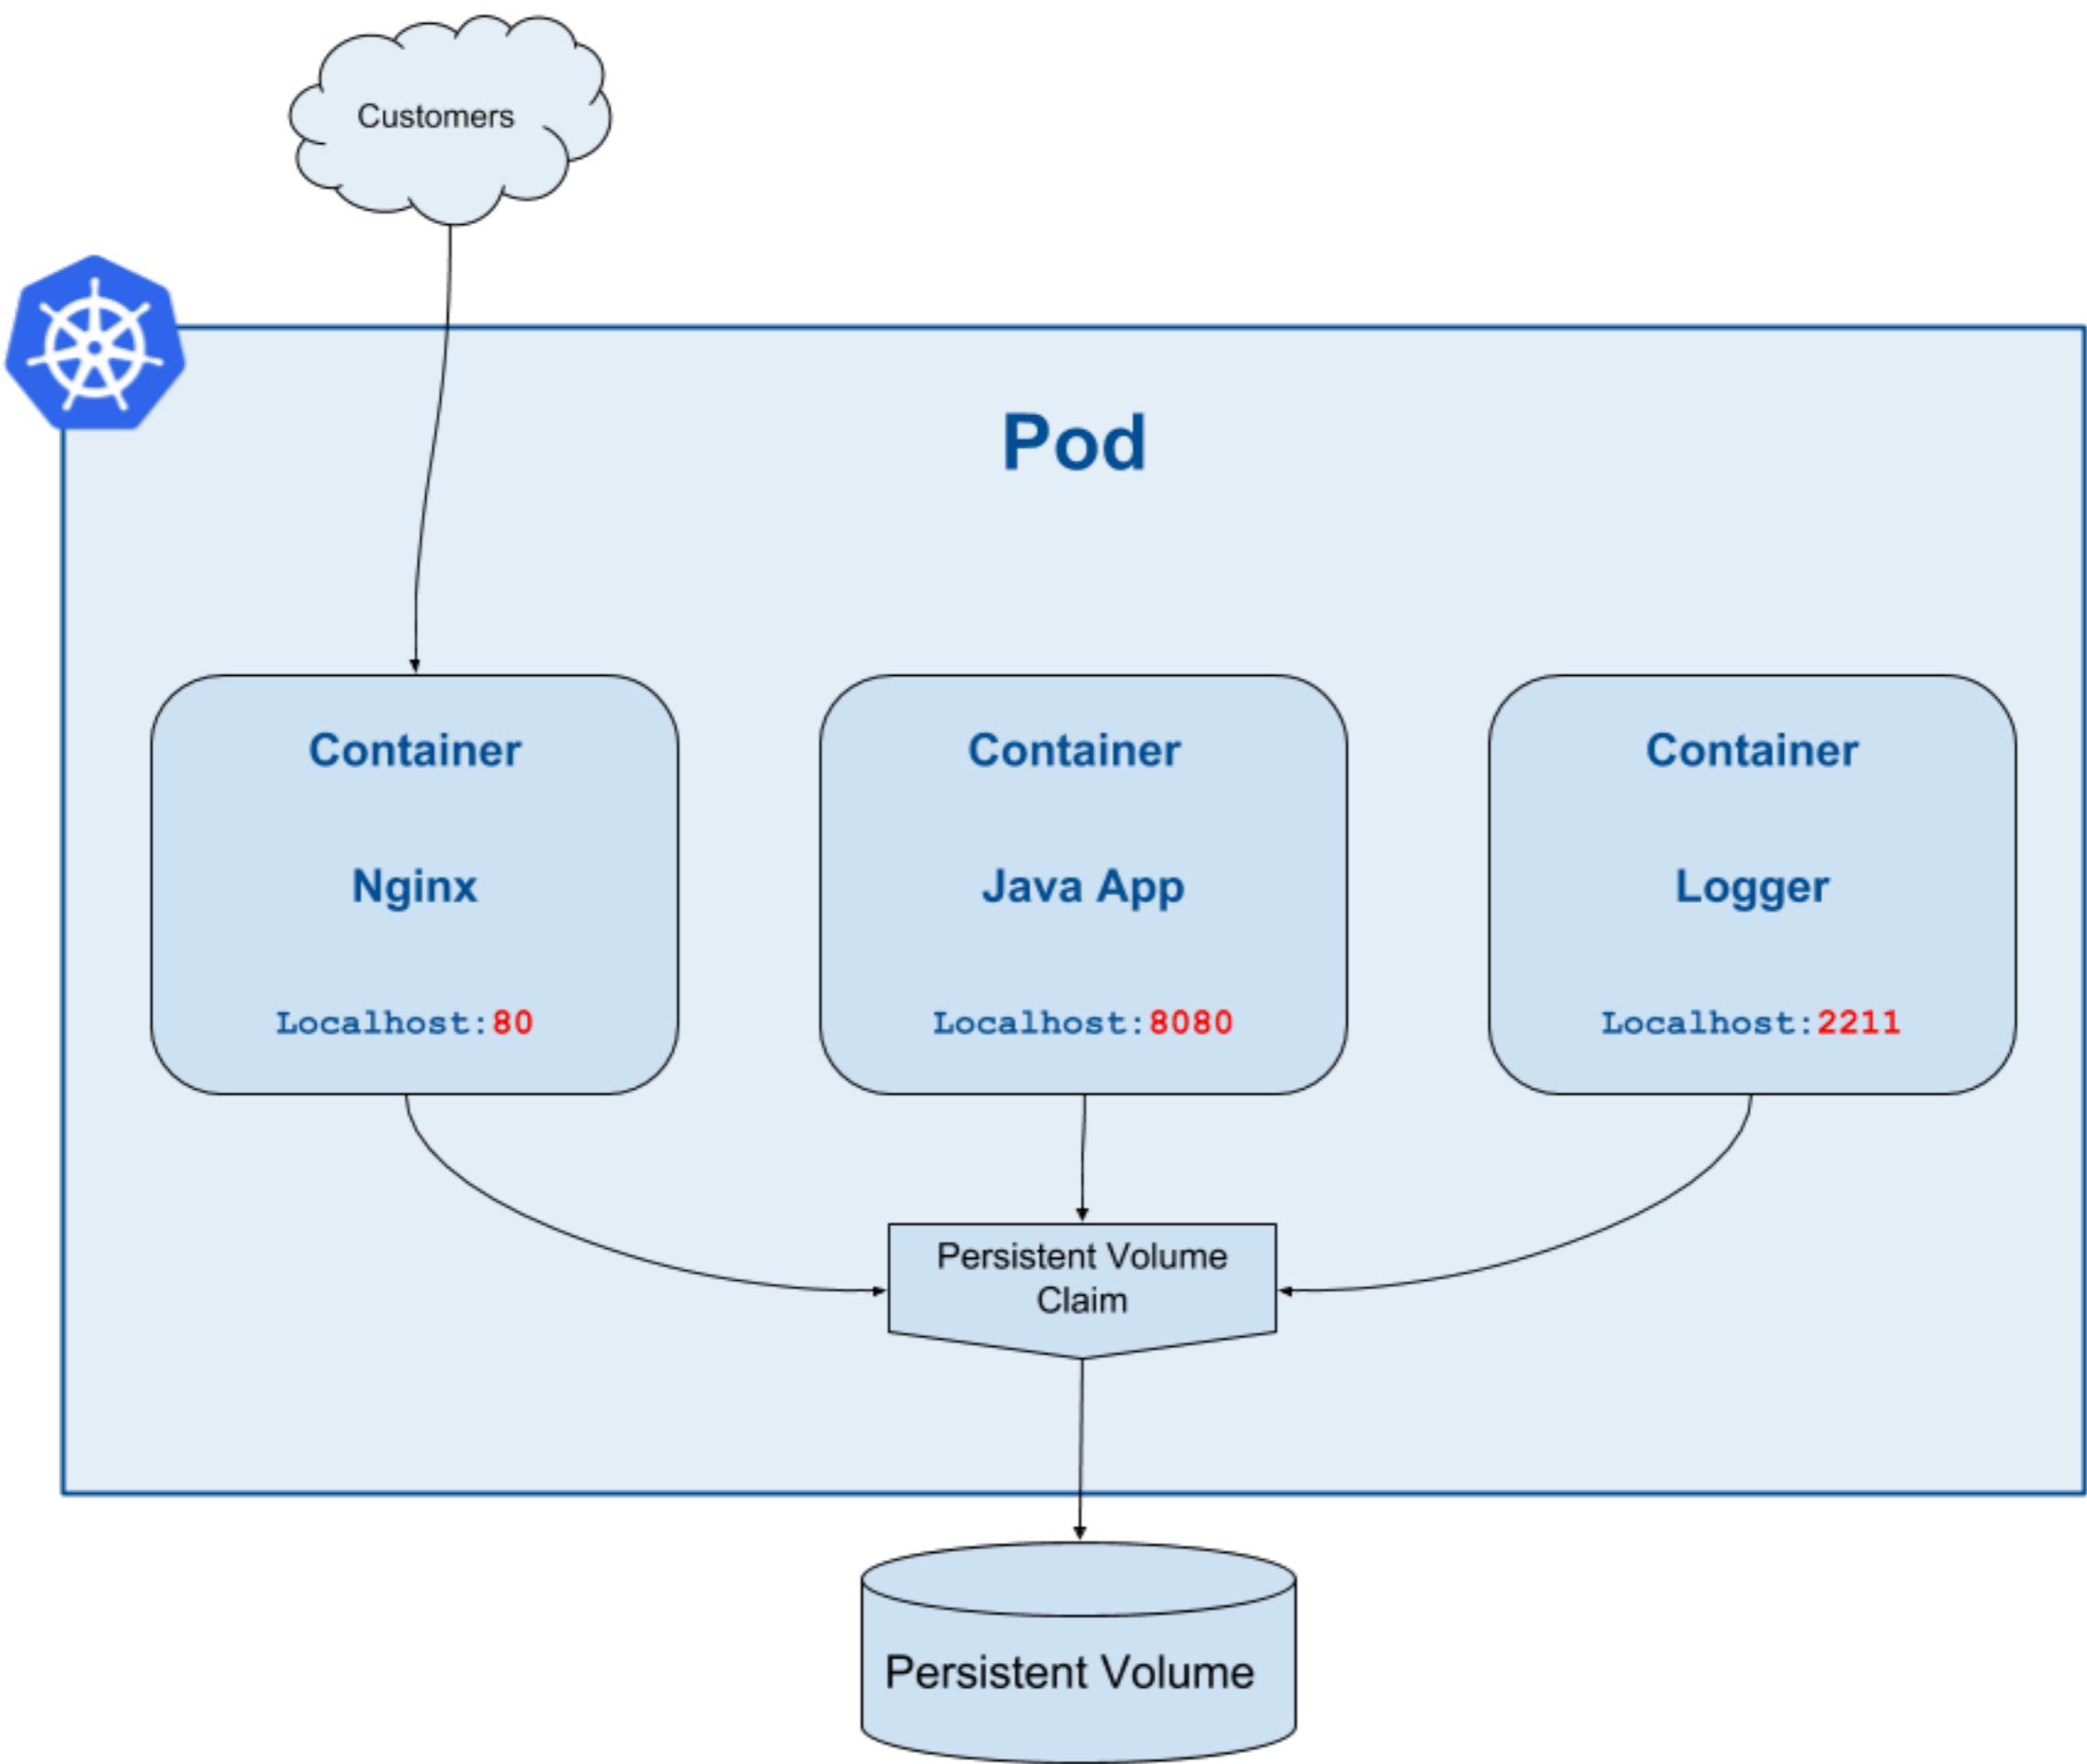


Branch Microservice



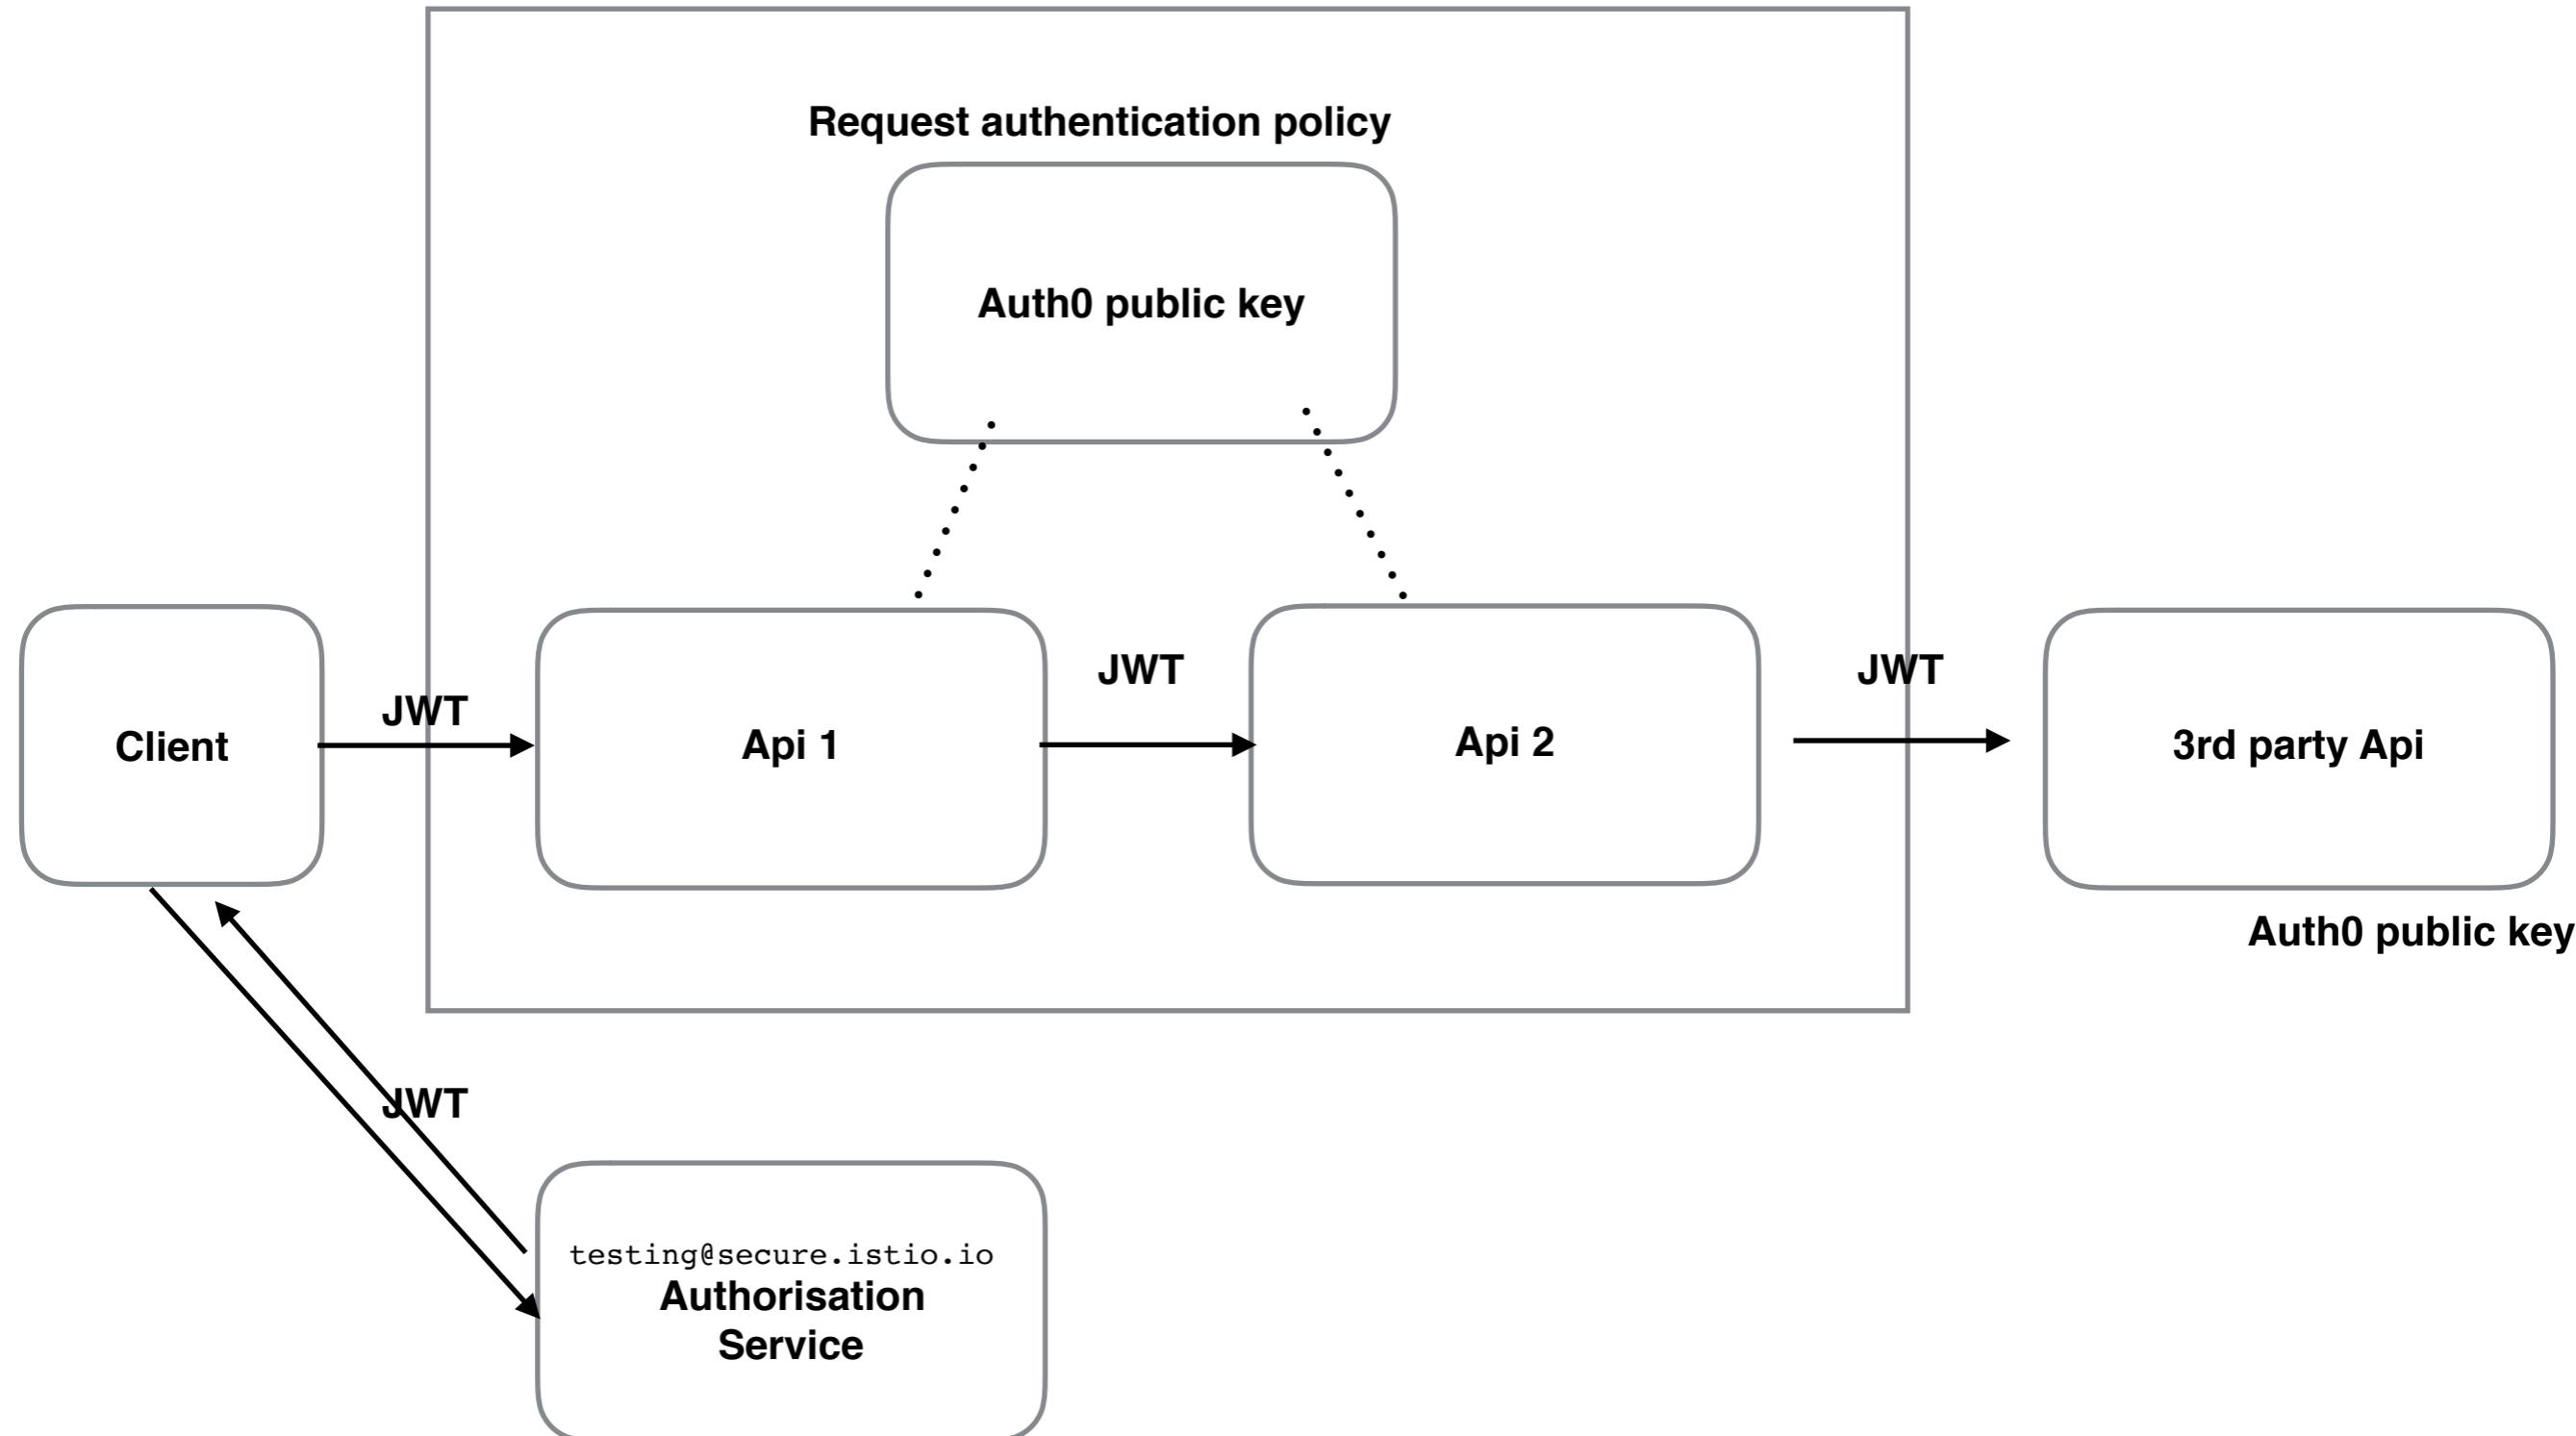
Aggregator Microservice





Concern	Approach
Package runtime dependencies Reproducible Environment “it works on my machine”	Docker image
Multiple Service Deployment Isolation	Containers
Simplify Deployment of many services	yaml, Helm
Distributed Deployment, scale	Kubernetes
Observability	Kubernetes dash board, weave scope, Kiali, graffana, kibana
Service Discovery	Cluster-IP, Node Port, Load Balancer, Ingress Gateway
Traffic Management (timeouts, retry, circuit breaker, throttle, canary, ...)	Service Mesh (Istio)
Log Management (debug & error)	EFK
Traceability	Prometheus, jager
Transaction	SAGA
Join Query (report query)	CQRS , Materialized View
Configuration	ConfigMap, Consul, KeyValut, ..
Performance	Shared db, grpc, Async, branch pattern
Authentication	OAuth2, JWT, OpenID connect

Oath 2



Appendix

- Applications run in *containers*, which in turn run inside of *Pods*.
- All Pods in your Kubernetes cluster have their own IP address and are attached to the same *network*.
- This means all Pods can talk directly to all other Pods.
- However, Pods are unreliable and come and go as scaling operations, rolling updates, rollbacks, failures and other events occur

- Each pod has its own IP address. And each pod thinks it has a totally normal ethernet device called `eth0` to make network requests through.
- When a pod makes a request to the IP address of another node, it makes that request through its own `eth0` interface. This tunnels to the node's respective virtual `vethX` interface.
- A network bridge connects two networks together. When a request hits the bridge, the bridge asks all the connected pods if they have the right IP address to handle the original request.
- In Kubernetes, this bridge is called `cbr0`. Every pod on a node is part of the bridge, and the bridge connects all pods on the same node together.
-

- when the network bridge asks all the pods if they have the right IP address, none of them will say yes.
- After that, the bridge falls back to the default gateway. This goes up to the cluster level and looks for the IP address.
- At the cluster level, there's a table that maps IP address ranges to various nodes. Pods on those nodes will have been assigned IP addresses from those ranges.
- For example, Kubernetes might give pods on node 1 addresses like 100.96.1.1, 100.96.1.2, etc. And Kubernetes gives pods on node 2 addresses like 100.96.2.1, 100.96.2.2, and so on.
- Then this table will store the fact that IP addresses that look like 100.96.1.xxx should go to node 1, and addresses like 100.96.2.xxx need to go to node 2.

- A Service is bound to a ClusterIP, which is a virtual IP address, and no matter what happens to the backend Pods, the ClusterIP never changes, so a client can always send requests to the ClusterIP of the Service.
- A Kubernetes Service object creates a stable network endpoint that sits in front of a set of Pods and load-balances traffic across them.
- You always put a Service in front of a set of Pods that do the same job. For example, you could put a Service in front of your web front-end Pods, and another in front of your authentication Pods. You never put a Service in front of Pods that do different jobs.
- Every service in a cluster is assigned a domain name like my-service.my-namespace.svc.cluster.local.
- when a request is made to a service via its domain name, the DNS service resolves it to the IP address of the service.
- Then kube-proxy converts that service's IP address into a pod IP address.

- You POST a new Service definition to the API Server
- The Service is allocated a ClusterIP (virtual IP address) and persisted to the cluster store
- The cluster's DNS service notices the new Service and creates the necessary DNS A records

```
apiVersion: v1
kind: Service
metadata:
  name: web-svc
labels:
  blog: svc-discovery
spec:
  type: LoadBalancer
```

It's important to understand that the name registered with DNS is the value of **metadata.name** and that the ClusterIP is dynamically assigned by Kubernetes.

- With service ClusterIP and Kubernetes DNS, service can be easily reached inside a cluster
- If you want more advanced features, such as flexible routing rules, more options for LB, reliable service communication, metrics collection and distributed tracing, etc., then you will need to consider Istio.
- The communication between services is no longer through Kube-proxy but through Istio's sidecar proxies. Istio sidecar proxy works just like Kube-proxy.

- ClusterIP is only reachable inside a Kubernetes cluster, but what if we need to access some services from outside of the cluster?
- With NodePort, Kubernetes creates a port for a Service on the host, which allows access to the service from the node network.
-

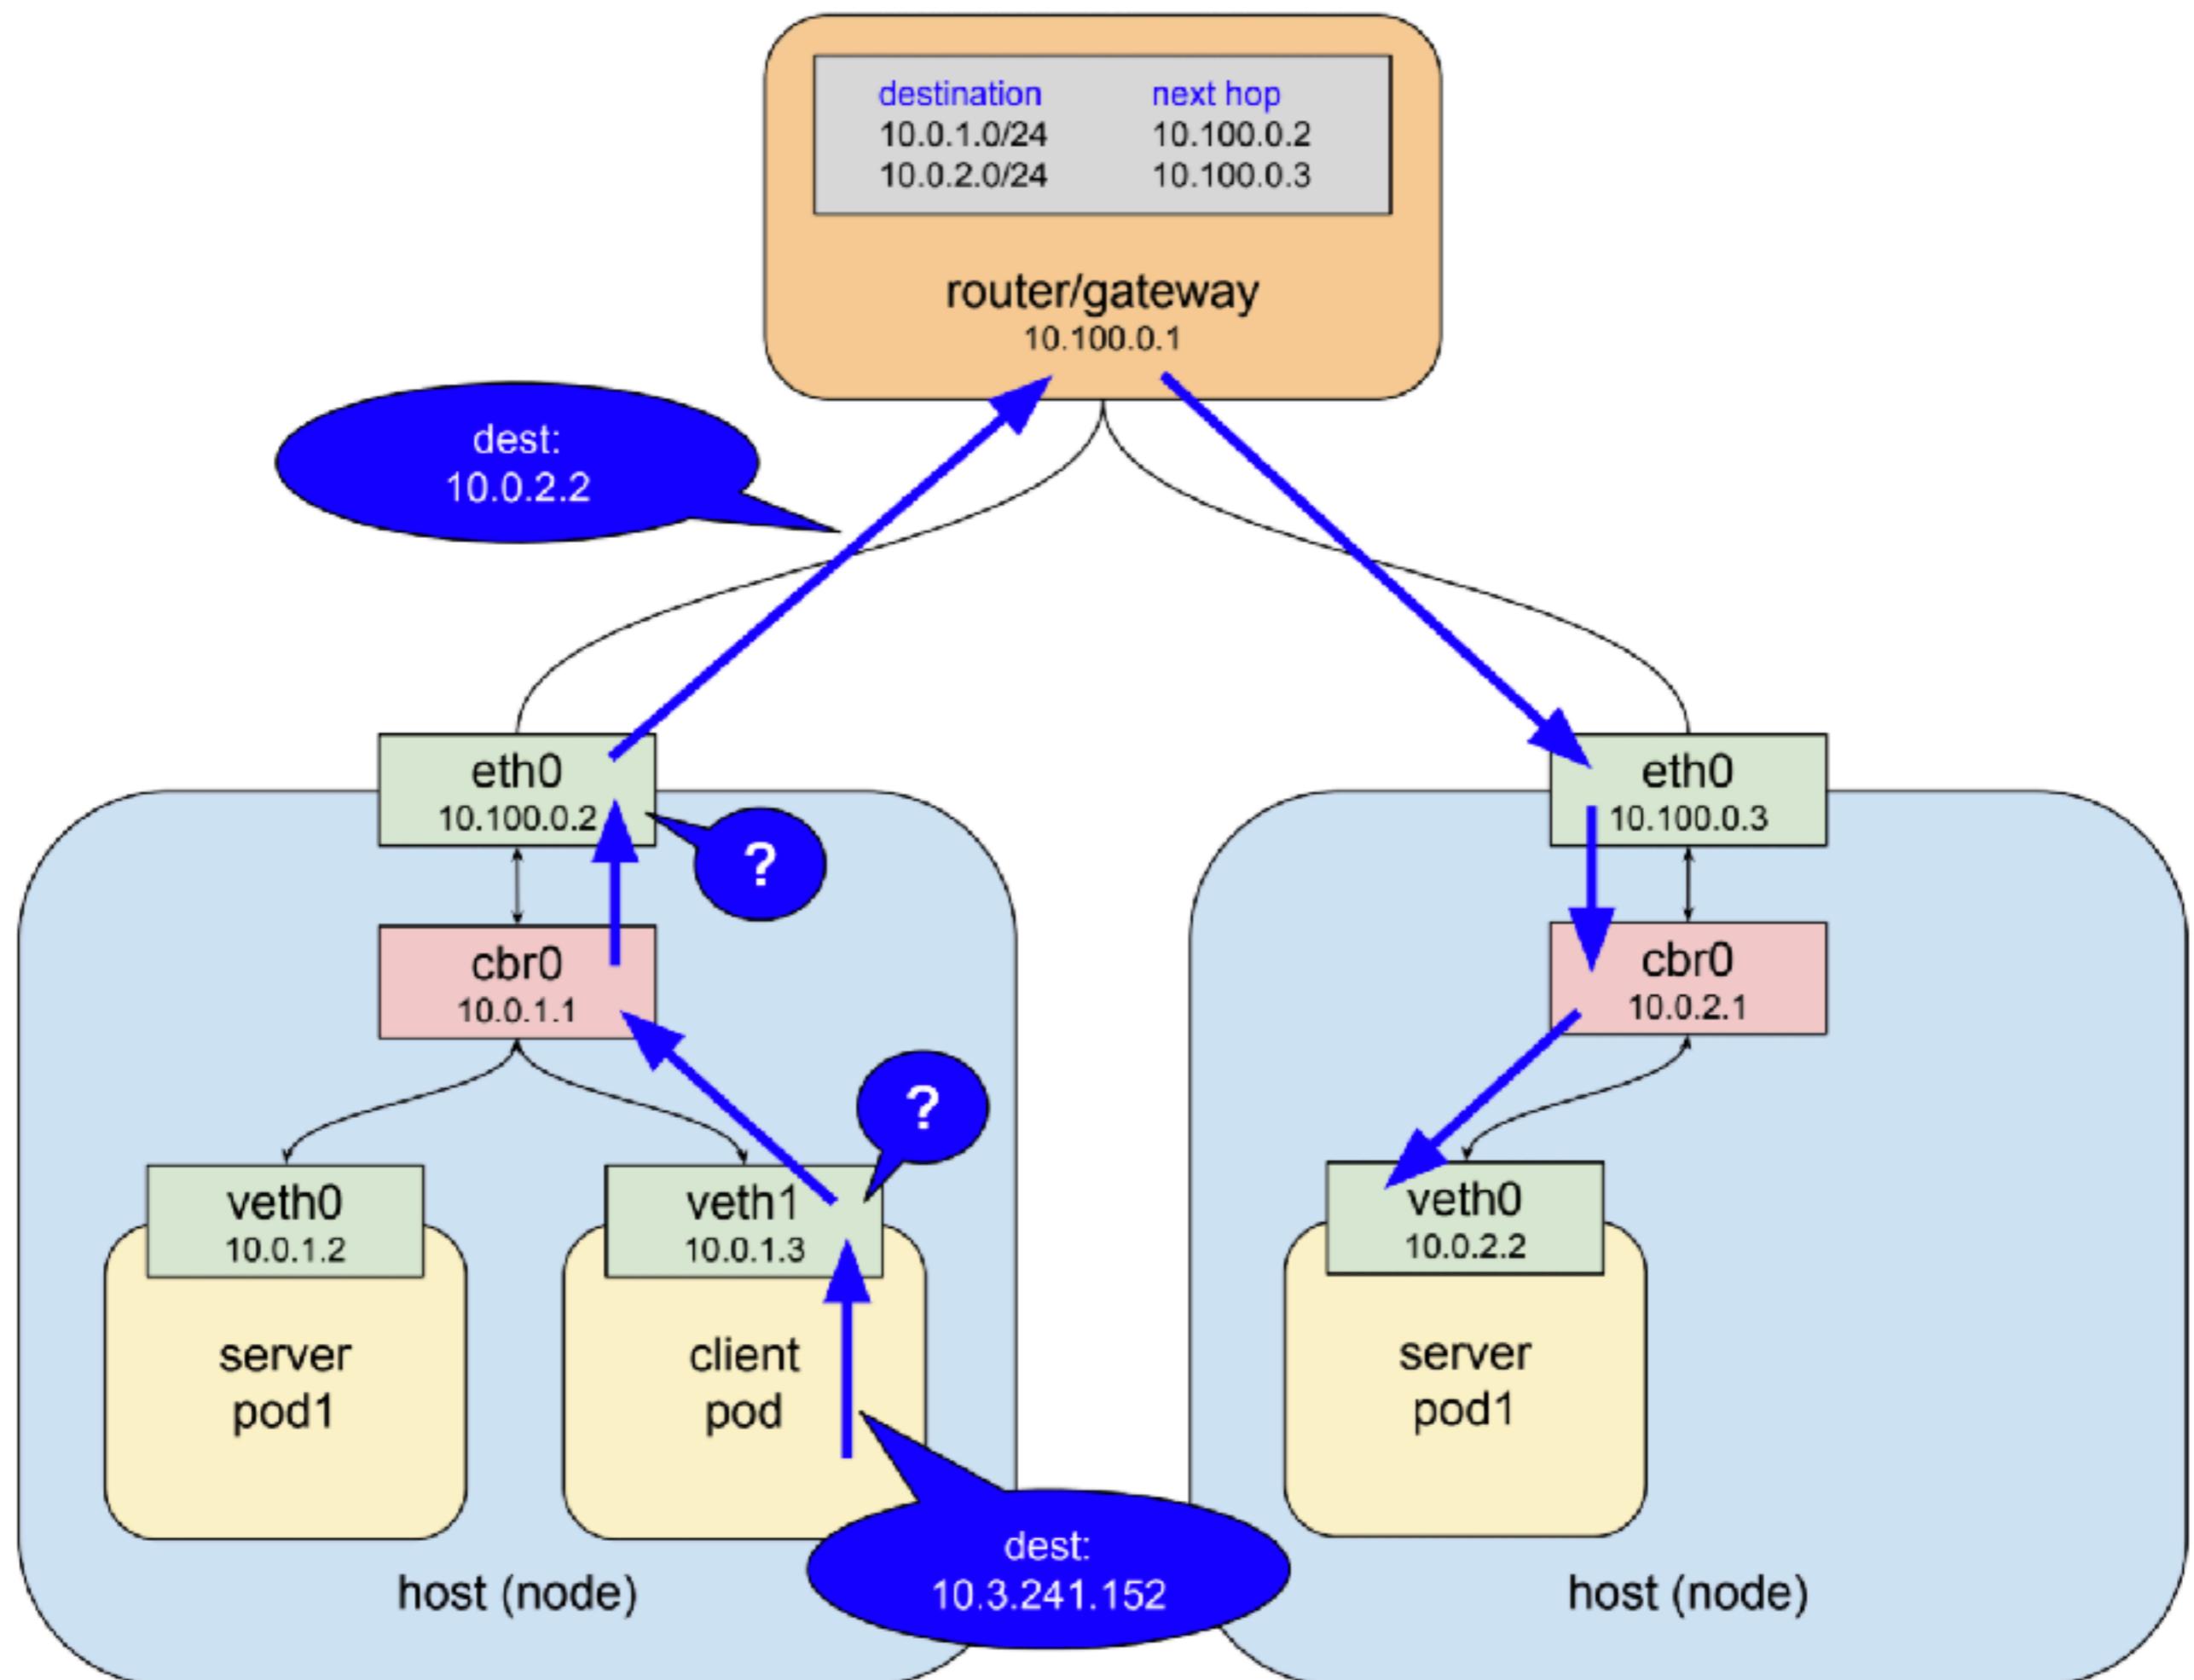
- NodePort is a convenient tool for testing in your local Kubernetes cluster, but it's not suitable for production because of these limitations.
- A single node is a single point of failure for the system. Once the node is down, clients can't access the cluster any more.
- A service can be declared as LoadBalancer type to create a layer 4 load balancer in front of multiple nodes. As this layer 4 load balancer is outside of the Kubernetes network, a Cloud Provider Controller is needed for its provision. This Cloud Provider Controller watches the Kubernetes master for the addition and removal of Service resources and configures a layer 4 load balancer in the cloud provider network to proxy the NodePorts on multiple Kubernetes nodes.

if we need to expose multiple services to the outside of a cluster, we must create a LoadBalancer for each service. However, creating multiple LoadBalancers can cause some problems:

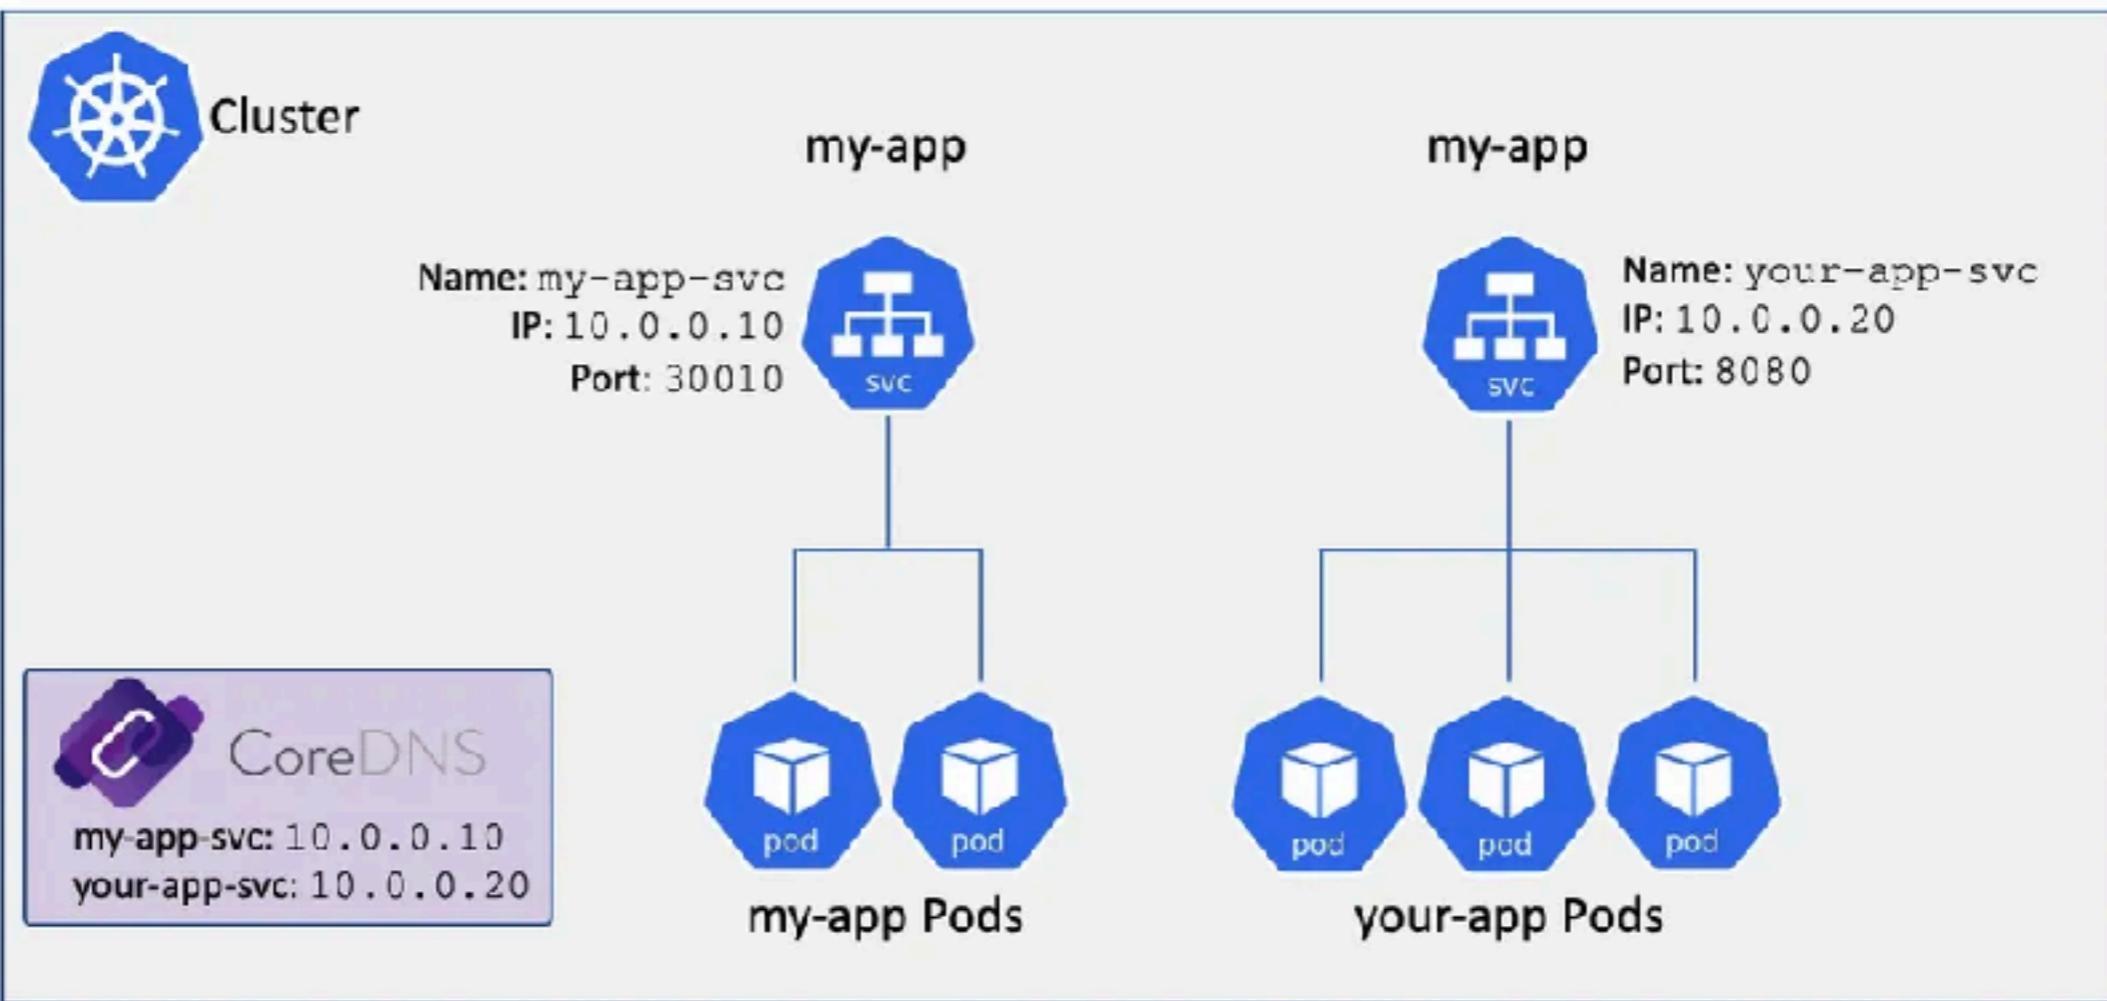
- Needs more public IPs, which normally are limited resources.
- Introduces coupling between the client and the server, making it hard to adjust your backend services when business requirements change.

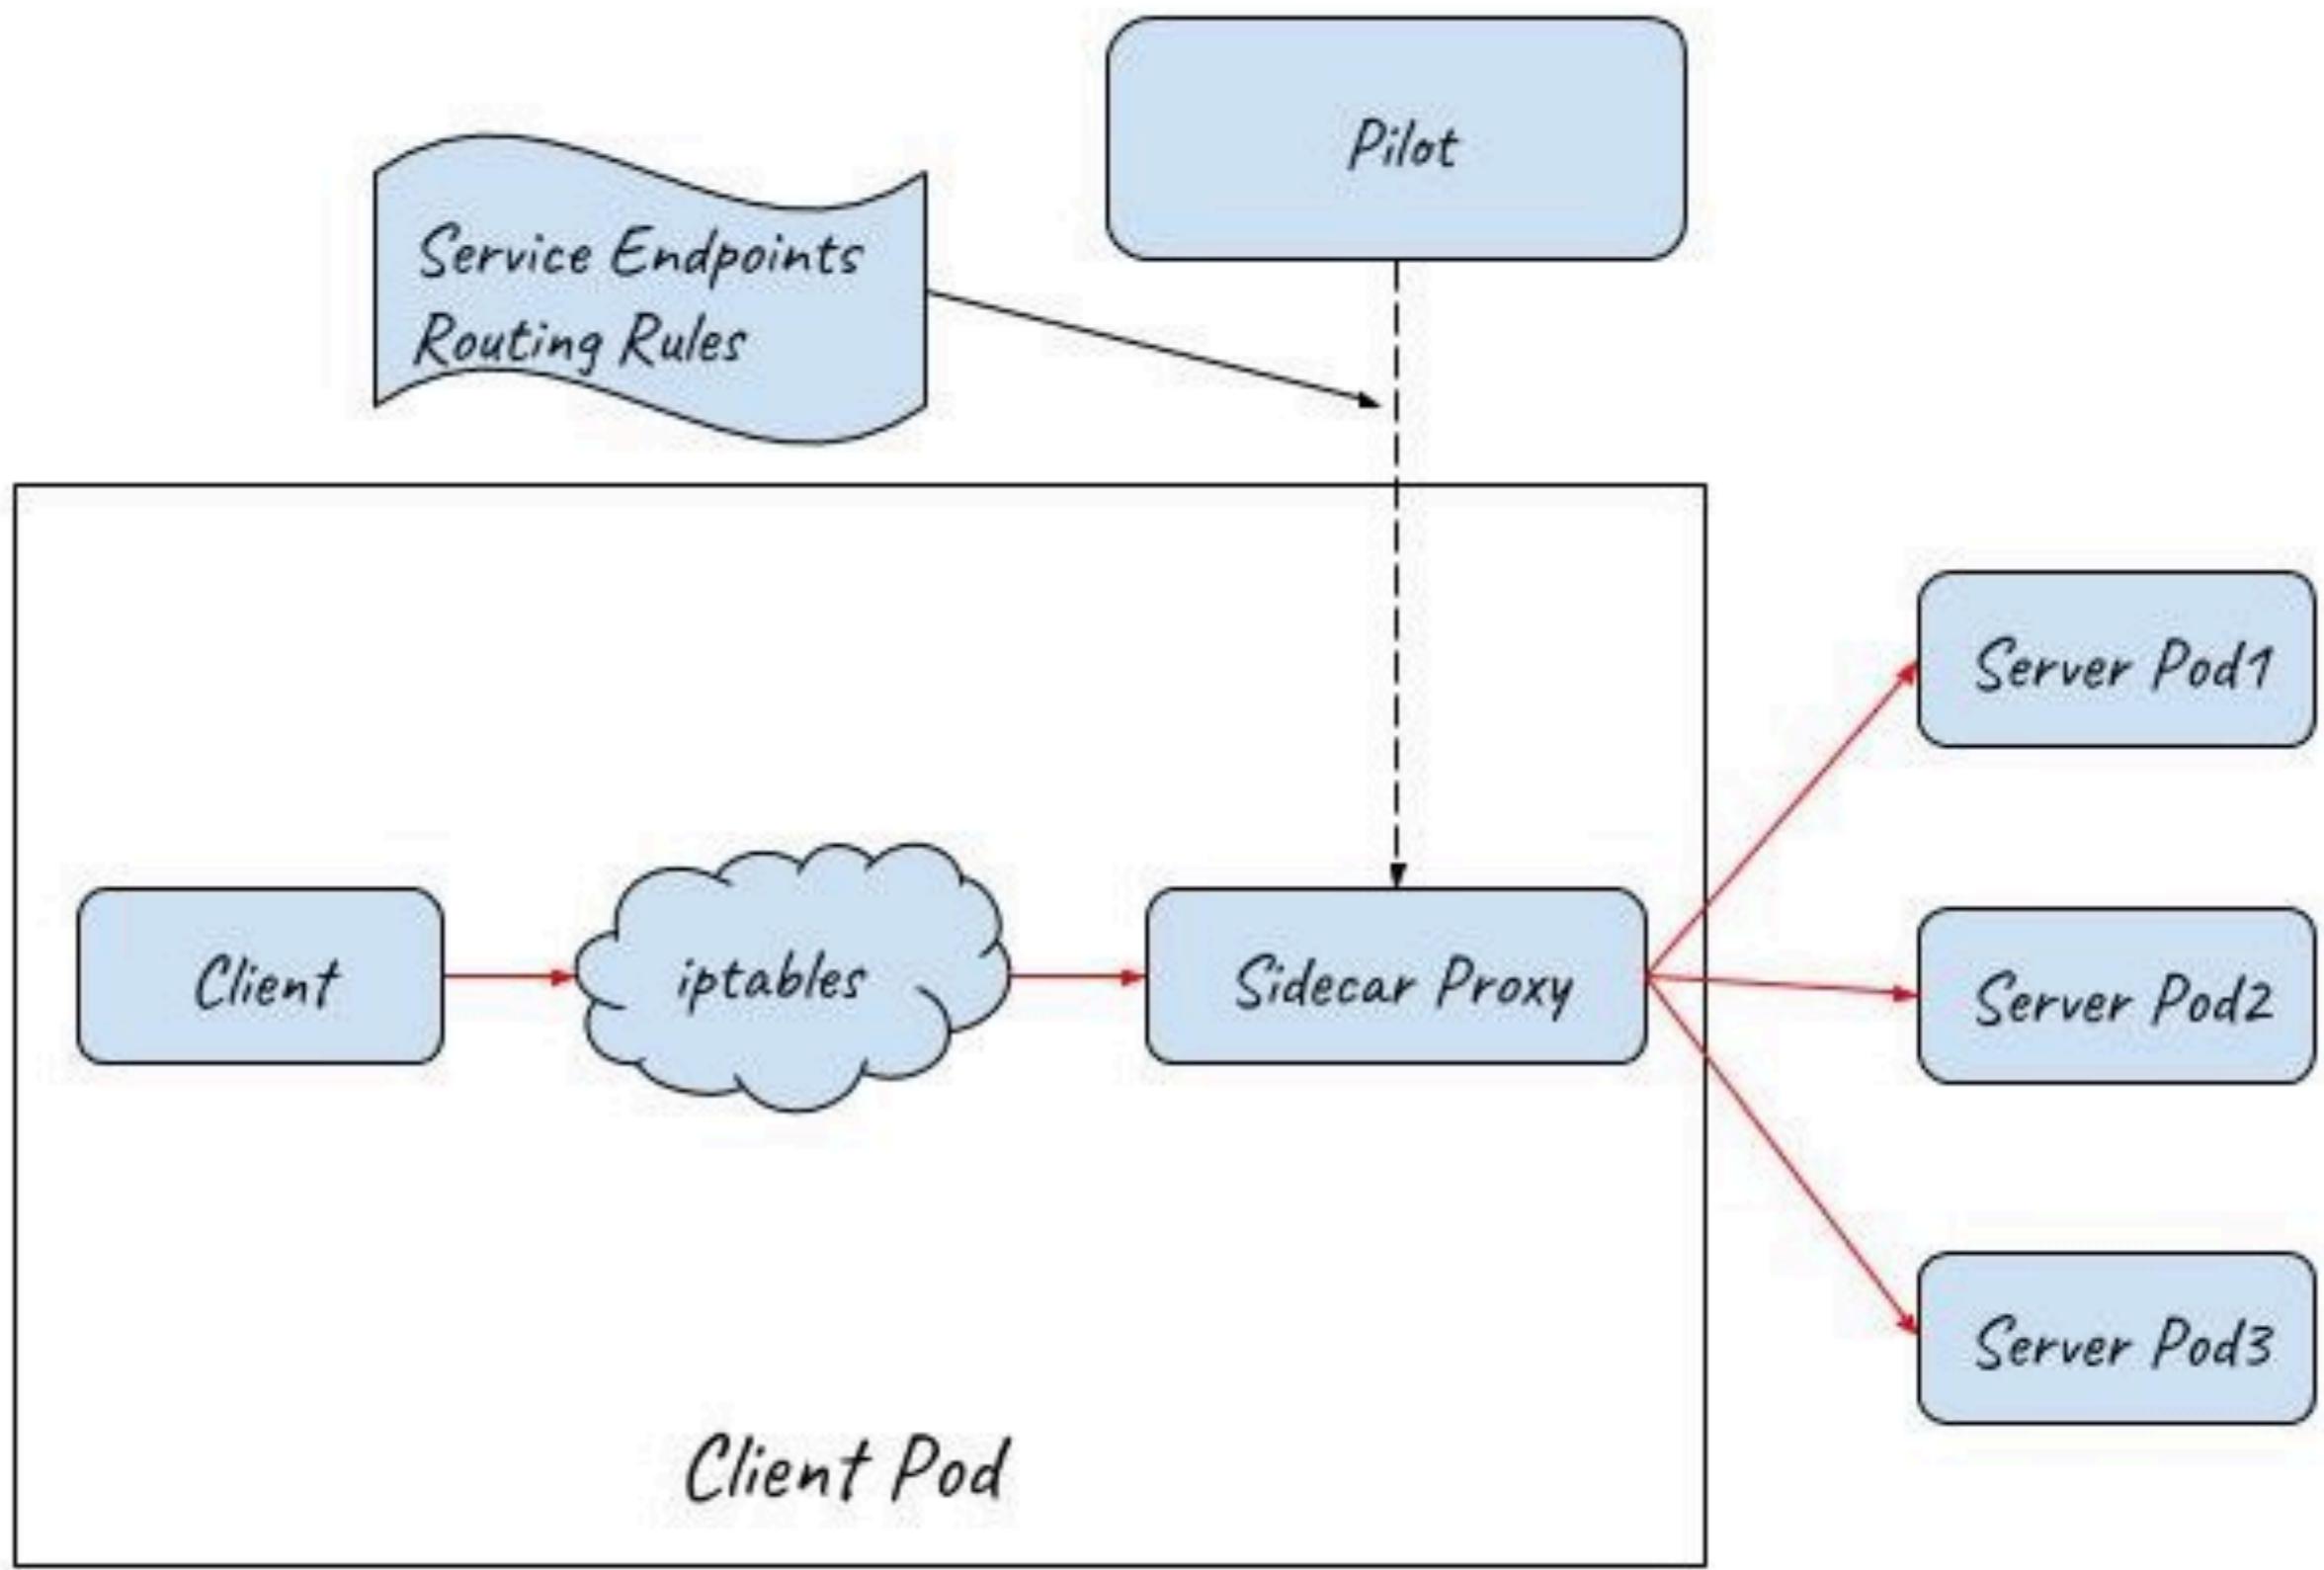
Kubernetes Ingress resource is used to declare an OSI layer 7 load balancer, which can understand HTTP protocol and dispatch inbound traffic based on the HTTP URL or Host.

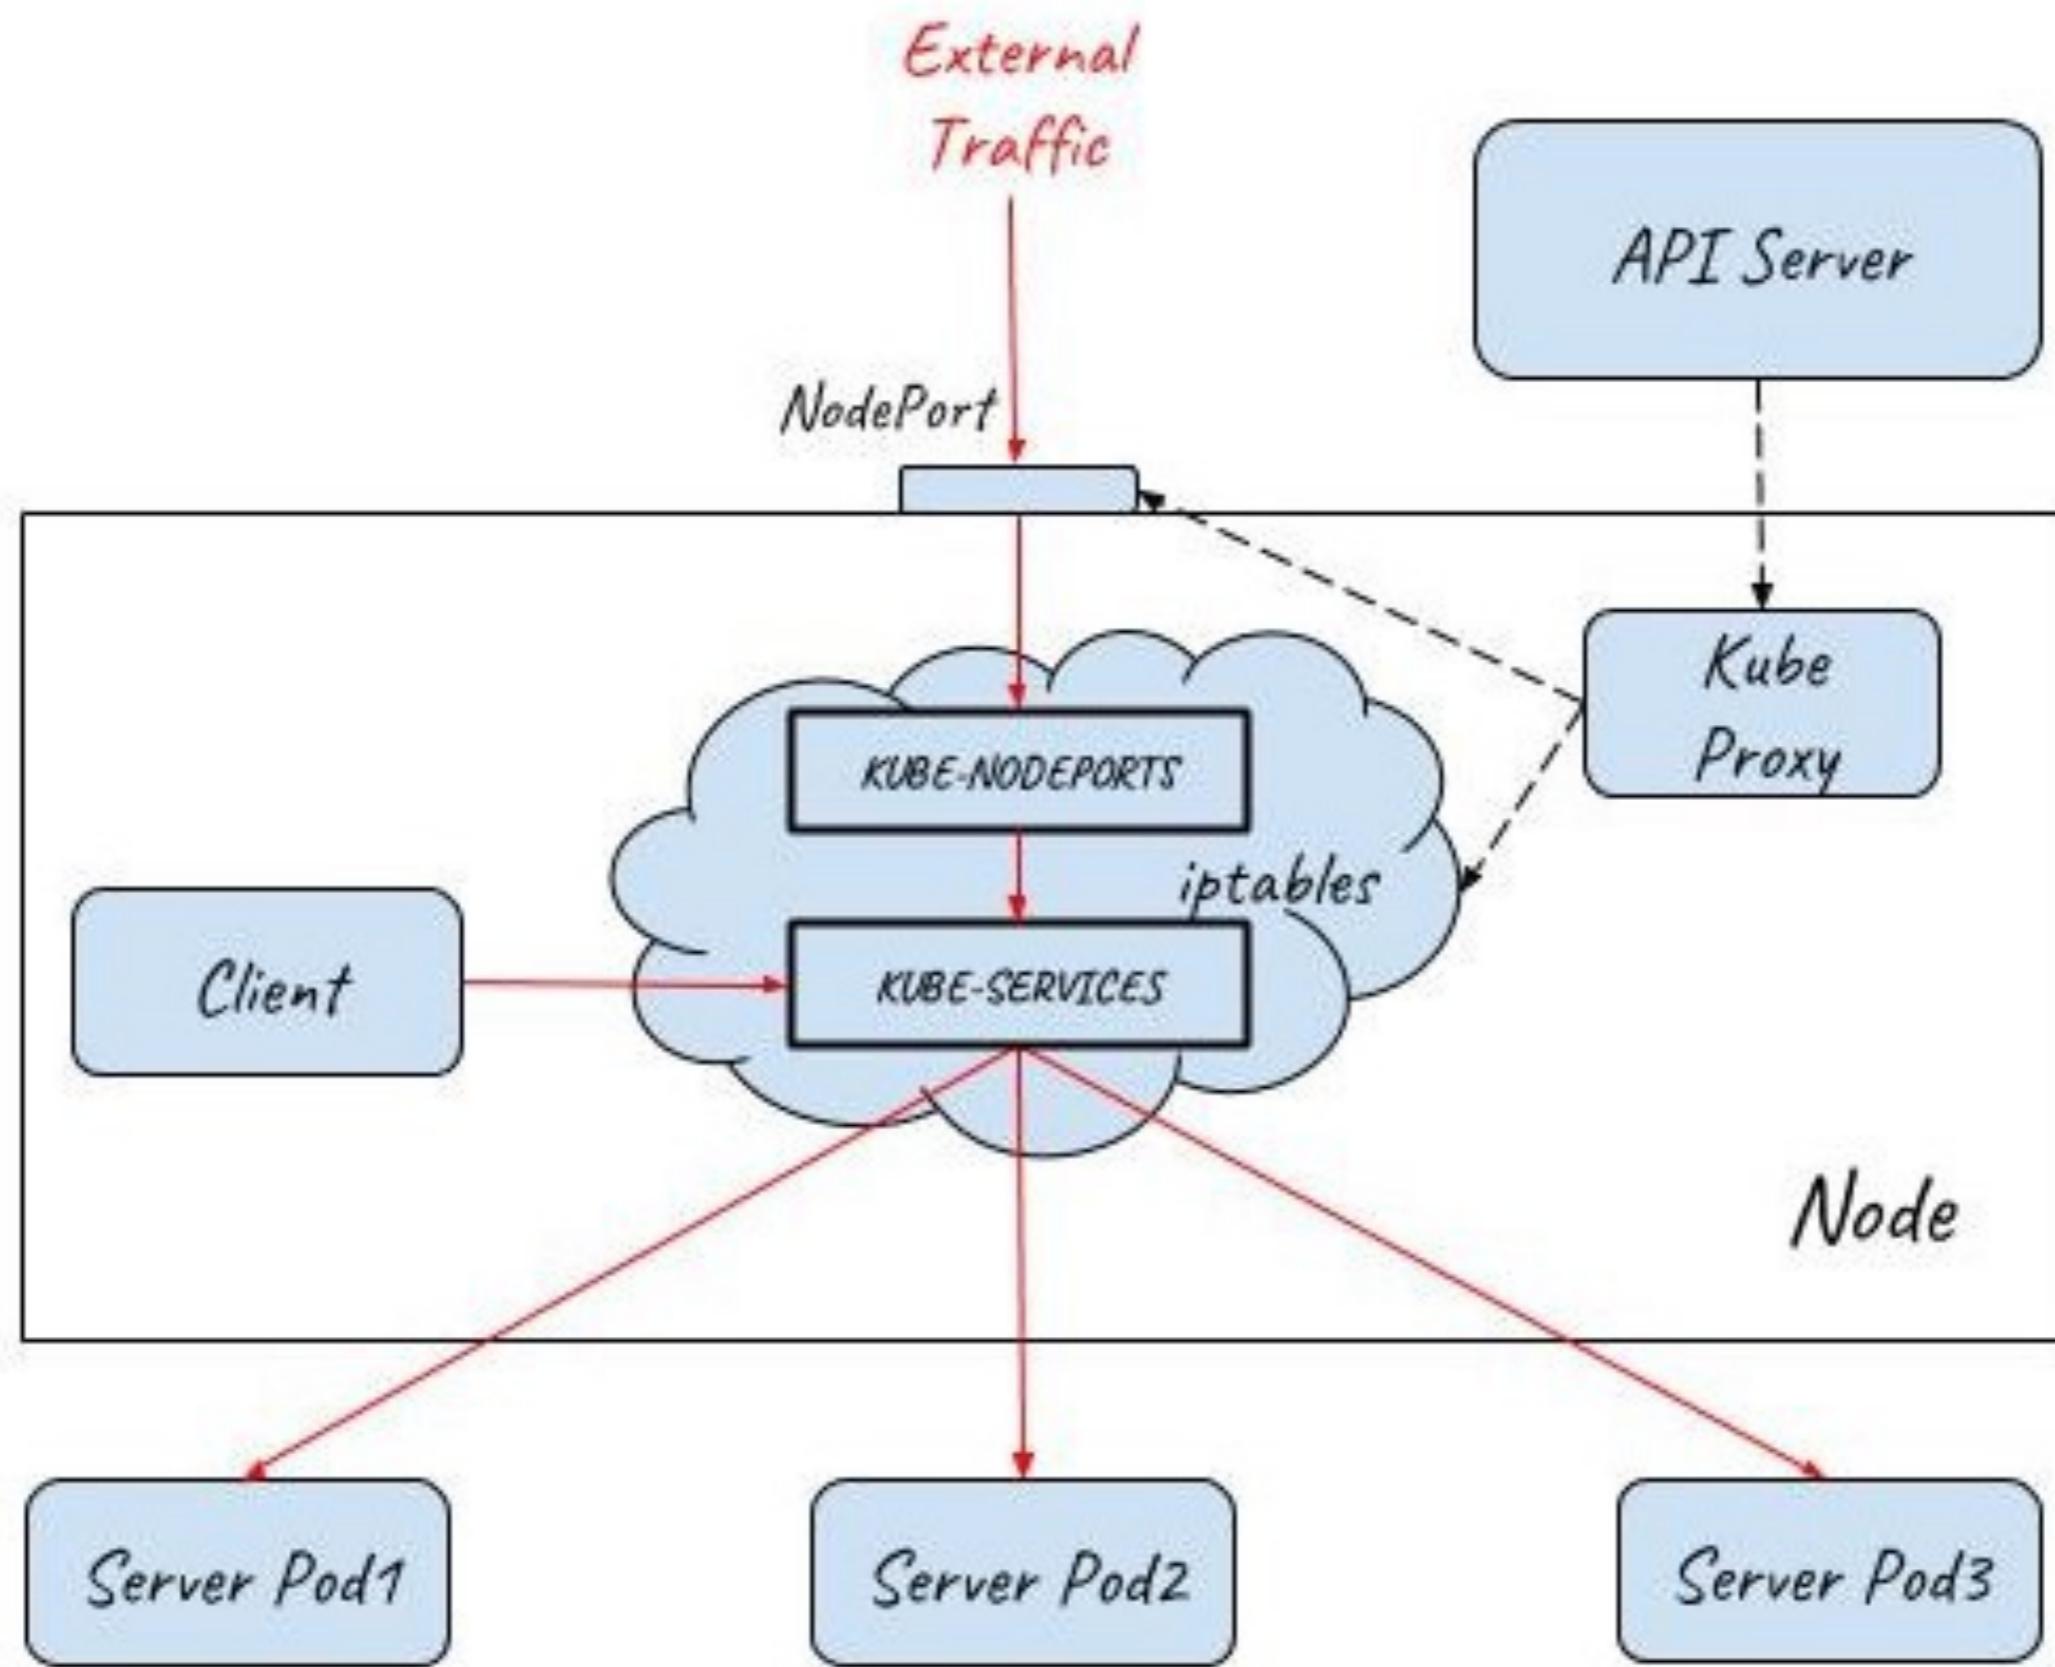
- Ingress controller itself is also deployed as Pods inside the cluster.
- Ingress controller provides a unified entrance for the HTTP services in a cluster, but it can't be accessed directly from outside because the ingress controller itself is also deployed as Pods inside the cluster.
- Ingress controller must work together with NodePort and LoadBalancer to provide the full path for the external traffic to enter the cluster.

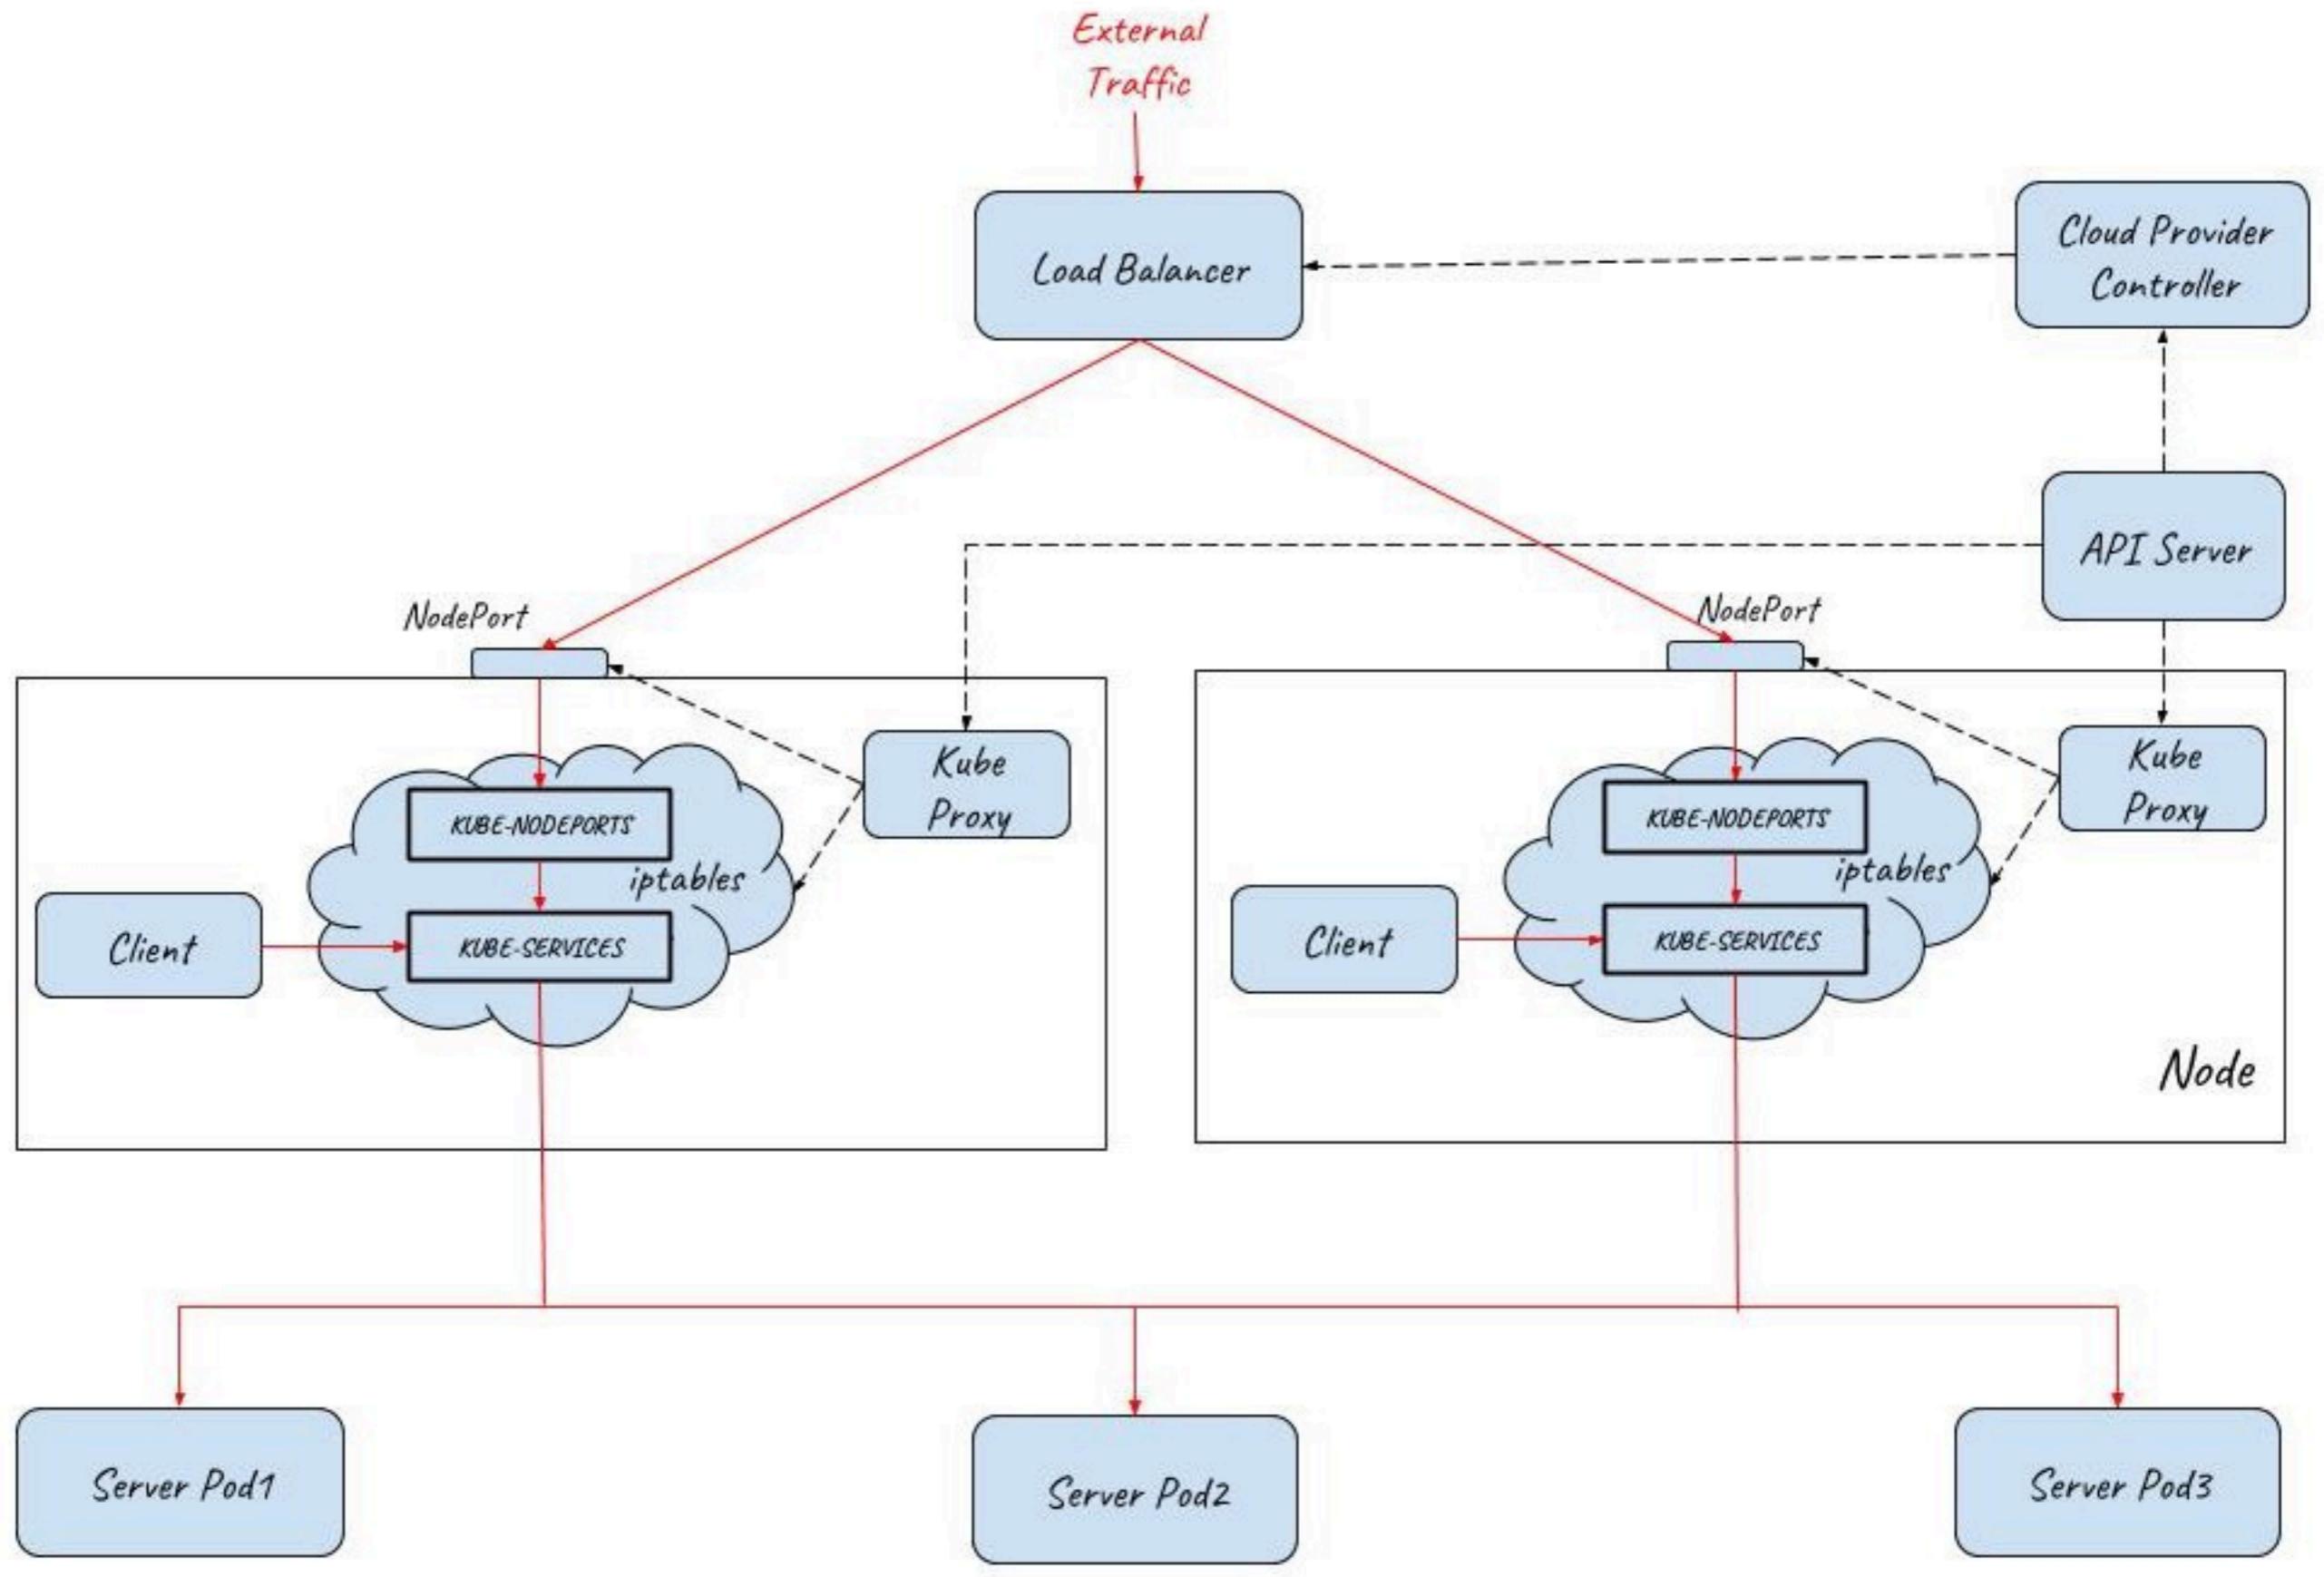


	OSI Layer	TCP/IP	Datagrams are called
Software	Layer 7 Application	HTTP, SMTP, IMAP, SNMP, POP3, FTP	Upper Layer Data
	Layer 6 Presentation	ASCII Characters, MPEG, SSL, TSL, Compression (Encryption & Decryption)	
	Layer 5 Session	NetBIOS, SAP, Handshaking connection	
	Layer 4 Transport	TCP, UDP	Segment
	Layer 3 Network	IPv4, IPv6, ICMP, <u>IPSec</u> , MPLS, ARP	Packet
Hardware	Layer 2 Data Link	Ethernet, 802.1x, PPP, ATM, <u>Fiber</u> Channel, MPLS, FDDI, MAC Addresses	Frame
	Layer 1 Physical	Cables, Connectors, Hubs (DLS, RS232, 10BaseT, 100BaseTX, ISDN, T1)	Bits

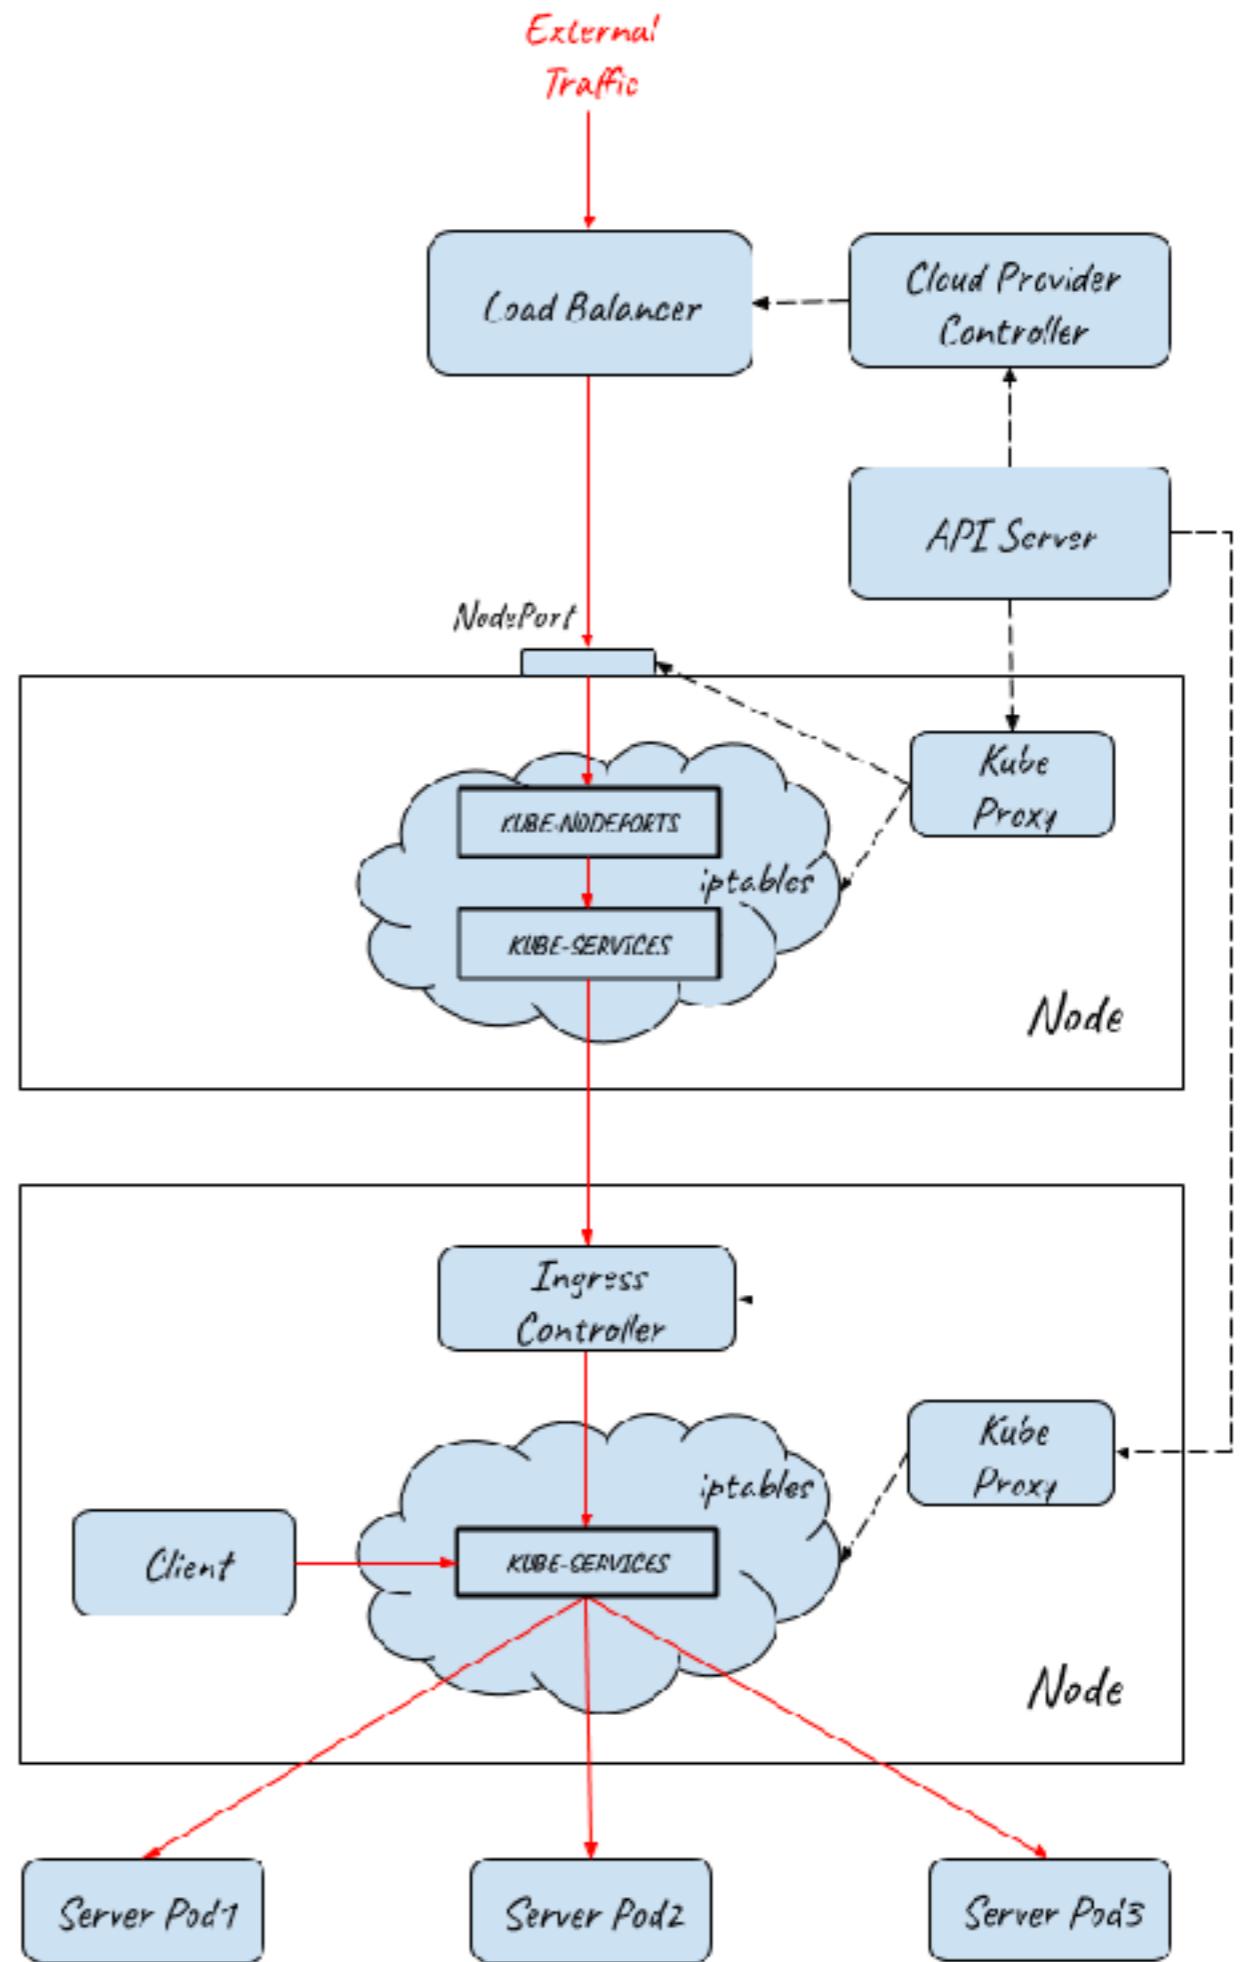








1. Internet/External traffic reaches the layer 4 load balancer.
2. Load balancer dispatches traffic to multiple NodePorts on the Kubernetes.
3. Traffic is captured by iptables and redirected to ingress controller Pods.
4. Ingress controller sends traffic to different Services according to ingress rules.
5. Finally, traffic is redirected to the backend Pods by iptables.



Client Request

- Load Balancer(External IP)
- Load Balancer (Node IP)
- Ingress Controller Service(ClusterIP)
- Ingress Controller Pod(Pod IP)
- Backend Service(ClusterIP)
- Backend Pod(Pod IP)

A searchable catalog of Customers

A searchable catalog of Freelancer

A searchable catalog of Projects

Timesheets for the Freelancers

Make Payments to the Freelancers

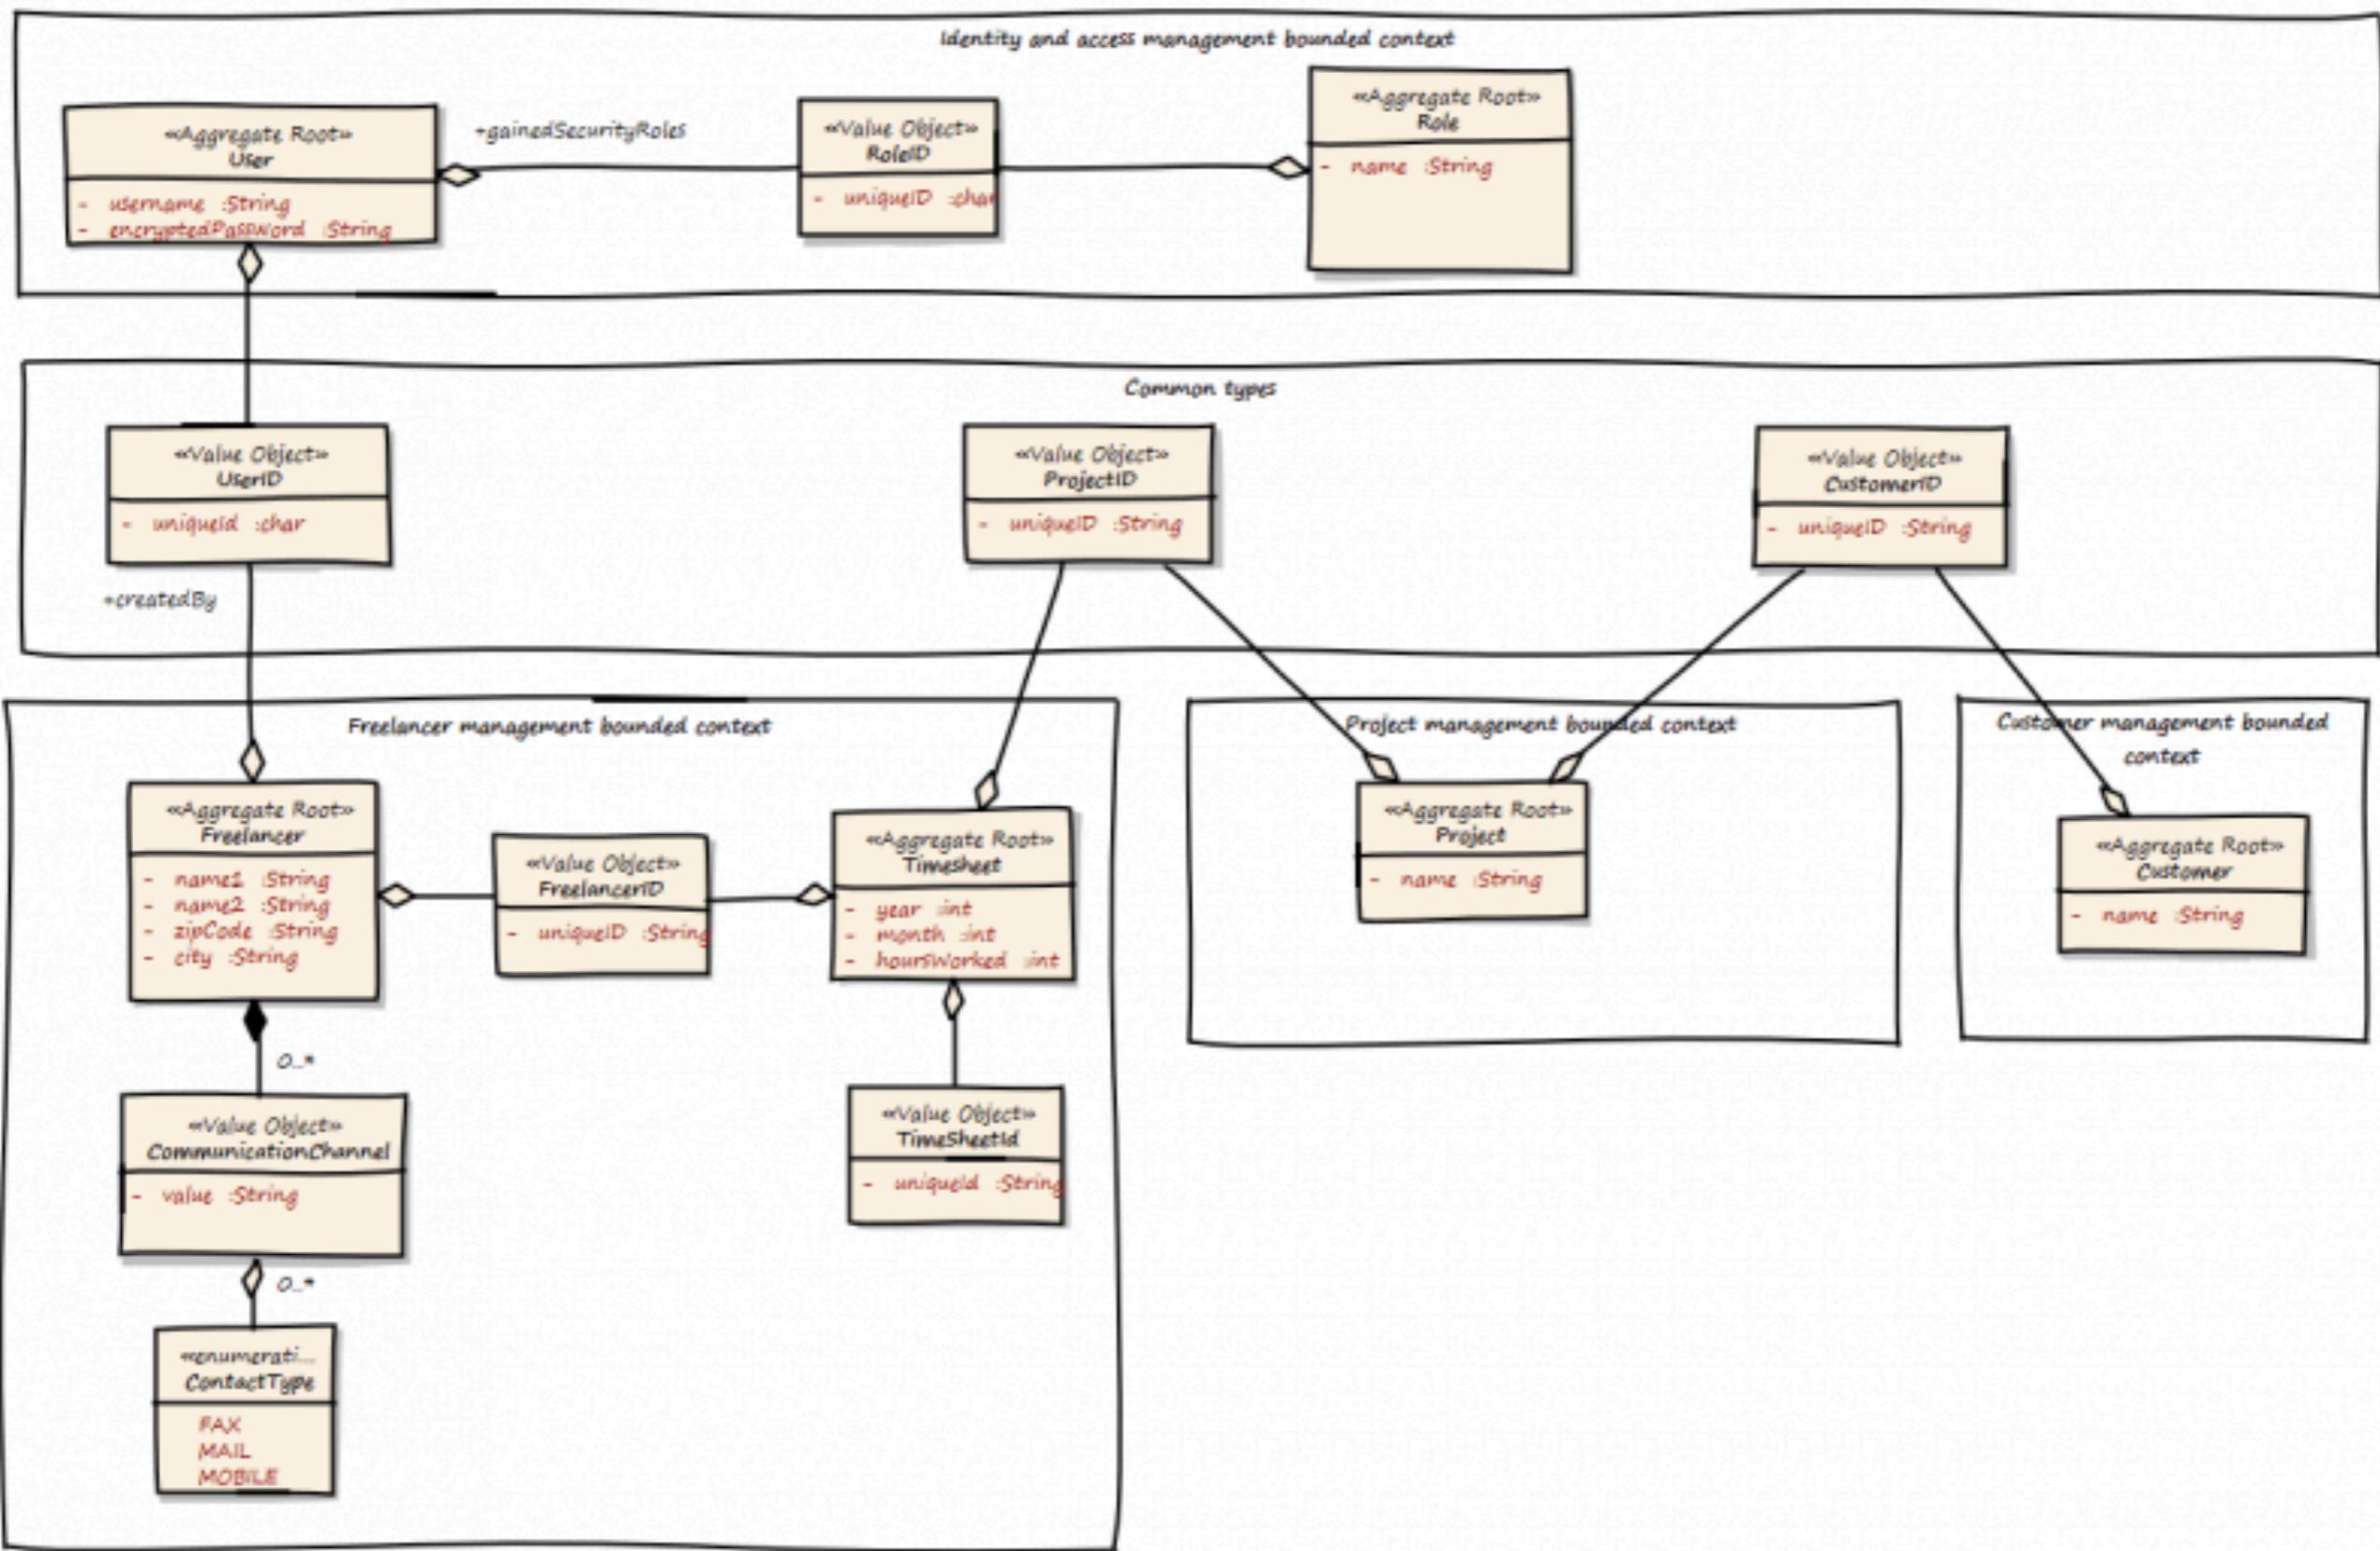
Body Leasing Domain

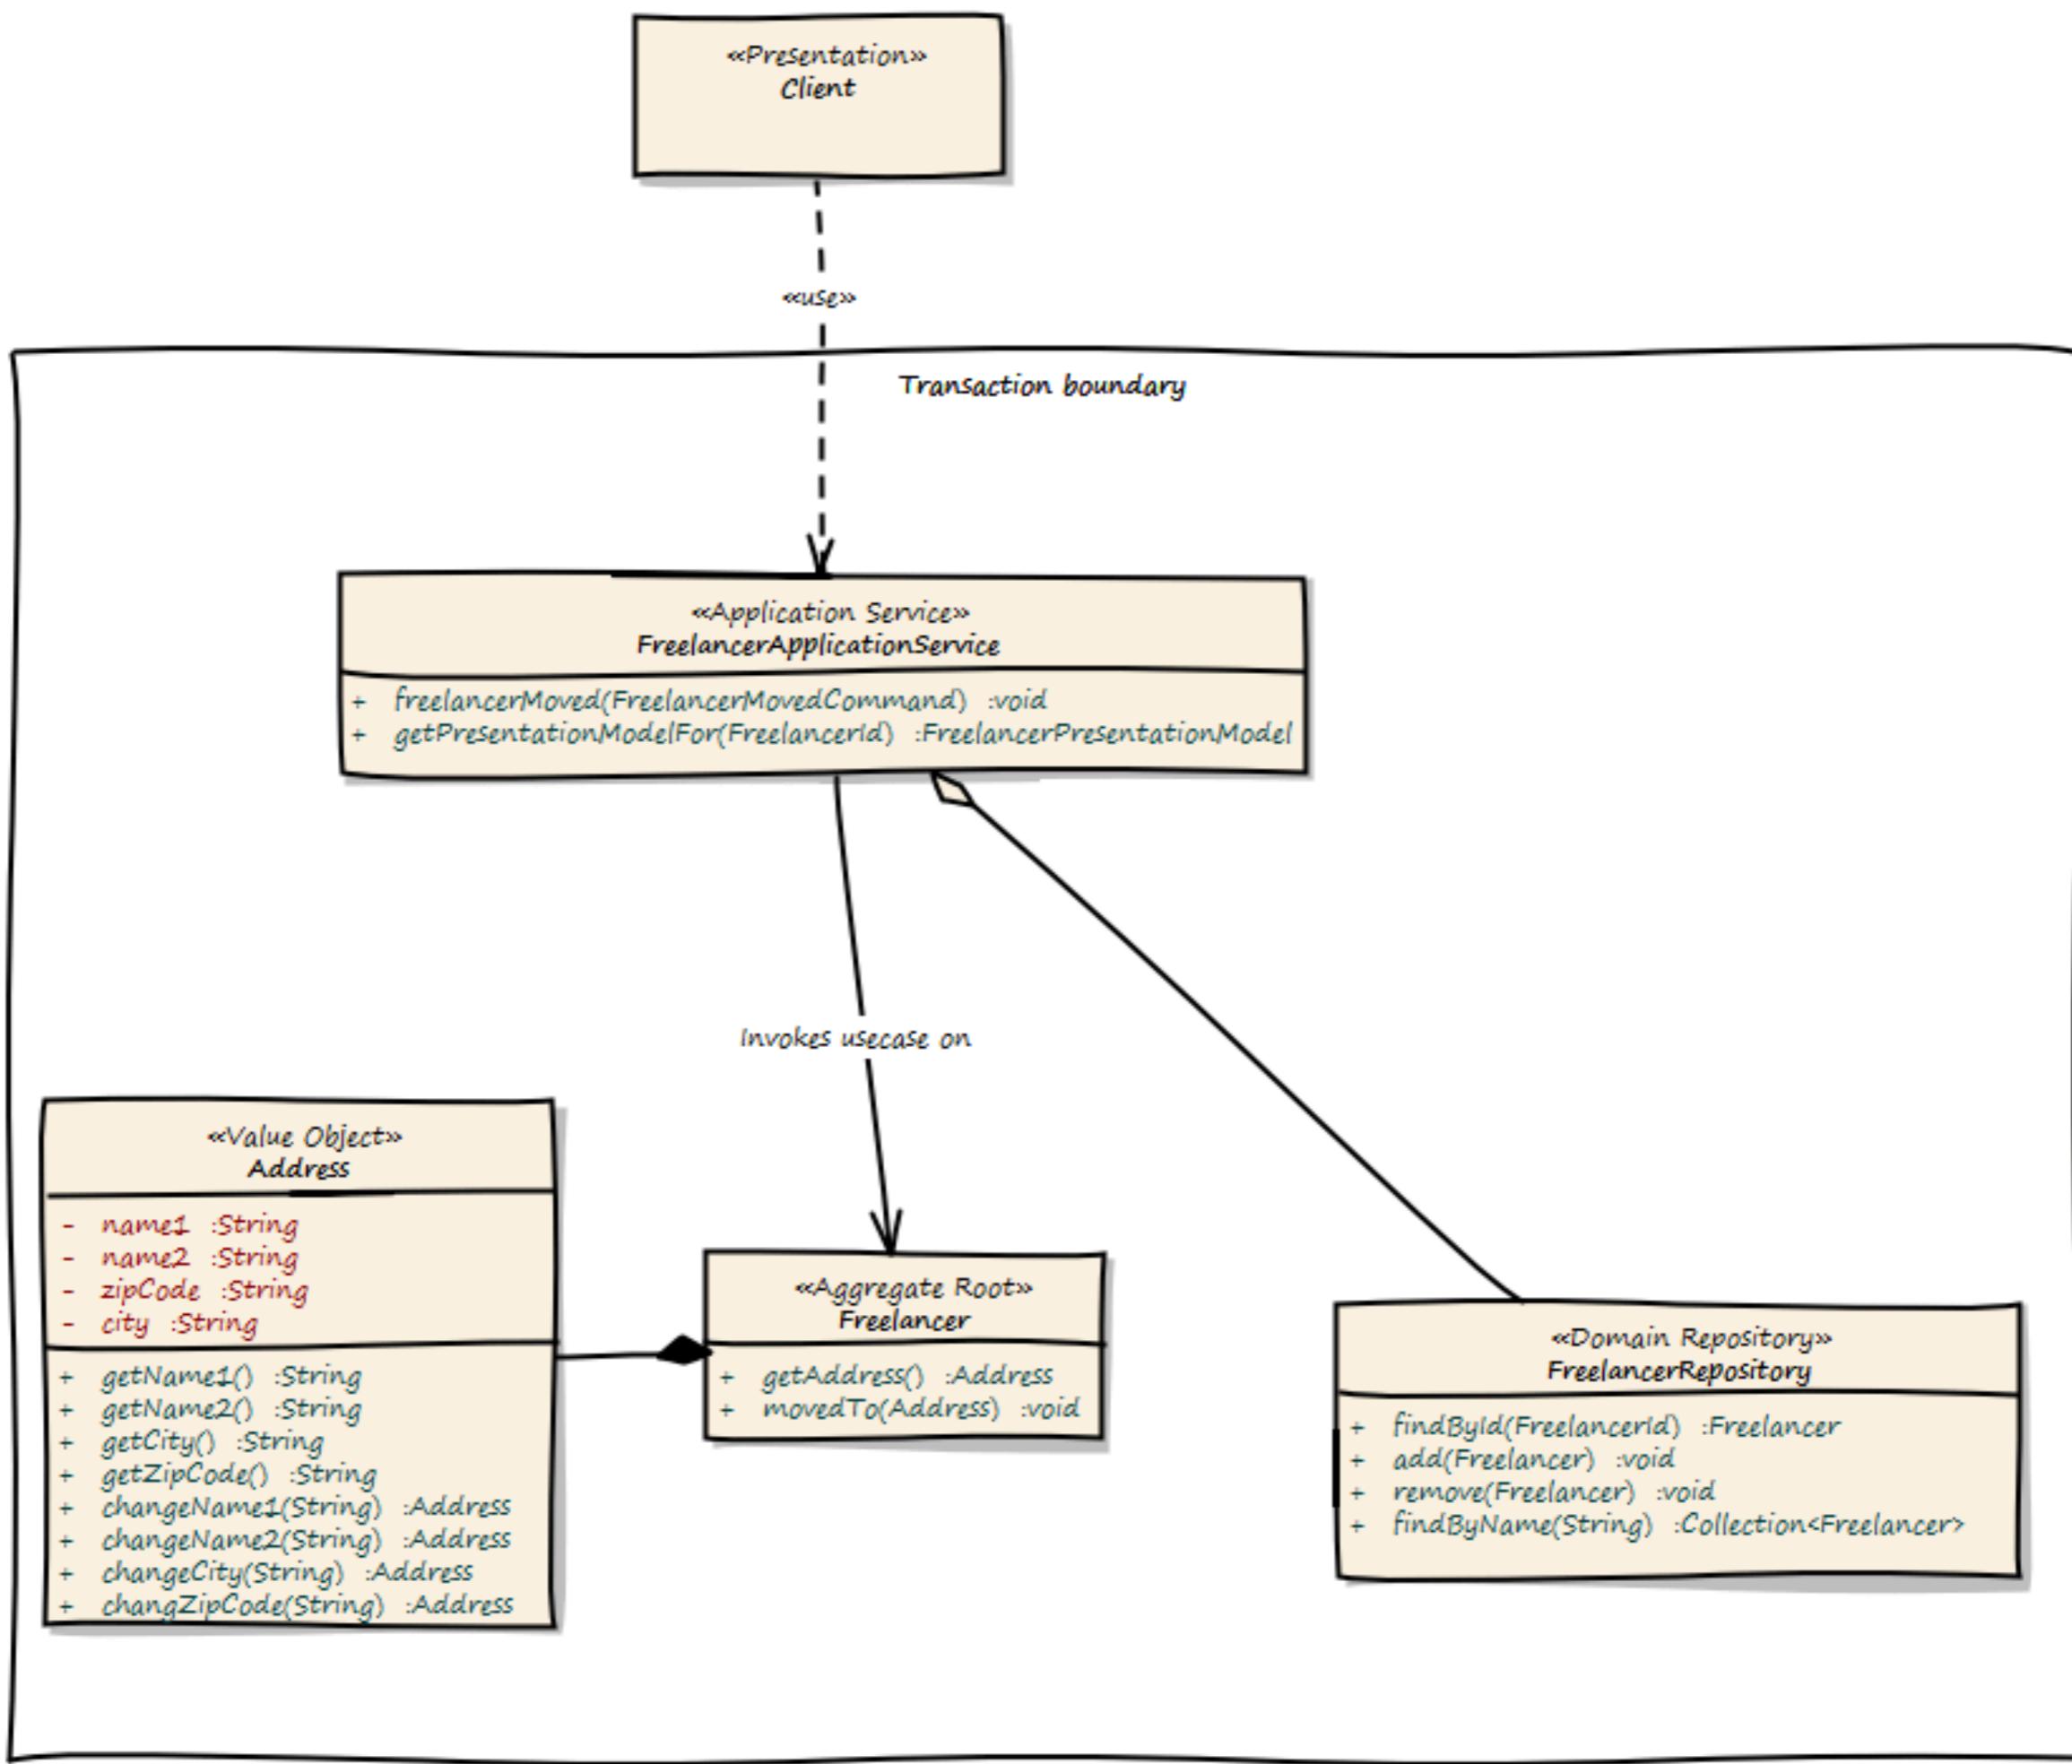
Identity and
Access
Management
Subdomain

Freelancer
Management
Subdomain

Customer
Management
Subdomain

Project
Management
Subdomain





<https://www.mirkosertic.de/blog/2013/04/domain-driven-design-example/>

Python flask todo api

C1

27017

27017

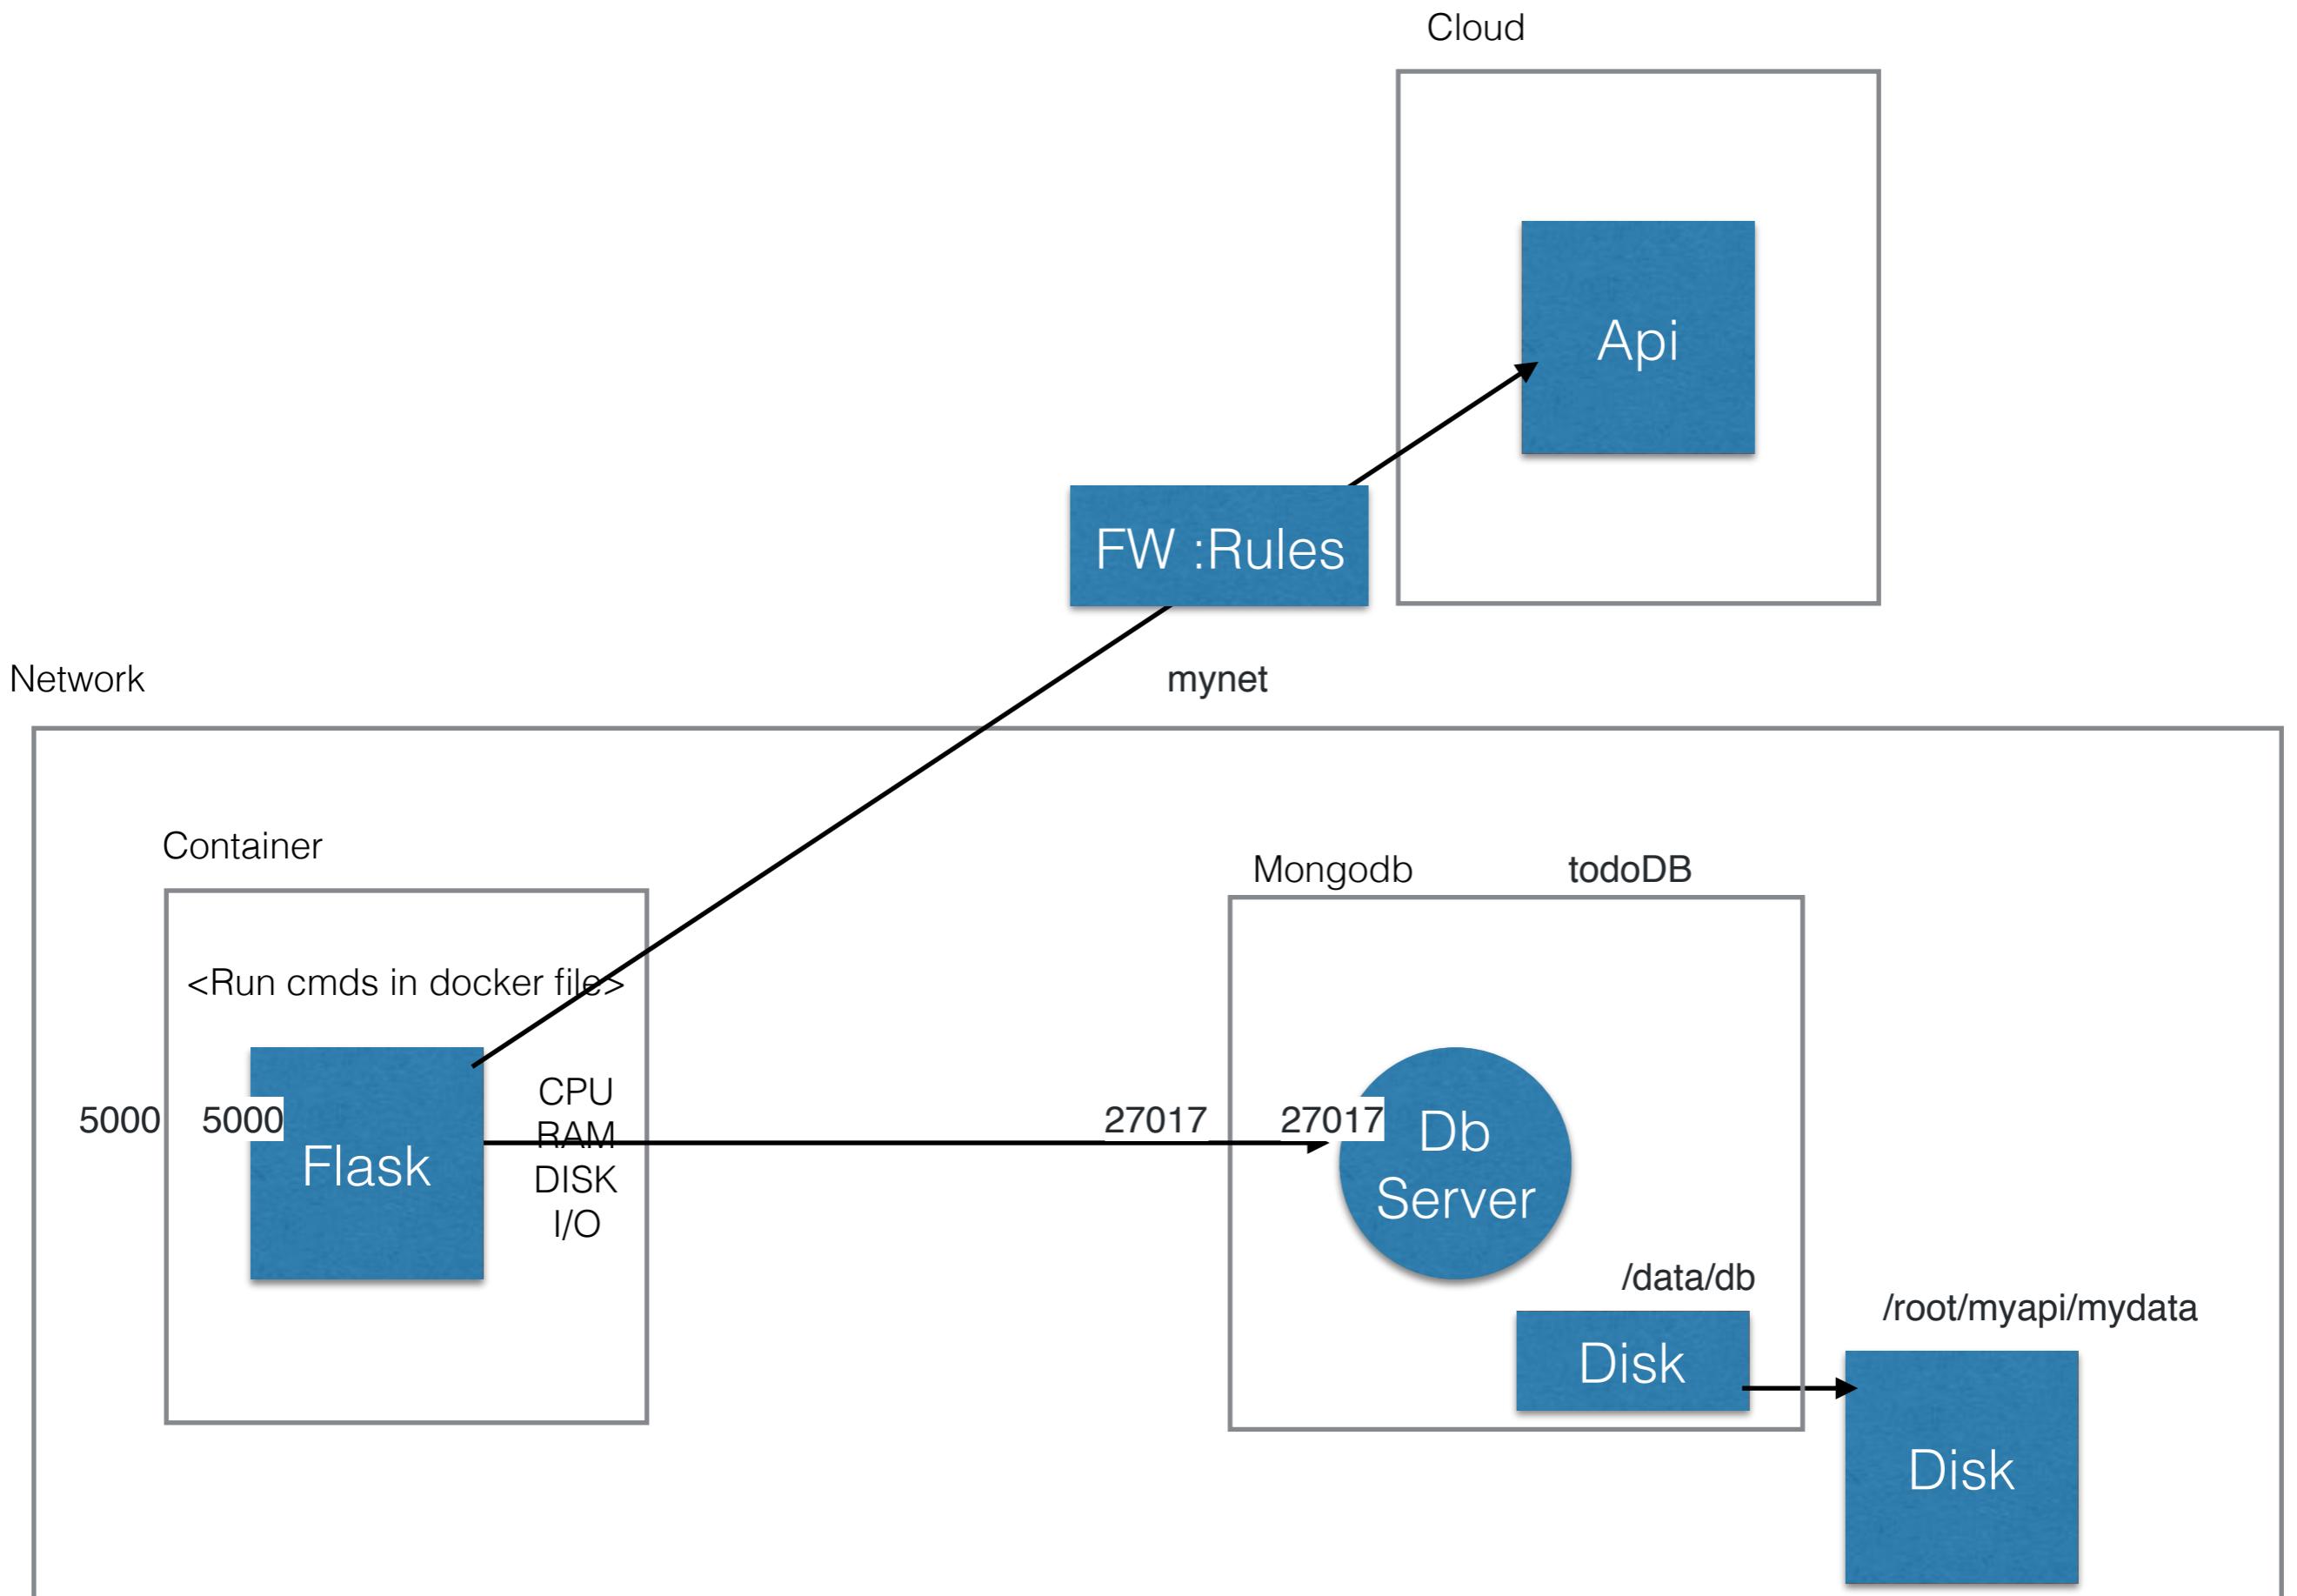
Mongodb

C2

/data/db

/root/myapi/mydata

mynet



- Docker File (cmd to execute in the deployed env)
- Image (installer/ setup/war/ package)
- Container (small vm)
-

